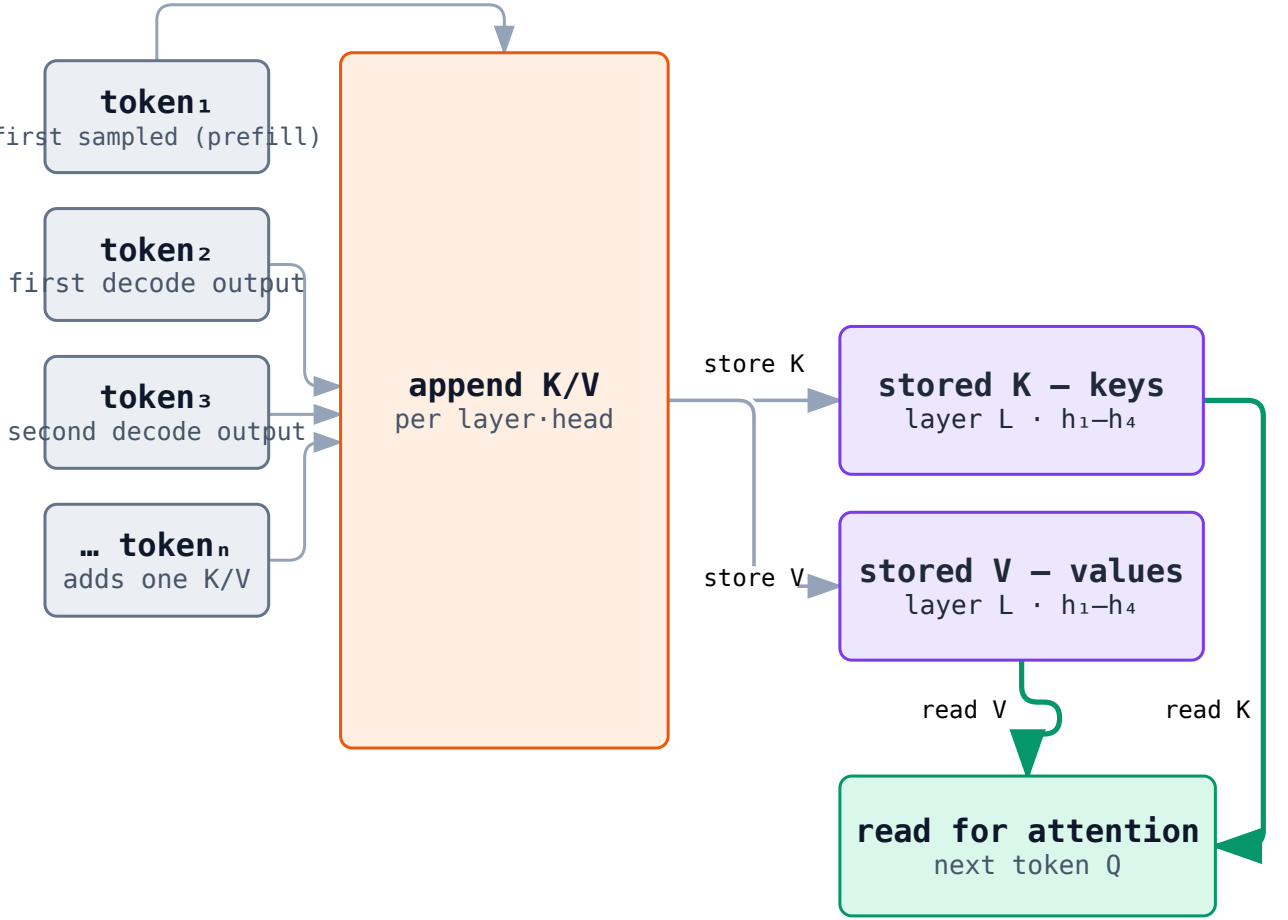
token₁
first sampled (prefill)
token₂
first decode output
token₃
second decode output
… tokenₙ
adds one K/V
append K/V
per layer·head
store K
store V
stored K — keys
layer L · h₁–h₄
stored V — values
layer L · h₁–h₄
read V
read K
read for attention
next token Q