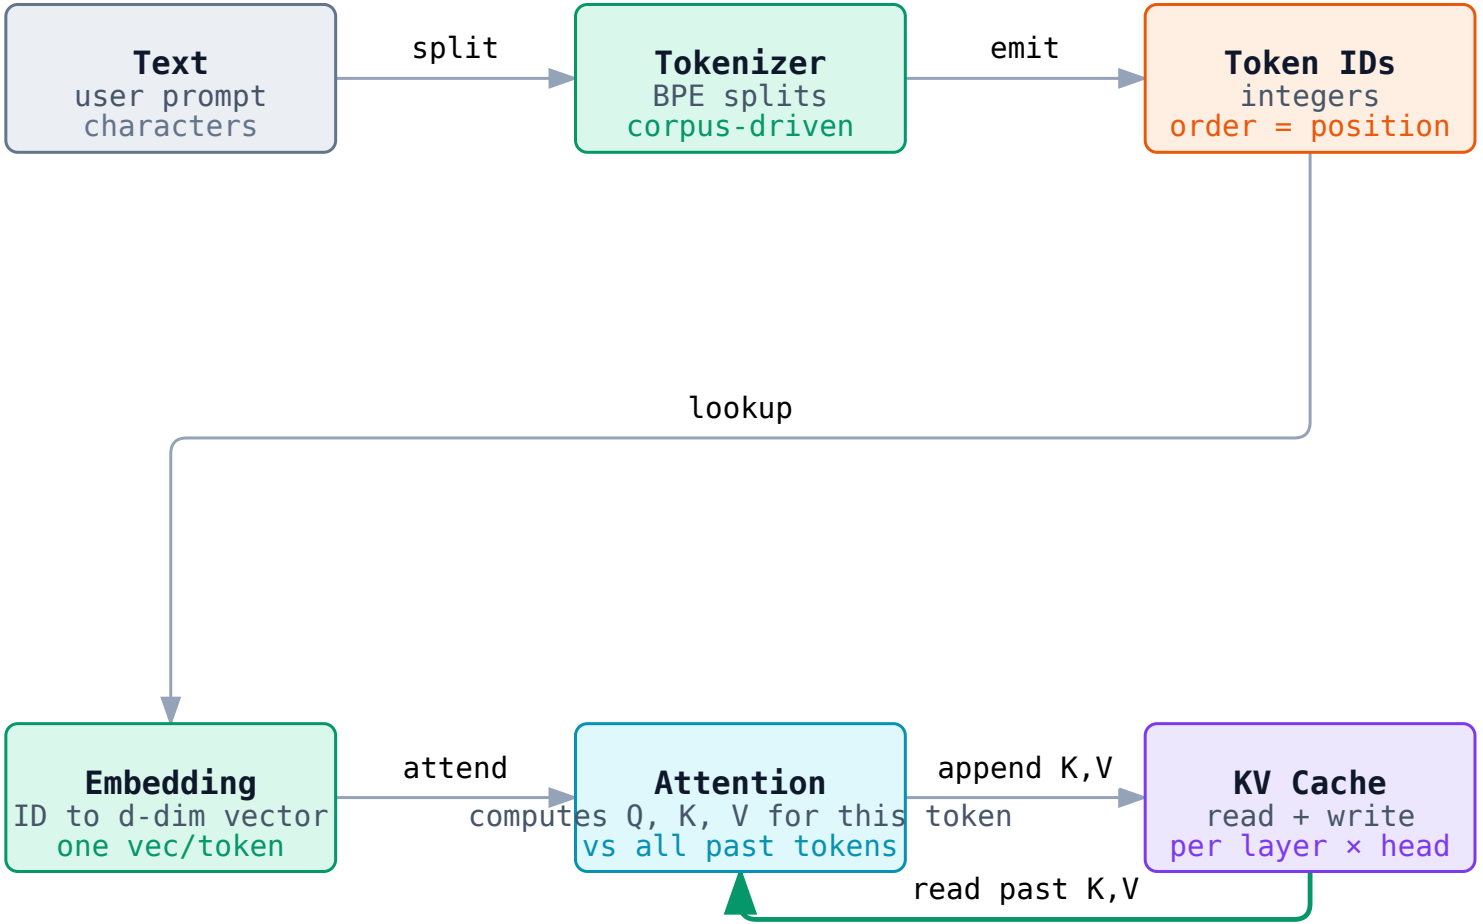
Text
user prompt
characters
split
Tokenizer
BPE splits
corpus-driven
emit
Token IDs
integers
order = position
lookup
Embedding
ID to d-dim vector
one vec/token
attend
Attention
computes Q, K, V for this token
vs all past tokens
append K,V
KV Cache
read + write
per layer × head
read past K,V