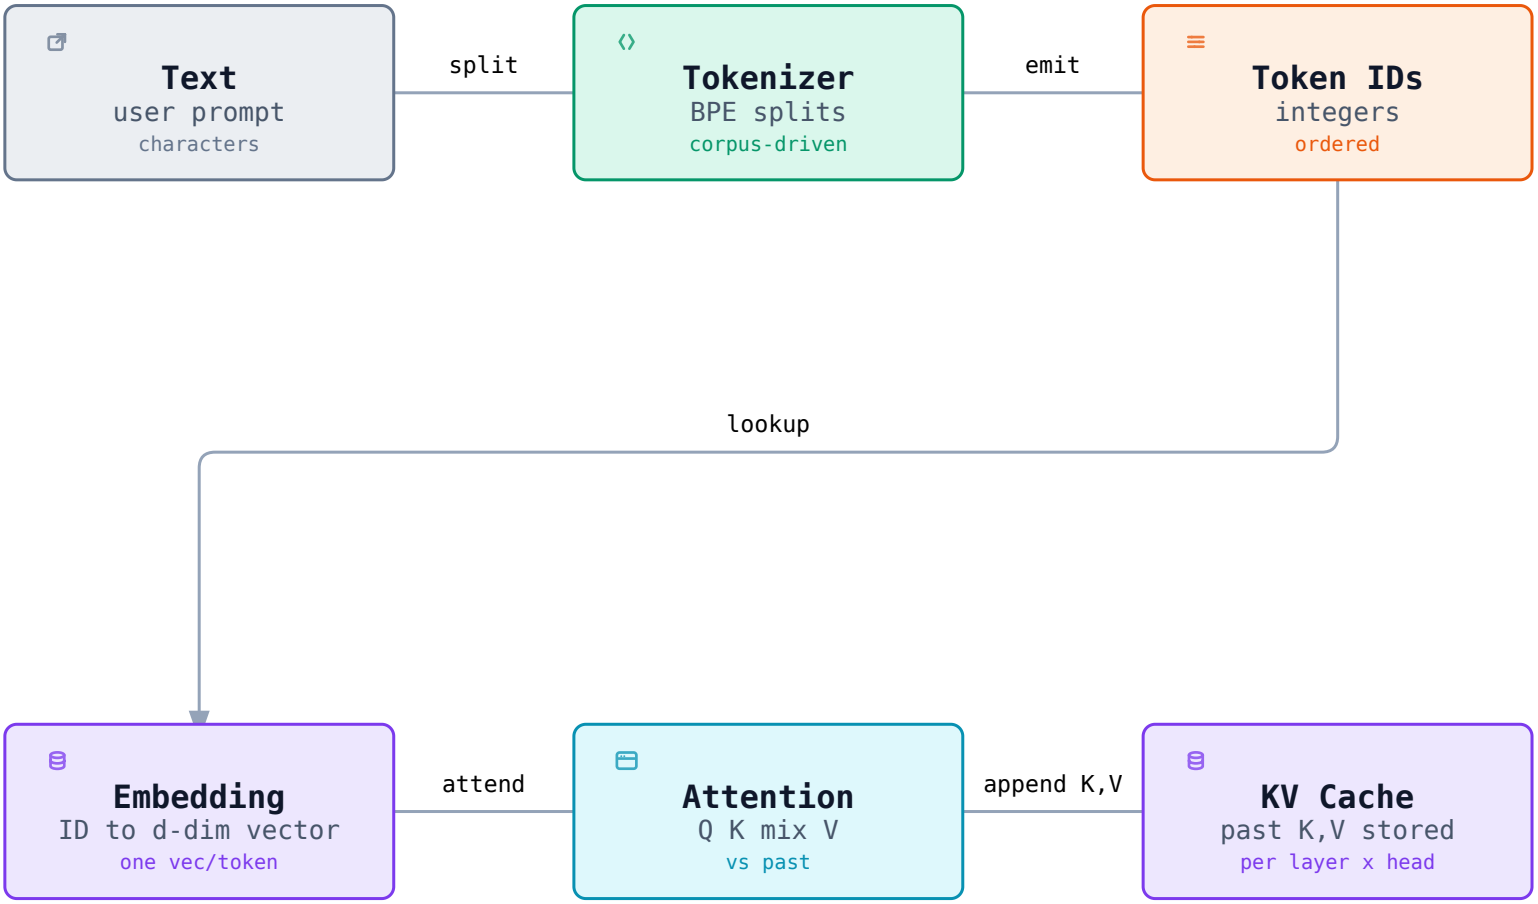
Text
user prompt
characters
split
Tokenizer
BPE splits
corpus-driven
emit
Token IDs
integers
ordered
lookup
Embedding
ID to d-dim vector
one vec/token
attend
Attention
Q K mix V
vs past
append K,V
KV Cache
past K,V stored
per layer x head