

01 / Source

📄 **Text**
user prompt
characters

02 / Tokenize

🔤 **Tokenizer**
BPE splits
corpus-driven

split
raw string

emit
integer IDs

🏷️ **Token IDs**
integers from vocab
ordered

lookup
dense vector

03 / Embed

📌 **Embedding**
ID → d-dim vector
one vec/token

attend
queries past K,V

04 / Attend

🔍 **Attention**
Q·K mix V
vs past

append K,V
remember prefix

05 / Remember

📌 **KV Cache**
past K,V stored
per layer × head

