



PREFILL — process the whole prompt, populate the KV cache, emit the first token



DECODE — generate one token at a time, reusing the cached K/V



visual grammar:

- grey box data / tokens
- blue box model / compute
- orange box KV state (persistent)