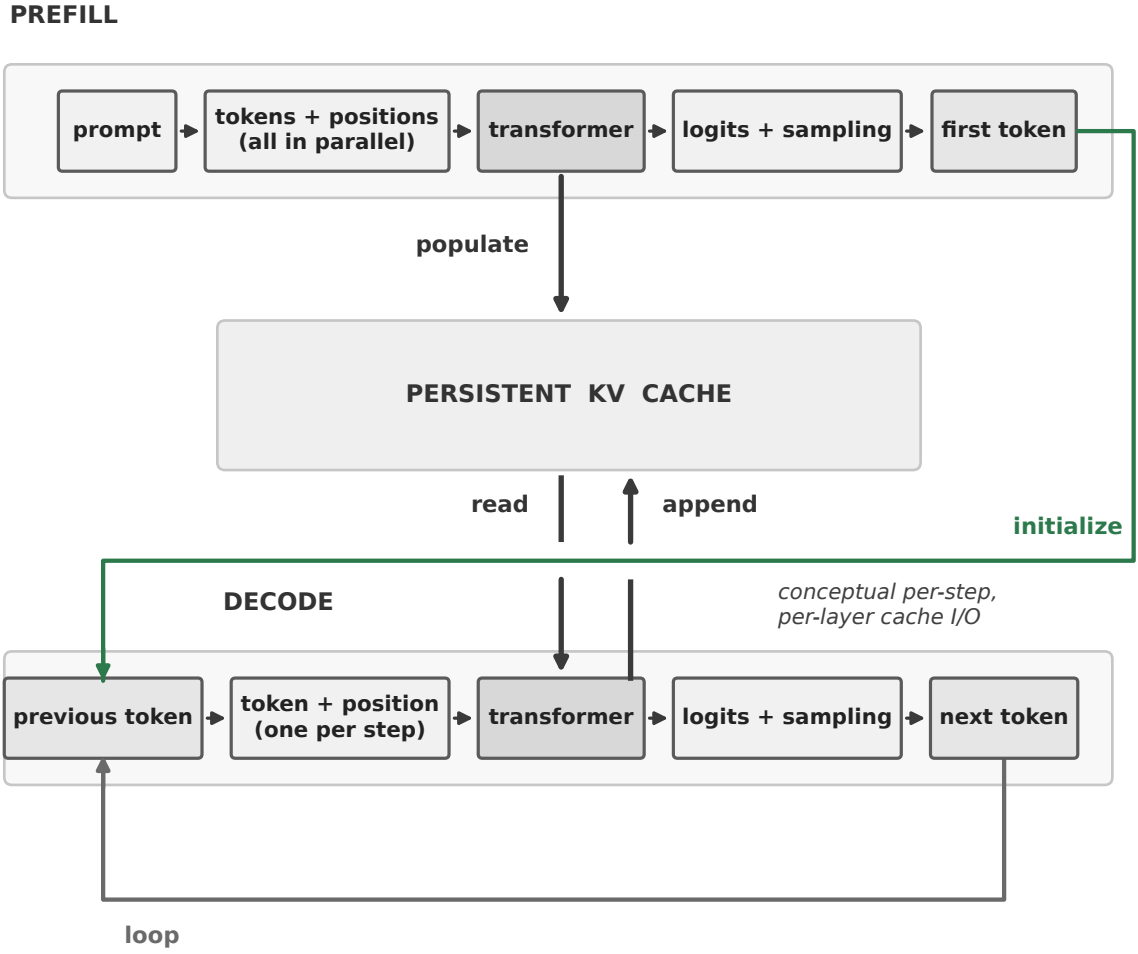
PREFILL
prompt
tokens + positions (all in parallel)
transformer
logits + sampling
first token
populate
PERSISTENT KV CACHE
read
append
initialize
DECODE
conceptual per-step, per-layer cache I/O
previous token
token + position (one per step)
transformer
logits + sampling
next token
loop