


Dense 70B



ALL 70B params active

every token


resident weights $\approx$ 140 GB

active per token $\approx$ 140 GB

KV grows linearly with context

KV same as dense

MoE (8 experts)



top-2 active $\rightarrow$ $\sim$ 14B

resident weights $\approx$ full expert set

active per token $\approx$ 28 GB

KV same as dense