


## Dense 70B

ALL active



■ resident  $\approx$  140 GB

■ active per token  $\approx$  140 GB


every token activates all 70B params

KV grows linearly with context

same as dense (matched atten

## MoE (8 experts)

top-2 active



■ resident  $\approx$  full expert set

■ active per token  $\approx$  28 GB

only top-2 of 8 experts compute per token

KV grows linearly with context

same as dense (matched atten