Dense 70B

ALL active

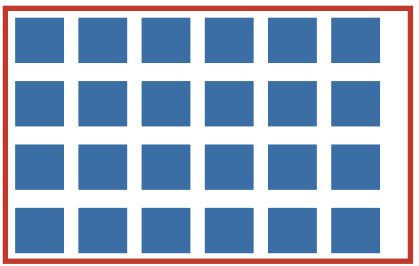

all 140 GB resident & active

every token activates all 70B params

KV grows linearly with context

KV: identical (dense = MoE)

controlled: same attention geometry, context length, and cache dtype

MoE (8 experts)

top-2 active

resident weights ≈ 94 GB

per-token active-weight bytes ≈ 25.6 GB (12.8B × 2 bytes)

per token: 11.4B top-2 + 1.4B always-on

KV grows linearly with context

KV: identical (dense = MoE)

controlled: same attention geometry, context length, and cache dtype