


Dense 70B



ALL 70B params active

every token

resident weights ≈ 140 GB

active per token ≈ 140 GB

KV grows linearly with context

same KV cache (attention is still dense)

MoE (8 experts)



top-2 active $\rightarrow \sim 14B$

resident weights $\approx$ full expert set

active per token ≈ 28 GB

KV grows identically to dense

same KV cache (attention is still dense)