

**Workload  
Characterization**



Capability / quality



Latency / compute



Context / KV cost



Cost / token



RAG vs end-to-end

**Model selection surface (driven by workload, not vice-versa)**