

Memory floor: inference vs fine-tuning (70B)

full fine-tune exceeds an 8×H100 host:
needs multi-node or offload

