

Memory floor: inference vs fine-tuning (70B)

full fine tune exceeds an 8xH100 host:

needs multi-node or offload

GPU memory floor (GB)

