


Memory floor: inference vs fine-tuning (70B)

full fine tune exceeds an 8xH100 host:

needs multi-node or offload

GPU memory floor (GB)



~1120 GB

~165 GB

~60 GB

QLoRA fits a single H100

Inference (weights+KV) Full fine-tune (weights+grad+opt+mem) QLoRA (weights+adapters)

