

Memory floor: inference vs fine-tuning (70B)

