


Memory floor: inference vs fine-tuning (70B)

full fine tune exceeds an 8xH100 host:

needs multi-node or offload

GPU memory floor (GB)



QLoRA fits a single H100

Inference (weights+KV) Full fine-tune (weights+grad+optimizer) QLoRA (weights+adapters)