


# Grouped-Query Attention (GQA): why the cache is 8× smaller

64 query heads (8 groups of 8) → 8 shared KV heads



MHA (full): 64 K/V per token → 2.62 MB/token

**GQA: 8 K/V per token → ~0.33 MB/token — 8× smaller cache**