


Grouped-Query Attention (GQA): why the cache is 8× smaller

64 query heads (8 groups of 8) → 8 shared KV heads



MHA (full): 64 K/V per token → 2.5 MB/token (worst case)

GQA: 8 K/V per token → ~0.33 MB/token — an 8× smaller cache (Ch. 7 note)