


Memory Tetris: how the 8×H100 host (640 GB) fills with context

640 GB pool = 8×H100 HBM



9.5K max = 229 GB used **32K context** = 288 GB used **128K context** = 539 GB used

baseline (weights + runtime) is context-independent; the KV cache is the lever that grows with context. FP16 ~2.5 MB/token [ILLUSTRATIVE][DERIVED].

-  KV cache (grows w/ context)
-  weights 140 GB
-  runtime / NCCL ~64 GB