


Memory Tetris: how the 8×H100 host (640 GB) fills as context grows



baseline (weights + runtime) is context-independent; the KV cache is the lever that grows with context. FP16 ~2.5 MB/token [2° DERIVED].