


Memory Tetris: how the 8×H100 host (640 GB) fills with context



baseline (weights + runtime) is context-independent;
the KV cache is the lever that grows with context.
FP16 ~2.5 MB/token [ILLUSTRATIVE][DERIVED].

- KV cache (grows w/ context)
- weights 140 GB
- runtime / NCCL ~64 GB