

The serving stack: batching trades latency, prefix cache skips prefill, P/D split trades memory & throughput



batching: latency↔throughput

KV transfer
(prefill→decode)

P/D split: memory↔throughput
KV passes between pools