



P/D Disaggregation: the two pools want opposite resources

input prompt →
↓



KV CACHE TRANSFER
over NVLink / fabric
prefill writes per-token KV; decode reads it
(shaded-dominant with Mooncake)



← tokens out to client

fill ~1.19 PFLOPS required vs 0.989 peak (FLOP-starved); decode 5.6 TB/s vs 3.35 TB/s (bandwidth-starved). One pool forces a compromise for both