



Single GPU

one accelerator · local memory

GPU → host



Single host

GPU rack node · NVLink fabric

memory floor



Multi-host

several nodes · cluster fabric

meet demand



Cluster / fleet

many nodes · network spine