

⟨⟩

## Single GPU

one accelerator · local memory

outgrows 1 GPU → host

forces local

⟨⟩

## Single host

GPU rack node · NVLink fabric

overflows memory floor

privacy /

⟨⟩

## Multi-host

several nodes · cluster fabric

input / fleet demand

cost pr

⟨⟩

## Cluster / fleet

many nodes · network spine