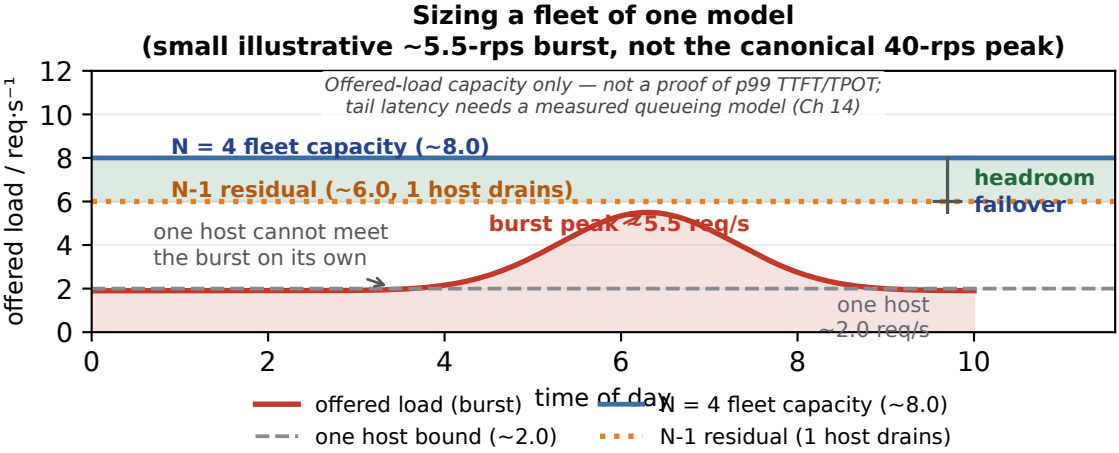
Sizing a fleet of one model
(small illustrative ~5.5-rps burst, not the canonical 40-rps peak)
Offered-load capacity only — not a proof of p99 TTFT/TPOT;
tail latency needs a measured queueing model (Ch 14)
N = 4 fleet capacity (~8.0)
N-1 residual (~6.0, 1 host drains)
headroom
failover
one host cannot meet
the burst on its own
burst peak ≈5.5 req/s
one host
~2.0 req/s
offered load / req·s⁻¹
time of day
offered load (burst)
N = 4 fleet capacity (~8.0)
one host bound (~2.0)
N-1 residual (1 host drains)