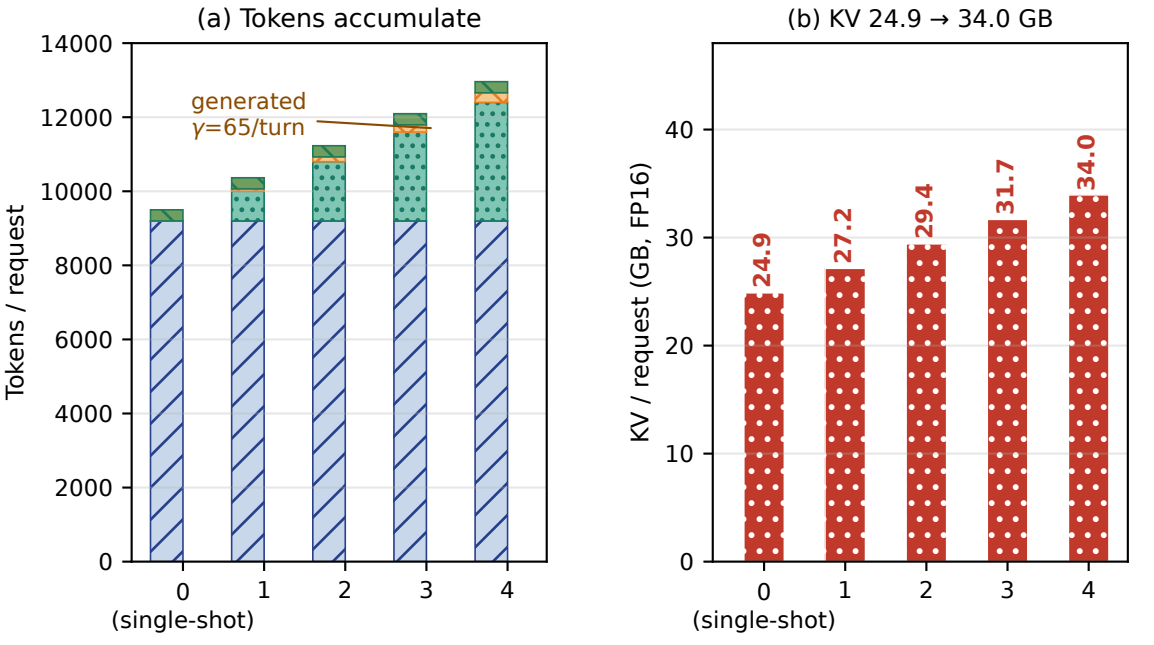
Token growth vs resulting KV growth across agent turns

append-only illustrative model: context only ever accumulates across turns — nothing is evicted



counterfactual: each column = a separate completion after T turns,
including one final 300-token answer — not successive states of one live agent.
Bars are RETAINED context/KV, not summed per-call processed or billed tokens;
tokens per submitted request grow ~2.5x under no prefix reuse (100-request ledger, Table 19-1).

initial context (l_0) · generated (γ=65) · KV / request (FP16)
retrieved (δ=800) · output (O=300)