


Agentic context accumulation: the KV block grows every turn



agentic depth is a fleet-sizing problem: KV grows ~36% (23.8 → 32.4 GB)

canonical example: KV = 2 MB/token (prompt) + 2 MB/token (retrieved context) + 2.5 MB/token (generated reasoning) = 6.5 MB/token