


Fleet throughput vs host count (KV/latency bound 2.1 req/s/host)



More hosts → lower per-host utilization → more tail headroom

