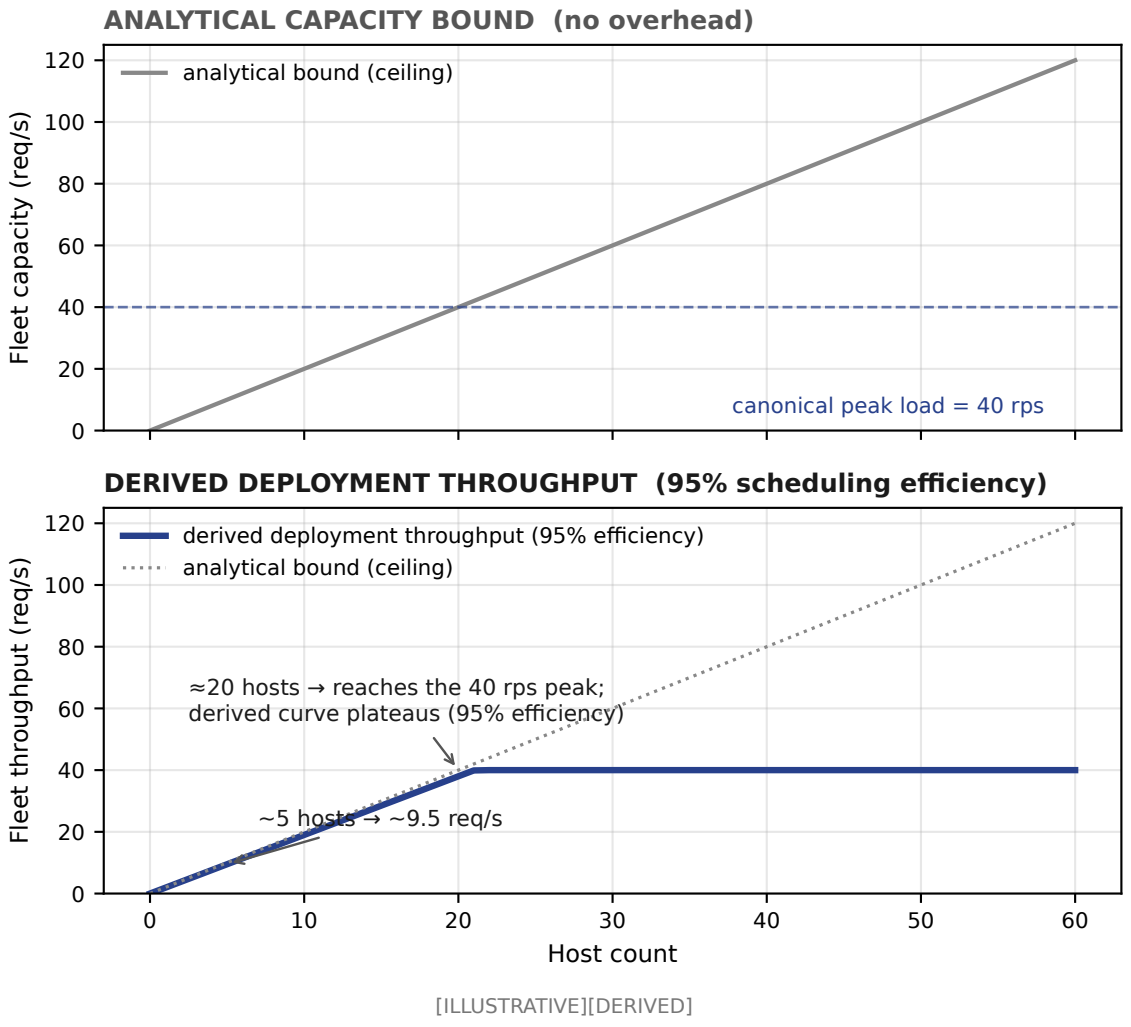
ANALYTICAL CAPACITY BOUND (no overhead)
analytical bound (ceiling)
Fleet capacity (req/s)
canonical peak load = 40 rps
DERIVED DEPLOYMENT THROUGHPUT (95% scheduling efficiency)
derived deployment throughput (95% efficiency)
analytical bound (ceiling)
Fleet throughput (req/s)
≈20 hosts → reaches the 40 rps peak;
derived curve plateaus (95% efficiency)
~5 hosts → ~9.5 req/s
Host count
[ILLUSTRATIVE][DERIVED]