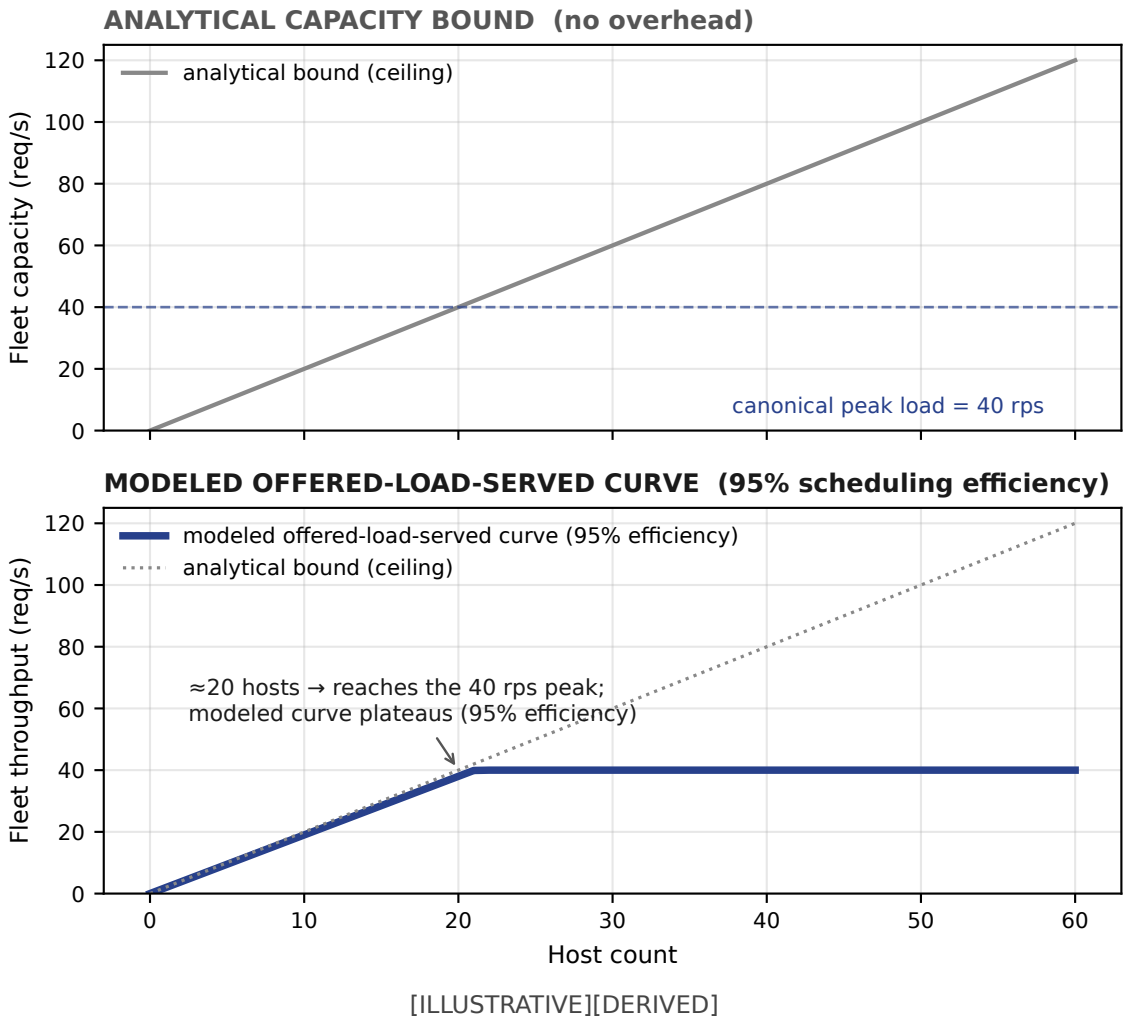
ANALYTICAL CAPACITY BOUND (no overhead)
analytical bound (ceiling)
Fleet capacity (req/s)
canonical peak load = 40 rps
MODELED OFFERED-LOAD-SERVED CURVE (95% scheduling efficiency)
modeled offered-load-served curve (95% efficiency)
analytical bound (ceiling)
Fleet throughput (req/s)
≈20 hosts → reaches the 40 rps peak;
modeled curve plateaus (95% efficiency)
Host count
[ILLUSTRATIVE][DERIVED]