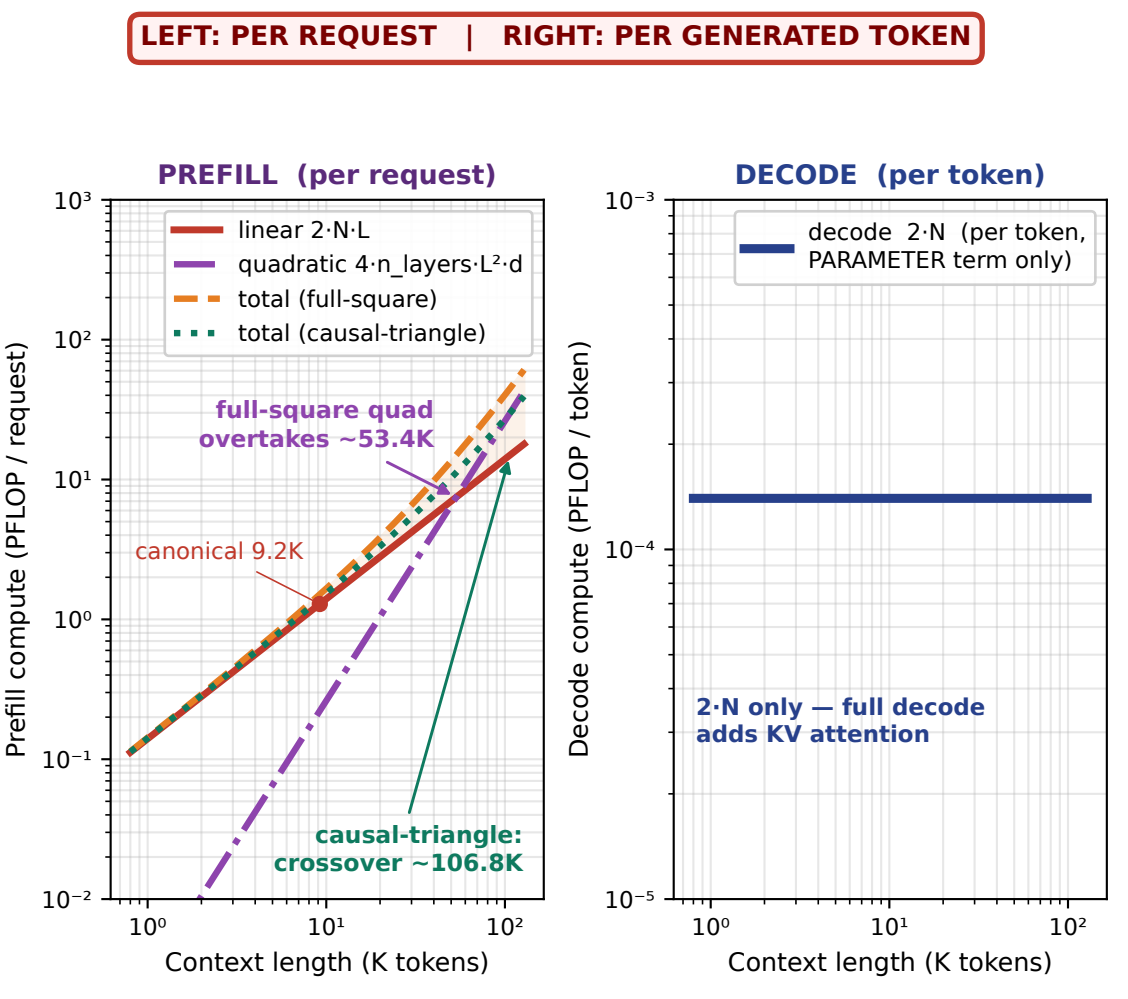
LEFT: PER REQUEST | RIGHT: PER GENERATED TOKEN
PREFILL (per request)
linear 2·N·L
quadratic 4·n_layers·L²·d
total (full-square)
total (causal-triangle)
full-square quad overtakes ~53.4K
canonical 9.2K
causal-triangle: crossover ~106.8K
Prefill compute (PFLOP / request)
Context length (K tokens)
DECODE (per token)
decode 2·N (per token, PARAMETER term only)
2·N only — full decode adds KV attention
Decode compute (PFLOP / token)
Context length (K tokens)