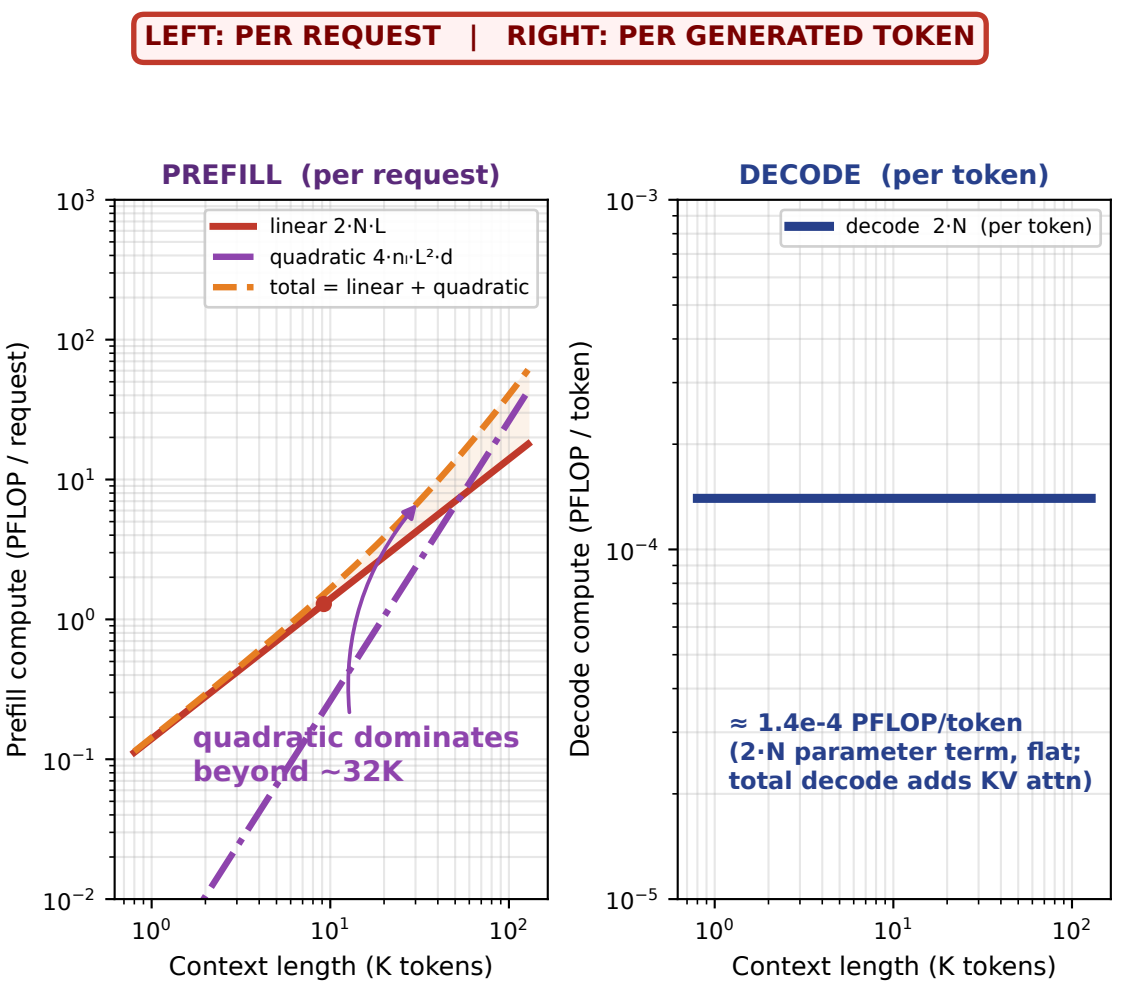
LEFT: PER REQUEST | RIGHT: PER GENERATED TOKEN
PREFILL (per request)
linear 2·N·L
quadratic 4·n_l·L²·d
total = linear + quadratic
quadratic dominates beyond ~32K
Prefill compute (PFLOP / request)
Context length (K tokens)
DECODE (per token)
decode 2·N (per token)
≈ 1.4e-4 PFLOP/token
(2·N parameter term, flat;
total decode adds KV attn)
Decode compute (PFLOP / token)
Context length (K tokens)