

Ch22 — prefill compute grows super-linearly; decode stays flat

