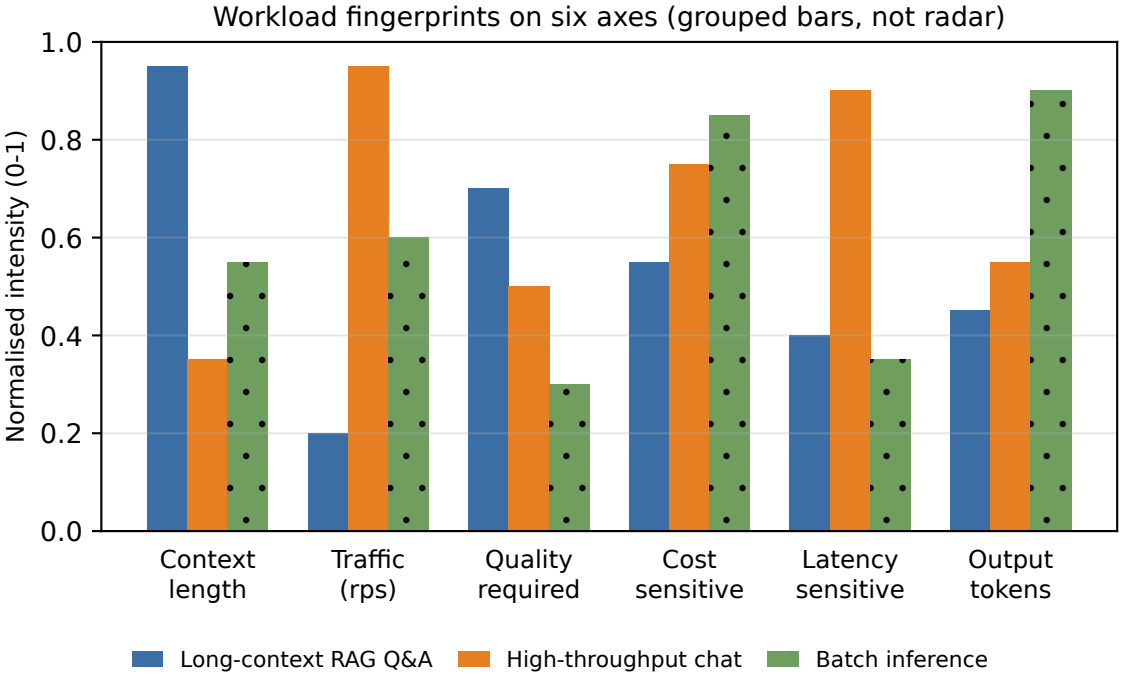
Workload fingerprints on six axes (grouped bars, not radar)
Normalised intensity (0-1)
1.0
0.8
0.6
0.4
0.2
0.0
Context length
Traffic (rps)
Quality required
Cost sensitive
Latency sensitive
Output tokens
Long-context RAG Q&A
High-throughput chat
Batch inference
per-axis 0-1 scale; compare within an axis, not across axes.