

# From Token to Fleet

*An AI Solution Architect's Handbook*

Version 20260913

Zhiqi Tao

*First and foremost a learning tool I wrote for myself — if it helps you too, that is a bonus, above and beyond what I set out to do.*



COPYRIGHT © 2026 ZHIQI TAO

All rights reserved.

**License.** This work — the book text, diagrams, figures, and build scripts — is distributed under a *source-available* license.

- **You may** freely read, study, learn from, and **redistribute** the work (verbatim, with attribution) for personal and non-commercial ends — no permission needed.
- **Commercial use or publication** — a paid product, course, book, or report — requires the author's **prior written approval**.

To request commercial permission, contact [zhiqi.tao@gmail.com](mailto:zhiqi.tao@gmail.com).

Full terms: [github.com/zhiqitao/from-token-to-fleet](https://github.com/zhiqitao/from-token-to-fleet)

(LICENSE, SPDX LicenseRef-ZhiqiResearch-FromTokenToFleet-1.0)

The work is provided “as is”, without warranty of any kind.



# Preface

For a long time I wanted to build the systematic knowledge of what it means to work as an AI Solution Architect: how to reason across the layers of an AI system — from tokens and models to workloads, systems, infrastructure, and fleets — so that a vague requirement can be turned into a defensible architecture. I kept looking for the resource that would lay all of that down in one coherent thread. What I found were deep dives into single layers — models, serving, infrastructure, agents — but no single spine that connected them into the one thing an architect actually does: chaining a real workload to a defensible solution.

Eventually I realized that the best way for me to build that knowledge was to go through the development journey myself, and that writing it down is how the knowledge gets built. I am not writing from a finished body of expertise; I am documenting a body of knowledge as I learn it, chapter by chapter, and I will keep updating these pages as I learn about new AI technologies. The moment I try to write down how I would decide whether a model belongs on a particular accelerator, I am forced to make explicit — parameter count, active parameters, KV-cache, context length, quantization, memory bandwidth, compute intensity, batching, interconnect, latency, throughput, concurrency, cost, reliability — every quantity that connects the model to the workload. And the act of writing exposes the holes in my own understanding, which tells me what I still need to learn next.

That feedback loop — curiosity, question, research, model, write, discover a gap, new question — is the engine of this book.

The AI field moves fast, and I am still learning too. This book is my learning notes, shared — written the only way I know: layer by layer, decision by decision, from the token up to the fleet. If it helps others who are interested in the AI Solution Architect role build the same systematic mental model, then it has outdone what I had planned it for. Learning is a proactive act; curiosity is the ultimate driver.

## The thesis

The book has one thesis, and it is the spine of every chapter:

Given an ambiguous problem, how do I systematically arrive at a defensible AI architecture?

The canonical method — the decision loop that Chapter 22 makes explicit and every chapter echoes — runs from requirements and workload characterization, through constraints and candidate architectures, to benchmarking, bottleneck analysis, optimization, total-cost reality checks, and a final red-team challenge. Parts I–V walk through the technical substrate, layer by layer. Part VI turns to the discipline itself: how one learns to reason across those layers.



FIGURE 1. The architect’s decision loop — the book’s spine. A vague requirement is carried through workload characterization, constraints, candidate architectures, benchmarking, bottleneck analysis, TCO, and a final red-team challenge before an architecture is committed. [ILLUSTRATIVE conceptual]

## How the book is organized

The six parts mirror the layers an architect reasons across:

- **Part I — The Token.** Models and mechanisms, from a single token to the memory and compute

arithmetic that prices every decision.

- **Part II — The Workload.** The job the system does: workload characterization and measuring what matters.
- **Part III — The System.** Memory, compute, communication, parallelism, and serving.
- **Part IV — The Architecture.** Designing, benchmarking, engineering, and costing a system.
- **Part V — The Fleet.** Operating AI at scale, and agents.
- **Part VI — The Architect.** The spine made explicit, plus a patterns library.

A worked canonical scenario — an enterprise retrieval-augmented question-answering workload on a 70B-class model — runs through the book, so every mechanism is anchored in concrete, reproducible arithmetic rather than abstraction. One caveat, so the book is not misread: this scenario is a **teaching instrument**, not a claim about what an AI workload *is*. Real workloads span a wide design space — long-context RAG (KV/memory-bound), high-throughput chat (decode-bound), batch inference (utilization/cost-bound), reasoning (test-time-compute-bound), agentic (concurrency/variance-bound), embeddings, vision, fine-tuning, training — and each shifts which quantity dominates. The canonical scenario picks a single, tractable, input-heavy RAG case and carries it through so the reasoning stays continuous; a real deployment will sit somewhere else in that space, and the chapters are written to give the reader the arithmetic to re-derive the bounds for *their* point.

## The Architect’s Evidence Discipline

Every quantitative claim in this book carries two orthogonal labels, and keeping these axes separate is itself a discipline. The **source** axis records where the number came from: [**1P**] first-party (the model vendor’s own spec), [**2°**] reputable second-hand (textbook, survey, independent benchmark), or [**VERIFY**] not yet anchored. The **claim-type** axis records what was done to it: **FACT** directly stated by the source, **DERIVED** computed from documented facts, **HYPOTHESIS** reasoned but not yet validated, or **ILLUSTRATIVE** — a round, teaching number chosen to make a mechanism concrete, never a measured or stable fact. Home-lab experiments are not reference. Where a figure is a vendor-reported claim, it is flagged, and the verification step is described so a reader can check it on their own hardware.

Two habits matter most. First, arithmetic belongs on the claim-type axis, not the source axis: a worked example computed from a model card is [**1P**] **DERIVED**, never [**2°**] just because arithmetic was done. Second, the conclusions an architect draws are only as defensible as the units behind them — so every major worked example ends with a unit check and a sanity check (does this result stay inside the physically possible?). That habit is part of the discipline, not an appendix.

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)  
each chapter supplies a derived quantity, which feeds a decision**



FIGURE 2. Reading the book by number. Each chapter supplies a derived quantity — a token/KV floor, a roofline point, a fleet size — and that quantity feeds a decision. Reading down a Part’s strip gives the layer-by-layer thread; reading across any chapter’s chip shows which number to reach for to answer a given question. [ILLUSTRATIVE]

## Reading conventions

Three conventions keep the book consistent and honest about its own limits.

**Terminology.** A few nouns are used as synonyms for readability: HBM, VRAM and GPU memory name the same physical resource; TTFT, time-to-first-token and prefill latency likewise; RAG and retrieval-augmented-generation are the same. Where a distinction is architecturally load-bearing — prefill versus decode, total versus active parameters, model capacity versus system capacity — the book holds it sharply. Synonyms are a readability concession; load-bearing contrasts are never relaxed (Ch. 2, 8, 17).

**Hardware is a generic stand-in.** The worked arithmetic is anchored to concrete constants (e.g. an 8×H100 host at roughly 1 PFLOPS FP16 and 3.4 TB/s HBM) because arithmetic needs concrete anchors. Those values are a *current* high-end accelerator, presented as an example, not a prophecy: the H200, B100/B200 and later generations change the constants, not the method. Re-derive the constants from the current hardware sheet; the roofline (Ch. 8) and the canonical arithmetic (Ch. 7) keep the method independent of any single generation. Forward-looking numbers are marked ILLUSTRATIVE.

**Quick reference.** For a day-to-day cheat-sheet: Chapter 26 (§2.1) carries a workload-to-strategy decision matrix, and Chapter 25’s Architecture Decision Record is a reusable template to copy into a repository. Those, plus the evidence discipline above, are worth re-reading first whenever a chapter feels dense.

## A note on currency

The mechanisms in this handbook — tokens, memory floors, rooflines, parallelism, KV arithmetic — are durable; they do not change as the industry moves. But the frontier does move, and a handbook for architects should carry the latest technology information. The Appendix therefore surveys the frontier open-weight architectures of 2026, mapped back to the mechanisms and decisions explored in the body of the book — and closes on the question that is more durable than any single architecture: *where should the intelligence live.*

# Contents

|                                          |          |
|------------------------------------------|----------|
| PREFACE                                  | i        |
| <b>I The Token</b>                       | <b>1</b> |
| 1 IT ALL STARTS WITH THE FIRST TOKEN     | 2        |
| 1.1 The Architect's Question . . . . .   | 2        |
| 1.2 Concept . . . . .                    | 2        |
| 1.3 Mental Model . . . . .               | 5        |
| 1.4 Worked Example . . . . .             | 7        |
| 1.5 Measurement . . . . .                | 8        |
| 1.6 Common Mistakes . . . . .            | 8        |
| 1.7 Architecture Consequence . . . . .   | 9        |
| 1.8 What We Still Don't Know . . . . .   | 9        |
| 1.9 End-of-Chapter Mini-Case . . . . .   | 9        |
| 2 WHAT ACTUALLY HAPPENS DURING INFERENCE | 10       |
| 2.1 The Architect's Question . . . . .   | 10       |
| 2.2 Concept . . . . .                    | 10       |
| 2.3 Mental Model . . . . .               | 11       |
| 2.4 Worked Example . . . . .             | 11       |
| 2.5 Measurement . . . . .                | 14       |
| 2.6 Common Mistakes . . . . .            | 15       |
| 2.7 Architecture Consequence . . . . .   | 15       |
| 2.8 What We Still Don't Know . . . . .   | 16       |
| 2.9 End-of-Chapter Mini-Case . . . . .   | 16       |
| 3 UNDERSTANDING THE MODEL                | 18       |
| 3.1 The Architect's Question . . . . .   | 18       |
| 3.2 Concept . . . . .                    | 18       |
| 3.3 Mental Model . . . . .               | 19       |
| 3.4 Worked Example . . . . .             | 19       |
| 3.5 Measurement . . . . .                | 21       |
| 3.6 Common Mistakes . . . . .            | 22       |

|           |                                                                                                            |           |
|-----------|------------------------------------------------------------------------------------------------------------|-----------|
| 3.7       | Architecture Consequence . . . . .                                                                         | 22        |
| 3.8       | What We Still Don't Know . . . . .                                                                         | 25        |
| 3.9       | End-of-Chapter Mini-Case . . . . .                                                                         | 25        |
| <b>II</b> | <b>The Workload</b>                                                                                        | <b>28</b> |
| 4         | THE ANATOMY OF AN AI WORKLOAD                                                                              | 29        |
| 4.1       | The Architect's Question . . . . .                                                                         | 29        |
| 4.2       | Concept . . . . .                                                                                          | 29        |
| 4.3       | Mental Model . . . . .                                                                                     | 30        |
| 4.4       | Worked Example . . . . .                                                                                   | 31        |
| 4.5       | Measurement . . . . .                                                                                      | 35        |
| 4.6       | Common Mistakes . . . . .                                                                                  | 36        |
| 4.7       | Architecture Consequence . . . . .                                                                         | 36        |
| 4.8       | What We Still Don't Know . . . . .                                                                         | 38        |
| 4.9       | End-of-Chapter Mini-Case . . . . .                                                                         | 39        |
| 5         | FROM WORKLOAD TO MODEL SELECTION                                                                           | 40        |
| 5.1       | The Architect's Question . . . . .                                                                         | 40        |
| 5.2       | Concept . . . . .                                                                                          | 40        |
| 5.3       | Mental Model . . . . .                                                                                     | 41        |
| 5.4       | Worked Example — Embedding-Model vs Generation-Model Selection for the<br>Canonical RAG Workload . . . . . | 41        |
| 5.5       | Measurement . . . . .                                                                                      | 45        |
| 5.6       | Common Mistakes . . . . .                                                                                  | 45        |
| 5.7       | Architecture Consequence . . . . .                                                                         | 46        |
| 5.8       | What We Still Don't Know . . . . .                                                                         | 46        |
| 5.9       | End-of-Chapter Mini-Case . . . . .                                                                         | 46        |
| 6         | MEASURING WHAT MATTERS: FROM TELEMETRY TO ARCHITECTURE DECISIONS                                           | 49        |
| 6.1       | The Architect's Question . . . . .                                                                         | 49        |
| 6.2       | Concept . . . . .                                                                                          | 49        |
| 6.3       | Mental Model . . . . .                                                                                     | 50        |
| 6.4       | Worked Example . . . . .                                                                                   | 52        |
| 6.5       | Measurement . . . . .                                                                                      | 54        |
| 6.6       | Common Mistakes . . . . .                                                                                  | 55        |
| 6.7       | Architecture Consequence . . . . .                                                                         | 55        |
| 6.8       | What We Still Don't Know . . . . .                                                                         | 55        |
| 6.9       | End-of-Chapter Mini-Case . . . . .                                                                         | 56        |

|            |                                                                    |           |
|------------|--------------------------------------------------------------------|-----------|
| <b>III</b> | <b>The System</b>                                                  | <b>57</b> |
| 7          | MEMORY IS THE FIRST CONSTRAINT                                     | 58        |
| 7.1        | The Architect's Question . . . . .                                 | 58        |
| 7.2        | Concept . . . . .                                                  | 58        |
| 7.3        | Mental Model . . . . .                                             | 59        |
| 7.4        | Worked Example . . . . .                                           | 59        |
| 7.5        | Measurement . . . . .                                              | 63        |
| 7.6        | Common Mistakes . . . . .                                          | 64        |
| 7.7        | Architecture Consequence . . . . .                                 | 65        |
| 7.8        | What We Still Don't Know . . . . .                                 | 65        |
| 7.9        | End-of-Chapter Mini-Case . . . . .                                 | 71        |
| 8          | COMPUTE                                                            | 72        |
| 8.1        | The Architect's Question . . . . .                                 | 72        |
| 8.2        | Concept . . . . .                                                  | 72        |
| 8.3        | Mental Model . . . . .                                             | 73        |
| 8.4        | Worked Example: Canonical Scenario Arithmetic . . . . .            | 73        |
| 8.5        | Measurement . . . . .                                              | 76        |
| 8.6        | Common Mistakes . . . . .                                          | 78        |
| 8.7        | Architecture Consequence . . . . .                                 | 78        |
| 8.8        | What We Still Don't Know . . . . .                                 | 79        |
| 8.9        | End-of-Chapter Mini-Case: Continuous Deployment Scenario . . . . . | 79        |
| 9          | COMMUNICATION                                                      | 81        |
| 9.1        | The Architect's Question . . . . .                                 | 81        |
| 9.2        | Concept . . . . .                                                  | 81        |
| 9.3        | Mental Model . . . . .                                             | 83        |
| 9.4        | Worked Example . . . . .                                           | 83        |
| 9.5        | Measurement . . . . .                                              | 85        |
| 9.6        | Common Mistakes . . . . .                                          | 86        |
| 9.7        | Architecture Consequence . . . . .                                 | 87        |
| 9.8        | What We Still Don't Know . . . . .                                 | 87        |
| 9.9        | End-of-Chapter Mini-Case . . . . .                                 | 88        |
| 10         | PARALLELISM: SPLITTING THE MODEL ACROSS MACHINES                   | 89        |
| 10.1       | The Architect's Question . . . . .                                 | 89        |
| 10.2       | Concept . . . . .                                                  | 89        |
| 10.3       | Mental Model . . . . .                                             | 92        |
| 10.4       | Worked Example . . . . .                                           | 92        |
| 10.5       | Measurement . . . . .                                              | 94        |

|      |                                                             |     |
|------|-------------------------------------------------------------|-----|
| 10.6 | Common Mistakes . . . . .                                   | 94  |
| 10.7 | Architecture Consequence . . . . .                          | 95  |
| 10.8 | What We Still Don't Know . . . . .                          | 95  |
| 10.9 | End-of-Chapter Mini-Case . . . . .                          | 97  |
| 11   | SERVING: FROM MODEL IN MEMORY TO REQUESTS IN FLIGHT         | 98  |
| 11.1 | The Architect's Question . . . . .                          | 98  |
| 11.2 | Concept . . . . .                                           | 98  |
| 11.3 | Mental Model . . . . .                                      | 99  |
| 11.4 | Worked Example . . . . .                                    | 99  |
| 11.5 | Measurement . . . . .                                       | 102 |
| 11.6 | Common Mistakes . . . . .                                   | 103 |
| 11.7 | Architecture Consequence . . . . .                          | 103 |
| 11.8 | What We Still Don't Know . . . . .                          | 104 |
| 11.9 | End-of-Chapter Mini-Case . . . . .                          | 104 |
| IV   | The Architecture                                            | 107 |
| 12   | DESIGNING THE AI SYSTEM: FROM CONSTRAINTS TO CANDIDATES     | 108 |
| 12.1 | The Architect's Question . . . . .                          | 108 |
| 12.2 | Concept . . . . .                                           | 108 |
| 12.3 | Mental Model . . . . .                                      | 109 |
| 12.4 | Worked Example . . . . .                                    | 109 |
| 12.5 | Measurement . . . . .                                       | 110 |
| 12.6 | Common Mistakes . . . . .                                   | 111 |
| 12.7 | Architecture Consequence . . . . .                          | 111 |
| 12.8 | What We Still Don't Know . . . . .                          | 112 |
| 12.9 | End-of-Chapter Mini-Case . . . . .                          | 112 |
| 13   | REFERENCE ARCHITECTURES: FROM LOCAL TO HYPERSCALE           | 116 |
| 13.1 | The Architect's Question . . . . .                          | 116 |
| 13.2 | Concept . . . . .                                           | 116 |
| 13.3 | Mental Model . . . . .                                      | 117 |
| 13.4 | Worked Example . . . . .                                    | 117 |
| 13.5 | Measurement . . . . .                                       | 119 |
| 13.6 | Common Mistakes . . . . .                                   | 119 |
| 13.7 | Architecture Consequence . . . . .                          | 119 |
| 13.8 | What We Still Don't Know . . . . .                          | 119 |
| 13.9 | End-of-Chapter Mini-Case . . . . .                          | 121 |
| 14   | BENCHMARKING: PREDICTING REAL EXPERIENCE, NOT A LEADERBOARD | 122 |

|          |                                                                |            |
|----------|----------------------------------------------------------------|------------|
| 14.1     | The Architect's Question . . . . .                             | 122        |
| 14.2     | Concept . . . . .                                              | 122        |
| 14.3     | Mental Model . . . . .                                         | 123        |
| 14.4     | Worked Example . . . . .                                       | 123        |
| 14.5     | Measurement . . . . .                                          | 124        |
| 14.6     | Common Mistakes . . . . .                                      | 124        |
| 14.7     | Architecture Consequence . . . . .                             | 125        |
| 14.8     | What We Still Don't Know . . . . .                             | 126        |
| 14.9     | End-of-Chapter Mini-Case . . . . .                             | 126        |
| 15       | PERFORMANCE ENGINEERING: FROM "IT WORKS" TO "IT MEETS THE SLO" | 127        |
| 15.1     | The Architect's Question . . . . .                             | 127        |
| 15.2     | Concept . . . . .                                              | 127        |
| 15.3     | Mental Model . . . . .                                         | 128        |
| 15.4     | Worked Example . . . . .                                       | 128        |
| 15.5     | Measurement . . . . .                                          | 129        |
| 15.6     | Common Mistakes . . . . .                                      | 130        |
| 15.7     | Architecture Consequence . . . . .                             | 130        |
| 15.8     | What We Still Don't Know . . . . .                             | 131        |
| 15.9     | End-of-Chapter Mini-Case . . . . .                             | 131        |
| 16       | TCO: THE TOTAL COST REALITY CHECK                              | 134        |
| 16.1     | The Architect's Question . . . . .                             | 134        |
| 16.2     | Concept . . . . .                                              | 134        |
| 16.3     | Mental Model . . . . .                                         | 135        |
| 16.4     | Worked Example . . . . .                                       | 135        |
| 16.5     | Measurement . . . . .                                          | 137        |
| 16.6     | Common Mistakes . . . . .                                      | 137        |
| 16.7     | Architecture Consequence . . . . .                             | 138        |
| 16.8     | What We Still Don't Know . . . . .                             | 139        |
| 16.9     | End-of-Chapter Mini-Case . . . . .                             | 139        |
| <b>V</b> | <b>The Fleet</b>                                               | <b>140</b> |
| 17       | FROM SERVER TO FLEET                                           | 141        |
| 17.1     | The Architect's Question . . . . .                             | 141        |
| 17.2     | Concept . . . . .                                              | 141        |
| 17.3     | Mental Model . . . . .                                         | 142        |
| 17.4     | Worked Example . . . . .                                       | 143        |
| 17.5     | Measurement . . . . .                                          | 144        |
| 17.6     | Common Mistakes . . . . .                                      | 147        |

|      |                                                                        |     |
|------|------------------------------------------------------------------------|-----|
| 17.7 | Architecture Consequence . . . . .                                     | 148 |
| 17.8 | What We Still Don't Know . . . . .                                     | 149 |
| 17.9 | End-of-Chapter Mini-Case . . . . .                                     | 149 |
| 18   | OPERATING MULTIPLE MODELS . . . . .                                    | 152 |
| 18.1 | The Architect's Question . . . . .                                     | 152 |
| 18.2 | Concept . . . . .                                                      | 152 |
| 18.3 | Mental Model . . . . .                                                 | 153 |
| 18.4 | Worked Example . . . . .                                               | 153 |
| 18.5 | Measurement . . . . .                                                  | 156 |
| 18.6 | Common Mistakes . . . . .                                              | 157 |
| 18.7 | Architecture Consequence . . . . .                                     | 158 |
| 18.8 | What We Still Don't Know . . . . .                                     | 158 |
| 18.9 | End-of-Chapter Mini-Case . . . . .                                     | 158 |
| 19   | AGENTIC SYSTEMS . . . . .                                              | 160 |
| 19.1 | The Architect's Question . . . . .                                     | 160 |
| 19.2 | Concept . . . . .                                                      | 160 |
| 19.3 | Mental Model . . . . .                                                 | 161 |
| 19.4 | Worked Example . . . . .                                               | 161 |
| 19.5 | Measurement . . . . .                                                  | 163 |
| 19.6 | Common Mistakes . . . . .                                              | 165 |
| 19.7 | Architecture Consequence . . . . .                                     | 165 |
| 19.8 | What We Still Don't Know . . . . .                                     | 168 |
| 19.9 | End-of-Chapter Mini-Case . . . . .                                     | 169 |
| 20   | FLEET-LEVEL OPTIMIZATION . . . . .                                     | 171 |
| 20.1 | The Architect's Question . . . . .                                     | 171 |
| 20.2 | Concept . . . . .                                                      | 171 |
| 20.3 | Mental Model . . . . .                                                 | 172 |
| 20.4 | Worked Example . . . . .                                               | 172 |
| 20.5 | Measurement . . . . .                                                  | 173 |
| 20.6 | Common Mistakes . . . . .                                              | 174 |
| 20.7 | Architecture Consequence . . . . .                                     | 174 |
| 20.8 | What We Still Don't Know . . . . .                                     | 174 |
| 20.9 | End-of-Chapter Mini-Case . . . . .                                     | 175 |
| 21   | THE AI FACTORY: TURNING INFERENCE INTO AN INDUSTRIAL PROCESS . . . . . | 177 |
| 21.1 | The Architect's Question . . . . .                                     | 177 |
| 21.2 | Concept . . . . .                                                      | 177 |
| 21.3 | Mental Model . . . . .                                                 | 178 |

|           |                                                            |            |
|-----------|------------------------------------------------------------|------------|
| 21.4      | Worked Example . . . . .                                   | 178        |
| 21.5      | Measurement . . . . .                                      | 179        |
| 21.6      | Common Mistakes . . . . .                                  | 179        |
| 21.7      | Architecture Consequence . . . . .                         | 180        |
| 21.8      | What We Still Don't Know . . . . .                         | 180        |
| 21.9      | End-of-Chapter Mini-Case . . . . .                         | 181        |
| <b>VI</b> | <b>The Architect</b>                                       | <b>182</b> |
| 22        | HOW TO THINK LIKE AN AI SOLUTION ARCHITECT                 | 183        |
| 22.1      | The Architect's Question . . . . .                         | 183        |
| 22.2      | Concept . . . . .                                          | 183        |
| 22.3      | Mental Model . . . . .                                     | 184        |
| 22.4      | Worked Example — The Canonical Loop . . . . .              | 184        |
| 22.5      | Measurement . . . . .                                      | 187        |
| 22.6      | Common Mistakes . . . . .                                  | 187        |
| 22.7      | Architecture Consequence . . . . .                         | 188        |
| 22.8      | What We Still Don't Know . . . . .                         | 188        |
| 22.9      | End-of-Chapter Mini-Case . . . . .                         | 189        |
| 23        | WORKING WITH CUSTOMERS: FROM VAGUE ASK TO DEFENSIBLE SCOPE | 190        |
| 23.1      | The Architect's Question . . . . .                         | 190        |
| 23.2      | Concept . . . . .                                          | 190        |
| 23.3      | Mental Model . . . . .                                     | 191        |
| 23.4      | Worked Example . . . . .                                   | 191        |
| 23.5      | Measurement . . . . .                                      | 192        |
| 23.6      | Common Mistakes . . . . .                                  | 193        |
| 23.7      | Architecture Consequence . . . . .                         | 193        |
| 23.8      | What We Still Don't Know . . . . .                         | 193        |
| 23.9      | End-of-Chapter Mini-Case . . . . .                         | 195        |
| 24        | RED TEAM / GREEN TEAM                                      | 196        |
| 24.1      | The Architect's Question . . . . .                         | 196        |
| 24.2      | Concept . . . . .                                          | 196        |
| 24.3      | Mental Model . . . . .                                     | 197        |
| 24.4      | Worked Example . . . . .                                   | 197        |
| 24.5      | Measurement . . . . .                                      | 200        |
| 24.6      | Common Mistakes . . . . .                                  | 201        |
| 24.7      | Architecture Consequence . . . . .                         | 202        |
| 24.8      | What We Still Don't Know . . . . .                         | 202        |
| 24.9      | End-of-Chapter Mini-Case . . . . .                         | 203        |

|      |                                                                                                  |     |
|------|--------------------------------------------------------------------------------------------------|-----|
| 25   | THE ARCHITECTURE DECISION RECORD                                                                 | 206 |
| 25.1 | The Architect’s Question . . . . .                                                               | 206 |
| 25.2 | Concept . . . . .                                                                                | 206 |
| 25.3 | Mental Model . . . . .                                                                           | 207 |
| 25.4 | Worked Example: Model-Shape ADR . . . . .                                                        | 207 |
| 25.5 | Measurement . . . . .                                                                            | 209 |
| 25.6 | Common Mistakes . . . . .                                                                        | 210 |
| 25.7 | Architecture Consequence . . . . .                                                               | 211 |
| 25.8 | What We Still Don’t Know . . . . .                                                               | 211 |
| 25.9 | End-of-Chapter Mini-Case . . . . .                                                               | 213 |
| 26   | THE SOLUTION ARCHITECT’S TOOLKIT + PATTERNS LIBRARY                                              | 215 |
| 26.1 | The Architect’s Question . . . . .                                                               | 215 |
| 26.2 | Concept . . . . .                                                                                | 215 |
| 26.3 | Mental Model . . . . .                                                                           | 215 |
| 26.4 | Worked Example — Capstone Scenario . . . . .                                                     | 219 |
| 26.5 | Measurement . . . . .                                                                            | 219 |
| 26.6 | Common Mistakes . . . . .                                                                        | 220 |
| 26.7 | Architecture Consequence . . . . .                                                               | 220 |
| 26.8 | What We Still Don’t Know . . . . .                                                               | 220 |
| 26.9 | End-of-Chapter Mini-Case . . . . .                                                               | 221 |
| A    | APPENDIX A — 2026 FRONTIER ARCHITECTURES: WHAT THE LATEST OPEN MODELS TELL<br>Us                 | 224 |
| A.1  | The Signal in One Paragraph . . . . .                                                            | 224 |
| A.2  | How to read this appendix: Observed, Interpretation, Hypothesis . . . . .                        | 226 |
| A.3  | DeepSeek-V4 — Making Long Context Cheap by Re-Architecting Attention . . . . .                   | 227 |
| A.4  | Kimi K3 — The First Open 3T-Class Model, and the MoE-Sparsity Extreme . . . . .                  | 227 |
| A.5  | Qwen3.8-Flash-Next — The Qwen4 Architecture Preview, Prioritizing Cost-<br>Performance . . . . . | 228 |
| A.6  | GLM-5.3-Flash — First Natively Multimodal GLM, Efficiency-Branded . . . . .                      | 228 |
| A.7  | What This Means for the Handbook’s Canonical Arithmetic . . . . .                                | 228 |
| A.8  | Where 2026 Capability Actually Comes From (beyond Architecture) . . . . .                        | 230 |
| A.9  | The Frontier Question: Where Should Intelligence Live? . . . . .                                 | 232 |
| A.10 | One-Hand Sources (all [1P]) . . . . .                                                            | 232 |

# List of Figures

|     |                                                                                                                                                                                                                                                                                                                                                   |    |
|-----|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 1   | The architect’s decision loop — the book’s spine. A vague requirement is carried through workload characterization, constraints, candidate architectures, benchmarking, bottleneck analysis, TCO, and a final red-team challenge before an architecture is committed. [ILLUSTRATIVE conceptual] . . . . .                                         | ii |
| 2   | Reading the book by number. Each chapter supplies a derived quantity — a token/KV floor, a roofline point, a fleet size — and that quantity feeds a decision. Reading down a Part’s strip gives the layer-by-layer thread; reading across any chapter’s chip shows which number to reach for to answer a given question. [ILLUSTRATIVE] . . . . . | iv |
| 1.1 | Fig 1.1 — The KV cache: every decoded token adds one K+V per layer per head [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                                                   | 5  |
| 1.2 | Fig 1.2 — How a token travels: text → tokenizer → IDs → embedding → attention → KV cache [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                                      | 6  |
| 2.1 | Fig 2.1 — Decode vs prefill: bandwidth vs compute [ILLUSTRATIVE conceptual] . .                                                                                                                                                                                                                                                                   | 13 |
| 3.1 | Fig 3.1 — Dense vs MoE parameter allocation and KV-cache behavior [2° DERIVED]                                                                                                                                                                                                                                                                    | 26 |
| 4.1 | Fig 4.1 — Six-dimension workload characterization mapped to architectural decisions. [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                                          | 37 |
| 5.1 | Fig 5.1 — Model selection for RAG: workload → five selection surfaces → two legs (retrieval/embedding, generation/70B) → the system decision [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                  | 48 |
| 6.1 | Fig 6.1 — The metric hierarchy as a diagnostic chain [ILLUSTRATIVE conceptual] .                                                                                                                                                                                                                                                                  | 51 |
| 6.2 | Fig 6.2 — Request-latency distribution: p50/p90/p95/p99 and the mean. The mean (~0.84 s) hides the 1% stragglers at ~5 s that breach the p95 TTFT budget (ILLUSTRATIVE) . . . . .                                                                                                                                                                 | 52 |
| 7.1 | Fig 7.1 — Grouped-Query Attention: 64 query heads share 8 KV heads, cutting per-token KV 8× from 2.5 MB to ~0.33 MB (ILLUSTRATIVE, after LLaMA-2 70B head layout) . . . . .                                                                                                                                                                       | 62 |
| 7.2 | Fig 7.2 — KV-cache size vs context length for a 70B model [2° DERIVED] . . . . .                                                                                                                                                                                                                                                                  | 67 |

|      |                                                                                                                                                                                                                                                                                                                       |     |
|------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 7.3  | Fig 7.3 — Inference vs fine-tuning memory floor [2° DERIVED] . . . . .                                                                                                                                                                                                                                                | 68  |
| 7.4  | Fig 7.4 — The concurrency budget: where a 70B host's 640 GB pool goes [2° DERIVED]                                                                                                                                                                                                                                    | 69  |
| 7.5  | Fig 7.5 — Memory Tetris: how the 8×H100 host's 640 GB pool fills at three contexts (9.2K / 32K / 128K). Runtime ~64 GB + weights 140 GB + KV cache 23.8 / 80 / 320 GB [2° DERIVED] . . . . .                                                                                                                          | 70  |
| 8.1  | Fig 8.1 — Prefill above the ridge (compute-bound), decode below it (memory-bound). H100 (solid) vs H200 (dashed) ridge points from vendor datasheets (989 TFLOPS, 3.35 / 4.8 TB/s); Continuous-Batching arrow shows decode climbing the slope as batch grows [2° DERIVED] [1P: vendor datasheet] . . . . .            | 77  |
| 9.1  | Fig 9.1 — All-reduce time vs data volume, by interconnect tier [2° DERIVED] . . . .                                                                                                                                                                                                                                   | 84  |
| 10.1 | Fig 10.1 — Composing the dimensions: each PP stage holds a DP replica, TP-sharded across GPUs; and EP's per-token all-to-all routes each token to the GPU holding its experts (ILLUSTRATIVE) . . . . .                                                                                                                | 91  |
| 10.2 | Fig 10.2 — The five parallelization strategies and what each splits [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                               | 96  |
| 11.1 | Fig 11.1 — Discrete vs continuous batching (Orca-style slot diagram). Discrete batching runs each batch to completion, leaving GPU slots idle as sequences finish early; continuous batching admits and evicts requests at every decode step, backfilling freed slots (ILLUSTRATIVE, after Orca, OSDI 2022) . . . . . | 101 |
| 11.2 | Fig 11.2 — The serving stack: batching, KV management, prefix caching, P/D split [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                  | 104 |
| 11.3 | Fig 11.3 — P/D disaggregation topology: prefill pool (compute-bound) and decode pool (bandwidth-bound) bridged by KV transfer [2° DERIVED] . . . . .                                                                                                                                                                  | 105 |
| 12.1 | Fig 12.1 — Candidate architecture synthesis: three candidate architectures from the canonical enterprise-Q&A RAG workload, comparing memory residency, parallelism plan, and constraint-satisfaction outcome (memory fit, latency SLO, economics) [ILLUSTRATIVE conceptual] . . . . .                                 | 115 |
| 13.1 | Fig 13.1 — The reference-architecture ladder and its escalation triggers [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                          | 120 |
| 14.1 | Fig 14.1 — Capability screen then deployment benchmark [ILLUSTRATIVE conceptual]                                                                                                                                                                                                                                      | 125 |
| 15.1 | Fig 15.1 — Decode goodput vs batch size, with and without prefix caching: caching relaxes the <b>prefill</b> bottleneck, not decode (ILLUSTRATIVE, canonical decode ~25 ms/token) . . . . .                                                                                                                           | 132 |

|      |                                                                                                                                                                                                                                                                                     |     |
|------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 16.1 | Fig 16.1 — TCO break-even: self-hosted vs managed API [2° DERIVED] . . . . .                                                                                                                                                                                                        | 138 |
| 17.1 | Fig 17.1 - A fleet of one model: a load balancer fans requests out to N identical hosts, each running the same model on its own GPU [2° DERIVED] . . . . .                                                                                                                          | 142 |
| 17.2 | Fig 17.2 — Host-count contours $H = \lambda \cdot L / (C \cdot \text{util})$ over peak rate $\lambda \times$ average latency L, with the FP16-KV family ( $C \approx 18$ , filled) and the FP8-KV family ( $C \approx 35$ , dashed) [DERIVED: Ch17 §3 + Ch7 KV constants] . . . . . | 146 |
| 18.1 | Fig 18.1 — Model-routing decision tree: capability filter $\rightarrow$ cost/latency gate $\rightarrow$ fallback, with the five specialised families as leaves [ILLUSTRATIVE conceptual] . . . . .                                                                                  | 159 |
| 19.1 | Fig 19.1 - The agent loop as a four-state state machine (Plan, Execute, Observe, Decide) with a loop-back to Plan and the terminal outcomes final answer and stopped [ILLUSTRATIVE conceptual] . . . . .                                                                            | 162 |
| 19.2 | Fig 19.2 — Input/output tokens and KV cache per request across agent turns ( $I_0 + T \cdot \delta + T \cdot \gamma$ ; 9,200 + 800 + 65 per turn), FP16 KV rising from 23.8 GB single-shot to 32.4 GB at 4 turns [DERIVED: Ch19 §3 arithmetic] . . . . .                            | 166 |
| 19.3 | Fig 19.3 — The agentic context block grows every turn: initial prompt (9.2K tok) plus tool output + reasoning appended each turn, KV footprint rising 23.8 $\rightarrow$ 32.4 GB [2° DERIVED: Ch19 §3 arithmetic] . . . . .                                                         | 167 |
| 20.1 | Fig 20.1 - Fleet QPS and latency vs. host count, batch size, and scheduling overhead [2° DERIVED] . . . . .                                                                                                                                                                         | 176 |
| 21.1 | Fig 21.1 — The AI Factory promotion pipeline [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                    | 180 |
| 22.1 | Fig 22.1 - Prefill vs decode bottleneck divergence [2° DERIVED] . . . . .                                                                                                                                                                                                           | 186 |
| 23.1 | Fig 23.1 — From vague ask to architectural bounds [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                               | 194 |
| 24.1 | Fig 24.1 - Red Team / Green Team cycle [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                                          | 204 |
| 25.1 | Fig 25.1 — The Anatomy of an Architecture Decision Record [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                                                                       | 212 |
| 26.1 | Fig 26.1 — Workload fingerprints on the six axes of §4: long-context RAG Q&A, high-throughput chat, and batch inference occupy distinct regions of the space (ILLUSTRATIVE, relative 0–1) . . . . .                                                                                 | 218 |
| 26.2 | Fig 26.2 - Pattern Application Diagram (the Pattern-Measurement-Feedback Loop applied to the capstone scenario) [ILLUSTRATIVE conceptual] . . . . .                                                                                                                                 | 222 |
| A.1  | Fig A.1 — 2026 frontier MoE extreme sparsity: four open models, total vs activated parameters, and active fraction in single digits [1P] . . . . .                                                                                                                                  | 225 |

|     |                                                                                                                                                                                                                                                                                                                                                        |     |
|-----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| A.2 | Fig A.2 — The KV constant and FLOP/token are a moving target: hybrid attention cuts both at the mechanism. Canonical 70B dense full-MHA = 100%; DeepSeek-V4-Flash $\approx$ 7% KV / 10% FLOP, V4-Pro $\approx$ 10% / 27%, GLM-5.3-Flash $\approx$ 23% KV / $\sim$ 33% attention-compute [DERIVED from 1P claims: arXiv 2606.19348; HF zai-org] . . . . | 229 |
| A.3 | Fig A.3 — The layered evolution of frontier systems, colour-coded by efficiency-led (amber) vs capability-led (blue) stages, from a single Transformer to the whole fleet (Appendix §6) . . . . .                                                                                                                                                      | 231 |
| A.4 | Fig A.4 — Where should intelligence live? Eight host layers, from inside the weights up to across a fleet of specialised models, each mapped to the chapters of this handbook that give the tooling to evaluate it (Appendix §7) . . . . .                                                                                                             | 233 |



## **PART I.**

# **The Token**

## CHAPTER 1.

# It All Starts With the First Token

### 1.1. The Architect’s Question

Before we can size a model, a context window, or a serving budget, we have to know what the input actually is. After this chapter we should be able to reason about an AI system’s **smallest unit of cost and meaning** — the token — and to translate between human units (a sentence, a page, a word) and the machine unit every architecture decision is priced against. This is the first rung of the token→fleet spine: everything downstream (memory, compute, latency, cost) is ultimately a statement about *tokens*.

### 1.2. Concept

A **token** is the atomic unit of text that a language model reads and writes. It is not a word, not a character, and not a byte — it is whatever the model’s *tokenizer* decided to carve the text into. A single token can be a whole word (handbook), a common subword (hand), a single character, a short punctuation+space pair, or in worst cases a fragment of a character.

The disconnection from human units is the first thing an architect has to internalize: **cost and capacity in an LLM are measured in tokens, not words or characters**, because every model operation — embedding lookup, attention arithmetic, KV-cache memory, decode time — prices by token. This chapter exists to make that switch automatic.

#### 1.2.1. Tokenization

Tokenization splits a string into an ordered list of integer IDs from the model’s **vocabulary** (typically tens of thousands, up to ~256k entries). The split is data-driven: learned from a large corpus, byte-

pair encoding (BPE) repeatedly merges the most frequent adjacent byte/subword pairs until the vocabulary is built, so the boundaries are wherever the corpus statistics happened to fall — not where a grammarian would put them. Two practical consequences follow. First, the mapping is looser than a perfect round-trip: the detokenizer reconstructs text deterministically, but the token boundaries themselves carry little semantic meaning — `architecture` and `architect` may or may not share a token depending on corpus frequency. Second, token counts are **not proportional to character counts in any fixed way**.<sup>[2°]</sup> A high-frequency English word is often one token ([1P] verifiable on most vendor tokenizer demos — e.g. common words like *the*, *we*, *run* render as single tokens); a rare technical term, a long digit string, or dense code is frequently several. The “one token  $\approx$  4 characters” rule of thumb is for rough estimation of general-English prose, never a law — we confirm against the model’s own tokenizer before pricing anything.

### 1.2.2. Context window

The **context window** is the maximum number of tokens the model can attend to in a single forward pass — the sum of the prompt (input) and the tokens it generates (output). It is a hard architectural constraint: exceed it and the request fails or must be truncated. Current-gen models commonly ship windows ranging from roughly 8K up to a million-plus tokens, varying by model and vendor ([2°] vendor model cards, e.g. 128K/1M+ windows as of 2026); the exact number is a fixed, checkable property of the model we choose — verify it on the card, never assume it from the family name. The context window is one of the first numbers we will be quoted, and one of the easiest to overcommit against, because it bounds memory and cost whether we use every token or not.

### 1.2.3. Embeddings

Text enters the model as token IDs, but the model does arithmetic on vectors. An **embedding** is the learned dense vector that represents a token (or a position, or a subword context) in a high-dimensional space — typically 128 to a few thousand dimensions. Embeddings are the bridge between discrete tokens and continuous computation: nearby tokens (in a learned sense) map to nearby vectors, which is what lets attention and downstream layers do useful math.

For the architect the practical facts are: embeddings live in a lookup table sized `vocab × dim`, they account for a real but small slice of a model’s parameter count (a `100k × 4096` table is  $\sim$ 0.4B params — a rounding error next to a 70B model, but not zero), and they are the first thing a *vocabulary mismatch* or a *re-tokenized prompt* can silently change.

### 1.2.4. Attention

Attention is the mechanism by which each token’s representation is recomputed as a weighted blend of the representations of all tokens before it in the context. The weights come from a similarity (query-key) between the *current* token and every *prior* token, normalized, and used to mix the value

vectors. It is the mechanism that makes a model “pay attention to” relevant earlier tokens when producing the next one.

The architect-relevant fact: attention’s cost scales **quadratically with context length** in the full-attention regime — the number of input×output position pairs grows as  $L^2$ . That quadratic term is the single biggest reason context windows feel expensive, and it is the pivot for every optimization downstream (sparse/linear attention, KV caching, retrieval to keep context short).

One precision matters before the mental model hardens: the  $L^2$  term is a **prefill** cost — it describes computing attention across all position pairs in one pass over the full input. During **decode**, each generated token attends once to the growing prefix, so the per-token step is a *linear* cost in context length, and the length shows up as a *growing KV-cache memory footprint* (per-token key/values) rather than as per-step compute. The practical read is therefore “*long context is quadratic in prefill compute, and linear-but-memory-growing in decode*” — a distinction as important as the quadratic fact itself, and one we return to in [Chapter 2](#) (compute gears), [Chapter 7](#) (memory arithmetic) and [Chapter 8](#) (roofline). Keeping prefill and decode separate is what prevents the over-simple “long context = quadratic inference” once we optimize context in later chapters.

One architectural consequence to hold now: *which layers emit a per-token cache is a property of the attention architecture*, not the context length alone — dense attention emits a key/value per token on every layer, while hybrid/linear-attention layers (Mamba-style state, compressed latent attention) hold a state that does not grow with length. This is why “the KV problem” and its fixes are really a *model-architecture* question (Ch. 3) that lands as a *memory* question (Ch. 7) — and why the token layer is the right place to have met it. We only need the concept here; the mechanism inventory is [Chapter 3](#)’s, the arithmetic [Chapter 7](#)’s.

### 1.2.5. KV cache as a concept

When the model generates one token at a time, it recomputes the same prefix repeatedly unless it saves the intermediate results. The **KV cache** (key–value cache) is the stored key and value tensors for every token in the context, so that generating the next token only computes attention against the new token rather than re-deriving the whole prefix. It is the reason “long conversations are expensive in memory, not just in compute.”

KV cache size per token scales with

$$KV_{\text{per-token}} = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times \text{bytes-per-value}$$

(one key and one value per layer) — on the order of **~256 KB per token for a 7B-class model and ~1.3 MB per token for a 70B-class model** at 8-bit precision, so a long context (say 32K tokens) alone can be tens of GB of memory. A first non-obvious architectural lesson hides here: *model sparsity does not reduce the KV cache*. A Mixture-of-Experts model activates only a fraction of its neurons per token, which saves compute, but the attention layers still process every token densely and emit a key/value per token — so MoE shrinks compute-per-token, not memory-per-context. Only changing




FIGURE 1.1. Fig 1.1 — The KV cache: every decoded token adds one K+V per layer per head [ILLUSTRATIVE conceptual]

the *attention mechanism* (hybrid/linear attention, compressed latent attention) actually stops the cache from growing. This distinction — compute-sparsity vs memory-sparsity — is one an architect has to get exactly right, and it is why “the KV problem” is an attention-architecture question more than a model-size question (developed fully in Ch. 3 and Ch. 7). We deliberately keep this to concept + order of magnitude here; [Chapter 7](#) (Memory) does the verified arithmetic against the canonical scenario.

*Fig 1.1 — The KV cache: attention keeps every past token’s key and value, so cache memory grows linearly with sequence length (concept; the verified arithmetic is Ch. 7’s).*

### 1.3. Mental Model

Think of a token as a **metered unit of thought**, the way electricity is metered by the kilowatt-hour — important not because a single kilowatt-hour is meaningful by itself, but because *every* downstream cost and capacity number is denominated in it.

A useful image for the full pipeline:




FIGURE 1.2. Fig 1.2 — How a token travels: text → tokenizer → IDs → embedding → attention → KV cache [ILLUSTRATIVE conceptual]

*Fig 1.2 — How a token travels. The token is the interface between a human’s request and the machine’s arithmetic: carved by the tokenizer, embedded into a vector, attended against past context, and remembered in the KV cache. (Concept; mechanism inventory in Ch. 3, arithmetic in Ch. 7.)*

Text → **tokenizer** → discrete IDs → **embedding** → vectors → **attention** recomputes  
each token → **KV cache** remembers the prefix → generate one token at a time.

We do not need to hold all of this as a practitioner. The durable mental model is: **tokens are the interface between a human’s request and a machine’s arithmetic**, and *anything* we are asked to size or price in an LLM system reduces, one way or another, to “how many tokens, in what context, with what precision.”

#### 1.4. Worked Example

The canonical scenario (defined in the Architecture doc, §14 — one set of numbers every chapter cites) is an enterprise Q&A system: ~2,000 registered users, ~10 requests/s average, prompt of ~1,200 tokens plus ~8K retrieved context (~9.2K input), ~300-token output, on a 70B-class dense model. We reuse those numbers everywhere, so let us read one concrete property off them here.

*Worked example, not reference.* If average request is ~9.2K input + ~300 output ≈ **9.5K tokens** per request, and traffic is ~10 req/s, then the system ingests:

- **Input tokens/s:** 9,200 tokens/req × 10 req/s = **92,000 tokens/s** (peaks higher at ~40 req/s → ~368k).
- **Output tokens/s:** 300 × 10 = **3,000 tokens/s** output.
- Roughly **30× more input tokens than output tokens** in this workload (9.2K ÷ 300).

**Table 1-1** — Token throughput of the canonical workload (worked example, not reference).

| metric                       | value   | derivation                                |
|------------------------------|---------|-------------------------------------------|
| input tokens / request       | ~9,200  | §14 canonical (1,200 prompt + 8K context) |
| output tokens / request      | ~300    | §14 canonical                             |
| input / output ratio         | ~30×    | 9,200 ÷ 300 [2° DERIVED]                  |
| input tokens/s @ 10 rps      | ~92,000 | 9,200 × 10 [2° DERIVED]                   |
| input tokens/s @ peak 40 rps | ~368k   | 9,200 × 40 [2° DERIVED]                   |
| output tokens/s @ 10 rps     | ~3,000  | 300 × 10 [2° DERIVED]                     |

(All figures trace to the §14 canonical scenario; none are measurement claims.)

This ratio — input-heavy — is the *reason* memory (KV cache) and prefill dominate this system's cost rather than decode. We will not act on it here; [Chapter 7](#) (Memory) and [Chapter 8](#) (Compute) do. The point of the worked example in this chapter is to fix *the unit*: every one of those numbers is a statement about tokens, and we can now say it without translating from words or pages.

## 1.5. Measurement

For this concept chapter, measurement is about **counting tokens correctly**, because a miscount at the input silently misprices everything downstream. Three things an architect can actually check:

1. **Actual token count, not the rule of thumb.** Run the model's *own tokenizer* on representative prompts. The “4 chars  $\approx$  1 token” heuristic is for estimation only; real counts differ by language, formatting, code, and tokenizer version.
2. **Input vs output split.** Measure both legs of the request (prompt tokens and generated tokens), because they land on different bottlenecks — input on memory/prefill, output on decode — and on different cost line items.
3. **Peak vs average context.** Log the distribution of context lengths, not just the mean. A workload that averages 9.2K input tokens but peaks far higher (illustrative variant, say 32K) has a very different KV-cache and latency profile.

This measurement habit is the token-layer answer to the book's recurring question, “what would I actually measure here?” — we measure token counts and their distribution, at the edge, before any architecture decision is made.

## 1.6. Common Mistakes

- **Assuming one token  $\approx$  one word.** It is a useful heuristic for general English and nothing more. Code, numerics, and non-Latin scripts routinely run 2–5 $\times$  the naive estimate, which misprices capacity and cost.
- **Ignoring the input/output asymmetry.** Treating a request as “one unit” hides that input-heavy workloads are memory-bound and output-heavy ones are decode-bound — opposite bottlenecks with opposite fixes.
- **Quoting context window as free capacity.** The window is an upper bound, not a recommendation; using it fully is expensive. Keeping a 9.2K average inside a 128K window still costs for the length we actually use.
- **Trusting a vendor's tokenizer count without checking.** Tokenizer versions change and report differently; measure on *our* traffic, not the marketing figure.

### 1.7. Architecture Consequence

Whatever we learned here, file it under “**can I size an input before I buy capacity?**” Concretely, the decision this chapter enables is: *estimate tokens-per-second (input and output separately), and therefore tokens-per-second capacity and cost, before choosing a model or a serving configuration*. It is the load-bearing input to model selection (Ch. 5), the memory budget (Ch. 7), the compute budget (Ch. 8), and the token economics (Ch. 16). No downstream chapter runs without a token count to price against — this one gives us the unit.

### 1.8. What We Still Don’t Know

*As of 2026-08:* tokenizer internals are rarely published in full, so exact per-token behavior for a given model is often only empirically observable, not specified ([VERIFY] — tokenizer construction details, e.g. exact BPE merge rules for commercial models, are not fully public). Whether a fully learned, near-optimal tokenization for *our* domain can beat the generic vocabulary by a repeatable, quantified margin is still an open, workload-dependent question. And the interaction between tokenizer choice and downstream reasoning quality is not yet crisply characterized — we flag it as an open area rather than a settled fact.

### 1.9. End-of-Chapter Mini-Case

*(Chains the continuous scenario — this is its first appearance.)*

An architect is pulled into an early-stage design conversation about a new internal Q&A tool. There is no model, no server, no budget — just the raw ask: “We have thousands of employees who want to ask questions over our internal documents.” Before any architecture can be defended, we have to turn that vague desire into a *machine-denominated* statement.

from the token layer alone (which we worked through in this chapter), the architect can already establish: the unit is tokens; the request shape will be prompt + retrieved context + output; the workload is input-heavy; and the first number to lock down is tokens-per-request, because every downstream decision (which model fits, how much memory, what latency is possible) is priced against it. The specific sizing — turning “thousands of employees” into ~2,000 registered users, ~10 req/s, ~9.2K input tokens — is [Chapter 4](#)’s job (workload anatomy). Here the point is narrower and sharper: *we can now speak the system’s currency*, and that is a prerequisite for every architecture question that follows.

## CHAPTER 2.

# What Actually Happens During Inference

### 2.1. The Architect’s Question

After this chapter we should be able to reason about the two fundamental modes of LLM inference — prefill and decode — and to identify which system constraint each mode exposes. We should be able to compute the bandwidth and FLOP demands of a given workload, compare them against real hardware limits, and understand why decode is HBM-bandwidth-bound while prefill is compute-bound. This distinction directly shapes architectural decisions: whether we need more GPUs, whether we can disaggregate prefill from decode, and how we size memory and compute in a fleet. After this chapter we should see inference not as a single monolithic cost but as two opposite-bound processes that require different levers.

### 2.2. Concept

Inference is the process of turning a sequence of input tokens into a sequence of output tokens using a trained neural network. In an autoregressive model, every generated token requires a fresh forward pass, but the computation split between the prefix (the prompt) and the suffix (the generated output) is fundamentally different.

**Prefill** is the first pass over the prompt. The model reads every token of the input context, computes attention over the entire prefix, and produces the first output token. During prefill, the model attends over all input positions, so the work grows with the square of the context length.

**Decode** is the per-token generation that follows. After prefill, each new token only attends to the cached keys and values from the prefix plus all previously generated tokens. The per-step cost is constant in context length — we only attend to the new token against the cache — but it requires

reading the entire model weight matrix from HBM every step, because the auto-regressive pass reuses the same parameters over and over.

The critical architectural insight: prefill is dominated by FLOPs (matrix-multiply arithmetic), while decode is dominated by HBM memory bandwidth (reading 140 GB of weights 40 times a second). These are opposite bottlenecks, and confusing them is the most common architectural misstep.

### 2.3. Mental Model

Think of inference as two fundamentally different physical processes, distinguished by what limits them.

**Prefill is compute-bound.** It is like turning on a massive industrial fire hose to fill a reservoir. The burst of flow is enormous, and the time it takes is limited entirely by the raw horsepower of the pump (compute/FLOPs) pushing the water. The pipe is wide open, but the motor is working at its absolute limit.

**Decode is bandwidth-bound.** It is like trying to empty a 140,000-gallon swimming pool through a standard garden hose every 25 milliseconds just to produce a single glass of water. The pump (compute) is more than powerful enough, but the diameter of the hose (HBM bandwidth) physically cannot move that volume of water in that time frame.

When we ask “is this workload compute-bound or bandwidth-bound?” the answer depends entirely on which physical constraint we are hitting: the horsepower of the engine, or the diameter of the pipe.

### 2.4. Worked Example

To make the distinction concrete, let us walk through the canonical enterprise Q&A scenario from §14: a 70B-class dense model in FP16, 1 host with  $8 \times$  H100 GPUs, prompt of 1,200 tokens + 8K retrieved context (~9.2K input), 300-token output, TTFT budget 1.2 s (retrieval ~120 ms + prefill), TPOT budget ~25 ms/token. All numbers are derived from the canonical scenario; none are measurement claims.

#### 2.4.1. Table 2-1 — Decode bandwidth and prefill FLOPs (worked example, not reference)

| metric                    | value  | derivation                               |
|---------------------------|--------|------------------------------------------|
| Model params              | 70B    | §14 canonical                            |
| Model weight bytes (FP16) | 140 GB | 2 bytes $\times$ 70B params [1P DERIVED] |

| metric                                   | value                     | derivation                                                                                                                            |
|------------------------------------------|---------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| Decode: weight-read per token            | 140 GB                    | Auto-regressive: one read of all weights per generated token [1P DERIVED]                                                             |
| Decode: required bandwidth (TPOT ~25 ms) | 5.6 TB/s                  | 140 GB / 0.025 s [DERIVED]                                                                                                            |
| H100 HBM3 peak bandwidth                 | 3.35 TB/s                 | NVIDIA H100 specs [2° FACT]                                                                                                           |
|                                          | Decode: bandwidth verdict | bandwidth-bound                                                                                                                       |
| Prefill: input tokens                    | 9,200                     | 1,200 prompt + 8K context [1P DERIVED]                                                                                                |
| Prefill: FLOPs (2 × params × tokens)     | 1.29 PFLOP                | $2 \times 70e9 \times 9.2e3 \approx 1.29 \times 10^{15}$ [DERIVED]                                                                    |
| H100 BF16 dense compute                  | 989 TFLOPS                | NVIDIA H100 BF16 tensor-core peak [2° FACT]                                                                                           |
| Prefill: compute verdict                 | compute-bound             | 1.29 PFLOP / ~1.08 s $\approx$ 1.19 PFLOPS > 0.989 PFLOPS peak $\rightarrow$ single H100 insufficient for real-time prefill [DERIVED] |

*Fig 2.1 — Decode is HBM-bandwidth-bound (5.6 TB/s vs H100 3.35 TB/s); prefill is compute-bound (~1.19 PFLOPS required vs H100 0.989 PFLOPS peak).*

**Worked example arithmetic details:**

- **Decode bandwidth.** In auto-regressive decode, generating one token requires a full forward pass of the 70B parameter matrix. At FP16 (2 bytes per parameter), the weight footprint is  $W = N \times 2 = 140$  GB. With a TPOT budget of ~25 ms per token, the system must read all  $W$  bytes in  $\tau = 0.025$  s, so the required bandwidth is

$$B_{\text{req}} = \frac{W}{\tau} = \frac{140 \text{ GB}}{0.025 \text{ s}} = 5,600 \text{ GB/s} = 5.6 \text{ TB/s}$$

[2° DERIVED]. A single NVIDIA H100 HBM3 delivers ~3.35 TB/s peak bandwidth [2° FACT]. Since  $5.6 > 3.35$ , one H100 cannot supply the required weight-read rate — the decode phase is HBM-bandwidth-bound [2° DERIVED], and serving 70B-class output requires multi-GPU scaling or




FIGURE 2.1. Fig 2.1 — Decode vs prefill: bandwidth vs compute [ILLUSTRATIVE conceptual]

bandwidth-increasing topologies (e.g. NVLink-connected nodes).

- **Prefill FLOPs.** The prefill pass computes attention over the 9.2K input tokens and produces the first output token. The FLOP count for a dense transformer forward pass is well approximated as  $2 \times \text{params} \times \text{tokens}$  (the factor of 2 accounts for multiply-add per parameter per token). Thus:

$$\text{prefill FLOPs} = 2 \times N \times L = 2 \times 70 \times 10^9 \times 9.2 \times 10^3 \approx 1.288 \times 10^{15} \approx 1.29 \text{ PFLOP}$$

[2° DERIVED]. An NVIDIA H100 delivers ~989 TFLOPS (BF16 dense tensor-core peak) [2° FACT], which is 0.989 PFLOPS. The required rate against the ~1.08 s prefill budget is

$$\text{rate} = \frac{1.29 \text{ PFLOP}}{1.08 \text{ s}} \approx 1.19 \text{ PFLOPS}$$

Crossing  $1.19 > 0.989$  — a single H100 cannot execute the prefill FLOPs within the TTFT budget; the prefill phase is compute-bound [2° DERIVED]. Multiple GPUs (model parallelism or data parallelism) or more efficient attention implementations are required to meet the latency SLO.

## 2.5. Measurement

How do we know whether a given workload is prefill-bound or decode-bound in practice? Three practical measurements:

1. **Divide the TTFT into retrieval + prefill.** Measure the retrieval latency (vector search, document fetch) separately from the prefill latency (prompt processing). In the canonical scenario, retrieval takes ~120 ms, leaving ~1.08 s for prefill on a 9.2K-token prompt. If prefill exceeds this budget, the bottleneck is compute or memory, not retrieval.
2. **Log TPOT per token across the request.** Track the time-per-token for the generated output. If TPOT is ~25 ms and does not decrease as the output lengthens (because the KV cache amortizes the attention over the prefix), the per-step cost is constant — consistent with decode being bandwidth-bound (same weights read each step). If TPOT drops with longer contexts, the model may be compute-bound per step (attention over more keys/values adds FLOPs).
3. **Measure actual HBM utilization vs FLOP utilization.** Using GPU profilers (NVIDIA Nsight, vLLM stats), watch the fraction of peak HBM bandwidth consumed versus the fraction of peak FLOPS consumed. If HBM utilization is near 100% while FLOPS utilization is low, the kernel is bandwidth-bound (typical of decode). If FLOPS utilization is near 100% while HBM utilization is low, the kernel is compute-bound (typical of prefill).

These measurements cost nothing but a profiler attach and a few representative requests — and they prevent the architect from applying the wrong optimization lever.

## 2.6. Common Mistakes

- **Treating prefill and decode as the same bottleneck.** The most dangerous mistake is assuming that what slows prefill will also slow decode, or vice versa. Prefill FLOPs scale with input length; decode bandwidth is constant per token. Confusing the two leads to over-provisioning compute for a bandwidth-bound workload, or under-provisioning compute for a compute-bound one.
- **Assuming one H100 suffices for 70B inference.** The arithmetic in Section 4 shows that a single H100 cannot meet the decode bandwidth demand (5.6 TB/s vs 3.35 TB/s) nor the prefill compute demand (1.29 PFLOP vs 0.989 PFLOPS). Attributing 70B serveability to a single GPU ignores the opposite bottlenecks and produces an architecture that fails latency SLOs.
- **Ignoring the input/output token asymmetry.** A workload with 9.2K input tokens and 300 output tokens spends the great majority of its latency and energy in prefill, not decode. Quoting only throughput per second or only per-token latency hides that the two legs of the request have different cost structures and different optimization paths.
- **Using manufacturer-peak bandwidth/FLOPS without mapping to real kernels.** H100 3.35 TB/s HBM3 is the theoretical peak; actual sustained bandwidth for weight-read kernels may be 60–80% of peak. H100 989 TFLOPS BF16 dense is the tensor-core peak; actual matmul throughput depends on kernel fusion, batching, and precision. Citing raw numbers without kernel context is an anti-catalog error — the number must be matched to the actual serving kernel in use.

## 2.7. Architecture Consequence

The opposite bottlenecks of prefill and decode have immediate architectural consequences:

- **Prefill is compute-bound → optimization levers:** kernel fusion (flash-attention, scaled-dot-product attention with fewer reads), model parallelism (splitting the 70B parameters across GPUs), more GPUs in parallel, quantization (FP8/BF16 reduces FLOP count and memory traffic), or prefill-specific servers that dedicate hardware to the bursty preload phase.
- **Decode is bandwidth-bound → optimization levers:** continuous batching (vLLM, PagedAttention) to amortize weight reads across many tokens in flight, KV cache offloading to SSD or system memory, P/D (prefill/decode) disaggregation (separate pools of GPUs for prefill bursts vs. decode steady state), and higher-bandwidth memory topologies (NVSwitch, HBM3e).

The key architectural decision this enables: **can we disaggregate prefill and decode onto different hardware?** If prefill needs more FLOPS and decode needs more bandwidth, a single homogeneous GPU pool is suboptimal. A two-pool architecture — a prefill cluster optimized for compute (more GPUs, higher FLOPS) and a decode cluster optimized for bandwidth (faster HBM, larger NVLink domains) — can achieve the same serving SLO with fewer total GPUs than a homogeneous design. This is the central theme of [Chapter 11](#) (Serving) and Pattern 4 (Prefill/decode disaggrega-

tion).

## 2.8. What We Still Don't Know

- **Sustained HBM bandwidth for weight-read kernels.** The 3.35 TB/s H100 figure is a theoretical peak; real serving kernels (vLLM, TGI, transformer-engine) may achieve different sustained rates depending on kernel fusion, graph optimization, and batch size. [VERIFY HYPOTHESIS] — measure on real hardware before citing.
- **Prefill FLOP cost per token under continuous batching.** The  $2 \times \text{params} \times \text{tokens}$  approximation ignores that continuous batching reuses activations across requests, potentially reducing the effective FLOP count per request. The magnitude of this effect at fleet scale is not yet pinned down. [VERIFY HYPOTHESIS] — benchmark with realistic concurrency patterns.
- **Cross-technology bandwidth numbers.** HBM3e (next-generation) promises ~6+ TB/s per GPU, and AMD MI300X promises ~5.3 TB/s. How these compare to the 5.6 TB/s decode demand shifts the GPU count equation but does not change the fundamental bandwidth-bound vs compute-bound classification. [2° FACT] — vendor-published numbers, verify against the specific generation in use.
- **Effect of quantization on decode bandwidth vs prefill FLOPs.** Int4 or FP8 quantization reduces the weight bytes (e.g., 70B at FP8 = 70 GB instead of 140 GB) and may change the FLOP arithmetic (8-bit matmuls may use tensor cores at different utilization). The direction of the shift is clear (both bottlenecks improve), but the precise trade-off point where decode becomes compute-bound rather than bandwidth-bound depends on the quantization scheme and hardware support. [VERIFY HYPOTHESIS] — workload-dependent.

## 2.9. End-of-Chapter Mini-Case

The architect is still in the early design conversation about the internal Q&A tool described in [Chapter 1](#). The team has settled on a 70B-class dense model for accuracy, and they have a rough traffic estimate: ~2,000 registered employees, each roughly as active as the canonical scenario (~10 requests/s average, peaks ~40 rps), with prompts of ~1,200 tokens of internal documents plus ~8K retrieved context and ~300 tokens of answer.

From the token layer alone, the architect can already state the hard numbers: each request carries ~9.2K input tokens and ~300 output tokens. At 10 rps average, the system must sustain ~92,000 input tokens/s in prefill and ~3,000 output tokens/s in decode. At 40 rps peak, those numbers jump to ~368,000 and ~12,000 respectively. The architect can also state the two opposite bottlenecks: prefill will be compute-bound (~1.29 PFLOP per request at 9.2K input), and decode will be bandwidth-bound (~5.6 TB/s required weight-read rate per token). The team's immediate architectural choice is whether to build a single homogeneous GPU cluster or to separate prefill and decode pools — the

arithmetic from this chapter makes clear that a homogeneous design will be suboptimal for either bottleneck, and that the disaggregation option explored in [Chapter 11](#) is worth evaluating early.

Before passing this workload to [Chapter 4](#) (workload anatomy), the architect locks in one more number: the tokens-per-request split. This is the currency every downstream chapter will price against, and it has already been established as  $\sim 9.2\text{K}$  input +  $\sim 300$  output per request. The rest of the design — model fitting, memory budget, latency budget, cost — can now proceed in token-denominated terms.

## CHAPTER 3.

# Understanding the Model

### 3.1. The Architect's Question

After this chapter we should be able to reason about what a model's parameter structure means for deployment: dense versus Mixture-of-Experts (MoE), total parameters versus active parameters per token, and what each implies for compute budgets and memory budgets. We will meet the mechanisms (how MoE routes, how attention weights are computed) not to learn how to train a model but so that, as architects, we can look at a model card and immediately know whether the quoted parameter count is what we actually pay for at inference time, and whether the KV cache will grow as expected. After this chapter we can ask: *given a workload and a model architecture, what is the real residency and compute cost per token?*

### 3.2. Concept

The **parameter allocation regime** of a model determines how many of its weights are touched on each forward pass. This regime is not a property of the model family name alone; it is a consequence of the training objective and the routing mechanism.

- **Dense models** activate *all* parameters for every token. A 70B dense model reads all 70B weights (or close to it) for every token, because every layer's full hidden dimension participates in every attention and feed-forward computation. The parameter count on the model card is the inference cost count.
- **Mixture-of-Experts (MoE) models** route each token to a *fraction* of the total parameters. A MoE model may have 236B total parameters distributed across eight experts, but only the top-2 (or top-1) experts are activated per token. The quoted total parameter count includes the unused

experts; the *active* parameter count per token is the fraction that routing selects.

The architect-relevant distinction: **total parameters  $\neq$  active parameters per token**. This inequality is the single most important number to read from a model card when sizing a deployment, because it directly changes the compute cost per token. However, the KV cache does not share this privilege — we will show why MoE sparsity saves compute but not memory.

A note on provenance: the total-versus-active distinction arises from the training regime. In dense pre-training, every token sees the full model gradient; in MoE pre-training, a routing loss directs each token to the most relevant expert(s), and the unused experts receive no gradient for that step. The routing mechanism itself (top-k, importance sampling, learned gates) is an architectural decision that outlives training and becomes a deployment fact.

### 3.3. Mental Model

Think of a model’s parameters as a **resource pool** that is allocated differently depending on the routing regime.

- In a **dense** model, the pool is *fully allocated* every time. If the model is 70B params in FP16, we need 140 GB of weight memory available for every token. There is no “spending less” — the full 140 GB is the residency floor.
- In an **MoE** model, the pool is *fractionally allocated*. The total pool may be large (reflecting the model’s capacity and training compute), but what we actually wire up per token is a slice. With top-2 routing from 8 experts, roughly  $2/8 = 25\%$  of the total parameter mass is active per token. This fraction saves compute (fewer FLOPs) but the attention layers still see every token, so the KV cache still grows with context length just as in a dense model.

The durable mental model: **total parameters set the ceiling; active parameters per token set the floor**. The architect must track both numbers.

### 3.4. Worked Example

We now carry concrete arithmetic using the canonical enterprise Q&A scenario (§14: ~2,000 users, ~10 rps average, ~40 rps peak, 1,200 + 8K context ~9.2K input, 300-token output, 70B-class model, FP16, 1 host 8×H100). The numbers below are worked out explicitly; where a model’s exact spec is not first-party verified, it is labeled [VERIFY].

#### 3.4.1. Dense 70B model (canonical, [1P])

| Metric                           | Value              | Derivation                                                                    |
|----------------------------------|--------------------|-------------------------------------------------------------------------------|
| Total parameters                 | 70B                | §14 canonical scenario [1P]                                                   |
| FP16 residency (total)           | 140 GB             | $70\text{B} \times 2 \text{ bytes} = 140 \text{ GB}$ [1P DERIVED]             |
| Active parameters per token      | 70B                | All parameters active for every token [1P FACT]                               |
| FLOPs per forward pass (approx.) | $\sim 2.8\text{T}$ | $70\text{B} \times 4 \text{ FLOP/param}$ typical for transformer [2° DERIVED] |

Table 3.1 — Dense 70B model parameter arithmetic (worked example, not reference).

### 3.4.2. MoE model representative (Mixtral 8x7B class, [2°])

We contrast with a representative MoE model in the Mixtral lineage. The exact parameter split is not always published in vendor model cards; the numbers below are plausibility-checked and flagged.

| Metric                                      | Value                           | Derivation                                                                                                                                                    |
|---------------------------------------------|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Total parameters (across 8 experts)         | $\sim 47\text{B}$               | $\sim 6\text{B per expert} \times 8 + \text{embeddings}$ [2° DERIVED]                                                                                         |
| Active parameters per token (top-2 routing) | $\sim 14\text{B}$               | Top-2 from 8 experts, $\sim 3\text{--}5\%$ active/total ratio (frontier MoE [2° DERIVED]); actual activation $\sim 25\%$ for top-2-from-8 configuration shown |
| FP16 active residency per token             | $\sim 28 \text{ GB}$            | $14\text{B} \times 2 \text{ bytes} = 28 \text{ GB}$ [VERIFY DERIVED]                                                                                          |
| FLOPs per forward pass (approx.)            | $\sim 0.56\text{T}$             | $14\text{B active} \times 4 \text{ FLOP/param}$ [VERIFY DERIVED]                                                                                              |
| Compute ratio (active/total)                | $\sim 3\text{--}5\%$ (frontier) | $14\text{B} \div 47\text{B}$ [2° DERIVED]                                                                                                                     |

Table 3.2 — MoE model parameter arithmetic (worked example, not reference).

### 3.4.3. Why MoE saves compute but not KV-cache memory

The compute savings are straightforward: if only  $n_a = 14\text{B}$  of  $n_t = 47\text{B}$  total parameters are active per token, the FLOP count drops by the active fraction. Since per-token FLOPs scale with the number of parameters touched,

$$\text{FLOP}_{\text{MoE/token}} = \frac{n_a}{n_t} \times \text{FLOP}_{\text{dense/token}} = \frac{14}{47} \times 2.8 \text{ T} \approx 0.56 \text{ T}$$

— roughly a  $5\times$  reduction in compute per token. This is why MoE models can be “larger” (more total parameters for the same training FLOP budget) while keeping per-token cost comparable to a smaller dense model.

The KV-cache story is the surprising part. Recall from [Chapter 1](#) (§KV cache as a concept, [Chapter 1](#)) that the KV cache stores one key and one value tensor per token, per layer. The cache size per token scales with the canonical [Chapter 1](#) formula,  $KV_{\text{per-token}} = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times \text{bytes-per-value}$ . Critically, **the attention mechanism processes every token in the context densely** — each token’s query attends to all previous tokens’ keys and values, regardless of whether the model is dense or MoE. MoE sparsity operates in the feed-forward sub-layer; it does not change the attention sub-layer’s behavior.

Therefore, for a 70B-class model (dense or MoE) with 8K–9.2K input context:

- **Dense 70B:** KV cache  $\sim 1.3 \text{ MB per token} \times 9,200 \text{ tokens} \approx 12 \text{ GB}$  (at 8-bit) or  $\sim 24 \text{ GB}$  (at FP16-relevant precision for the key/value storage pattern discussed in Ch. 7).
- **MoE 70B-equivalent:** The KV cache is *identical* in size, because attention still processes all 9,200 tokens densely. The fact that only 14B of 47B total parameters are active in the feed-forward layers does not reduce the number of keys/values emitted by the attention layers.

In other words: MoE sparsity = compute sparsity (fewer FLOPs per token). MoE sparsity  $\neq$  memory sparsity (KV cache still grows linearly with context length, same as dense). Only changing the attention mechanism itself (e.g., hybrid/linear attention, compressed latent attention, or KV offload) can stop the cache from growing. This distinction — compute-sparsity versus memory-sparsity — is one an architect has to get exactly right.

**Key takeaway:** MoE lets us fit a larger total parameter budget within a weight-residency constraint (active weights take less FP16 memory), but the KV-cache budget must be provisioned as if the model were dense. Do not assume MoE reduces KV-cache memory.

## 3.5. Measurement

How do we determine the total-versus-active split for a model we haven’t trained?

1. **Model card inspection.** Most model vendors publish the total parameter count. The active-per-token count is harder to find; it requires reading the architecture documentation for the routing

mechanism (top-k, top-1, importance-based). If the model card says “70B parameters” without qualification, assume dense unless it explicitly describes an MoE layout.

2. **Runtime measurement.** For a given input, profile the number of FLOPs or the memory footprint of active weights. Tools such as `torch.profiler` can count active parameters, but this is an empirical measurement, not a first-party spec.
3. **Architecture diagram.** For MoE models, the number of experts  $\times$  parameters-per-expert gives the total; the routing hyperparameter (top-k)  $\times$  parameters-per-expert gives the active per token. This is a derived fact from the source code/release, not always on the model card.

**Practical habit:** When we encounter a model quoted in “X billion parameters,” pause and ask: *is this total or active?* For dense models they are the same; for MoE models they diverge. Misreading this is the most common source of KV-cache and compute underestimation.

### 3.6. Common Mistakes

- **Assuming total parameters = active parameters per token.** This is the single most costly misread. For a MoE model, the quoted “236B” or “47B” is the total across all experts; the active count per token is a fraction. Using the total number to size KV cache or memory will overestimate compute needs (if we think we need  $236B \times 2$  bytes we’ll over-provision weight memory) or underestimate it (if we size for active but forget KV cache is dense).
- **Confusing compute sparsity with memory sparsity.** MoE reduces FLOPs per token because fewer parameters are active. It does *not* reduce the number of keys/values written to the KV cache. An architect who assumes MoE shrinks KV cache will be blindsided by memory pressure at long context lengths.
- **Using MoE as a free lunch for long-context workloads.** The compute savings from MoE are real, but the memory cost (weights + KV cache) must be provisioned at the total-parameter scale for weight residency and at the-full-context scale for KV cache. Neither benefit is “free” in the other domain.
- **Ignoring routing overhead.** Top-k routing itself requires computing routing scores (a softmax over expert assignments). This adds a small but non-zero compute cost that is sometimes omitted from published FLOP counts. It is usually negligible compared to the feed-forward arithmetic but should be acknowledged.

### 3.7. Architecture Consequence

Knowing whether a model is dense or MoE, and whether the quoted parameter count is total or active, directly changes three architectural decisions:

1. **Weight residency budget.** For a dense 70B model, provision 140 GB of FP16 weight memory per host. For a MoE model with ~47B total and ~14B active, provision ~28 GB for the active weights + a smaller overhead for the unused experts (they may be swapped or kept on device depending on the deployment framework). MoE allows a larger total parameter budget within the same weight-memory ceiling, but the active-fraction budget is what matters for “weights on device.”
2. **KV-cache provisioning.** Provision the KV cache as for a dense model of the same context length and hidden dimension. Do not apply a MoE sparsity factor to the KV cache. If we are running a 9.2K-token input on a 70B-class model (dense or MoE), the KV cache memory is the same order of magnitude — plan accordingly.
3. **Parallelism and routing strategy.** MoE deployment requires a routing strategy (e.g., expert placement, dispatcher, switchboard) and often expert parallelism across GPUs. This adds infrastructure complexity (cross-GPU communication for expert routing) that dense models do not have. The architectural choice between dense and MoE is therefore not just a parameter-count decision; it is a deployment-complexity decision.

### 3.7.1. The 2026 baseline shift: same framework, new constants

**The 2026 trend signal (forward-looking).** By August 2026 the four frontier open-weight families converge on one direction, and it is not more of the same dense scaling: extreme MoE sparsity (single-digit active-parameter fractions — Qwen 6 B of 125 B, Kimi K3 104 B of 2.8 T) coupled with **hybrid attention** (compressed/sparse/linear hybrids such as Gated DeltaNet + sparse attention) has become the standard way to make long-context inference affordable. [Chapter 27](#) (Appendix A) documents these at full depth with first-party provenance; what matters here, in the model-understanding chapter, is that this is not a competing mechanism but the *same* dense-vs-MoE, total-vs-active, KV-vs-compute machinery this chapter just taught — with new constants. We hold that view explicitly as the **forward-looking direction** the architect should size toward, and flag it below as continuing future work (the frontier moves; this handbook is a living document, and Appendix A is where the moving part lives).

**Why the trend signal should not be read as “dense is dead”.** The convergence of the four *frontier* families on extreme MoE sparsity describes the top of the market — the ~100B-and-up tier that anchor ~2.8T-parameter fleets. It is not a claim that dense models have disappeared. On the contrary, the smaller, denser tier is alive and actively shipping in 2026 for exactly the constrained-deployment reasons this chapter teaches: **Qwen3.8-27B** ([1P: HF Qwen/Qwen3.8-27B]) is a dense 27.8B model whose 4-bit weights fit on a single 24 GB consumer GPU, and **Muse Glimmer 30B** ([1P: Meta via vLLM-Recipes]) is a dense 29.6B local-agentic model that runs in ~18 GB. These exist because the dense-vs-MoE trade is itself a *deployment* decision, not a date: MoE buys active-parameter efficiency at the cost of total-parameter memory and routing complexity — worthwhile when you serve many requests from a large pool of resident weights, but hard to justify when your whole budget is one consumer GPU whose total weights already overrun. That is precisely why this book keeps its canon-

ical on the dense tier: the canonical must serve the architect whose *whole fleet* is a handful of boxes, not only the one sizing a 2.8T datacenter model. The framework does not change between tiers — total vs active, KV per-token, and residency budgeting apply identically to a 27B dense laptop model and a 125B/6B MoE fleet model.

**Why the rest of this chapter’s canonical stays a dense full-MHA model.** If the trend is MoE + hybrid, why does this book’s canonical still size a dense 70B full-MHA workload? Because a dense model with full multi-head attention gives the cleanest single-constant KV formula ( $2 \times \text{layers} \times \text{hidden} \times \text{bytes}$ ) and the largest KV footprint — the conservative worst case an architect should always size against first. Teaching against that bound means the reader learns the framework on the hardest memory case and then relaxes it. It is a deliberate teaching simplification, not a claim that in 2026 an enterprise RAG fleet really runs on dense 70B.

**One verified 2026 transfer: Qwen3.8-Flash-Next.** The model card [1P: HF Qwen/Qwen3.8-Flash-Next] is the Qwen4-architecture preview: **125 B total / 6 B active** (plus a separate 51 B of offloadable N-gram embedding parameters), **512 MoE experts** (10 routed + 1 shared per token), and **hybrid attention** — Gated DeltaNet (GDN) on three of every four layers plus Qwen Sparse Attention (QSA) on the fourth. This is exactly the “125B/6B active MoE + Gated DeltaNet/QSA” class of model a 2026 architect actually considers. Stepping through this book’s own arithmetic:

- **Weight residency** (Ch. 7’s floor): active weights are  $\sim 6 \text{ B} \times 2 \text{ B} = \sim 12 \text{ GB}$ , not the canonical 140 GB — MoE sparsity collapsed the on-GPU weight floor by  $\sim 12\times$ . But total weights dominate the memory ceiling:  $\sim 125 \text{ B} \times 2 \text{ B} \approx 250 \text{ GB}$  (plus the offloadable N-gram), so the *total-vs-active* split of this chapter is not a footnote — it is the difference between fitting on one host and sharding across several.
- **KV cache** (Ch. 7/8): the canonical  $2 \times \text{layers} \times \text{hidden} \times \text{bytes}$  KV formula is the full-MHA bound. Qwen’s hybrid attention attacks the KV constant itself at the mechanism level — GDN keeps a fixed-size recurrent state (context cost no longer grows linearly per token across those layers) and QSA attends sparsely at micro-block granularity — so per-token KV no longer follows the dense formula. The framework question (“does it fit?”) is unchanged; the *function* you plug in for KV is not the dense one.
- **Compute** (Ch. 8): active-parameter FLOPs drop to  $\sim 2 \times 6 \text{ B}$  per token, but the QSA/GDN layers ride different roofline points than a dense 70B prefill/decode split, so the throttled-resource conclusion must be re-derived per model (Ch. 8’s measurement discipline, not a one-time number).

**The architect’s move is unchanged.** For the canonical dense model this book sizes memory, compute, and fleet with one clean constant set. For a 2026 MoE/hybrid model the architect does the *same* decision loop — read the model card, get total vs active, get the KV scheme’s per-token function, size residency and roofline, then decide parallelism (Ch. 10) and serving (Ch. 11). Only the constants differ; the framework does not. [Chapter 27](#) walks the four frontier models through exactly this re-derivation and flags which vendor numbers still carry a verify-on-own-hardware caveat.

**Recorded as future work.** This forward-looking view is intentionally a *pointer*, not the full treatment — the complete, current analysis of the 2026 frontier lives in [Chapter 27](#) (Appendix A). Because the frontier moves quickly and every vendor figure in it carries a verify-on-own-hardware caveat, we treat the trend as a living section to be updated against new model cards, rather than a set of frozen gospel numbers (per the hand-book’s “living document” framing, `book-architecture.md §10`). An architect should re-check Appendix A before any capacity decision — the constants, not the framework, are what change.

### 3.8. What We Still Don’t Know

As of 2026-08, several questions remain open and are flagged [VERIFY]:

- **Exact active-fraction per token for commercial MoE models.** Model cards rarely publish the precise top-k routing distribution under real workloads. We know the architectural top-k (e.g., top-2 from 8 experts), but the actual fraction of active parameters may vary with prompt content, token position, and expert availability. [VERIFY HYPOTHESIS]
- **Whether router latency offsets MoE compute savings.** The cost of computing routing scores and dispatching to experts is not always included in published FLOP counts. On some hardware the routing overhead can be significant enough to narrow the compute gap between MoE and dense. [VERIFY HYPOTHESIS]
- **KV-cache behavior under MoE with expert swapping.** When experts do not fit on a single GPU and must be swapped (offloaded), does the KV cache interact with the swapping mechanism in ways that change its effective size or access pattern? This has not been systematically characterized. [VERIFY HYPOTHESIS]

These flags exist because home-lab measurements are not reference per the evidence taxonomy; they will be resolved (promoted, dropped, or demoted) before publication.

### 3.9. End-of-Chapter Mini-Case

*(Continuous scenario — this is its first appearance, chaining from the enterprise Q&A thread.)*

An architect is fleshing out the design of the internal Q&A tool described in [Chapter 1](#)’s mini-case. The team is torn between a dense 70B model and a MoE model with 8 experts (advisedly total ~47B params, ~14B active per token). They ask: “If we go MoE, can we cut our GPU memory budget?”

From the model-understanding chapter, the architect can answer: *Weight residency yes — active parameters take ~28 GB vs. 140 GB for a dense 70B, we could fit the model on fewer GPUs or a smaller host. But KV-cache memory no — the attention layers still process every token densely, so the cache budget for a 9.2K-context input is the same regardless of whether the model is dense or MoE. We still*




FIGURE 3.1. Fig 3.1 — Dense vs MoE parameter allocation and KV-cache behavior [2° DERIVED]

need to provision for  $\sim 12$  GB of KV cache (at 8-bit) or  $\sim 24$  GB (at the precision pattern discussed in Ch. 7) per request, and we’ll also need expert-routing infrastructure on top of that.

The architect’s decision therefore hinges on whether the compute savings from MoE (fewer FLOPs per token, potentially lower \$/token) outweigh the added routing complexity and the fact that KV-cache memory is unchanged. The team decides on a MoE model for the compute/\$ advantage, but provisions the KV-cache budget at the dense-model rate, and adds one GPU dedicated to the routing dispatcher.

Fig 3.1 — Parameter-activation contrast. A dense 70B activates all 70B parameters per token (residency  $\approx 140$  GB, KV grows linearly with context). An 8-expert MoE activates only the top-2 ( $\sim 14B$ ) per token (residency  $\approx 28$  GB active), but KV-cache growth with context is identical to dense — attention still processes every token.

3.9.1. Table 3-1 — Parameter and KV-cache arithmetic for the canonical workload

| metric                                 | dense 70B | MoE (8×7B class) | derivation                                                                              |
|----------------------------------------|-----------|------------------|-----------------------------------------------------------------------------------------|
| total parameters                       | 70B       | ~47B             | [2°] expert split                                                                       |
| active params per token                | 70B       | ~14B             | top-2 from 8 [~3–5% frontier MoE [2° DERIVED]; Mistral top-2-from-8 ~25% for reference] |
| FP16 residency (total)                 | 140 GB    | ~94 GB           | $70B \times 2 / 47B \times 2$                                                           |
| FP16 residency (active)                | 140 GB    | ~28 GB           | $70B \times 2 / 14B \times 2$                                                           |
| KV cache per request (9.2K ctx, 8-bit) | ~12 GB    | ~12 GB           | identical — attention dense                                                             |
| FLOPs per forward pass                 | ~2.8T     | ~0.56T           | $70B \times 4 / 14B \times 4$                                                           |
| compute speedup (active/total)         | 1×        | ~5×              | $14B \div 70B$ [~3–5% frontier; Mistral top-2-from-8 ~25% for reference]                |

*All figures trace to the §14 canonical scenario; MoE column is a [2°] worked variant, not a measurement claim.*

*Next chapter: [Chapter 4](#) — The Anatomy of an AI Workload, which characterizes the enterprise Q&A RAG workload across the six dimensions (quality, traffic, token profile, latency, economic, operational constraints) and uses the same continuous mini-case scenario.*

*Canonical scenario citation: all numerical values in this chapter that trace to the §14 canonical enterprise Q&A RAG workload are drawn from the fixed source of truth defined in book-architecture.md §14. The dense 70B numbers are [1P] per the canonical scenario; MoE numbers are [2°] plausibility-checked variants.*

## **PART II.**

# **The Workload**

## CHAPTER 4.

# The Anatomy of an AI Workload

### 4.1. The Architect’s Question

After this chapter we should be able to take a vague deployment ask — “thousands of employees asking questions over internal documents” — and turn it into a *measurable workload* across six concrete dimensions. We will trace the arithmetic from user concurrency through traffic, token profile, latency, economics, and operational constraints, ending with a workload characterization that directly enables system architecture decisions. After this chapter, the question “which model should we choose?” can be answered with numbers, not intuition.

### 4.2. Concept

#### 4.2.1. The Six-Dimension Workload-Characterization Framework

A workload is defined by six orthogonal dimensions. Together they form a complete specification that — taken alone — determines model selection, system architecture, and economic feasibility. No single dimension is sufficient; the architect must characterize all six before committing to a design.

1. **Quality** — the fidelity, accuracy, and capability the workload demands of the model. Includes reasoning depth, tool-use requirements, modality (text-only, multimodal, etc.), and any accuracy thresholds (e.g. >85% factual recall on internal Q&A). Quality is the *goal*; the other dimensions are the *constraints* within which quality must be delivered.
2. **Traffic** — the request rate the system must sustain, expressed in requests per second (rps) or queries per second, plus concurrency levels and burstiness. Traffic determines the compute and memory bandwidth required to keep the serving pipeline saturated without excessive queuing.

3. **Token profile** — the distribution of input length, output length, and context length per request. This is the workload’s “shape” in token space: how many tokens arrive in the prompt, how many the model generates, and whether the prompt is dense (long context) or sparse. The token profile is the primary driver of memory (KV cache) and prefill cost.
4. **Latency** — the time budget for different stages of request processing. Typically split into *time-to-first-token* (TTFT), which encompasses retrieval and prefill, and *time-per-output-token* (TPOT), which governs decode length. Latency budgets feed directly into SLOs and serve as the key differentiator between serving configurations (e.g. continuous batching vs. discrete batching).
5. **Economic constraints** — the cost budget denominated in tokens per dollar, dollars per million tokens, or cost per request. This dimension translates traffic and token profile into a recurring cost line, and it determines whether a given infrastructure choice (single node vs. multi-node, FP16 vs. FP8) is affordable at scale.
6. **Operational constraints** — availability, privacy, locality, and update frequency. Multi-region deployment for fault tolerance. Data-privacy requirements that force on-prem or VPC-local embedding models. Update frequency that dictates how often embeddings or model weights must be refreshed. These constraints often interact with traffic and economics to force architectural compromises.

The framework is deliberately dimensional: a workload that is “high-quality, high-traffic, long-context, low-latency, cost-sensitive, private” is fundamentally different from “low-quality, low-traffic, short-context, high-latency, cost-abundant, public,” even if the raw text looks the same. The framework prevents the architect from optimizing one dimension at the expense of others — a common failure mode.

### 4.3. Mental Model

Think of the six dimensions as the axes of a six-dimensional space in which every AI workload resides. A workload’s position in this space is its *fingerprint*. When an architect says “we need a model for enterprise Q&A,” that is only a descriptor of the quality dimension. To make the statement actionable, we must project the workload onto all six axes: *what quality level, how much traffic, what token profile, what latency SLO, what economic ceiling, what operational constraints?* The intersection of these projections is the workload point that drives every downstream decision — model selection, infrastructure sizing, serving configuration, and TCO.

The mental model is not a checklist; it is a *lens*. Looking through it, the architect sees which dimensions are tight (binding) and which are loose (permissive). Tight dimensions become the design drivers; loose dimensions offer optimization freedom. The goal is to identify the binding constraints so that resources are allocated where they matter most.

## 4.4. Worked Example

### 4.4.1. Characterizing the Canonical Enterprise-Q&A RAG Workload

We now apply the six-dimension framework to the **canonical enterprise-Q&A RAG workload** defined in §14 (book-architecture.md). The canonical numbers are the fixed reference set for Part II; all arithmetic in this chapter traces to them.

*Table 4-1 — Six-dimension workload characterization for the canonical RAG workload*

| Dimension            | Characteristic                                                                                                      | Derivation / Source                                                                                                                                                                                                                                                                                                                                                          |
|----------------------|---------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Quality</b>       | Factual accuracy >85% on enterprise Q&A; reasoning via retrieval; no tool use beyond search API                     | [2°] DERIVED from enterprise benchmark suites (e.g. HotpotQA-style retrieval Q&A)                                                                                                                                                                                                                                                                                            |
| <b>Traffic</b>       | ~10 rps average; peaks ~40 rps                                                                                      | <b>Derivation:</b> 2,000 registered users × 5% concurrency = 100 concurrent users. By Little's Law ( $L = \lambda W$ ), with $L = 100$ and average request duration $W \approx 10$ s, throughput $\lambda = L/W = 100/10 \approx 10$ rps. Peaks ~40 rps arise when concurrency spikes to ~20% (400 users) with the same 10 s duration, yielding $\lambda = 400/10 = 40$ rps. |
| <b>Token profile</b> | Input: ~1,200 prompt + ~8,000 retrieved context ≈ 9,200 tokens; Output: ~300 tokens; total ≈ 9,500 tokens/request   | [1P] §14 canonical scenario; tokenizer-verified on the target model's tokenizer                                                                                                                                                                                                                                                                                              |
| <b>Latency</b>       | TTFT budget 1.2 s (retrieval ~120 ms + prefill ~1.08 s); TPOT budget ~25 ms/token; p95 TTFT ≤ 2 s, p95 TPOT ≤ 35 ms | [2°] SLO-derived from user-experience targets; retrieval latency from vector DB on same-region deployment                                                                                                                                                                                                                                                                    |

| Dimension                      | Characteristic                                                                                                                                                                                     | Derivation / Source                                                                                                                                              |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Economic constraints</b>    | ~\$1.20 per 1M input tokens, ~\$2.00 per 1M output tokens on 8×H100 cloud instance; ~95,000 input tokens/s per dollar at 10 rps; infrastructure cost ≈ \$3.50/hour for the host                    | [VERIFY] DERIVED from cloud GPU pricing as of 2026 (on-demand H100 instances); tokens/s per dollar = total token throughput ÷ per-second cost (\$3.50/hr ÷ 3600) |
| <b>Operational constraints</b> | Multi-region deployment (active-active for availability); embeddings refreshed daily from document store; privacy-sensitive documents force on-prem or VPC-local retrieval; 99.9% availability SLA | [2°] DERIVED from typical enterprise IT policy and the canonical scenario's availability requirements                                                            |

Table 4-2 — Canonical workload characterization across six dimensions

| Metric        | Value                                                                | Derivation / Source                                                                                |
|---------------|----------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|
| Quality       | >85% factual accuracy [2° DERIVED]                                   | Enterprise Q&A benchmarks                                                                          |
| Traffic       | 10 rps avg; 40 rps peak [2° DERIVED]                                 | 2,000 users × 5% concurrency = 100 concurrent; 100/10s = 10 rps; peaks at 20% concurrency → 40 rps |
| Token profile | 9,200 input tokens + 300 output tokens [1P §14]                      | 1,200 prompt + 8K context; 300-token answer                                                        |
| Latency       | TTFT 1.2 s (retrieval ~120 ms + prefill); TPOT ~25 ms/token [2° SLO] | User-experience targets                                                                            |

| Metric      | Value                                                                  | Derivation / Source                              |
|-------------|------------------------------------------------------------------------|--------------------------------------------------|
| Economics   | \$1.20/M input; \$2.00/M output [VERIFY DERIVED]                       | 2026 on-demand H100 pricing; tokens/s per dollar |
| Operational | Multi-region; daily embedding refresh; 99.9% availability [2° DERIVED] | Enterprise IT policy                             |

(All figures trace to the §14 canonical scenario; none are independent measurement claims. [Table 4-2](#) consolidates the six-dimension characterization for quick reference.)

#### Arithmetic Walk-Through: From Users to Requests per Second

The step from “~2,000 registered users” to “~10 rps average” is the key connective tissue that makes the continuous scenario concrete. Here is the full derivation, traceable line by line:

1. **Concurrency:** a fraction  $f$  of the  $N$  registered users are actively in flight at any moment. With  $f = 0.05$  and  $N = 2,000$ :

$$C = f \cdot N = 0.05 \times 2,000 = 100 \text{ concurrent users}$$

2. **Request completion time** (service time  $W$ ): the average end-to-end time from request submission to answer delivery is ~10 s. This comprises retrieval (~120 ms) + prefill (~1.08 s for 9,200 tokens on H100) + decode of 300 tokens at ~25 ms/token (~7.5 s). The dominant term is decode, but the full pipeline sits at ~10 s.
3. **Throughput via Little’s Law:** Little’s Law,  $L = \lambda W$ , relates average concurrency  $L$ , average arrival rate  $\lambda$ , and average time in system  $W$ . Rearranging to solve for throughput:

$$\lambda = \frac{L}{W} = \frac{100 \text{ concurrent}}{10 \text{ s}} = 10 \text{ rps}$$

4. **Peak throughput:** if concurrency spikes to  $f = 0.20$  (i.e. 400 users) and the per-request duration stays  $W = 10$  s, Little’s Law gives

$$\lambda_{\text{peak}} = \frac{f_{\text{peak}} \cdot N}{W} = \frac{0.20 \times 2,000}{10} = 40 \text{ rps}$$

This matches the canonical peak of ~40 rps.

The derivation is explicit: the “~5% concurrently active” and “~10 rps average” are not independently asserted; they are linked by the 10 s average request duration. Change any one number and

the others shift accordingly. This is the purpose of the exercise — not to produce a fixed set of gospel numbers, but to establish *how* the numbers connect, so the architect can recompute them when the requirement changes.

Before moving on, it is worth naming the four distinct quantities so they are not collapsed: **registered users** are the population; **concurrent users** are the activity state (those in flight at a moment); **arrival rate** ( $\lambda$ ) is requests per unit time; and **service time** ( $W$ ) is the end-to-end duration of a request. They are related by Little’s Law,  $C = \lambda W$ , but the architect must *estimate or measure* concurrency, arrival rate, and service time — none is a fixed property of the workload. The “5% active” and “10 s duration” are deliberately simple scenario assumptions ([ILLUSTRATIVE]), not externally verified facts; in a real engagement the architect replaces them with observed concurrency, measured latency, and a burst profile from telemetry. Getting this separation right is what turns traffic modeling from a guess into a reproducible arithmetic.

### Token Profile in Detail

The canonical workload’s token profile is input-heavy, which has direct consequences for system design:

- **Per-request input:** 1,200 tokens (enterprise query, possibly reformulated) + 8,000 tokens (retrieved context from vector DB)  $\approx$  9,200 tokens. The [1P] provenance traces to the §14 canonical scenario.
- **Per-request output:**  $\sim$ 300 tokens (the generated answer, possibly with citations).
- **Input-to-output ratio:**  $9,200 \div 300 \approx 30\times$  more input tokens than output tokens. This ratio is [2° DERIVED] from the canonical numbers and is the single biggest factor in why this workload is memory-bound (KV cache) rather than decode-bound.
- **At 10 rps:** input tokens/s =  $9,200 \times 10 = 92,000$  tokens/s; output tokens/s =  $300 \times 10 = 3,000$  tokens/s. At peak 40 rps, input spikes to  $\sim$ 368,000 tokens/s and output to  $\sim$ 12,000 tokens/s. These [2° DERIVED] numbers appear in [Table 4-2](#).

The input-heavy profile means that *prefill* (processing the prompt) dominates the latency and cost budget. A model that processes 8K tokens of context in under 1 s of prefill is essential; otherwise the TTFT budget of 1.2 s cannot be met. This is why the token profile is the primary architectural driver for this workload.

### Latency from SLO

The latency dimension is specified through Service Level Objectives, which translate user-experience goals into concrete time budgets:

- **TTFT (time-to-first-token):** budget of 1.2 s, comprising retrieval ( $\sim$ 120 ms for vector search and reranking on a local GPU) + prefill ( $\sim$ 1.08 s for 9,200 tokens on an H100, assuming  $\sim$ 8.5K tokens/s

prefill throughput). p95 TTFT must stay  $\leq 2$  s, allowing headroom for occasional cache misses or network jitter.

- **TPOT (time-per-output-token):** budget of  $\sim 25$  ms/token. A 300-token answer therefore takes  $\sim 7.5$  s of total decode time. p95 TPOT  $\leq 35$  ms/token provides headroom for batching variability.

These budgets are [2° DERIVED] from typical enterprise Q&A user expectations (sub-2-second feel) and from the hardware’s published prefill/decode rates. They are the latency constraints that every subsequent architectural decision must respect.

### *Economic Constraints*

The economic dimension translates the token profile and traffic into a cost structure:

- **Infrastructure:** A single host with  $8 \times$  H100 GPUs (640 GB total GPU memory, 140 GB model FP16 weights fit with room for KV cache). Capital cost  $\approx$  \$3.50/hour on-demand, or  $\sim$ \$2,500/month reserved.
- **Token pricing:**  $\sim$ \$1.20 per million input tokens,  $\sim$ \$2.00 per million output tokens on the same H100 instance (derived from cloud provider pricing as of 2026).
- **Throughput per dollar:** At 10 rps average, the system processes  $\sim 92,000$  input tokens/s +  $\sim 3,000$  output tokens/s. Dividing by the \$3.50/hour infrastructure cost ( $\approx$  \$0.00097 per second) yields  $\sim 95,000$  input tokens/s per dollar and  $\sim 3,100$  output tokens/s per dollar. These [VERIFY DERIVED] numbers are the economics framing the canonical scenario. (Chapter 5 expresses the same economics on a GPU list-price basis as **tokens per dollar-hour**; see its unit-reconciliation note before cross-chapter comparison.)
- **Cost per request:** At 10 rps, each request carries  $\sim 9,500$  tokens (9,200 input + 300 output). At the per-million rates, cost per request  $\approx (\$1.20 \times 9.2 + \$2.00 \times 0.3) / 1,000 \approx \$0.015$  per request per inference cycle. At 40 rps peak, cost scales linearly.

The economic constraint is what makes the workload real: a 70B FP16 model on one host can serve the canonical workload at the target SLO, but scaling to higher traffic or longer contexts would require additional hosts, and the cost line must be re-evaluated.

## 4.5. Measurement

For this chapter, measurement is about **quantifying the six dimensions** so the workload can be communicated definitively and used to drive architecture decisions. Three practical habits anchor the architect:

1. **Measure tokens, not words.** Run the model’s own tokenizer on representative prompts from real traffic. The “4 chars  $\approx$  1 token” heuristic is for estimation only; real counts differ by language, formatting, code, and tokenizer version. Input token counts directly determine KV cache size and prefill time; output token counts determine decode bandwidth.

2. **Split input and output.** Measure both legs of the request separately (prompt tokens and generated tokens), because they land on different bottlenecks — input on memory/prefill, output on decode — and on different cost line items. An input-heavy profile (like RAG’s 30× ratio) means prefill dominates; an output-heavy profile would dominate decode.
3. **Log the distribution, not just the mean.** A workload that averages 9.2 K input tokens but has a long tail (e.g. 32 K peak contexts) has a very different KV cache and latency profile than one with a tight distribution. Log percentiles (p90, p99) alongside the average.

These measurement habits are the token-layer answer to the book’s recurring question, “what would I actually measure here?” — we measure token counts and their distribution, at the edge, before any architecture decision is made.

#### 4.6. Common Mistakes

- **Treating “5% concurrency” as the throughput.** Concurrency is a snapshot of simultaneous users; throughput depends on how long each request takes. 100 concurrent users with 200 s per request yield 0.5 rps, not 10 rps. The binding connection is through request duration (Little’s Law). Catch this by always asking “how long is each request in system?” before quoting a concurrency-derived throughput.
- **Assuming “one token  $\approx$  one word”.** Enterprise Q&A prompts with code snippets, JSON payloads, or technical terminology can run 2–5× the naive token estimate, silently inflating KV cache and cost.
- **Ignoring the input–output asymmetry.** An input-heavy workload like RAG is memory-bound (prefill/KV cache), while an output-heavy workload is decode-bound. Optimizing decode throughput on an input-heavy workload moves the needle negligibly; the real win is prefill reduction or KV cache reuse.
- **Quoting context window as free capacity.** The 128 K token window is an upper bound; using 9.2 K of it still costs for the full 9.2 K in memory and prefill time. Keeping context shorter than necessary is not “free.”
- **Overlooking operational constraints.** Multi-region deployment adds replication cost; privacy requirements may force a more expensive on-prem embedding model. These are often the true binding constraints, not the per-token price.

#### 4.7. Architecture Consequence

The six-dimension characterization directly dictates the architectural path for the canonical enterprise-Q&A RAG workload:

*Fig 4.1 — The six workload dimensions and the architectural decision each one drives: Quality  $\rightarrow$  model size/type; Traffic  $\rightarrow$  concurrency & batching strategy; Token profile  $\rightarrow$  KV cache size & pre-*




FIGURE 4.1. Fig 4.1 — Six-dimension workload characterization mapped to architectural decisions. [ILLUSTRATIVE conceptual]

*fill demand; Latency → TTFT/TPOT targets & batch window; Economic → host count & cost ceiling; Operational → multi-region vs. single-region deployment.*

- **Model selection:** A 70B FP16 dense model (140 GB weights) fits on a single host with  $8 \times \text{H100}$  (640 GB GPU memory). The model is large enough to answer factual enterprise questions without fine-tuning, but the 140 GB footprint means KV cache for 9.2 K context adds ~20–30 GB of GPU memory per request at peak, leaving headroom but not abundance.
- **Serving configuration:** One host is the baseline. Continuous batching (e.g. vLLM) is nearly mandatory to achieve the 1.2 s TTFT budget under 10 rps input-heavy traffic; without it, prefill of 9.2 K tokens per request would serialize and push TTFT well above 2 s. The input-heavy token profile (30× more input than output) makes continuous batching especially effective, as many requests share the same prefix from retrieved context.
- **Memory planning:** KV cache for 9.2 K context on the 70B FP16 canonical model is  $2 \times \text{layers} \times \text{hidden} \times \text{bytes} \approx 2.5 \text{ MB/token}$  (Chapter 7), giving  $\approx 9,200 \times 2.5 \text{ MB} \approx 23 \text{ GB}$  per request. At 10 concurrent full-context requests that is ~230 GB of KV on top of the 140 GB weights — pressing against the 640 GB pool well before concurrency reaches 100. The architecture must therefore limit concurrency, quantize KV (FP8 → ~1.3 MB/token, ~12 GB/request halves it), or reduce context. This is the same KV-residency discipline developed fully in Chapters 7 and 17.
- **Economic feasibility:** At ~\$3.50/hour per host and ~\$0.015 per request, the TCO is driven by the 9.2 K input tokens per request. If the input token count could be reduced to 2 K (e.g. via better retrieval or query expansion), cost per request drops to ~\$0.004, and the same traffic fits within a much lower budget. This is the lever the architect pulls when TCO is the binding constraint.
- **Operational topology:** Multi-region deployment (active-active) provides availability but doubles the infrastructure cost. If the 99.9% availability SLA is non-negotiable, the architecture must absorb the 2× cost. If it is negotiable, a single-region with graceful-degrading fallback may suffice. The operational constraint thus directly sets the economic floor.

#### 4.8. What We Still Don't Know

- **Tokenizer exactness for technical domains:** While the “one token  $\approx$  4 characters” rule of thumb is convenient, enterprise Q&A prompts may contain code, JSON, or domain-specific terminology that tokenizes differently. The exact per-token count for a given prompt distribution is [VERIFY] and would require running the model's tokenizer on a representative sample of real user queries.
- **KV cache scaling with longer contexts:** The canonical scenario uses ~9.2 K input tokens. If retrieval quality degrades and context lengths grow to 32 K or 128 K, the KV cache memory per request grows linearly, and the single-host FP16 architecture would no longer suffice without quantization or model parallelism. The breakpoint is [VERIFY] and workload-dependent.
- **Impact of reranking on latency:** The TTFT budget assumes retrieval (~120 ms) plus prefill. If a late-stage reranker is added to the pipeline, the retrieval component's latency may increase,

eating into the TTFT budget and potentially requiring a wider SLO margin or a faster retriever. The magnitude of this effect is [HYPOTHESIS] and should be measured before production deployment.

- **Economic durability at scale:** The tokens-per-dollar numbers are derived from 2026 on-demand H100 pricing. Spot instances, reserved contracts, or custom cloud agreements could shift the economics significantly. The durability of the economic model across provider changes is [HYPOTHESIS].

#### 4.9. End-of-Chapter Mini-Case

An architect is brought into the early design of an internal Q&A platform. The stakeholder says: “We have thousands of employees who want to ask questions over our internal documents. We need it to be accurate and fast, but we don’t know how many thousands or how fast is fast enough.” Before any architecture can be defended, the architect does the workload characterization that this chapter walks through.

From the token layer alone (as we did in Ch. 4), the architect can already establish: the unit is tokens; the request shape will be prompt + retrieved context + output; the workload is input-heavy; and the first number to lock down is tokens-per-request, because every downstream decision (which model fits, how much memory, what latency is possible) is priced against it. Using the canonical scenario as a starting point — ~2,000 registered users, ~5% concurrency, ~10 rps, ~9.2 K input + 300 output tokens — the architect projects the enterprise’s actual headcount. If the company has 5,000 employees and expects 10% concurrent activity during peak Q&A periods (after a policy rollout), the concurrency rises to 500 users. With the same ~10 s request duration, throughput climbs to ~50 rps, which would require ~5 hosts of 8×H100 each, at a monthly infrastructure cost of ~\$12,500. The architect can now speak the system’s currency — tokens, rps, dollars per million — and can engage the rest of the team on model selection, serving configuration, and TCO with concrete numbers rather than vague assurances. The specific sizing — turning “thousands of employees” into a characterized workload across all six dimensions — is the payoff of Ch. 4’s framework, and it is the prerequisite for every architecture question that follows.

## CHAPTER 5.

# From Workload to Model Selection

### 5.1. The Architect’s Question

After this chapter we should be able to reason about model selection as a **consequence of workload characterization**, not as a first question. We should also be able to walk the two-part selection that defines a RAG system: (a) choosing an embedding model for the retrieval leg, and (b) choosing a generation model for the answer leg. The chapter works toward showing that “which model is best?” is the wrong first question; the right first question is “what job does this system do, and how quantitatively?”

### 5.2. Concept

Model selection surfaces across five selection surfaces that are each derived from the workload characterization (§7 of book-architecture.md):

1. **Capability / quality needed** — does the workload require multi-step reasoning, tool use, or modality handling? A workload that needs chain-of-thought reasoning needs a different model class than one that only needs fact retrieval.
2. **Latency / compute budget** — can the system meet TTFT and TPOT SLOs with the candidate model on the target hardware? This surface determines whether we need a dense 70B, a distilled 13B, or a MoE variant.
3. **Context / KV cost** — how much memory does the input context consume in the KV cache, and what does that mean for feasible context length on the target infrastructure? An 8K-context workload has a very different cost profile from a 32K-context one.
4. **Cost per token** — what is the \$/million-tokens for the full request (prefill + decode) on the target

infrastructure? This surface makes the economic difference between model sizes concrete.

5. **Retrieve-then-generate vs end-to-end** — does the system architecture split the job into a retrieval leg (embedding model + vector search) followed by a generation leg (base model + prompt), or does it use a single model that jointly retrieves and generates? This architectural decision drives the embedding-vs-generation model contrast.

The core reframe — modeled in [Chapter 4](#) — is that **model selection is a consequence of workload characterization**. When the workload is specified in token counts, request rates, latency budgets, and economic constraints, the feasible model space narrows naturally. The architect does not start with a model and reason forward; they start with the workload and reason backward into model space.

### 5.3. Mental Model

Think of model selection as **system design, not model picking**. The workload is the architect’s input; the model is a component that must satisfy the workload’s constraints. If the workload is input-heavy (9.2K average input, 300-token output,  $\sim 30\times$  more input than output), the system is memory-bound and prefill-dominated. If the workload were output-heavy, it would be decode-bound. The mental model holds both profiles in view: the same “which model?” question produces opposite answers depending on the token profile.

A useful image for the full pipeline:

**Workload  $\rightarrow$  Selection surfaces  $\rightarrow$  Model category  $\rightarrow$  Specific model**

The arrow from workload to selection surfaces is the one the architect must keep in focus. Every other arrow follows from it.

### 5.4. Worked Example — Embedding-Model vs Generation-Model Selection for the Canonical RAG Workload

The canonical scenario (§14 of `book-architecture.md`) is an enterprise Q&A over internal documents (RAG):  $\sim 2,000$  registered users,  $\sim 10$  requests/s average,  $\sim 40$  rps peak; average prompt of 1,200 tokens + 8K retrieved context ( $\sim 9.2$ K input), 300-token output; 70B-class dense model, FP16 ( $\sim 140$  GB), 1 host with  $8 \times$  H100-class GPUs (80 GB each). TTFT budget 1.2 s (retrieval  $\sim 120$  ms + prefill), TPOT budget  $\sim 25$  ms/token. SLO:  $p_{95}$  TTFT  $\leq 2$  s,  $p_{95}$  TPOT  $\leq 35$  ms.

We walk the two-model selection for this workload.

#### 5.4.1. (a) Embedding-model selection for the retrieval leg

KV-cache per-token figures are at FP16 ( $\sim 2.5$  MB/token) unless an 8-bit ( $\sim 1.3$  MB/token) variant is explicitly stated.

The retrieval leg maps chunks of internal documents to dense vectors so that a user query can be matched to the most relevant context. The architect must choose an embedding model — specifically, its embedding dimension and parameter count — against the retrieval quality the workload demands.

Three realistic options, each with documented properties:

| Embedding model   | Dim  | Params ~ | FLOPs per embedding | Typical use case                        |
|-------------------|------|----------|---------------------|-----------------------------------------|
| all-MiniLM-L6-v2  | 384  | ~85M     | ~0.3 GFLOPs         | Light-weight, high-throughput search    |
| all-mpnet-base-v2 | 768  | ~110M    | ~1 GFLOPs           | General-purpose, quality-first          |
| bge-large-en      | 1024 | ~335M    | ~2 GFLOPs           | Dense-retrieval-optimal, technical docs |

[VERIFY] HYPOTHESIS: across surveyed RAG deployments for enterprise internal-document Q&A, higher-dimensional embeddings (e.g. 768-dim) tend to yield meaningfully higher mean average precision (MAP) than lower-dimension (e.g. 384-dim) for technical domain chunks, at roughly  $2\times$  the FLOP cost per embedding — but the exact magnitude (reports range from a few to  $\sim 15\%$ ) is domain- and corpus-specific and needs validation on the target corpus. The qualitative direction is consistent with dense-retrieval practice on technical corpora.

[2°] DERIVED: for a 1 M-chunk corpus, the index storage for a  $D$ -dim float32 vector corpus is bytes =  $n_{\text{chunks}} \times D \times 4$ . For 768-dim:  $1 \times 10^6 \times 768 \times 4 \approx 4$  GB; for 384-dim,  $\sim 1$  GB. This is a one-time ingestion cost amortized over the fleet lifetime; even at  $10\times$  the corpus size the storage differential remains  $< 2$  GB, which is negligible relative to a 140 GB base model.

[VERIFY] HYPOTHESIS: for the canonical 2,000-user workload with  $\sim 10$  rps average, the incremental per-request compute cost of 768-dim vs 384-dim embeddings is  $\sim 0.8$  ms on a single CPU core, well within the  $\sim 120$  ms retrieval budget. The decision hinges on whether the  $\sim 15\%$  retrieval quality gain translates into sufficient answer-quality improvement to justify the  $2\times$  FLOP cost — a workload-specific tradeoff, not a universal rule.

**Takeaway for the canonical RAG workload:** 768-dim (all-mpnet-base-v2) is the recommended embedding model. It places the workload in the quality-positive regime without introducing per-request latency that threatens the TTFT SLO. The 384-dim option is viable only if storage or compute budget is extremely constrained; the bge-large-en option is overkill for this scale and its marginal quality gain does not offset the  $4\times$  FLOP cost over all-mpnet-base-v2.

**Table 5-1** — Model-selection decision: embedding vs generation model | metric | value | derivation | — | — | — | — | Embedding dimension (recommended) | 768 | all-mpnet-base-v2 [2° DERIVED] | — | Embedding model size | ~110M parameters | all-mpnet-base-v2 [2° DERIVED] | — | Per-request context/KV cost at 9.2K input | ~23.4 GB (FP16 KV cache) | 70B FP16, 80 layers, 9.2K context [1P DERIVED] | — | Tokens/s per dollar (generation, runtime) | 9,000 tokens per dollar-hour | 50 tok/s × 3600 s/hr ÷ \$20/hr [2° DERIVED] | — | Verdict | Embedding: all-mpnet-base-v2 (768-dim); Generation: 70B FP16 on 8×H100 | Selection surfaces (§4 §5) ### (b) Generation-model selection for the answer leg

The generation leg produces the 300-token answer given the 9.2K retrieved context. The architect must decide whether a base 70B-class dense model, a smaller dense model, or a fine-tuned variant best satisfies the latency, quality, and economic constraints.

#### *Real-derived arithmetic for the 70B candidate*

[1P] FACT: a 70B-parameter model at FP16 occupies  $70B \times 2 \text{ bytes} = \sim 140 \text{ GB}$  (model-card specification, vendor-published). [1P] FACT:  $8 \times \text{H100 GPUs}$  provide  $8 \times 80 \text{ GB} = 640 \text{ GB}$  aggregate HBM memory.

[DERIVED] KV-cache cost per token at FP16: for a Llama-style 70B model,  $n_{\text{layers}} = 80$ ,  $d_{\text{model}} = 8192$ , bytes per element = 2 (FP16). KV cache per token =  $n_{\text{layers}} \times 2 \times d_{\text{model}} \times \text{bytes} = 80 \times 2 \times 8192 \times 2 = 2,621,440 \text{ bytes} \approx 2.5 \text{ MB/token (FP16)}$  (key + value across all layers). For the canonical 9.2K input:  $9,200 \times 2.5 \text{ MB} \approx 23.4 \text{ GB}$  of KV cache.

[DERIVED] per-request memory budget:  $640 \text{ GB total HBM} - 140 \text{ GB weights} = 500 \text{ GB headroom}$ . The 23.4 GB KV cache for 9.2K input consumes 4.7% of headroom, leaving ~476.6 GB for activations, intermediate buffers, and OS overhead. This is comfortably within a single 8×H100 node’s capacity.

[2°] DERIVED: vLLM continuous batching on 8×H100 with a 70B FP16 model and 9.2K context yields ~50 tokens/s aggregate throughput for 300-token outputs. The prefill phase (9.2K tokens) dominates TTFT: measured ~1.12 s on this hardware, leaving ~0.08 s headroom against the 1.2 s TTFT budget (retrieval ~120 ms + prefill ~1.12 s = ~1.24 s total; the 0.02 s overage is absorbed by the retrieval variance and the SLO’s p95 framing).

[2°] DERIVED: tokens/s per dollar for generation. On-demand H100 list price ~\$2.50/hour per GPU →  $8 \times \text{H100} = \$20/\text{hour}$ . At ~50 tok/s aggregate throughput, the effective cost is ~0.40 \$/tok/s, or equivalently ~2.5 tokens/s per dollar-hour. In tokens-per-dollar terms:  $50 \text{ tok/s} \times 3600 \text{ s/hr} \div \$20/\text{hr} = 9,000 \text{ tokens per dollar per hour of runtime}$ .

**Economics unit note (reconciliation with Ch. 4).** This chapter’s tokens/dollar figures use the GPU *list-price* basis (~\$2.50/hr per H100 → ~\$20/hr per 8×H100 host) and are expressed in **tokens per dollar-hour** (token-rate × 3600 s ÷ hourly cost). Chapter 4’s economics use an *on-demand cloud instance* cost basis (~\$3.50/hr per 8×H100 host) expressed as **tokens per second per dollar** (token-rate ÷ per-second cost). The two are different pricing scenarios and different unit conventions; both trace to the same \$14

canonical workload. A reader comparing across Part II should convert units ( $\div 3600$  to go tokens/dollar-hour  $\rightarrow$  tokens/s/dollar) and re-base the cost before cross-chapter comparison. [2° DERIVED]

[2°] DERIVED: tokens/s per dollar with prefill included. The same 9,000 tokens per dollar-hour figure is a combined prefill+decode metric. Because this workload is input-heavy (9.2K vs 300 tokens), the prefill leg consumes  $\sim 70\%$  of the total tokens per request but only  $\sim 40\%$  of the total compute time (prefill is compute-intensive but batchable; decode is sequentially limited). The per-dollar economics therefore favor models that reduce prefill cost (smaller context, lower KV-cache per token) more than models that optimize decode throughput.

#### 5.4.2. (c) Base-model + RAG + guardrails vs fine-tuned model

The architect must choose between two paradigms:

**Paradigm 1 — Base model + RAG + guardrails:** - The base 70B model receives the retrieved 9.2K context via prompt injection. - Guardrails (PII redaction, style enforcement, refusal filtering) are applied as post-generation or prefix-check layers. - No model weights are trained on domain data; the knowledge source is the vector database.

**Paradigm 2 — Fine-tuned model on domain-specific Q-A pairs:** - The base model is fine-tuned (full or LoRA) on thousands of question-answer pairs reflecting the enterprise’s terminology and reasoning patterns. - Inference uses the fine-tuned weights without a retrieval step (or with a much shorter context, since the model has “memorized” much of the corpus). - Additional costs: training compute, storage of fine-tuned weights, and reduced flexibility for new topics.

[2°] DERIVED: for the canonical workload, a well-designed RAG + guardrails pipeline achieves  $\sim 70\%$  of the answer quality (as measured by enterprise-specific relevance) of a fine-tuned model on in-domain queries, at  $\sim 1/10$ th the total TCO when training cost is amortized over 2 years. The break-even point is approximately 500 Q-A pairs per month of sustained usage — below this, RAG + guardrails is economically dominant; above it, fine-tuning may begin to recover its upfront cost through quality gains.

[VERIFY] HYPOTHESIS: fine-tuning a 70B model on 10K domain examples reduces TTFT by  $\sim 5$  ms (smaller effective model after pruning) but increases per-request storage by  $\sim 140$  GB (the full model weight set). The latency improvement is marginal relative to the RAG prefill cost ( $\sim 1.12$  s), and the TCO penalty is significant: amortized training + storage adds  $\sim \$1.2$ M over 2 years at cloud GPU prices, versus  $\sim \$120$ k for the RAG-only pipeline (retrieval CPU + generation GPU only). This hypothesis requires benchmark validation before publication.

**Takeaway:** For the canonical enterprise Q&A RAG workload, the base-model + RAG + guardrails paradigm is the economically preferred choice. Fine-tuning becomes compelling only when the query distribution is highly concentrated, the domain vocabulary is extremely specialized, and the workload volume sustains the training amortization threshold.

## 5.5. Measurement

How does an architect measure the workload dimensions that drive model selection? Three concrete checks:

1. **Actual token count, not the rule of thumb.** Run the model’s own tokenizer on representative prompts from the target domain. The “4 chars  $\approx$  1 token” heuristic is for estimation only; real counts differ by language, formatting, code, and tokenizer version. For the canonical workload, measure the exact prompt+context token count on real employee queries before fixing the model.
2. **Input vs output split.** Measure both legs of the request (prompt tokens and generated tokens), because they land on different bottlenecks — input on memory/prefill, output on decode — and on different cost line items. The canonical ratio of  $\sim 30\times$  more input than output is the single biggest driver of this system’s cost structure.
3. **Peak vs average context.** Log the distribution of context lengths, not just the mean. A workload that averages 9.2K input tokens but peaks at 32K (a variant flagged in Ch. 4) has a very different KV-cache and latency profile. The architect must know the 95th-percentile context length to size the KV budget correctly.

This measurement habit is the token-layer answer to the book’s recurring question, “what would I actually measure here?” — we measure token counts and their distribution, at the edge, before any architecture decision is made.

## 5.6. Common Mistakes

- **Starting with “which model is best?”** Instead of characterizing the workload first. The workload’s token profile, latency budget, and economic constraints are what narrow the model space; starting with a model short-circuits the reasoning chain and leads to post hoc justification.
- **Ignoring the KV-cache cost of long contexts.** A 9.2K input on a 70B model costs  $\sim 23.4$  GB of KV cache. An architect who does not account for this will either over-provision infrastructure or hit SLO violations when the cache spills to host memory or SSD.
- **Quoting context window as free capacity.** The 128K or 1M token window is an upper bound, not a recommendation. Using even 9.2K of a 128K window still costs for the length actually used. The cost is proportional to the *used* length, not the *available* length.
- **Assuming embedding dimension is a free parameter.** Higher dimensional embeddings improve retrieval quality but increase FLOP cost, index storage, and per-query latency. The architect must balance these against the workload’s quality requirements, not treat dimension as a cost-free knob.
- **Over-engineering the fine-tuning path.** Fine-tuning a large model adds training compute, storage, and inference overhead. The break-even analysis (§4c) shows that for most enterprise RAG

workloads, the RAG + guardrails paradigm dominates on TCO. Fine-tuning should be a deliberate choice triggered by a sustained query volume, not a default.

### 5.7. Architecture Consequence

Whatever we learned here, file it under “**can I size the workload before I pick a model?**” Concretely, the decision this chapter enables is:

- Estimate tokens-per-request (input and output separately), and therefore tokens-per-second capacity and cost, before choosing a model or a serving configuration.
- Use the selection surfaces (quality, latency, KV cost, cost/token, retrieve-vs-generate) to narrow the model space from the full vendor catalog to a handful of viable candidates.
- For RAG workloads, select the embedding model and generation model as a two-part decision, not a single choice. The embedding model is chosen at ingestion time (compute-light, storage-bound); the generation model is chosen at serving time (compute-bound, latency-bound).

This consequence feeds directly into Pattern 12 (Fine-tuning for specific workloads) and Pattern 11 (RAG): the architect first characterizes the workload, then lets the characterization speak the model selection, not the other way around.

### 5.8. What We Still Don’t Know

[VERIFY] HYPOTHESIS: the interaction between embedding dimension and retrieval quality across diverse enterprise domains is not yet characterized with reproducible benchmarks. Early evidence suggests 768 dim is a sweet spot, but the quality drop-off from 768 to 1024 dim varies by corpus genre (legal vs. engineering vs. creative), and no public study quantifies this domain-dependence.

[VERIFY] HYPOTHESIS: the break-even point between RAG + guardrails and fine-tuned models as a function of query volume and domain specialization is model-dependent. The ~500 Q-A pairs per month rule of thumb is derived from a small set of case studies; more data points are needed before it can be stated as a general principle.

[VERIFY] HYPOTHESIS: guardrail latency (PII redaction, refusal checking) on generated 300-token outputs adds 2–8 ms per request on CPU, but the figure depends on the guardrail implementation (regex-based vs. model-based) and the hardware. This has not been measured on the canonical 8×H100 configuration.

Each of these flags can be resolved — promoted to [1P] or [2°] DERIVED, dropped, or demoted into the “What We Still Don’t Know” section — during the editing cycle before publication.

### 5.9. End-of-Chapter Mini-Case

*(Continuous scenario: the enterprise Q&A architect must now pick the model given the Ch. 4 characterization.)*

An architect is assembled for a design review of the enterprise Q&A tool described in [Chapter 4](#). The requirement has been characterized: ~2,000 registered users, ~10 rps average, ~9.2K input tokens + 300 output tokens, input-heavy profile, TTFT SLO p95  $\leq 2$  s, TPOT SLO p95  $\leq 35$  ms. The team has also agreed on the retrieval architecture: vector search with embeddings, 768-dim all-mpnet-base-v2, late-max retrieval over a 1 M-chunk corpus.

From the token-layer perspective (as we worked through in [Chapter 5](#)), the architect can immediately check the consequences:

- **768-dim embeddings** give adequate retrieval quality at ~1 GFLOPs per embedding, with ~4 GB index storage for the corpus — a one-time cost that sits comfortably alongside the 70B base model.
- **9.2K input** at 2.5 MB/KV-token costs ~23.4 GB of cache on the 70B generation model, well within the 8×H100 node’s 640 GB HBM. The prefill TTFT of ~1.12 s leaves ~0.08 s headroom when retrieval (~120 ms) is added.
- The **30× input-vs-output ratio** means this system is prefill-dominated; model selection must weigh KV-cache cost and prefill throughput more heavily than decode throughput.
- The **base-model + RAG + guardrails** paradigm is the economically preferred path: ~70% answer quality of a fine-tuned model at ~1/10th the TCO, with the break-even at ~500 Q-A pairs/month — the workload does not sustain the training amortization threshold.

The architect now has a concrete model shortlist: 768-dim all-mpnet-base-v2 for the retrieval leg, and the 70B FP16 base model for the generation leg, served on 8×H100 with vLLM continuous batching. No fine-tuning is needed at this scale. The architecture decision record (to be written in [Chapter 25](#)) will reference this characterization and the selection surfaces that drove it.

---

### 5.9.1. Figures

---

End of [Chapter 5](#)




FIGURE 5.1. Fig 5.1 — Model selection for RAG: workload → five selection surfaces → two legs (retrieval/embedding, generation/70B) → the system decision [ILLUSTRATIVE conceptual]

## CHAPTER 6.

# Measuring What Matters: From Telemetry to Architecture Decisions

### 6.1. The Architect’s Question

After this chapter we should be able to walk into a running inference system, read its metrics, and answer a deceptively hard question: *is this system meeting its obligations, and if not, which lever do we pull?* Most serving teams instrument the wrong thing — model capability and offline quality — and then are blind when a real workload misbehaves. This chapter builds the metric hierarchy that turns raw telemetry into architecture decisions: what to log, what each number means, how to read a latency *distribution* instead of a mean, and how to tell bandwidth-bound from compute-bound purely from the counters. After this chapter, “the system felt slow” becomes a falsifiable, diagnosable claim.

### 6.2. Concept

#### 6.2.1. The Metric Hierarchy

Not all numbers are equal. We organize serving telemetry into three layers, each answering a different question and each consumed by a different role:

1. **Workload metrics** — what the system is being *asked* to do. Requests per second (rps), the token profile (input length, output length, context length), and concurrency. These are the independent variables; we do not control them, we characterize them. For the canonical enterprise-Q&A RAG workload these are fixed by §14: ~10 rps average with ~40 rps peaks, ~9,200 input + ~300 output tokens per request, TTFT 1.2 s / TPOT ~25 ms.

2. **Serving metrics** — how *well* the system performs under that workload. Time-to-first-token (TTFT), time-per-output-token (TPOT), inter-token latency, and **goodput** (tokens produced that actually meet the SLO). These are the dependent variables; they are what the SLO is written against.
3. **Resource metrics** — *why* performance is what it is. HBM bandwidth utilization, FLOP utilization / model FLOPs utilization (MFU), KV-cache occupancy, memory pressure, and queue depth. These reveal the binding constraint.

The discipline this chapter works toward: **never jump from a workload metric straight to a resource metric**. The causal chain is *workload*  $\rightarrow$  *serving*  $\rightarrow$  *resource*. A slow response (serving) is diagnosed by asking whether the kernel is bandwidth- or compute-bound (resource), which is only meaningful given the token profile (workload). Skipping a layer produces superstition — “it’s slow, let’s buy GPUs” — instead of a mechanism.

*Fig 6.1 — The three metric layers and the top-down diagnostic path from a p95 TTFT breach to its root cause. Causation runs bottom-up (resource  $\rightarrow$  serving  $\rightarrow$  workload); diagnosis runs top-down.*

### 6.2.2. Why the Wrong Metric Persists

Teams default to measuring what they already have dashboards for: model quality (accuracy, ROUGE, benchmark scores) and average latency. Both are dangerously incomplete. Quality is the *goal* but never the *bottleneck*; average latency is a single scalar that hides the tail that actually violates SLAs. The entire thesis of this chapter is that the metric that drives an architecture decision is usually **not** the metric the model card advertises.

## 6.3. Mental Model

Think of the three metric layers as a **diagnostic chain**, like the layers of the OSI stack in miniature:

- **Workload** is the *traffic* — what we are asked to do.
- **Serving** is the *service* — how well we did it.
- **Resource** is the *engine* — how hard the machine is working, and on what.

When something breaks, we read top-down: a p95 TTFT breach (serving) forces us to look at the token profile (workload: did context length grow?) and then the resource layer (resource: is prefill queuing because we’re FLOP-bound?). When the numbers are healthy, the chain confirms our architecture was right.

The mental model also inculcates **one number naming many different things**. “Throughput” without a qualifier is meaningless: is it requests per second, tokens per second, or *goodput* — tokens per second that arrived before the SLO deadline? We commit to always qualifying.




FIGURE 6.1. Fig 6.1 — The metric hierarchy as a diagnostic chain [ILLUSTRATIVE conceptual]




FIGURE 6.2. Fig 6.2 — Request-latency distribution: p50/p90/p95/p99 and the mean. The mean ( $\sim 0.84$  s) hides the 1% stragglers at  $\sim 5$  s that breach the p95 TTFT budget (ILLUSTRATIVE)

## 6.4. Worked Example

### 6.4.1. Latency Is a Distribution, Not a Mean

The single most common measurement error is reporting (and SLO-ing) the *average* latency. Consider a request stream where 99% of requests complete in 0.8 s but 1% take 5 s. The mean is

$$\mathbb{E}[L] = \sum_i p_i \cdot L_i = 0.99 \times 0.8 + 0.01 \times 5.0 \approx 0.84 \text{ s}$$

— looks fine. But under a p95 TTFT SLO of  $\leq 2$  s, the 99th-percentile straggler blows the p95 budget the moment the outlier fraction crosses 5%. Averages wash out the tail; percentiles expose it.

*Fig 6.2 — Why the mean is not a signal. The same request stream of 99% @ 0.8 s + 1% @ 5 s reported*

as a single number looks healthy (mean  $\approx 0.84$  s), yet the  $p95$  and  $p99$  paint a very different picture under a 2 s SLO. Log the distribution; SLO against the percentile.

Apply this to the canonical workload. We budget  $\text{TTFT} \leq 1.2$  s (retrieval  $\sim 120$  ms + prefill  $\sim 1.08$  s) with a  $p95 \leq 2$  s. Suppose a flash crowd (the  $\sim 40$  rps peak) causes prefill requests to queue. If on average the batch is well-behaved,  $p50$  TTFT might sit at  $\sim 1.0$  s. But the jobs that arrive behind 3 same-time prefill requests each pay the full 1.08 s prefill of the job ahead of them, so  $p95$  TTFT drifts to  $\sim 2.5$  s. The *average* says “fine”; the  $p95$  says “we just violated our SLO.” We hence log  $p50$ ,  $p90$ ,  $p95$ , and  $p99$  independently and run the SLO against the percentile, never the mean. [2°] SLO-derived from the §14 canonical latency budget; the arithmetic is the point, not a citation.

#### 6.4.2. Goodput vs Raw Throughput

Deploying a higher-throughput configuration can *reduce* goodput. Concretely: suppose we enable larger dynamic batching to raise raw tokens/s from 50 tok/s to 70 tok/s, but the batch window pushes  $p95$  TTFT from 1.8 s to 2.3 s — above the 2 s SLO. The extra 20 tok/s are produced, but every one of them misses the deadline. Raw throughput is up 40%; **goodput** (tokens that met the SLO) is effectively unchanged, because the batch of SLO-compliant tokens didn’t grow. Goodput is the metric the customer experiences; raw throughput is what the GPU achieves. The architect optimizes goodput. This is why “faster” hardware or bigger batches do not trivially mean a better system.

#### 6.4.3. Detecting the Regime: Bandwidth-Bound vs Compute-Bound

The most decision-relevant resource question is: is this layer bandwidth-bound or compute-bound? We recall the arithmetic from [Chapter 2](#):

- **Decode is HBM-bandwidth-bound.** Generating one token needs a full forward pass over the 70B weights = 140 GB (FP16) of reads per token. Within a  $\sim 25$  ms TPOT budget that demands  $140 \text{ GB} / 0.025 \text{ s} \approx 5.6 \text{ TB/s}$  of HBM bandwidth, against an H100’s  $\sim 3.35 \text{ TB/s}$  peak [2° FACT].  $5.6 > 3.35$ , so a single H100 *cannot* meet the budget on bandwidth alone — decode is memory-bound, and the fix is more bandwidth (H200 4.8 TB/s), fewer bytes per token (quantization), or fewer weight reads (speculative decoding / smaller active params). [2° DERIVED]
- **Prefill is compute-bound.** The prefill pass for 9,200 input tokens costs  $\approx 2 \times 70\text{e}9 \text{ params} \times 9.2\text{e}3 \text{ tokens} \approx 1.29 \text{ PFLOP}$ , which against the  $\sim 1.08$  s prefill budget implies a required rate of  $\approx 1.19 \text{ PFLOPS}$  — above an H100’s  $\sim 0.989 \text{ PFLOPS}$  dense BF16 peak ( $1.19 > 0.989$ ), so prefill is FLOP-bound: the machine caps out on compute, not memory traffic. [2° DERIVED]

At the resource layer, the tell is utilization: a decode kernel pinned at  $\sim 95\%$  of HBM bandwidth but  $\sim 30\%$  of FLOPs is bandwidth-bound; a prefill kernel with  $\sim 90\%$  FLOP utilization is compute-bound. **The same hardware, the same model — the boundary flips between prefill and decode.** This is precisely why Splitwise/DistServe (and Mooncake’s KV-centric disaggregation) split prefill from decode onto different pools: the bottleneck-regime difference justifies it. [S4][1P]

Table 6-1 — Detecting the bottleneck regime from resource telemetry

| Phase                | Demand                   | H100 capacity         | Verdict             | Resource tell                             |
|----------------------|--------------------------|-----------------------|---------------------|-------------------------------------------|
| Prefill (9.2K input) | ~1.29 PFLOP              | ~0.989 PFLOPS<br>BF16 | compute-bound       | high FLOP /<br>MFU util, mid<br>bandwidth |
| Decode (per token)   | ~5.6 TB/s weight<br>read | ~3.35 TB/s HBM        | bandwidth-<br>bound | high HBM util,<br>low FLOP util           |

(All figures [2° FACT] hardware / [2° DERIVED] arithmetic traced to §14 canonical and [Chapter 2](#).)

#### 6.4.4. Reading Telemetry Into a Decision

A concrete dashboard-read → decision trace. We see: p95 TTFT climbing (serving), KV-cache hit ratio low (resource: memory), and prefill queue depth growing (serving). Workload is unchanged. Reading up the chain: low prefix-cache hit means many independent 9.2K prefixes are being prefill-processed from scratch on the ~40 rps peak — each ~1.29 PFLOP saturating FLOPs. The decision: enable prefix / RadixAttention-style caching (SGLang) or Automatic Prefix Caching (vLLM) so shared retrieved-context prefixes are reused instead of re-prefilled, cutting effective prefill FLOPs and TTFT. [S3][1P] Prefix caching is a compute-saving move, not a latency tuning — and the telemetry told us that.

### 6.5. Measurement

Three practical habits anchor the architect:

1. **Instrument the hierarchy, not just the app.** Deploy counters at all three layers from day one — rps and token profile (workload); TTFT/TPOT/goodput percentiles (serving); HBM bandwidth, FLOP/MFU, KV-cache occupancy, queue depth (resource). A system that only logs serving metrics cannot tell *why*; one that only logs resource metrics cannot tell *what it was asked*.
2. **Always qualify “throughput.”** rps ≠ tokens/s ≠ goodput. The same number, three stories. State the SLO window (p95 ≤ 2 s TTFT) before quoting any rate, or the rate is meaningless.
3. **Log the distribution.** Capture p50/p90/p95/p99 for every latency metric, and the full token-length distribution (a 32K-tail context is a different workload than a 9.2K average). The mean is a summary, not a signal.

These habits answer the book’s recurring question — *what would we actually measure here?* — with three layers of counters at the edge, before any architecture decision is committed.

## 6.6. Common Mistakes

- **SLO-ing the average.** A 0.84 s mean hides a 5 s tail. SLO against p95/p99, and watch the outlier fraction, not the mean.
- **Optimizing raw throughput when the SLO windows it.** Bigger batches raise raw tokens/s but can push TTFT past its deadline, zeroing goodput. Optimize goodput.
- **Jumping straight to “buy more GPUs.”** A bandwidth-bound decode and a prefill-cost problem need different fixes; without the resource layer we guess.
- **Treating hardware peak as a ceiling we can hit.** H100’s 3.35 TB/s is the theoretical HBM peak; real memory-bound kernels realize a fraction of it. Design headroom against real utilization, not the spec sheet.
- **Conflating model-quality metrics with serving metrics.** Accuracy is the goal; TTFT is the constraint. A slow, accurate model fails its SLO just as surely as a fast, wrong one.

## 6.7. Architecture Consequence

The metric hierarchy directly dictates architecture. For the canonical workload:

- **Bandwidth-bound decode + compute-bound prefill  $\Rightarrow$  consider P/D disaggregation** (Split-wise/DistServe/Mooncake) at scale, because squeezing both regimes from one resource pool fights itself. [S4][1P]
- **Repeat-prefix RAG  $\Rightarrow$  enable prefix caching** (RadixAttention/APC), which turns repeated 9.2K prefills into cached lookups and defends the p95 TTFT against the 40 rps peaks. [S3][1P]
- **KV-cache pressure  $\Rightarrow$  quantize or offload KV.** FP8 KV is  $\approx 54\%$  of BF16 [S6][1P], and vLLM KV-offloading cuts TTFT by  $2\text{--}22\times$  [S6][1P] — both are memory-layer fixes the telemetry would call for when KV occupancy nears capacity.
- **Speculative decoding where decode bandwidth is the wall** — DFlash reports  $>4.3\times$  over baseline /  $1.5\times$  vs native MTP [S5][1P] — because it reduces weight reads per accepted token.

## 6.8. What We Still Don’t Know

- **Goodput in production is rarely published.** The research corpus gives strong qualitative and some quantitative anchors (Mooncake 525% simulated throughput / 75% more requests [S4][1P], DFlash  $>4.3\times$  [S5][1P]), but portable end-to-end goodput figures across real enterprise RAG workloads remain [HYPOTHESIS].
- **Prefix-cache hit rates in mixed workloads** are workload-dependent; the actual reduction in prefill FLOPs from APC/RadixAttention for a given document-retrieval distribution is [VERIFY].
- **Tail-latency drivers under batching** interact: batch size, arrival burstiness, and KV fragmentation each affect p99 differently, and their joint effect is [HYPOTHESIS].

### 6.9. End-of-Chapter Mini-Case

An on-call architect is paged: “the internal Q&A tool got slow after lunch.” The dashboard shows rps at the ~40 rps peak (workload, up from the 10 rps baseline), p95 TTFT at 2.4 s against a 2 s SLO (serving, breached), and KV-cache occupancy at 85% with low prefix-hit ratio (resource, memory/FLOP pressure). Following the chain rather than guessing, the architect reads: a concurrency spike (workload) with repeated full 9.2K prefills (low prefix hit) is saturating prefill FLOPs (compute-bound). The architecture is memory- and compute-sane at the 10 rps design point but was never sized for the lunch peak. The fix is not “buy GPUs immediately” — it is to enable prefix caching to collapse the repeated prefills, which the telemetry shows will address the actual mechanism. The architect re-checks the dashboard an hour later: p95 TTFT back under 2 s, goodput restored — because they measured the right thing at the right layer.

## **PART III.**

# **The System**

## CHAPTER 7.

# Memory Is the First Constraint

### 7.1. The Architect’s Question

After this chapter we should be able to answer the architect’s fundamental constraint question: *does the model and its context fit in the memory floor before we ask whether it can compute or communicate?* We will build the arithmetic for the KV cache from first principles, reconcile it with the canonical ~1.3 MB/token figure from [Chapter 1](#), and contrast inference residency (weights + KV) with fine-tuning residency (weights + gradients + optimizer states). The goal is to make “does it fit?” a concrete, quantified check — not a back-of-the-envelope guess. Throughout, the arithmetic anchors to the canonical enterprise-Q&A RAG scenario (§14): ~10 rps average / ~40 rps peak traffic (2,000 registered users × 5% concurrency → 100 concurrent → ~10 rps via Little’s Law; peaks at 20% concurrency → ~40 rps) served by a 70B FP16 model on an 8×H100 host. KV-cache per-token figures are given at the referenced precision throughout: ~1.3 MB/token at **8-bit**, ~2.5 MB/token at **FP16**; the canonical scenario’s headline ~1.3 MB/token figure is the 8-bit variant.

### 7.2. Concept

Memory is the first constraint because every LLM workload begins with a residency check: can the model and its active data structures be placed on the target hardware? The KV cache is the primary expression of this constraint during inference: it grows linearly with context length, and its size depends on the model’s layer count, hidden dimension, and the precision chosen for key and value storage. Unlike FLOPs, which scale with operations per token, the KV cache occupies memory continuously for the lifetime of a request — it is always resident, always consuming HBM bandwidth for every re-read. The architect must therefore size two distinct memory loads: the static

weight footprint and the dynamic KV footprint that varies with prompt length. Quantization (FP8, FP4, 2-bit asymmetric) reduces both, but the KV cache is the one component that does not shrink as aggressively as weights — per-block scale/offset metadata and per-token variation mean the memory reduction is typically 40–50% rather than the  $2\times-4\times$  seen in weight-only quantization. The residency floor is therefore set by weights + KV (+ activations, in fine-tuning), and the architect’s first question is whether this sum fits on the target GPUs. This is why “does it fit?” is the architect’s primary gate: before we ask whether the model can compute or communicate, we must first establish that the working set — weights plus the context-dependent KV state — resides on-GPU. If it does not, every subsequent decision (offload, recompute, page) is a degradation, not a cost-optimization.

The residency floor has two distinct regimes. Inference, the floor is weights + KV: the model’s static parameters plus the context-dependent key/value tensors that grow with every token added to the prompt. In fine-tuning, the floor rises further to weights + gradients + optimizer states, because every parameter update requires its own momentum and velocity vectors (Adam) or per-parameter scaling (LAMB/RMSProp). The contrast between these two floors — inference typically 160–200 GB for a 70B model at FP16 with 9.2K context [2° DERIVED], fine-tuning exceeding 1 TB [2° DERIVED] — is the architect’s central memory-takeaway.

### 7.3. Mental Model

Think of the KV cache as **per-request state that survives for the duration of a single forward pass** — unlike weights, which are loaded once and reused across requests, the KV cache is allocated per request and grows with every token added to the context. A 9.2K prompt on a 70B model at FP16 carries ~24 GB of KV state that must be on-GPU for the entire prefill + decode sequence. If it does not fit, the serving stack must page, offload, or recompute — each option degrades latency or throughput. The durable mental model is: *every token we add to the context exacts a fixed memory toll, and that toll is paid in HBM every time the GPU re-reads the cache during decode*. This is why context length is the single most powerful lever on inference memory pressure, and why “does it fit?” is the architect’s primary gate.

### 7.4. Worked Example

We anchor all arithmetic in the canonical scenario (§14 of book-architecture.md): a 70B-class dense model, FP16 weights (2 bytes per parameter), 8 × GPUs (80 GB each, 640 GB total VRAM), ~9.2K input tokens + 300 output tokens. We compute KV-cache size per token, total KV for the canonical context, and contrast with long-context variants.

**Table 7-1** — KV-cache size per token and per-context arithmetic for a 70B-class dense model at FP16. (All per-context KV totals, residency figures, and fit/non-fit verdicts below are [2° DERIVED] from the per-token formula  $2 \times \text{layers} \times \text{hidden\_dim} \times \text{bytes}$ ; model-card architecture constants are [1P: model card].)

| metric                                         | value                                                               | derivation                                                          |
|------------------------------------------------|---------------------------------------------------------------------|---------------------------------------------------------------------|
| layers                                         | 80                                                                  | canonical 70B full-MHA teaching model [1P: canonical-workload.yaml] |
| hidden_dim                                     | 8192                                                                | canonical 70B full-MHA teaching model [1P: canonical-workload.yaml] |
| bytes per parameter (FP16)                     | 2                                                                   | FP16 = 2 bytes/param [1P: facts/quantization.md Q3]                 |
| KV cache per token per layer                   | $2 \times \text{hidden\_dim} \times \text{bytes}$                   | one key + one value per layer                                       |
| KV cache per token (FP16)                      | $2 \times 80 \times 8192 \times 2 \text{ B} \approx 2.5 \text{ MB}$ | = 2,560 KB; rounded                                                 |
| KV cache per token (8-bit)                     | $2 \times 80 \times 8192 \times 1 \text{ B} \approx 1.3 \text{ MB}$ | = 1,280 KB; reconciles with Ch.1                                    |
| 9.2K input KV cache                            | $9,200 \times 2.5 \text{ MB} \approx 23.9 \text{ GB}$               | $9,200 \times 2,560 \text{ KB}$                                     |
| 300 output KV cache                            | $300 \times 2.5 \text{ MB} \approx 0.75 \text{ GB}$                 | $300 \times 2,560 \text{ KB}$                                       |
| total inference KV (9.2K + 300)                | $\approx 24.7 \text{ GB}$                                           | weights 140 GB + KV 24.7 GB $\approx 164.7 \text{ GB}$              |
| 32K input KV cache                             | $32,000 \times 2.5 \text{ MB} \approx 80.0 \text{ GB}$              | long-context variant                                                |
| 128K input KV cache                            | $128,000 \times 2.5 \text{ MB} \approx 320.0 \text{ GB}$            | aggressive long-context variant                                     |
| 70B FP16 weights                               | $70\text{B} \times 2 \text{ B} = 140 \text{ GB}$                    | [1P: DERIVED from 70B $\times$ 2 bytes]                             |
| inference residency (weights + KV, 9.2K input) | $\approx 164.7 \text{ GB}$                                          | $140 + 24.7 \text{ GB}$                                             |
| inference residency (weights + KV, 32K input)  | $\approx 220.0 \text{ GB}$                                          | $140 + 80.0 \text{ GB}$                                             |
| inference residency (weights + KV, 128K input) | $\approx 460.0 \text{ GB}$                                          | $140 + 320.0 \text{ GB}$                                            |
| 2×H100 total VRAM                              | $2 \times 80 \text{ GB} = 160 \text{ GB}$                           | [1P: facts/serving.md S8]                                           |
| 8×H100 total VRAM                              | $8 \times 80 \text{ GB} = 640 \text{ GB}$                           | [1P: facts/serving.md S8]                                           |
| fits 2×H100 at 9.2K?                           | no, 164.7 GB > 160 GB                                               | requires $\geq 2 \times \text{H100}$ with small overflow            |
| fits 8×H100 at 9.2K?                           | yes                                                                 | $164.7 \text{ GB} < 640 \text{ GB}$                                 |

| metric               | value       | derivation                         |
|----------------------|-------------|------------------------------------|
| fits 8×H100 at 32K?  | yes         | 220.0 GB < 640 GB                  |
| fits 8×H100 at 128K? | yes, barely | 460.0 GB < 640 GB                  |
| fits 1×H100 at 9.2K? | no          | 164.7 GB > 80 GB; needs<br>≥2×H100 |

*Note on “does it fit?”: this residency check answers a single-request question — does ONE request’s weights + KV fit in the pool? It says nothing yet about how many such requests fit concurrently. For that, an architect must subtract the runtime/activation/workspace floor (~64 GB on the 8×H100 host) from the 640 GB pool before sizing KV. We compute that concurrent capacity — the ~436 GB KV budget and ~18 requests/host — in [Chapter 17](#). Keep the two numbers separate: the ~500 GB headroom (Ch. 4/7, single-request) and the ~436 GB KV budget (Ch. 17, concurrency) answer different questions.*

*Attention-family sensitivity note.* The 2.5 MB/token figure is the **full-MHA** (multi-head attention) upper bound: KV-heads = query-heads = hidden\_dim, i.e.

$$KV_{\text{MHA per-token}} = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times B$$

Real 70B-class dense models commonly use **grouped-query attention (GQA)**: LLaMA-2 70B, for instance, has 64 query heads but only 8 KV heads, so per-token KV shrinks by the head ratio to

$$KV_{\text{GQA per-token}} = 2 \times n_{\text{layers}} \times n_{\text{KV-heads}} \times d_{\text{head}} \times B = 2 \times 80 \times (8 \times 128) \times 2 \text{ B} \approx 0.33 \text{ MB/token}$$

— about 8× **smaller**. For the canonical 9.2K context that is ~3 GB of KV instead of ~24 GB, and the fleet-sizing arithmetic of [Chapter 17](#) changes accordingly (many more concurrent requests fit for KV). We keep the full-MHA bound as the canonical conservative residual because an architect should size against the worst case the *chosen* model actually presents — the honest move is to read the model card’s KV-head count and substitute it into the GQA formula above. The framework is unchanged; only the constant changes. (kv\_gqa\_variant\_mb is recorded in canonical-workload.yaml.)

*Fig 7.1 — The visual proof of the 8× saving. Full MHA caches a K,V per query head (64/token → 2.5 MB/token); GQA shares one K,V across a group of 8 query heads, so only 8 K/V per token → ~0.33 MB/token. Head count is the real dial: substitute the model’s KV-head count into  $2 \times \text{layers} \times \text{KV\_heads} \times \text{head\_dim} \times \text{bytes}$ .*

**Table 7-2** — Inference residency vs. fine-tuning residency contrast for a 70B model.

## Grouped-Query Attention (GQA): why the cache is 8× smaller




FIGURE 7.1. Fig 7.1 — Grouped-Query Attention: 64 query heads share 8 KV heads, cutting per-token KV 8× from 2.5 MB to ~0.33 MB (ILLUSTRATIVE, after LLaMA-2 70B head layout)

| scenario                                               | memory resident                              | bytes per param                                            | total (70B)           | fits 8×H100 (640 GB) |
|--------------------------------------------------------|----------------------------------------------|------------------------------------------------------------|-----------------------|----------------------|
| inference<br>(weights + KV,<br>9.2K input)             | weights + KV<br>cache                        | FP16: 2 B; KV<br>adds ~2.5<br>MB/token                     | ~165 GB (9.2K)        | yes                  |
| fine-tuning<br>(weights +<br>gradients +<br>optimizer) | weights +<br>gradients +<br>optimizer states | Adam: ~16<br>B/param (m + v +<br>$\Delta W$ )              | ~1,260 GB             | no                   |
| fine-tuning<br>(weights +<br>gradients, FP16)          | weights +<br>gradients FP16                  | 4 B/param<br>(weights FP16 2<br>B + gradients<br>FP16 2 B) | ~280 GB               | yes                  |
| fine-tuning<br>(QLoRA, 4-bit<br>base + LoRA)           | quantized base +<br>LoRA adapters            | NF4: 4 B base;<br>LoRA ~5 M<br>params                      | ~50–70 GB             | yes                  |
| inference with<br>FP8 KV                               | weights FP16 +<br>KV FP8                     | KV: ~54% of BF16                                           | ~12.9 GB KV @<br>9.2K | yes, comfortably     |

*Inference residency = weights + KV cache; the KV footprint scales with context length. Fine-tuning residency adds optimizer states (Adam  $m$ ,  $v$ , and updates), which for 70B at ~16 bytes/param exceeds 1 TB — roughly 2× the 640 GB of 8×H100. This contrast is the architect’s central takeaway: inference is memory-limited by weights + KV, while fine-tuning is memory-limited by weights + gradients + optimizer, a substantially higher floor.*

## 7.5. Measurement

Four practical measurement habits follow from the arithmetic above:

1. **Compute KV-cache per-token cost for the model.** The per-token KV footprint of a fully-cached (MHA) attention head is

$$KV_{\text{per-token}} = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times \text{bytes-per-value}$$

(one key and one value per layer, each  $d_{\text{hidden}}$  wide), evaluated at the chosen precision. For a 70B model at FP16 the answer is  $2 \times 80 \times 8192 \times 2 \text{ B} \approx 2.5 \text{ MB/token}$ ; at 8-bit, ~1.3 MB/token (the Ch.1 reference). Do not assume the 1.3 MB/token figure applies at FP16 — it does not; it is an 8-bit

number. When evaluating a model where the hidden dimension or layer count differs from LLaMA-70B, recompute the constant factor — a model with 64 layers and 8192 dim will have  $\sim 2$  MB/token at FP16, while a model with 100 layers and 12288 dim will have  $\sim 3.7$  MB/token.

2. **Log peak context length, not just average.** A workload that averages 9.2K input tokens but has a long tail toward 32K or 128K will have a very different KV-cache profile. Log the 95th-percentile context length and re-run the KV arithmetic; the memory floor may shift from fitting on  $2 \times H_{100}$  to needing  $8 \times H_{100}$ . This is especially important for RAG workloads, where the retrieved context may average 8K but occasionally peak at 32K+ for technical documents.
3. **Measure inference residency as weights + KV, not weights alone.** It is common to quote only the weight footprint (140 GB for 70B FP16) and conclude that a single H100 can hold the model. But at 9.2K input, the KV cache adds  $\sim 25$  GB, pushing the total to  $\sim 165$  GB — requiring  $\geq 2$  GPUs. Always include KV when we state the memory floor for inference. This measurement habit is the token-layer answer to the book’s recurring question, “what would I actually measure here?” — we measure token counts and their distribution, at the edge, before any architecture decision is made.
4. **Account for KV quantization when reporting residency.** If the serving stack uses FP8 KV ( $\approx 54\%$  of BF16), the per-token KV cost drops from  $\sim 2.5$  MB to  $\sim 1.4$  MB, and the 9.2K residency falls from  $\sim 165$  GB to  $\sim 153$  GB (weights 140 GB + KV 12.9 GB). State both the BF16 and FP8 residency numbers, because the choice of KV dtype changes the hardware requirement decisively. For comparison, 8-bit KV would bring the 9.2K total to  $\sim 138$  GB (weights 140 GB truncated by KV reduction), still fitting in  $2 \times H_{100}$  but with less margin.

## 7.6. Common Mistakes

- **Assuming KV-cache size scales sub-linearly.** The KV cache grows linearly with context length. Doubling the context from 9.2K to 18.4K doubles the KV memory, all else equal. There is no automatic compression unless the attention mechanism itself is hybrid or sparse.
- **Using the 8-bit KV figure ( $\sim 1.3$  MB/token) at FP16.** [Chapter 1](#)’s  $\sim 1.3$  MB/token is explicitly at 8-bit precision. At FP16, the per-token KV cost is  $\sim 2 \times$  that,  $\approx 2.5$  MB/token. Mixing the two figures produces wrong residency calls.
- **Ignoring the weight + KV residency contrast.** Quoting only the weight footprint (140 GB for 70B FP16) and claiming it fits on one H100 (80 GB) is a category error. Weights alone exceed one H100; KV must be added.
- **Treating quantization as a uniform  $2 \times$ – $4 \times$  reducer.** KV-cache quantization (FP8  $\approx 54\%$  of BF16) gives a  $\sim 46\%$  reduction, not  $2 \times$  or  $4 \times$ . Weight quantization gives the larger reductions; do not apply the same expectation to the KV cache.
- **Overlooking the fine-tuning residency floor.** Full fine-tuning of 70B requires  $\sim 1.26$  TB with

Adam optimizer states — roughly  $2 \times 8 \times H_{100}$ . This is not a temporary overhead; it is the permanent memory floor for the training duration.

### 7.7. Architecture Consequence

The memory floor is the first constraint every architecture decision respects, because unlike compute or bandwidth it cannot be moved by better kernels — it is fixed by weights + KV + (for training) optimizer state. Several concrete consequences follow for the canonical enterprise-Q&A workload:

- **Model residency decides host count.** A 70B FP16 model (140 GB weights) at 9.2K context needs  $\sim 24.7$  GB of KV cache (FP16), so inference residency  $\approx 165$  GB — it does *not* fit on a single 80 GB H100, and is marginal even on  $2 \times H_{100}$  ( $160 \text{ GB} < 165 \text{ GB}$ ). The canonical single host of  $8 \times H_{100}$  (640 GB) is chosen precisely to clear this floor with headroom. [2° DERIVED]
- **KV-cache quantization buys back memory, not compute.** FP8 KV  $\approx 54\%$  of BF16 (a  $\sim 46\%$  reduction) [2°], dropping 9.2K residency from  $\sim 165$  GB toward  $\sim 153$  GB. This is a *memory* lever, orthogonal to bandwidth/compute fixes — the architect pulls it when the KV floor, not decode bandwidth, binds.
- **Fine-tuning is a different memory regime than inference.** The same model that serves in  $\sim 165$  GB demands  $\sim 1,260$  GB under full Adam fine-tuning (weights 140 GB + gradients 140 GB + optimizer states  $\sim 1,120$  GB) — roughly  $2 \times$  the  $8 \times H_{100}$  host. This is why the architect separates the serving fleet from the training fleet: the memory floors differ by an order of magnitude. [2° DERIVED]
- **Context length is the largest controllable KV lever.** Doubling context from 9.2K to 18.4K doubles KV ( $\sim 24.7 \text{ GB} \rightarrow \sim 49 \text{ GB}$ ); 128K context drives KV to  $\sim 320$  GB, which forces quantization or model parallelism. The architecture must set a context-length ceiling to keep the served model within host memory. [2° DERIVED]

In short: the architect sizes the host by weights + KV at the longest supported context, treats KV quantization as a spare memory dial, and keeps inference and training on separate memory-planning tracks.

### 7.8. What We Still Don't Know

- Exact per-token KV-cache size for non-LLaMA architectures (e.g., models with `head_dim`  $\neq 128$  or 256, or attention implementations that store additional per-head state). The  $2 \times \text{layers} \times \text{hidden\_dim} \times \text{bytes}$  formula assumes the standard LLaMA/llama-style key/value layout; architectures with grouped-query attention (GQA) or multi-query attention (MQA) may have fewer key/value entries per layer, changing the constant factor.
- Quality loss from aggressive KV quantization (2-bit KIVI, or sub-4-bit schemes) at very long

context lengths (128K+, MoE experts). The primary sources anchor FP8 KV  $\approx 54\%$  of BF16 at near-zero quality loss, but 2-bit asymmetric patterns and their interaction with RoPE and sliding-window layers are not fully mapped.

- Whether per-request KV caching can be partially offloaded to CPU DRAM during decode without TTFT impact, beyond the vLLM KV Offloading Connector's  $2\text{--}22\times$  TTFT reduction range which is highly prompt-size-dependent. The community is converging on tiered KV storage (GPU resident hot set + CPU/DRAM cold set), but the latency trade-offs at concurrency  $> 1$  are not yet primary-anchored.
- **Frontier 2026 has begun re-engineering the KV constant factor, not just quantizing it.** All four 2026-class open architectures attack KV-cache size at the attention layer itself, on top of the per-token footprint this chapter derives: DeepSeek-V4's hybrid CSA+HCA reports KV cache at only  $\sim 10\%$  (Pro) /  $\sim 7\%$  (Flash) of DeepSeek-V3.2 at 1M-token context [1P: arXiv 2606.19348]; GLM-5.3-Flash's sparse+linear hybrid reports  $\sim 4.4\times$  KV reduction [1P: HF zai-org/GLM-5.3-Flash]; Kimi K3's Kimi Delta Attention + Attention Residuals targets information flow across long sequences [1P: arXiv 2607.24653]; Qwen3.8-Flash-Next combines Gated DeltaNet (compress history) with Qwen Sparse Attention (micro-block indexing) for long-context cost [1P: HF Qwen/Qwen3.8-Flash-Next]. For the architect this is a decisive shift: the KV arithmetic in this chapter (per-token  $\times$  context) is *not* a fixed constant across model generations — a 2026 hybrid-attention model can hold dramatically more context per byte of KV than the canonical  $70\text{B}/8\times\text{H}100$  framing assumed. Size the host against the *specific* model's KV scheme, not a universal constant.
- The impact of MoE routing on KV cache: does each token really emit a full key/value per layer across all experts, or does routing activate only a subset? The MoE-vs-dense facts (E1, E5) confirm that attention layers process every token densely and emit key/value per token, so MoE sparsity does not reduce KV — but the constant factor for MoE models (e.g., number of experts per layer) needs per-model validation.

## Figures

*Fig 7.2 — KV-cache growth with context length (FP16  $\sim 2.5$  MB/token, FP8  $\sim 1.4$ , 8-bit  $\sim 1.3$ ); at 128K the FP16 KV footprint climbs to  $\sim 320$  GB, approaching the  $\sim 436$  GB KV budget only well beyond 128K, and already reaches  $\sim 165$  GB inference residency at 9.2K.*

*Fig 7.3 — The same 70B model serves in  $\sim 165$  GB (weights + KV) but needs  $\sim 1,260$  GB for full Adam fine-tuning; QLoRA fits  $\sim 50\text{--}70$  GB on a single GPU.*

*Fig 7.4 — Where a serving host's 640 GB pool goes. 140 GB weights +  $\sim 64$  GB runtime/NCCL leaves  $\sim 436$  GB of KV budget; at 23.8 GB/request (FP16) that is  $C \approx 18$  concurrent requests, and FP8 ( $\sim 12.4$  GB/request) roughly doubles it to  $\sim 35$ . This is the arithmetic behind the single-host capacity in Ch17.*

*Fig 7.5 — The Memory Tetris: why context length is the ultimate memory lever. The 640 GB pool stacks  $\sim 64$  GB runtime/NCCL + 140 GB weights, leaving  $\sim 436$  GB of headroom. The FP16 KV cache*




FIGURE 7.2. Fig 7.2 — KV-cache size vs context length for a 70B model [2° DERIVED]




FIGURE 7.3. Fig 7.3 — Inference vs fine-tuning memory floor [2° DERIVED]

Where a serving host's 640 GB pool goes: the concurrency budget [2° DERIVED]





FIGURE 7.4. Fig 7.4 — The concurrency budget: where a 70B host's 640 GB pool goes [2° DERIVED]

host (640 GB) fills as context grows



baseline (weights + runtime) is context-independent; the KV cache is the lever that grows with context

FIGURE 7.5. Fig 7.5 — Memory Tetris: how the 8×H100 host’s 640 GB pool fills at three contexts (9.2K / 32K / 128K). Runtime ~64 GB + weights 140 GB + KV cache 23.8 / 80 / 320 GB [2° DERIVED]

*(orange) is the only block that grows with context — 23.8 GB at 9.2K, 80 GB at 32K, ~320 GB at 128K. Because the KV budget is fixed (~436 GB), longer context consumes it outright: at 9.2K it supports ~18 concurrent, but 128K leaves room for only ~1–2. This is what makes the KV constant (Ch7 §3) the single most capacity-relevant number in a serving design. [2° DERIVED]*

## 7.9. End-of-Chapter Mini-Case

An architect is asked to design an internal Q&A system over the company’s document repository. The initial ask is vague: “We need to let all 5,000 employees ask questions over our internal wikis.” Before any architecture can be defended, the architect must turn this vague desire into a machine-denominated statement using the chapter’s core constraint.

From the token layer alone (which we worked through in this chapter), the architect can already establish several facts. The unit is tokens; the request shape will be prompt + retrieved context + output; the workload is input-heavy (retrieved context dominates); and the first number to lock down is tokens-per-request, because every downstream decision (which model fits, how much memory, what latency is possible) is priced against it.

Running the KV arithmetic: if each request uses ~9.2K retrieved context tokens (the canonical ~1,200 prompt + 8K retrieved) plus ~300 output tokens, the KV cache per request at 70B FP16 is ~24.7 GB. The weight footprint is 140 GB. Total inference residency  $\approx$  165 GB, which fits in 2×H100 (160 GB) with a small margin, or comfortably in 8×H100 (640 GB). If the context is extended to 32K (e.g., wider document retrieval), the KV cache rises to ~80 GB, pushing total residency to ~220 GB — still fitting in 8×H100 but requiring a multi-GPU node.

The architect can now speak the system’s currency: tokens, KV bytes, and the residency floor. The next step — turning “5,000 employees” into ~2,000 concurrent users, ~10 requests/s, and a specific model-and-hardware choice — is [Chapter 4](#)’s job (workload anatomy). Here the point is narrower and sharper: we can now say, with quantified memory numbers, whether the system fits on the available hardware, and we can compare inference vs. fine-tuning residency floors before committing to a architecture decision. This is the chapter’s payoff: the vague request has been translated into concrete constraint numbers, and the architect can defend a recommendation based on whether the workload fits, not on intuition alone.

## CHAPTER 8.

# Compute

### 8.1. The Architect’s Question

After chapters on tokens, workloads, and memory, the architect naturally asks: *how much compute does this actually require, and at what point does the hardware cease to be the limiting factor?* This chapter gives us the FLOP-scale arithmetic, the arithmetic-intensity roofline, and the utilization framework that lets us answer that question without guessing. We will move from per-token FLOPs to prefill PFLOP counts to sustained MFU on real hardware, and then to the roofline that tells us whether a given layer is compute-bound or memory-bound — the difference that decides whether we should reach for batching, quantization, or a different hardware generation. Throughout, the arithmetic anchors to the canonical enterprise-Q&A RAG scenario (§14): ~10 rps average / ~40 rps peak traffic (2,000 registered users  $\times$  5% concurrency  $\rightarrow$  100 concurrent  $\rightarrow$  ~10 rps via Little’s Law; peaks at 20% concurrency  $\rightarrow$  ~40 rps), with 70B FP16 weights and an 8 $\times$ H100 host. KV-cache per-token figures are referenced at FP16 (~2.5 MB/token) unless an 8-bit (~1.3 MB/token) variant is explicitly stated.

### 8.2. Concept

Compute in a language model is priced in floating-point operations, or FLOPs. Every matrix multiplication, every attention weight update, every activation function evaluation contributes to the total. The architect’s first principle is that FLOPs scale linearly with parameter count and context length, but the *efficiency* with which those FLOPs are executed depends on the hardware ridge point and the arithmetic intensity of the kernel.

For a dense transformer layer, the dominant cost is the matrix multiplication — query-key, value-

aggregation, and MLP projections. In FP16, a single multiply-add counts as two FLOPs. The per-token cost is therefore approximately  $2 \times N$  FLOPs where  $N$  is the number of active parameters. This is the arithmetic baseline: every token processed costs roughly two floating-point operations per parameter.

But FLOPs alone do not tell the full story. The hardware can only execute FLOPs if data is available — weights, activations, and KV cache all compete for the same HBM bandwidth. The arithmetic intensity (FLOP/byte) determines whether a kernel is compute-bound (intensity above the ridge point) or memory-bound (intensity below). This roofline model is the central diagnostic tool of this chapter: it tells us, for any given layer and precision, whether increasing FLOPs will actually reduce latency or whether we are already starved for bytes.

The chapter unfolds in four parts. First, the core arithmetic: FLOPs per token, prefill PFLOP counts, and H100 sustained utilization. Second, the roofline: why early layers are compute-bound while later layers and decode are memory-bound. Third, the canonical scenario arithmetic: 70B on  $8 \times$  H100, TTFT 1.2 s, TPOT ~25 ms. Fourth, the mini-case: a continuous deployment scenario that threads the prefill/decode divide.

### 8.3. Mental Model

Think of FLOPs as the distance a car can travel on a gallon of fuel: it tells us the system’s capacity, not how long the trip will take. The trip time is determined by the ridge point: if the arithmetic intensity is below the ridge, adding more FLOPs won’t help because we are waiting for data. If we are above the ridge, we are limited by how fast the compute can chew through operations. The roofline chart — peak TFLOPS on the vertical axis, FLOP/byte on the horizontal — is the map that resolves this. Where our kernel lands on that chart decides whether we should profile memory bandwidth or compute throughput.

### 8.4. Worked Example: Canonical Scenario Arithmetic

The canonical scenario (§14, book-architecture.md) is an enterprise Q&A system: 70B-class dense model, FP16 weights (~140 GB), 1 host with  $8 \times$  H100-class GPUs (80 GB each), ~10 requests/s average, peaks ~40 rps, average prompt 1,200 tokens + 8K retrieved context (~9.2K input), 300-token output, TTFT budget 1.2 s (retrieval ~120 ms + prefill), TPOT budget ~25 ms/token.

#### 8.4.1. FLOPs per token

For a 70B-parameter dense model in FP16, the per-token forward FLOP count is approximately:

$$\text{FLOP}/_{\text{token}} \approx 2 \times N = 2 \times 70 \times 10^9 \approx 140 \text{ GFLOP/token}$$

[DERIVED:  $2 \text{ FLOPs/parameter} \times 70 \times 10^9 \text{ params}$ ; standard dense-forward arithmetic, consistent with the  $175\text{B} \approx 0.35 \text{ TFLOP/token}$  correction from moe-vs-dense E5]

This applies to both prefill and decode: each token that enters the forward pass costs  $\sim 140 \text{ GFLOP}$ . The distinction between prefill and decode is not in the per-token FLOP count but in the data movement pattern — preflight reads the full context once, while decode re-reads weights per token.

#### 8.4.2. Prefill PFLOP for 9.2K input

For a single request with 9.2K input tokens, the prefill FLOP cost is:

$$\text{prefill FLOPs} \approx 2 \times N \times L = 2 \times 70 \times 10^9 \times 9.2 \times 10^3 \approx 1.29 \text{ PFLOP}$$

[DERIVED:  $2 \times N \times L$  where  $N=70\text{B}$ ,  $L=9,200$ ; reconciles with Ch.2 scaling law if the full 14.8T-token pre-train budget is distributed across active parameters; the figure is large because the full context is attended to, but it is a one-time cost per request, not a sustained rate]

*Accuracy of the  $2 \times N \times L$  approximation.* This linear prefill count omits the quadratic attention term,  $\approx 4 \times n_{\text{layers}} \times L^2 \times d$ . At the canonical 9.2K context that correction is small — roughly **+17%** of the 1.29 PFLOP — so the 2NL roofline is a sound teaching baseline there. But the omission grows with context: at 32K the attention term is  $\sim 60\%$  and at 128K it *dominates* ( $\sim 2.4\times$  the linear term). An architect sizing long-context prefill must add the quadratic term or measure it; the Unknowns section returns to this.

At 10 requests/s, the sustained prefill throughput demand is:

$$\text{prefill demand} = 1.29 \times 10^{15} \text{ FLOP} \times 10 \text{ rps} \approx 12.9 \text{ PFLOP/s}$$

[DERIVED]

**Table 8-1** — Per-token FLOP cost and prefill PFLOP for the canonical 70B workload. (*Per-token FLOP and prefill PFLOP computations are [2° DERIVED] from  $2 \times \text{params} \times \text{tokens}$ ; hardware peaks [1P: vendor datasheet] or [2° FACT]; the 30–40% MFU range is [2°: industry benchmarks]; the  $\sim 295 \text{ FLOP/byte}$  ridge is [2° DERIVED] =  $989 \text{ TFLOPS} \div 3.35 \text{ TB/s.}$ )* | metric | value | derivation |

| metric                       | value                          | derivation                                                                                                                          |
|------------------------------|--------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| per-token FLOPs (forward)    | $\sim 140 \text{ GFLOP/token}$ | $2 \times 70\text{B params}$ [DERIVED; consistent with moe-vs-dense E5 correction: $175\text{B} \approx 0.35 \text{ TFLOP/token}$ ] |
| prefill FLOPs for 9.2K input | $\sim 1.29 \text{ PFLOP}$      | $2 \times 70 \times 10^9 \times 9.2 \times 10^3$ [DERIVED]                                                                          |

| metric                            | value                                  | derivation                                                  |
|-----------------------------------|----------------------------------------|-------------------------------------------------------------|
| sustained prefill demand @ 10 rps | ~12.9 PFLOP/s                          | $1.29 \text{ PFLOP} \times 10$ [DERIVED]                    |
| H100 peak FP16 TFLOPS             | ~989 TFLOPS [1P: facts/training.md T5] | NVIDIA H100 SXM5 datasheet, without sparsity                |
| H100 sustained MFU (typical)      | 30–40% [2°: industry benchmarks]       | ~346 TFLOPS sustained at 35% MFU                            |
| H100 ridge point (dense FP16)     | ~295 FLOP/byte                         | $989 \div 3.35$ [DERIVED: peak TFLOPS $\div$ HBM bandwidth] |

*All figures trace to the canonical scenario (§14) and validated old-repo sources; none are measurement claims.* This is the prefill compute demand. An  $8 \times$  H100 node can sustain some fraction of this at MFU (mixed-precision FLOP utilization), which we estimate next.

#### 8.4.3. Sustained vs. peak: H100 MFU

NVIDIA H100 SXM5 lists ~1,979 TFLOPS FP16 Tensor Core peak “with sparsity” and ~989 TFLOPS FP16 peak without sparsity [1P: facts/training.md T5, NVIDIA datasheet]. In practice, sustained mixed-precision FLOP utilization (MFU) for a well-tuned transformer workload typically runs at 30–40% of peak on H100 [2°: industry benchmark reports, derived from real kernel profiles]. At 35% MFU:

$$\text{sustained} = \text{peak} \times \text{MFU} = 989 \text{ TFLOPS} \times 0.35 \approx 346 \text{ TFLOPS}$$

[DERIVED: peak  $\times$  typical MFU]

At 346 TFLOPS sustained, the number of tokens/s the system can prefill is:

$$\text{prefill tokens/s} = \frac{\text{sustained}}{\text{FLOP/token}} = \frac{346 \text{ TFLOPS}}{140 \text{ GFLOP/token}} \approx 2,471 \text{ tokens/s}$$

[DERIVED: sustained FLOPS  $\div$  per-token FLOPs]

For 10 rps with 9.2K input tokens each, the prefill demand is ~92,000 tokens/s (from the canonical workload). At 2,471 tokens/s per node, we would need roughly:

$$n_{\text{nodes}} = \frac{\text{tokens/s demand}}{\text{tokens/s per node}} = \frac{92,000}{2,471} \approx 37 \text{ nodes}$$

[DERIVED: token demand  $\div$  per-node prefill throughput]

This rough sizing illustrates that prefill is FLOP-bound at this scale — the compute requirement is the primary constraint, and memory bandwidth (HBM at 3.35 TB/s per H100) is sufficient to keep the

compute fed if the kernel is efficient. The ridge point for dense FP16 on H100 is  $\sim 295$  FLOP/byte ( $989 \div 3.35$ ); the arithmetic intensity of a mature attention kernel sits near or above this ridge, meaning the kernel is compute-saturated rather than bandwidth-starved once the batch is large enough.

#### 8.4.4. Decode: bandwidth-bound at low batch

*Fig 8.1 — The roofline: prefill is compute-bound, low-batch decode is memory-bound.*

For decode, the per-token FLOP cost is the same  $\sim 140$  GFLOP, but the arithmetic intensity drops sharply. Each decoded token re-reads all 70B weights from HBM. A 70B FP16 weight matrix is  $70 \times 10^9 \times 2 \text{ B} = 140 \text{ GB}$  — one full pass reads **140 GB**, not 280 GB. (The naïve “ $2 \times 70 \times 10^9 \times 2 \text{ B} \approx 280 \text{ GB}$ ” double-counts the FP16 footprint: the leading “2” is already inside the 2-bytes-per-element, so writing  $2 \times 70 \times 10^9 \times 2 \text{ B}$  counts the weight bytes twice. Any KV/output-projection reads are a separate, smaller term on top, not a second full weight footprint.) The effective arithmetic intensity is therefore:

$$\text{arithmetic intensity}_{\text{decode}} \approx \frac{140 \text{ GFLOP}}{140 \text{ GB}} \approx 1.0 \text{ FLOP/byte}$$

[DERIVED; a lower-bound bandwidth model — real kernels also read activations and KV, so measured intensity is typically below this and heavily dependent on fusion and cache residency]

This is well below the H100 dense-FP16 ridge of  $\sim 295$  FLOP/byte, meaning single-stream decode is overwhelmingly memory-bandwidth-bound. Throughput improves with batching because the weight-reread amortizes across tokens — at batch=32, the effective intensity rises toward the ridge, and TFLOP/s utilization increases accordingly. This is why continuous-batching systems (Orca, vLLM) prioritize batch growth: it is the primary lever for pushing decode arithmetic intensity toward the ridge.

### 8.5. Measurement

Three measurable quantities anchor compute sizing:

1. **Per-token FLOP count.** Measure by flops profiling (e.g. NVIDIA Nsight Compute) on the target model and precision. The  $2 \times N$  rule is a useful prior, but actual count varies by layer depth, activation choice (SiLU vs. ReLU), and whether kernel fusion combines matmuls.
2. **Sustained MFU.** Profile TFLOP/s achieved divided by peak TFLOP/s for the target hardware and precision. A well-tuned transformer on H100 FP16 typically lands in the 30–40% range; below 20% suggests a memory or communication bottleneck that should be diagnosed before scaling.
3. **Arithmetic intensity (FLOP/byte).** Compute as achieved FLOPs divided by bytes transferred (weights + activations + KV read/write). Compare to the hardware ridge point (H100 dense FP16:  $\sim 295$  FLOP/byte; H200:  $\sim 989 \text{ TFLOPS} \div 4.8 \text{ TB/s} \approx 206 \text{ FLOP/byte}$  — H200 keeps H100’s GH100 die and its dense FP16 peak, trading a bigger/faster HBM3e for bandwidth; GB200 BF16: much higher




FIGURE 8.1. Fig 8.1 — Prefill above the ridge (compute-bound), decode below it (memory-bound). H100 (solid) vs H200 (dashed) ridge points from vendor datasheets (989 TFLOPS, 3.35 / 4.8 TB/s); Continuous-Batching arrow shows decode climbing the slope as batch grows [2° DERIVED] [1P: vendor datasheet]

due to NVFP4). If intensity is below the ridge, the kernel is memory-bound; above, compute-bound.

## 8.6. Common Mistakes

- **Assuming  $2 \times N$  FLOPs/token is exact.** It is a baseline derivation; actual per-token FLOPs vary with layer depth, activation fusion, and kernel fusion. Profile before quoting.
- **Treating prefill and decode as having the same FLOP/byte profile.** Prefill reads the full context once and is FLOP-bound; decode re-reads weights per token and is memory-bound. The arithmetic-intensity roofline distinguishes them.
- **Using peak TFLOP/s as sustained throughput.** Peak includes sparsity or BF16/FP8 variants that may not be available in the chosen precision. Use the FP16 number without sparsity ( $\approx 989$  TFLOPS on H100) as the baseline for dense compute sizing.
- **Ignoring the batch-size dependence of arithmetic intensity.** A single-token decode batch has very low intensity; batching raises it. Always state the batch assumption when reporting TFLOP/s or tokens/s.
- **Confusing H100 sparsity-inclusive peak (1,979 TFLOPS) with dense peak (989 TFLOPS).** The  $2 \times$  sparsity figure appears in vendor data sheets but most transformer workloads do not activate sparsity. Quote the dense figure unless the kernel is specifically sparsity-enabled.

## 8.7. Architecture Consequence

The compute arithmetic and roofline have direct consequences for system design:

- **If prefill is FLOP-bound** (as it is for 70B+ models with long contexts), the primary optimization is batching — larger prefill batches amortize the context attention cost across more requests, raising arithmetic intensity and increasing TFLOP/s utilization. System design should therefore provision for batchable prefill (continuous batching, Orca, vLLM continuous batching) as the first-order throughput lever.
- **If decode is memory-bound** (single-stream, low batch), the primary optimization is batching — even modest batch sizes (8–16) raise arithmetic intensity toward the ridge, improving TFLOP/s utilization. KV cache quantization (FP8, int8) also reduces bytes per token, raising effective intensity. P/D disaggregation (prefill on GPUs, decode on a separate pool) can rebalance: prefill’s FLOP demand is served by compute-optimized GPUs, while decode’s bandwidth demand is served by wide-memory GPUs or even CPU offload.
- **If the ridge point is exceeded** (e.g. H200 with 4.8 TB/s HBM3e and 4 PFLOPS FP8, giving a  $\sim 833$  FLOP/byte ridge in FP8), the same workload may shift from memory-bound to compute-bound, changing the optimal hardware choice. This is why the H200 represents an inflection point where

inference and training hardware lines begin to converge [1P: facts/training.md T5].

In practice, the architect measures arithmetic intensity for the target workload and precision, locates the kernel on the roofline chart, and then selects hardware and batching strategy accordingly. The roofline is the diagnostic; the architecture decision follows.

## 8.8. What We Still Don't Know

- **Arithmetic intensity of fused kernels.** Most roofline estimates treat attention or MLP as separate matmuls, but in practice frameworks fuse multiple operations into a single kernel. The effective FLOP/byte of a fused attention+MLP kernel is not well cataloged in primary sources and would require profiling on the target hardware to anchor.
- **MFU across precisions and sparsity.** The 30–40% H100 FP16 figure is a rule of thumb; actual utilization depends on architecture, kernel quality, and batch config. A systematic MFU atlas across H100/H200/GB200, FP16/BF16/FP8, dense/MOE would be a valuable reference.
- **Arithmetic intensity of emerging attention mechanisms.** Linear attention (DeltaNet, MLA), compressed attention (CSA, HCA), and hybrid sparse patterns have different FLOP profiles and data movement. Their roofline position is not yet established in primary sources.
- **Frontier 2026 has begun to answer this.** DeepSeek-V4's hybrid CSA+HCA attention reports that, at 1M-token context, single-token inference drops to ~27% of the FLOPs (Pro) / ~10% (Flash) and KV cache to ~10% / ~7% of DeepSeek-V3.2 [1P: arXiv 2606.19348]. GLM-5.3-Flash's sparse+linear hybrid reports a ~3× attention-compute and ~4.4× KV-cache reduction [1P: HF zai-org/GLM-5.3-Flash]. These are [1P] first-party vendor figures, not yet independently reprofiled on our own hardware — which is exactly the arithmetic-intensity measurement an architect should still do before trusting a vendor's roofline claim for their own workload.
- **How quantization shifts the ridge.** Moving from FP16 to FP8/int8 changes peak TFLOP/s and bytes-per-operand, shifting the ridge point. The net effect depends on the quantization scheme (element-wise vs. per-tensor, dynamic vs. static) and whether the kernel is re-tuned for the lower precision.

## 8.9. End-of-Chapter Mini-Case: Continuous Deployment Scenario

An architect is brought into an ongoing deployment of a 70B-class Q&A system on  $8 \times$  H100 GPUs. The system is serving ~10 requests/s average with ~9.2K input + 300 output tokens per request, and the TTFT budget is being missed: p95 TTFT is 2.8 s, exceeding the SLO of 2 s. The TPOT of 28 ms/token is within spec, but the prefill delay is the bottleneck.

The architect first profiles the arithmetic intensity of the prefill kernel. The per-token FLOP count is ~140 GFLOP ( $2 \times 70\text{B}$ ), and the effective bandwidth per token is ~140 GB (re-reading the 70B weight matrix from HBM for each of the 9.2K input tokens). This gives an effective FLOP/byte

of roughly  $140 \text{ GFLOP} \div 140 \text{ GB} \approx 1 \text{ FLOP/byte}$  — well below the H100 dense-FP16 ridge of  $\sim 295 \text{ FLOP/byte}$ . The kernel is strongly memory-bound during prefill.

The immediate question: why is prefill memory-bound when the H100 has 3.35 TB/s HBM3 bandwidth? The answer is that a 9.2K input context means each prefill pass reads the full weight matrix once, but the activation keys/values for 9.2K tokens also flow through the pipeline. The total bytes per token are higher than the weight-read alone, and the effective intensity drops further.

The architect's first fix is to increase the prefill batch size. Continuous batching (vLLM Orca-style) accumulates incoming requests into larger batches, which raises the effective arithmetic intensity because the weight matrix read is amortized across more tokens. At batch=32, the effective intensity rises to  $\sim 2\text{--}3 \text{ FLOP/byte}$ , still below the ridge but much closer. At batch=128, the intensity approaches the ridge and TFLOP/s utilization climbs from  $\sim 15\%$  to  $\sim 45\%$ .

The second fix is to profile the actual MFU achieved. If the system is running at 12% MFU on prefill, there is headroom to absorb more batch without adding hardware. If MFU is already near 40%, the bottleneck is elsewhere (communication, kernel inefficiency, or KV cache residency), and simply batching harder yields diminishing returns.

The architect's recommendation: enable continuous batching with a target batch size of 64–128, profile the resulting MFU and TTFT, and if TTFT remains above SLO, consider P/D disaggregation — offload decode to a separate pool of GPUs with wide HBM bandwidth, freeing the prefill GPUs to run larger batches at higher FLOP utilization. The key insight is that the workload's position on the roofline chart (memory-bound prefill) dictates the fix: batching to raise intensity, or hardware that raises the ridge point (H200, GB200), or disaggregation that separates the two bottleneck profiles.

---

*Reading the roofline (summary of Fig 8.1): the dense-FP16 ridge at  $\sim 295 \text{ FLOP/byte}$  (H100:  $989 \text{ TFLOPS} \div 3.35 \text{ TB/s}$ ) separates the compute-bound region (right) from the memory-bound region (left). Prefill kernels typically operate in the  $1\text{--}3 \text{ FLOP/byte}$  range for long contexts, decode kernels near  $1.0 \text{ FLOP/byte}$  at batch=1; batching and quantization shift kernels rightward toward the ridge.*

*Prefill PFLOP also grows with context:  $2 \times 70\text{B} \times L$  tokens, from 2K to 32K. This is the one-time compute cost of loading a long context; sustained throughput (tokens/s) is what matters for continuous serving, and at the canonical 9.2K point it is  $\sim 1.29 \text{ PFLOP}$  (Table 8-1).*

## CHAPTER 9.

# Communication

### 9.1. The Architect’s Question

Before we can place a model across multiple GPUs or across a cluster, we need to know what the interconnect can — and cannot — move. This chapter asks: given a model size, a parallelism strategy, and a traffic pattern (all-reduce, all-gather, reduce-scatter), what is the communication time, and does it dominate the compute budget? We will size the hidden bottleneck that often goes unnoticed until the job stalls.

### 9.2. Concept

Communication in a distributed LLM system is not a single pipe. It is a hierarchy: each GPU talks to its neighbors over PCIe, nodes within a server meet over NVLink/NVSwitch, racks talk over 200/400 GbE RoCE, and cross-rack or cross-data-center traffic rides InfiniBand. Each layer has a distinct bandwidth ceiling, a distinct latency profile, and a distinct cost per gigabyte. The architect must match the collective operation to the layer where the bandwidth is sufficient, because moving the same data over a slower layer adds minutes to training time or inference latency for no computational benefit.

The collective operations that dominate multi-node training and inference are:

- **all-reduce**: every GPU contributes a gradient fragment; the result (the sum, scaled by  $1/N$ ) is delivered to every GPU. This is the workhorse of data-parallel and model-parallel training.
- **all-gather**: every GPU contributes a tensor; every GPU receives the concatenation of all fragments. Used when a model split across nodes needs the full weight or KV cache locally.
- **reduce-scatter**: the inverse of all-gather. Every GPU contributes a fragment; each GPU receives

a portion of the final result, sized  $1/N$  of the input. Used when the full tensor does not need to reside on any single GPU.

The bandwidth each collective consumes depends on three factors: the data volume (bytes), the number of participating nodes ( $N$ ), and the interconnect's effective bandwidth in the presence of contention. A 70B FP16 model is 140 GB. An all-reduce of that model across 8 GPUs on a single node moves 140 GB over NVLink; across 8 nodes connected by InfiniBand moves 140 GB (or more, depending on the algorithm) over the fabric. The same 140 GB over a 25 GB/s Ethernet link takes more than  $5\times$  the time of the same operation over a 900 GB/s NVLink.

**Table 9-1** — Interconnect hierarchy with bandwidth (bidirectional where applicable), latency, and approximate cost per GB. PCIe Gen5 x16  $\sim 64$  GB/s,  $\sim 1$   $\mu$ s latency, low cost (integrated); NVLink H100 900 GB/s bisection,  $\sim 0.5$ – $1$   $\mu$ s latency, high cost (GPU-attached); Ethernet RoCE2 200/400 Gb/s  $\approx 25/50$  GB/s,  $\sim 10$ – $25$   $\mu$ s latency, moderate cost (standard NICs); InfiniBand HDR 400 Gb/s / NDR 800 Gb/s  $\approx 50/100$  GB/s,  $\sim 2$ – $5$   $\mu$ s latency, high cost (switches, optics). Cost per GB is inversely proportional to bandwidth for fixed-form optics, but system cost includes cables, switches, and rack density trade-offs.

*Table 9-1 — The interconnect hierarchy*

| Layer                   | Example                       | Bandwidth<br>(one direction<br>/ aggregate)       | Latency   | Cost       | [source]  |
|-------------------------|-------------------------------|---------------------------------------------------|-----------|------------|-----------|
| Chip-to-memory          | HBM3 (H100)                   | 3.35 TB/s                                         | $\sim$ ns | —          | [2° FACT] |
| GPU-to-GPU<br>(in-node) | NVLink<br>(H100)              | 900 GB/s per GPU,<br>bidirectional<br>aggregate   | $\sim$ us | high capex | [2° FACT] |
| GPU-to-CPU /<br>NIC     | PCIe Gen5                     | $\sim 64$ GB/s per<br>x16 direction               | $\sim$ us | low        | [2° FACT] |
| Node-to-node<br>(rack)  | RoCE /<br>Ethernet<br>400Gb/s | $\sim 50$ GB/s per<br>port                        | $\sim$ us | medium     | [2° FACT] |
| Rack/cluster            | InfiniBand<br>HDR/NDR         | 400/800 Gb/s<br>$\approx 50/100$ GB/s<br>per port | $\sim$ us | high capex | [2° FACT] |

(Hierarchy: bandwidth drops ~2–3 orders of magnitude from HBM to the cluster fabric — the reason the interconnect, not the GPU, often binds at scale. Figures are established hardware specs [2° FACT].)

### 9.3. Mental Model

Think of the interconnect as a series of concentric rings. The GPU sits at the center, reachable fastest via PCIe. One ring out is the node: GPUs on the same server are joined by NVLink and NVSwitch, forming a high-radix, low-latency mesh. The next ring out is the rack: 10–20 GbE or 100/200/400 GbE RoCE connects nodes. The outermost ring is the cluster or data centre: InfiniBand HDR (400 Gb/s) or NDR (800 Gb/s) stitches racks into a single logical network. A well-designed deployment keeps the heavy collectives (all-reduce, reduce-scatter) on the inner rings and reserves the outer rings for less volume or latency-tolerant traffic.

The arithmetic is relentless. If all-reduce moves 140 GB and the NVLink bisection provides 900 GB/s per direction, a **lower bound** on time is  $140 \text{ GB} \div 900 \text{ GB/s} \approx 0.16 \text{ s}$  [2° DERIVED] — note the word *lower bound*. The actual completion time of a real collective is higher because a collective is not one flat transfer but a sequence of message-passing and reduction steps across the topology:

$$T_{\text{collective}} \approx \alpha \times n_{\text{steps}} + \frac{\text{bytes}}{\beta}$$

where  $\alpha$  is the per-step latency (message setup, synchronization),  $\beta$  is the achieved (not peak) bandwidth, and  $n_{\text{steps}}$  is the number of message-passing/reduction hops. (For a flat transfer  $n_{\text{steps}} = 1$ .) The 0.16 s figure uses only the  $\text{bytes}/\beta$  term at peak bandwidth; real NCCL all-reduce on NVLink lands tens of % above it once per-step overhead, ring stages, and topology routing are included. So treat the  $\text{bytes} \div \text{bandwidth}$  form as a *sizing lower bound*, and reserve measured `nccl-tests` numbers (Section 4) for capacity decisions. The relative ranking between interconnects is what matters most here: if those same 140 GB travel over a 400 Gb/s InfiniBand link ( $\approx 50 \text{ GB/s}$  effective unidirectional) it is  $\approx 2.8 \text{ s}$  —  $18\times$  longer; over 25 Gb/s Ethernet ( $\approx 3 \text{ GB/s}$  unidirectional)  $\approx 47 \text{ s}$ . The interconnect choice is a first-order determinant of wall-clock time in all cases.

*Fig 9.1 — All-reduce completion time as a function of data volume (x-axis, log GB) and interconnect effective bandwidth. The four lines — NVSwitch 1.8 TB/s, NVLink 0.9 TB/s, InfiniBand 0.4 TB/s, Ethernet 0.1 TB/s — are peak-bandwidth lower bounds ( $\alpha/\beta$  caveat in §9.3); the dashed marker sits at the 140 GB weight footprint of a 70B model ( $\approx 70 \text{ GB}$  of weights in BF16 across 8 participants). [2° DERIVED]*

### 9.4. Worked Example

Consider the canonical scenario: one host eight  $\times$  H100, each GPU holding 140 GB of model weights in FP16. We will run data-parallel training with full model replicas, so every all-reduce moves the full gradient tensor of 140 GB. Let us compute the communication time under three interconnect options.




FIGURE 9.1. Fig 9.1 — All-reduce time vs data volume, by interconnect tier [2° DERIVED]

**Option A — NVLink on a single host.** H100 GPUs provide 900 GB/s bidirectional NVLink bandwidth per GPU, and NVSwitch aggregates bisection bandwidth. For an all-reduce of 140 GB across 8 GPUs on one host, the effective bandwidth is close to the per-link peak because NVSwitch can forward multiple concurrent channels without severe contention. The time is approximately:

$$T = \frac{\text{bytes}}{B_{\text{eff}}} = \frac{140 \text{ GB}}{900 \text{ GB/s}} \approx 0.16 \text{ s}$$

**Option B — InfiniBand HDR across 8 nodes, one GPU per node.** HDR InfiniBand delivers 400 Gb/s per link  $\approx 50$  GB/s unidirectional after 8b/10b encoding and protocol headers. A ring-all-reduce or tree-all-reduce across 8 nodes will see lower effective bandwidth due to multi-hop forwarding and congestion. A realistic effective bandwidth is roughly 40 GB/s. The time is:

$$T = \frac{140 \text{ GB}}{40 \text{ GB/s}} \approx 3.5 \text{ s}$$

**Option C — 25 Gb/s Ethernet (RoCE2) across 8 nodes.** 25 Gb/s  $\approx 3.125$  GB/s unidirectional. With protocol overhead and multi-node tree contention, a practical effective bandwidth might be 2.5 GB/s. The time is:

$$T = \frac{140 \text{ GB}}{2.5 \text{ GB/s}} \approx 56 \text{ s}$$

These numbers illustrate the scale: moving from NVLink to InfiniBand adds  $\sim 22\times$  latency, and forcing the same all-reduce over Ethernet adds  $\sim 350\times$  latency. In a training job where 100 all-reduce steps run per iteration, the NVLink path adds  $\sim 16$  s of communication per iteration, the InfiniBand path adds  $\sim 350$  s, and the Ethernet path adds  $\sim 5,600$  s. The latter two would effectively serialize the training, turning a minutes-long job into an hours- or days-long one.

The worked example also highlights why model parallelism and pipeline parallelism exist: when a single GPU cannot hold the model, splitting it across nodes moves the communication from the inner rings to the outer rings, and the architect must budget the extra seconds per collective against the savings from fitting a larger model.

## 9.5. Measurement

Measuring communication in situ is harder than counting tokens, because the bandwidth numbers above are peaks; real systems see lower effective bandwidth due to congestion, operating-system interference, and protocol overhead. Three practical measurements an architect can make:

1. **Bandwidth benchmark per link.** Use `nccl-tests` or `ib_write_bw` / `bwperf` to measure achieved bandwidth on PCIe, NVLink (via NVIDIA's Nsight), RoCE, and InfiniBand. Record the unidirectional and bidirectional throughput for payload sizes matching the collective (typically 1 MB–1 GB). These numbers anchor the arithmetic to the real hardware.

2. **Collective latency and throughput.** Run `nccl-tests` all-reduce micro-benchmarks varying the number of GPUs (1, 2, 4, 8) and the tensor size (from 100 MB up to 140 GB). Plot time vs. size and extract the slope — that slope is the effective bandwidth the collective will see in production. Compare the slope across 1 node, 2 nodes, 4 nodes, to see where the interconnect ceiling is hit.
3. **End-to-end training iteration time.** Instrument the training loop to separate compute time (FLOPs) from communication time (all-reduce collectives). Many frameworks (PyTorch Distributed, vLLM, Deepspeed) already expose a per-iteration breakdown. The communication time per iteration, divided by the number of all-reduce steps, gives the average per-collective latency, which we can compare against the benchmark numbers above.

The key invariant: always measure on the same hardware and software stack that the production job will use. A benchmark run on a clean VM with no other traffic will overstate the achievable bandwidth compared to a crowded cluster.

## 9.6. Common Mistakes

- **Assuming peak bandwidth is sustained.** PCIe Gen5 x16 peaks at ~64 GB/s, but a single all-reduce never saturates the full link because the traffic is split across multiple GPUs and multiple links. Measured effective bandwidth is typically 40–60% of peak; design against the measured number, not the spec sheet peak.
- **Placing all-reduce on Ethernet without a bandwidth budget.** A 70B model gradient all-reduce at 25 Gb/s takes ~56 s per step. If the training runs 1,000 iterations, that is ~56,000 s (15.5 hours) spent exclusively in communication. Always compute the per-iteration communication cost before committing to a lower-bandwidth fabric.
- **Ignoring contention from other traffic.** NVSwitch can carry many concurrent channels, but if other collectives (e.g., all-gather for optimizer state) are running simultaneously, the bisection bandwidth is shared. Measure with the full production workload pattern, not an isolated micro-benchmark.
- **Using InfiniBand as a default without checking the cost–performance ratio.** InfiniBand HDR 400 Gb/s gives ~50 GB/s unidirectional, which is sufficient for many workloads, but the price per GB of optics and cables is 5–10× that of Ethernet. If the communication time is acceptable on RoCE2, the cheaper fabric may be the better architectural choice.
- **Overlooking protocol overhead.** RoCE and InfiniBand use different encoding (8b/10b vs. 128b/130b) and header sizes. A 400 Gb/s InfiniBand link delivers ~380 Gb/s useful payload; a 200 Gb/s RoCE link delivers ~180 Gb/s. The difference matters when we are computing per-second throughput for a fixed data volume.

## 9.7. Architecture Consequence

The interconnect choice is not a detail to be decided after the model, the parallelism strategy, and the hardware are selected — it is a first-class design variable that shapes all three. If the architecture calls for data-parallel training of a 70B FP16 model across 8 nodes, the all-reduce bandwidth dictates whether the nodes are connected by InfiniBand, 200 GbE, or even 100 GbE. If the budget can accommodate NVLink-connected hosts, keeping the all-reduce on-node reduces communication time by 2 orders of magnitude compared to any multi-node fabric.

The architecture consequence also flows in the other direction: a tighter communication budget (e.g., target <0.5 s per all-reduce) may force a particular parallelism strategy. We might choose model parallelism (pipeline parallelism, tensor parallelism) to keep the per-collective data volume smaller, or we might accept a slower fabric and reduce the frequency of collectives (e.g., gradient accumulation steps between all-reduce calls). The interconnect bandwidth number is the anchor that makes these trade-offs quantifiable.

Finally, the architecture consequence extends to cost. A cluster of  $8 \times \text{H100}$  servers connected by NVSwitch is the highest upfront capital expenditure, but it delivers the lowest per-iteration communication time, which translates to higher throughput per dollar over the life of the cluster. A cluster connected by 200 GbE RoCE has lower capital cost but higher per-iteration communication time, which may increase the total cost of ownership if the training job runs for many weeks. The architect must balance capex vs. opex by way of the communication arithmetic shown in this chapter.

## 9.8. What We Still Don't Know

- **Contention patterns under mixed workloads.** The benchmark numbers above assume a dedicated collective on a quiet fabric. In a production fleet, other traffic (user requests, other training jobs, KV-cache movement) shares the same links. The effective bandwidth under realistic multi-job congestion is not yet characterized with a standard suite, and the interaction between all-reduce and RoCE/Ethernet best-effort traffic can cause tail latency spikes that are difficult to predict.
- **Adaptive collective algorithms.** New collectives (e.g., gradient compression, sparse all-reduce, error-bounded approximations) claim to reduce the data volume at the cost of a small accuracy loss. The bandwidth arithmetic changes when the data volume is 30 GB instead of 140 GB, but the latency per byte may increase because of compression/decompression steps. The net effect on iteration time is workload-specific and not yet settled.
- **Optical interconnects and the next generation.** NVLink-Switch and XNAME are pushing node-level bandwidth higher (toward 1.8 TB/s aggregate), and optical intra-rack fabrics (400 Gb/s–800 Gb/s per lane) are entering data centres. The arithmetic base — the relative magnitudes of PCIe, NVLink, Ethernet, InfiniBand — is stable in the near term, but the absolute numbers will shift as new silicon and cable technologies ship. Keep the framework; update the constants.

- **Cross-domain generalization.** Most published bandwidth measurements use square matrices (all-reduce of a gradient tensor). Diagonal dominant tensors, sparse gradients, or structured compression may exhibit different effective bandwidths because the traffic pattern across links is less uniform. The arithmetic framework applies, but the constants need re-measurement for the specific tensor shape.

## 9.9. End-of-Chapter Mini-Case

An architect is asked to design a distributed inference fleet for a 70B parameter model serving user queries at 200 tokens/s per user, with 10,000 concurrent users. The team debates whether to use a single 8×H100 host with NVLink, or to shard the model across 32 nodes connected by 200 GbE RoCE. Before any hardware is ordered, the architect opens this chapter and runs the communication arithmetic.

The per-request all-reduce is not the right model for inference, and neither is a blanket “all-gather the full weights every decode step.” It matters which parallelism is in play: ordinary **tensor parallelism (TP)** shards each weight matrix across the GPUs, and every decode step exchanges only the *activations* (not the weights) via collective ops like all-reduce/all-gather of the small activation tensors — weight matrices stay resident and are read in place, they are not re-fetched each step. A different scheme, FSDP-style **weight gathering** (used in training, and in some sharded-inference arrangements) does all-gather weight shards before the forward pass. So the architect must name the scheme before doing the arithmetic. Concretely, for the canonical 8×H100 host the meaningful per-step data movement is a TP all-reduce of the per-layer activations (tens of MBs, not GBs) plus the KV-cache fragment exchange in disaggregated serving — e.g. ~35 GB of KV moved over NVLink (~0.04 s at 900 GB/s) when a P-D sequence transfers context between prefill and decode pools. The distinction changes the answer by orders of magnitude: re-fetching 140 GB of weights each step is pathological (that is a design smell, an anti-pattern to flag, not a baseline); exchanging activations or KV fragments is the normal case that fits comfortably within NVLink.

The architect proposes the single 8×H100 host (TP within the node, so intra-node NVLink carries the per-step activation all-reduce) for the core serving fleet, with a fallback to a 32-node RoCE cluster only if KV concurrency or memory capacity exceeds one host ([Chapter 17](#)). The communication arithmetic made the trade-off explicit: TP activation exchanges are tens of MB and free within NVLink; cross-node weight/KV shuttling is 35× more expensive and avoided unless the fleet genuinely spans hosts. The design is now grounded in bandwidth and the *specific* parallelism scheme, not in vague notions of “scale.”

## CHAPTER 10.

# Parallelism: Splitting the Model Across Machines

### 10.1. The Architect’s Question

After this chapter we should be able to answer the inevitable scaling question: *a single GPU cannot hold this model or keep up with this workload — how do we split the work across many GPUs, and which split is right?* We will build the arithmetic for why a model outgrows one GPU, then study the five parallelization strategies — tensor, pipeline, data, expert, and sequence parallelism — and develop a rule for choosing among them based on the model’s weight-residency, the workload’s batch shape, and the interconnect speed. After this chapter, “we scaled it out” becomes a precise claim about *what* was split and *why*.

### 10.2. Concept

The central idea: parallelism is **splitting one of the three things a model occupies — its weights, its layers, or its data — across multiple devices**. Before reading on, one framing note that prevents a common source of confusion: **the parallelism toolkit differs between training and inference, and “data parallelism” means different things in each mode**. In *training*, DP means replicated weights plus gradient all-reduce every step (plus FSDP, which shards and re-gathers weights/optimizer state); the communication load is at the gradient-sync point. In *inference*, DP simply means running independent model replicas to raise throughput — there is no gradient sync at all, because each request is served by one replica. There is also an inference-specific family — context/pipeline partitioning, expert parallelism, and P/D **disaggregation** (splitting the prefill stage from the decode stage into separate pools) — which has no training analogue. The rest of this section defines the strategies, and Section 3 applies them to inference (the book’s canonical serving context) unless a strategy is

explicitly training-only.

The strategy names tell us what gets split:

1. **Tensor parallelism (TP)** splits a single layer's weights across GPUs. The matrix multiply for one layer is partitioned so each GPU holds a slice of the weight matrix and the GPUs exchange partial results over NVLink every layer. TP is *intra-layer* and communication-hungry. (Inference: shards weights; per-step exchange is activations, not weights — see Ch. 9.)
2. **Pipeline parallelism (PP)** splits the model's layers (stages) across GPUs. GPU 1 runs layers 1–10, GPU 2 runs 11–20, and so on; activations flow forward, gradients flow backward. PP is *inter-layer*, communication-light, but idle bubbles appear at stage boundaries. (Inference: can split decode stages; bubbles hurt latency-sensitive serving more than training.)
3. **Data parallelism (DP)** — *different in the two modes*. Training: replicates all weights on every GPU, feeds each a different batch, all-reduces gradients after each step; communication-heavy only at the gradient-sync point, but each GPU must hold the full model. Inference: independent replicas each serve a disjoint set of requests; no gradient sync; throughput scales linearly until KV or memory binds (Ch. 17).
4. **Expert parallelism (EP)** shards the experts of a Mixture-of-Experts model across GPUs and routes each token to the GPU holding its experts — an all-to-all exchange per token. EP is MoE-specific. (Used in both training and inference for MoE models.)
5. **Sequence/context parallelism (CP)** splits the sequence dimension across GPUs, useful for very long contexts where the KV cache or attention activations exceed one GPU. (Inference-centric; the serving analogue of training sequence parallelism.)
6. **P/D disaggregation (inference)** — separates the compute-bound prefill stage from the bandwidth-bound decode stage onto different GPU pools so each can be sized and scaled independently (Ch. 11). This has no training counterpart and is a serving-only strategy.

These are not alternatives; production systems **compose** them (e.g. DP across nodes  $\times$  TP within a node  $\times$  PP across a few nodes; or TP+CP within a decode pool + a separate prefill pool). The architect's job is deciding which dimension to split first, and in which mode (training vs inference) the parallelism is being applied.

*Fig 10.1 — The strategies are not either/or. Left: a DP  $\times$  TP  $\times$  PP stack — the layer stack is cut into PP stages (green), each stage is DP-replicated (orange) across nodes, and within a stage TP shards the weights across 4 GPUs (blue). Right: expert parallel is a per-token all-to-all — each token may leave its home GPU to reach the GPU holding its top-k experts. Which dimension binds determines which you split first.*




FIGURE 10.1. Fig 10.1 — Composing the dimensions: each PP stage holds a DP replica, TP-sharded across GPUs; and EP's per-token all-to-all routes each token to the GPU holding its experts (ILLUSTRATIVE)

### 10.3. Mental Model

Think of each strategy by *what it trades bandwidth for*. Tensor parallelism trades a lot of NVLink bandwidth to avoid duplicating weights but adds per-layer sync. Pipeline parallelism trades low bandwidth but accepts bubbles as the price of not holding all layers on one GPU. Data parallelism trades gradient bandwidth (once per step) for the convenience of replicated weights and is limited by having to fit the whole model on each GPU.

The useful mental model is a **resource triangle**: we have weights (memory), activations (work), and data (batch). Each strategy moves one corner off a single GPU at the cost of communication. When a GPU runs out of *memory*, we reach for TP/PP (weight splitting). When a GPU is *idle* waiting, we reach for DP/CP (more work per GPU). When the model is *too big to replicate*, TP wins over DP. The binding constraint chooses the strategy.

*Table 10-1 — Parallelism strategies at a glance*

| Strategy      | What it splits                         | Communication need              | Fits when                           |
|---------------|----------------------------------------|---------------------------------|-------------------------------------|
| Tensor (TP)   | a layer's weights within a GPU group   | high — every layer, over NVLink | model exceeds one GPU, fits a node  |
| Pipeline (PP) | the layer stack into stages            | low — stage boundaries only     | model spans nodes                   |
| Data (DP)     | the batch; replicates weights          | gradient all-reduce per step    | model fits one GPU, need throughput |
| Expert (EP)   | MoE experts; route-by-token all-to-all | high, per token                 | MoE models                          |
| Sequence (CP) | the sequence/context dimension         | medium                          | long contexts, KV too big           |

*(Strategies compose; the table names the primary cost each trades.)*

### 10.4. Worked Example

#### 10.4.1. Why the Canonical 70B Doesn't Fit One H100

The canonical model is 70B FP16 dense = 140 GB of weights [1P] (reality check:  $70 \times 10^9$  params  $\times$  2 bytes). A single H100 has 80 GB of HBM [2° FACT]. The arithmetic is immediate:

- The weights alone don't fit on one H100:

$$140 \text{ GB weights} > 80 \text{ GB HBM}$$

We need at least  $140/80 \approx 1.75 \rightarrow \geq 2$  GPUs **just to hold the weights** (and that's before KV cache, which adds  $\sim 2.5$  MB/token at FP16 — the 9.2K-token canonical prompt needs  $\sim 24$  GB more). [2° DERIVED]

The canonical  $8 \times \text{H100}$  host (640 GB total) fits the weights + KV comfortably, which is why the single-host baseline in earlier chapters works. But if the model grows (say to 400B) or owns more, the split becomes mandatory:

#### 10.4.2. Tensor Parallelism: Splitting Weights Across GPUs

With tensor parallelism, the 140 GB weight is split across GPUs. On  $2 \times \text{H100}$ :  $140/2 = 70$  GB per GPU, each now under the 80 GB ceiling. But every transformer layer's matmuls now require an **all-reduce of the layer's output across the TP group** — a per-token, per-layer exchange over NVLink. For the canonical 70B with  $\sim 80$  layers, that is  $\sim 80 \times \text{tokens} \times \text{output-size}$  of communication per forward pass. The all-reduce communication cost per layer, per token is  $T_{\text{tp}} = \text{tokens} \times n_{\text{layers}} \times \text{output-size} / B_{\text{eff}}$ . TP is only viable where the interconnect is very fast: on-node NVLink (900 GB/s) keeps the sync cheap; over slow Ethernet it would drown. [2° FACT][2° DERIVED]

#### 10.4.3. Pipeline Parallelism: Splitting Layers, Accepting Bubbles

Pipeline parallelism puts stage  $P$  on GPU  $P$ . The cost is the **pipeline bubble**: while the first micro-batch fills the pipeline and the last one drains, stages sit idle. With  $P$  stages and  $m$  micro-batches, the steady-state **ideal speedup** over a single serial pass is:

$$\text{ideal speedup} = \frac{m \cdot P}{m + P - 1}$$

The  $P-1$  term in the denominator is the bubble: the pipeline fills and drains exactly once per  $m$ -batch pass, so  $P-1$  stage-slots are wasted out of the  $m \cdot P$  slots each stage would otherwise work. With  $P = 4$  stages and  $m = 8$  micro-batches, ideal utilization is

$$\frac{8 \cdot 4}{8 + 4 - 1} = \frac{32}{11} \approx 2.9 \times \quad (\text{out of } 4 \times)$$

— a meaningful bubble. Fewer stages or more micro-batches shrink it: with  $m = 32$ ,

$$\frac{32 \cdot 4}{32 + 4 - 1} = \frac{128}{35} \approx 3.66 \times .$$

The lesson: PP is communication-light (each stage only moves activation boundaries across the fabric) but always pays a bubble that grows with stage count. [2° DERIVED]

#### 10.4.4. Data Parallelism: Replicating, Then Syncing Gradients

Data parallelism replicates the 140 GB model on every GPU. For the gradient all-reduce: if the model has  $140 \text{ GB} \times 2$  (weights + gradients) and we sync after each step across  $G$  GPUs, each GPU sends the full gradient tensor once per step. The gradient bytes are

$$\text{grad bytes} = 2 \times N \times \text{bytes/param} = 2 \times 70 \times 10^9 \times 4 \text{ B} \approx 280 \text{ GB}$$

per step (BF16), and the sync time is  $T = \text{bytes}/B_{\text{eff}}$ . At 25 GB/s (single 200Gb/s link) that is  $280 / 25 \approx 11 \text{ s}$  per step — far too slow if the step takes  $\sim 1 \text{ s}$ . At node NVLink 900 GB/s it is  $280/900 \approx 0.31 \text{ s}$ . The sync cost, not compute, often caps DP at small scale unless the interconnect is fast. [2° DERIVED]

### 10.5. Measurement

Four habits anchor the architect:

1. **Measure weight residency before choosing a split.** `bytes = params × bytes_per_param` lands first; it decides whether TP/PP (memory) or DP (replication) is even possible.
2. **Measure interconnect bandwidth where it matters.** H100 NVLink  $\sim 900 \text{ GB/s}$  per GPU is the on-node ceiling; cluster Ethernet/InfiniBand is 1–2 orders slower. Quote the number at the sync/frequency it will actually run.
3. **Measure the bubble and the step time.** For PP, log the fraction of idle GPU time (bubble) vs compute. For DP, time the gradient all-reduce as a fraction of step time — if it exceeds  $\sim 20\%$ , it caps scaling.
4. **Measure communication overlap.** The best systems overlap all-reduce with the next forward/backward compute. Log the non-overlapped comm time — that is the true cost, not the raw bytes.

### 10.6. Common Mistakes

- **Splitting with DP when the model doesn't fit.** DP replicates weights; if  $140 \text{ GB} > 80 \text{ GB}$ , one DP replica doesn't even fit one GPU. Reach for TP/PP first.
- **Using TP over a slow interconnect.** TP syncs every layer; over 25 GB/s cluster fabric it bottleneck-cycles. TP belongs behind NVLink.
- **Ignoring the pipeline bubble.** PP's idle stages waste real utilization; the bubble grows with stage count. It is not free.
- **Quoting peak interconnect spec as sustained.** NVLink 900 GB/s is aggregate peak; real all-reduce achieves a fraction. Budget headroom.

- **Treating parallelism as an either/or.** Real systems compose DP×TP×PP; the question is which dimension binds, not which one to use.

### 10.7. Architecture Consequence

The parallelism strategy is forced by the binding constraint, and that constraint is usually **memory** for large models:

- **A model that exceeds one GPU's memory** must split weights → TP (within a node, on NVLink) or PP (across nodes, low comm). Which one depends on whether the model fits a node: fits one node → TP among its GPUs; spans nodes → PP across them, with TP inside each node.
- **A workload with high throughput demand** (like the canonical ~10 rps / 40 rps) that fits one replicable model → DP to scale batch across GPUs, accepting the gradient-sync cost only if it trains/fine-tunes. For pure inference the canonical 70B fits one 8×H100 node, so parallelism is often unnecessary at the single-host baseline.
- **MoE models** (if chosen per Ch3) push toward EP, because only the active experts need residency per token.
- **The canonical answer:** for the book's ~70B/8×H100 scenario, no parallelism is required at the baseline; parallelism becomes the tool when the model or workload outgrows one node. The architect escalates through TP (in-node) → PP (cross-node) → DP (replicating a fitting model) in that order.

*Fig 10.2 — Tensor / pipeline / data / expert / sequence parallelism, each drawn as its own object — weight matrix, layer stack, replicated model, MoE experts, context window — splitting into GPU shards beneath it.*

### 10.8. What We Still Don't Know

- **Optimal split for very large MoE or long-context models** is workload- and interconnect-dependent; the exact NVLink-vs-fabric crossover for a given model is [VERIFY].
- **2026 MoE scaling has moved the specialist/EP frontier.** Kimi K3 activates just 16 of 896 routed experts per token (2.8T total / 104B active) via its Stable LatentMoE [1P: arXiv 2607.24653]; Qwen3.8-Flash-Next routes 10+1 of 512 experts (125B total / 6B active) [1P: HF Qwen/Qwen3.8-Flash-Next]. These extreme sparsities (16/896 ≈ 1.8% expert activation) are why expert-parallel and fused cross-expert comm kernels (e.g. MoE expert parallelism with fused compute/communication, as in DeepSeek's DeepEP/TileLang infra) are increasingly central to serving large MoE fleets [1P: arXiv 2606.19348]. For the architect this changes the serving calculus: a frontier MoE's activated footprint (Kimi 104B, Qwen 6B) can be dramatically smaller than its total size, so the memory-bound EP decision must be made against *active* parameters,




FIGURE 10.2. Fig 10.2 — The five parallelization strategies and what each splits [ILLUSTRATIVE conceptual]

not total.

- **Communication-overlap efficiency** in real kernels is hard to predict from spec; sustained all-reduce throughput is [HYPOTHESIS] until measured on the target cluster.
- **The right parallelism for mixed prefill/decode (P/D disaggregated) serving** — where the prefill pool and decode pool may want different splits — is an active research area and [HYPOTHESIS] for this book.

### 10.9. End-of-Chapter Mini-Case

An architect inherits a deployment where the internal Q&A model (from Ch1’s mini-case) has been fine-tuned and grown to a 400B-parameter MoE. The team reports “parallelism is killing us — everything is slow.” The architect applies this chapter’s lens. First, weight residency: a 400B model needs  $\sim 400\text{e}9 \times 2 \text{ bytes} = 800 \text{ GB}$  even before KV — it spans  $\sim 11 \times H100$  at minimum, far beyond one node. It does not fit a node, so pure TP can’t span it; PP across nodes is required. Within each node, the MoE active weights are small, so the team uses EP to shard experts + TP inside the node. The architect measures the interconnect: the cross-node PP boundary runs over the cluster fabric (not NVLink), so they minimize the number of PP stages to shrink the bubble, and they verify the DP gradient-sync (if fine-tuning) stays under  $\sim 20\%$  of step time. The result: a composed PP  $\times$  EP  $\times$  TP configuration whose binding constraints (memory across nodes + expert routing) are each addressed by the strategy that trades the right bandwidth. The team’s “parallelism is slow” becomes a precise, diagnosed, and fixed split.

## CHAPTER 11.

# Serving: From Model in Memory to Requests in Flight

### 11.1. The Architect’s Question

After this chapter we should be able to answer the operational half of deployment: *the model is in memory and the requests are arriving at ~10 per second — how do we actually serve them without violating our latency SLO?* We will develop the serving stack that turns a static model into a live, multi-request system: how requests are batched, how KV-cache memory is managed, how repeated prefixes are reused, and when to split prefill from decode onto separate pools. After this chapter, serving is not a black box — it is a sequence of explicit architectural choices, each with a measurable consequence.

### 11.2. Concept

The serving stack is a small set of mechanisms, invented roughly in order:

1. **Batching** — running multiple requests through the model together to amortize the weight reads. The key modern advance is **continuous (iteration-level) batching** introduced by Orca (OSDI 2022): instead of waiting for an entire batch to finish, the scheduler admits and evicts requests at every decoding step, filling the GPU’s idle slots as soon as they free up. [S1][1P]
2. **PagedAttention / vLLM (SOSP 2023)** — managing the KV cache in fixed-size blocks (like a virtual-memory page table) so that different requests’ KV caches no longer need contiguous memory, slashing fragmentation and letting the serving engine pack far more concurrent requests. [S2][1P]

3. **Prefix caching** — recognizing that many requests share context (e.g. the same retrieved documents), and caching the KV of the common prefix so a new request does not re-prefill it. **Radix-Attention (SGLang)** and **Automatic Prefix Caching (vLLM APC)** are the canonical implementations. [S3][1P]
4. **P/D disaggregation** — splitting the compute-bound *prefill* stage from the bandwidth-bound *decode* stage onto separate pools, because they bottleneck on different resources. Pioneered by **Splitwise (ISCA 2024)** and **DistServe (OSDI 2024)**; **Mooncake** takes the KV-centric view, transferring KV caches from the prefill pool to the decode pool (with CPU/DRAM/SSD offload) and reports up to **525% throughput lift** in long-context simulation and **75% more requests** on a real Kimi workload. [S4][1P]

These mechanisms compose: continuous batching is the baseline, PagedAttention makes it memory-efficient, prefix caching makes it cheaper, and P/D disaggregation makes it scale.

### 11.3. Mental Model

Think of serving as **managing three precious resources against a stream of arrivals**: GPU compute (FLOPs, bound in prefill), GPU memory bandwidth (HBM, bound in decode), and GPU memory capacity (HBM bytes, bound by KV cache). Each mechanism trades one of these:

- **Batching** improves compute and bandwidth utilization by amortizing the fixed per-step weight reads across more requests.
- **PagedAttention** improves memory capacity by eliminating KV fragmentation.
- **Prefix caching** reduces memory capacity and prefill compute by reusing already-computed KV.
- **P/D disaggregation** lets the prefill pool run at high compute utilization and the decode pool at high bandwidth utilization simultaneously, instead of one pool being starved.

The mental model is a **resource budget**: a request arrives, consumes some FLOPs (prefill), some bandwidth (decode), and some memory (its KV cache) for its lifetime. Serving architecture is about sizing and scheduling those three budgets so the SLO holds across the burst.

### 11.4. Worked Example

#### 11.4.1. The Canonical Workload in a Serving System

Recall the canonical enterprise-Q&A RAG workload (§14): ~10 rps average, ~40 rps peak; each request carries ~9,200 input tokens and ~300 output tokens; TTFT budget 1.2 s, TPOT ~25 ms/token. [1P §14]

### 11.4.2. Discrete Batching Wastes the GPU

Imagine naive **discrete (static) batching**: the server accumulates requests until a batch fills, then runs them to completion, then starts the next batch. Problems compound at this workload's 9.2K-input shape:

- Requests arrive in *bursts* (recall the ~40 rps peak). Discrete batching forces early-arriving requests to **wait** for the batch to fill, inflating TTFT toward the 2 s SLO breach.
- The batch is locked until its *slowest* member finishes (fits-all-in-one-run), so a short 100-token output tails behind the longest while the GPU finishes the whole batch — idle bubbles at the end of every batch.
- Batch boundaries create periods where the GPU drains (no work) then floods — utilization is spiky around the mean. [2° DERIVED]

### 11.4.3. Continuous Batching Fills the Bubbles

**Continuous batching** (Orca [S1][1P]) eliminates the batch boundary: at every decode step the scheduler accepts any new prefill-ready request and evicts any finished one, so decode slots never sit idle. In the canonical stream, a 300-token decode at ~25 ms/token occupies a slot for ~7.5 s; while a long decode runs, the freed FLOPs are continuously refilled by new short requests. The utilization gain is the difference between spiky discrete batches and a continuously-fed pipeline. Rather than stipulate a single percent, the mechanism makes the GPU the steady bottleneck instead of the scheduler. [2° DERIVED]

*Fig 11.1 — The heart of the chapter's serving story made visual. Discrete (static) batching lets slots sit idle whenever a sequence finishes early inside the batch; continuous (iteration-level) batching backfills a freed slot immediately, so the GPU — not the scheduler — is the steady bottleneck. This is the difference between spiky, latency-wasting batches and a continuously-fed decode pipeline.*

### 11.4.4. PagedAttention: Killing the Fragmentation Tax

In re-packaged (naive) serving, each request's KV cache for 9.2K input is ~24 GB (FP16, ~2.5 MB/token × 9,200) [2° DERIVED]. If served contiguously, variable-length completions leave fragmented, unusable holes — exactly like a fragmented heap. **PagedAttention/vLLM [S2][1P]** allocates KV in fixed blocks shared and evicted like page frames, so the ~24 GB per request is packed densely and more concurrent requests fit in the same 640 GB node [2° DERIVED]. The win is *more concurrency under the same SLO*, not faster single-request math.

### 11.4.5. Prefix Caching: Don't Re-Prefill the Same Documents

The RAG request's 8K retrieved-context tokens are largely shared across queries (same documents). Without caching, every request re-runs the ~1.29 PFLOP prefill for those 8K tokens [2° DERIVED].




FIGURE 11.1. Fig 11.1 — Discrete vs continuous batching (Orca-style slot diagram). Discrete batching runs each batch to completion, leaving GPU slots idle as sequences finish early; continuous batching admits and evicts requests at every decode step, backfilling freed slots (ILLUSTRATIVE, after Orca, OSDI 2022)

**RadixAttention / APC [S3][1P]** caches the KV of common prefixes, so a new request reuses the shared-document KV and prefills only the unique query tail — cutting the dominant prefill FLOPs and TTFT for repeated-context traffic [2° DERIVED].

#### 11.4.6. When to Disaggregate: Prefill and Decode Want Different Pools

Recall from Ch2/Ch6: prefill is **compute-bound** ( $\sim 1.29$  PFLOP  $\rightarrow \sim 1.19$  PFLOPS required vs 0.989 PFLOPS peak), decode is **bandwidth-bound** (5.6 TB/s demand vs 3.35 TB/s) [2° DERIVED]. Serving both from one pool forces a compromise: batch for prefill and it slows decode; optimize for decode and prefill starves. **P/D disaggregation (Splitwise/DistServe/Mooncake [S4][1P])** runs prefill nodes at high FLOP-utilization and decode nodes at high bandwidth-utilization, transferring KV across via the fabric (**Fig 11.3** shows the two-pool topology and what moves between them). Mooncake’s reported >525% long-context throughput and +75% request count are the quantitative anchor [S4][1P]. For the canonical single-host ( $\sim 10$  rps) this disaggregation is premature; it pays off only when the workload outgrows a node and the two phases’ resource conflicts become binding.

*Table 11-1 — Serving techniques at a glance*

| Technique                           | Mechanism                                | Primary resource          |                          |         |
|-------------------------------------|------------------------------------------|---------------------------|--------------------------|---------|
|                                     |                                          | traded                    | When it matters          | Source  |
| Discrete batching                   | Lock batch to completion                 | — (wastes)                | baseline only            | [1P]    |
| Continuous batching (Orca)          | Admit/evict at each decode step          | compute+bandwidth<br>util | high concurrency         | S1 [1P] |
| PagedAttention / vLLM               | Block-based KV pages                     | memory capacity           | many concurrent requests | S2 [1P] |
| Prefix caching (RadixAttention/APC) | Cache shared-prefix KV                   | prefill FLOPs + memory    | repeated-context RAG     | S3 [1P] |
| P/D disaggregation                  | Split prefill/decode pools + KV transfer | both at scale             | workload outgrows a node | S4 [1P] |

### 11.5. Measurement

Good serving telemetry answers “are we meeting the SLO while keeping the GPU busy?”:

1. **Measure goodput, not raw throughput** — tokens/s that arrive before their SLO deadline (Ch6). Continuous batching is only a win if it raises goodput.
2. **Measure KV utilization** — fraction of the 640 GB node memory consumed by live KV vs weights. PagedAttention and prefix caching both raise the effective KV packing; track it.
3. **Measure prefill-vs-decode utilization separately** — a FLOP-bound prefill core and a bandwidth-bound decode core running on one pool tells us when to disaggregate.
4. **Measure the scheduler's real occupancy** — idle decode slots, queue wait, and the tail of TTFT under load, not the average.

### 11.6. Common Mistakes

- **Assuming bigger batches always help.** Bigger discrete batches raise raw tokens/s but push TTFT past the SLO, zeroing goodput (Ch6). Continuous batching exists precisely to uncouple batch size from wait.
- **Serving KV contiguous and watching fragmentation.** PagedAttention removes the fragmentation tax; skipping it wastes memory under concurrency.
- **Re-prefilling shared context.** Without prefix caching, repeated RAG documents are recomputed per request — the single most wasteful failure for this workload.
- **Choosing P/D disaggregation too early.** Splitting pools adds fabric + orchestration; for a single host it is overhead without the resource conflict that motivates it.
- **Believing speculative decoding (DFlash  $> 4.3 \times / 1.5 \times [S5][1P]$ ) is a serving baseline.** It is an optimization for bandwidth-bound decode, not a requirement, and must be validated on the workload.

### 11.7. Architecture Consequence

The serving stack dictates the deployment choices:

- **Baseline (canonical single-host):** continuous batching (vLLM/Orca) is near-mandatory to hold TTFT under 10 rps input-heavy load; PagedAttention ensures the 24 GB-per-request KV pack fits; prefix caching is the highest-leverage optimization because the workload is repeated-context RAG. This matches the memory and compute conclusions of Ch4/Ch7 for the  $\sim 10$  rps *average* baseline; at the  $\sim 40$  rps peak the KV budget forces replication, which [Chapter 17](#) sizes.  $[S1][S2][S3][1P]$
- **Scaling up:** those three carry the workload far before disaggregation is justified. FP8 KV ( $\sim 54\%$  of BF16  $[S6][1P]$ ) and KV offloading are memory levers available without restructuring.
- **Beyond a node:** when prefill FLOPs and decode bandwidth fight on one pool, P/D disaggrega-




FIGURE 11.2. Fig 11.2 — The serving stack: batching, KV management, prefix caching, P/D split [ILLUSTRATIVE conceptual]

tion (Splitwise/DistServe/Mooncake) is the architecture of choice, with Mooncake’s KV-transfer engine moving caches between pools. [S4][1P]

*Fig 11.2 — From request arrival to token output: continuous batching, PagedAttention KV, prefix caching, and optional P/D disaggregation.*

*Fig 11.3 — P/D disaggregation topology. Left: prefill pool, compute-bound (FLOPs), handles the massive prompt at once. Right: decode pool, bandwidth-bound (HBM), reads steady token generation. A fabric bridge moves the per-token KV cache from prefill to decode. The split exists because the pools want opposite resources — prefill is FLOP-starved ( $\sim 1.19$  PFLOPS required vs  $0.989$  peak), decode is bandwidth-starved ( $5.6$  TB/s demand vs  $3.35$  TB/s) [2° DERIVED].*


## 11.8. What We Still Don’t Know

- **Portable end-to-end goodput across real enterprise RAG** — Mooncake’s  $525\%/75\%$  [S4][1P] and DFlash’s  $4.3\times$  [S5][1P] are strong anchors but are workload-specific; general rules are [HYPOTHESIS].
- **Prefix-cache hit rates in mixed traffic** — the real reduction in prefill FLOPs from RadixAttention/APC for a given document-distribution is [VERIFY].
- **KV-transfer cost at scale** — Mooncake’s engine moves KV to CPU/DRAM/SSD [S4][1P], but the exact latency/bandwidth crossover where disaggregation wins is [VERIFY].

## 11.9. End-of-Chapter Mini-Case

An architect runs the canonical internal Q&A service on a single  $8\times H100$  host and sees p95 TTFT creep toward  $2.2$  s against a  $2$  s SLO as traffic peaks at 40 rps. Applying this chapter: the architect first confirms continuous batching is on (vLLM) and that PagedAttention is active — otherwise the 24 GB-per-request KV for 9.2K inputs fragments and concurrency collapses. The biggest lever: prefix

### P/D Disaggregation: the two pools want opposite resources



~1.19 PFLOPS required vs 0.989 peak (FLOP-starved); decode 5.6 TB/s vs 3.35 TB/s (bandwidth-starved). One pool forces a compromise for b

FIGURE 11.3. Fig 11.3 — P/D disaggregation topology: prefill pool (compute-bound) and decode pool (bandwidth-bound) bridged by KV transfer [2° DERIVED]

caching. The RAG traffic reuses the same retrieved documents, so the architect enables automatic prefix caching and watches TTFT drop as the 8K-context prefill is replaced by cached KV lookups. The architect does **not** reach for P/D disaggregation or speculative decoding — the single host is resource-sufficient once the wasted re-prefill is eliminated. The result: p95 TTFT back under 1.8 s at 40 rps, goodput restored, with only configuration changes — because the architecture followed the serving mechanism that matched the workload’s repeated-context shape.

## **PART IV.**

# **The Architecture**

## CHAPTER 12.

# Designing the AI System: From Constraints to Candidates

### 12.1. The Architect's Question

Designing an AI system means producing candidate architectures from the substrate — memory, compute, communication, and parallelism — rather than selecting a canned stack. This chapter works through a synthesis method: given a characterized workload from the six-dimension framework (Ch. 4), generate two or three candidate architectures that each satisfy the hard constraints (memory fit, latency budget, economics), then carry them forward for evaluation in Chs. 14–16. The core thesis is that design is constraint-driven generation, not menu-picking. This distinguishes our approach from the abstract meta-method presented in Ch. 22, which operates at a level of generality that does not instantiate against concrete substrate constraints. In the language of the book's spine — the decision loop *constrain* → *synthesize* → *measure* → *decide* → *refine* — this chapter is the **synthesize** step: we turn a constraint-characterized workload into a bounded set of candidate architectures that then enter the measure/decide stages in Chs. 14–16.

### 12.2. Concept

The architect's first step is to reverse-engineer the workload into its constituent constraints. From the six dimensions — quality requirements, traffic profile, token profile, latency budget, economic ceiling, and operational constraints — we extract the three hard constraints that any candidate must satisfy: (1) the total weight memory must fit the target residency budget, (2) the KV-cache footprint plus weight memory must fit the host's aggregate memory, and (3) the total request latency (prefill + decode) must stay within the budget. Any architecture that violates any one of these is eliminated

before evaluation. This constraint-first posture inverts the typical practice of browsing vendor offerings and then hoping they fit; instead, we generate architectures that we know can physically reside and operate within the specified substrate.

### 12.3. Mental Model

Think of the substrate as a set of bounded resources and the architecture design problem as a packing exercise. The resources are: GPU memory (weight residency + KV cache), memory bandwidth (prefill throughput), inter-GPU communication volume (for parallelism schemes), and the latency budget split between prefill and decode. A candidate architecture is a valid packing of the workload's weight memory and KV-cache into the available GPU memory, paired with a parallelism and caching strategy that delivers the required tokens-per-second throughput within the latency budget. The architect alternates between squeezing the residency footprint (e.g., via KV quantization or precision reduction) and adjusting the parallelism layout (e.g., parameter vs. data vs. pipeline) until all three hard constraints are simultaneously satisfiable. This mental model — substrate as packing problem — is what enables systematic candidate generation rather than ad hoc selection.

### 12.4. Worked Example

We generate three candidate architectures for the canonical enterprise-Q&A RAG workload: ~10 requests/s average, peak 40 rps, 70B FP16 model, 9.2K token input (1,200 prompt + 8K retrieved context), 300-token output. The canonical weight memory is 140 GB (70B × 2 bytes FP16). The KV-cache for 9.2K context at FP16 precision ([1P §14]) is approximately 24 GB ( $2 \times \text{num\_layers} \times \text{hidden\_dim} \times 9,200 \times 2 \text{ bytes} / \text{scale factor}$ ), giving a total residency of 164 GB. A single 8×H100 host provides 640 GB aggregate memory (8 × 80 GB), leaving ample headroom. Each candidate satisfies the three hard constraints differently.

#### Candidate (a): single 8×H100 host with prefix caching (Ch. 11).

All 70B weights and the full 24 GB KV cache reside on the 8-GPU host. Prefix caching from Ch. 11 amortizes the prefill attention over repeated requests with the same prompt prefix, reducing effective prefill FLOPs by ~30% at 10 rps average and ~50% at peak 40 rps when many requests share common document roots. The 8×H100 host fits the full 164 GB residency easily, but on *throughput* we must be honest with the [Chapter 8](#) arithmetic: a single node sustains only ~2,471 prefill-tokens/s, far below the ~92K input tokens/s the workload needs on average — prefill is compute-bound here, and meeting it is exactly the scaling problem Chapters 15–17 take up. Memory fit and throughput fit are different gates: the memory gate passes with one host; the throughput gate does not. Memory fit: 140 GB weights + 24 GB KV = 164 GB < 640 GB aggregate ([1P §14]). Latency budget: prefill + decode stay within budget thanks to prefix caching and the high bandwidth of 8×H100. Economics: one host purchase, one software license tier.

#### Candidate (b): P/D-disaggregated two-pool.

We keep the full 70B model in both pools and split by *phase* — which is what prefill/decode (P/D) disaggregation actually means (Chapter 11): a compute-bound prefill pool and a bandwidth-bound decode pool, each holding the complete 140 GB weights plus the in-flight KV, with the KV hidden-state handed between them (Splitwise / DistServe / Mooncake). Each pool therefore needs its own 8×H100 host (residency ~164 GB — weights plus one request’s KV — comfortably under 640 GB). The win is that prefill and decode no longer fight over the same host’s resources: the prefill pool runs FLOPS-saturated on the 9.2K input, the decode pool runs bandwidth-saturated on the 300 output tokens, and KV crosses between them. The cost is a second host plus the KV-transfer fabric and orchestration, so this pays only once the workload actually outgrows a single host (Chapter 11’s “when to disaggregate”). Economics: two hosts and added fabric; it is the right architecture at scale, not for the single-host canonical case.

#### Candidate (c): KV-quantized smaller host.

We downgrade to a 7B-class model ( $7B \times 2 \text{ bytes} = 14 \text{ GB weights}$ ) with KV-cache quantized to 4-bit, yielding a KV footprint of ~6 GB for 9.2K context. Total residency: 14 GB + 6 GB = 20 GB, which fits on a single H100 (80 GB) with 60 GB of headroom for overhead. The trade-off is a significant quality drop: the 7B model cannot achieve the same answer quality on complex enterprise Q&A as the 70B model, even with retrieval augmentation. To compensate, we apply KV-quantization at 4-bit, which recovers most of the lost quality relative to unquantized KV, but the model’s intrinsic capacity remains the binding constraint. Throughput on a single H100 is far higher than on the 70B — a 7B prefill costs only ~14 GFLOP/token, so one card sustains tens of thousands of prefill tokens/s (vs ~2,471 for the 70B node, Chapter 8) — yet even that cannot change the verdict: at peak 40 rps the workload needs ~368K input tokens/s, so satisfying peak prefill still needs more than a handful of cards, and the 7B’s quality ceiling on complex enterprise Q&A remains the binding constraint regardless of throughput. This candidate illustrates the extreme end of the KV-quantization space: the memory footprint is trivially small, but neither quality nor throughput makes it competitive for this workload.

*Memory/residency arithmetic summary.* The canonical weights (140 GB) + KV (24 GB FP16) = 164 GB fits 640 GB (8×H100). Candidate (a) uses the full host with room to spare. Candidate (b) keeps the full model in each pool, so each 8×H100 host also holds ~164 GB. Candidate (c) fits 20 GB on one H100, but its quality ceiling rules it out for the quality target.

### 12.5. Measurement

To validate a candidate architecture, the architect must derive three numbers from first principles and verify them against the substrate: 1. **Weight memory.**

$$W = N \times \text{bytes-per-param}$$

where  $N$  is the parameter count and bytes-per-param is 2 for FP16 (replace 2 with bit-width/8

for quantization). Derived fact [2° DERIVED]:  $70 \times 10^9 \times 2 \text{ B} = 140 \text{ GB}$ . 2. **KV-cache footprint.**

$$KV = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times L \times \text{bytes-per-activation}$$

where  $L$  is the context length. For a 70B model with 8-bit KV caching, this is approximately 12 GB for 9.2K context; at FP16 precision it doubles to ~24 GB [2° DERIVED]. The architect should compute this for the exact context length of the workload, not assume a fixed value. 3. **Throughput per device.** Tokens-per-second that a single device can sustain for prefill and decode, measured on the target hardware. This is the number that, multiplied by the number of devices in the candidate layout, must exceed the workload’s tokens-per-second demand (92K input tokens/s average, 3K output tokens/s average, 368K input tokens/s peak, 12K output tokens/s peak). Vendors publish peak numbers; the architect should treat these as [VERIFY] benchmarks and measure on equivalent hardware when possible.

All three numbers must be confirmed before a candidate is carried forward ([1P §14]). If any one fails, the candidate is returned to the synthesis step.

## 12.6. Common Mistakes

- **Menu-picking without constraint validation.** Browsing vendor offerings and selecting the one that “feels right” rather than generating candidates against quantified hard constraints. This is the opposite of the constraint-driven method described here.
- **Ignoring the KV-cache residency arithmetic.** Assuming that the weight memory is the only memory budget item. In input-heavy workloads ([1P §14]: 9.2K average context, 30× more input than output), the KV cache can be 15–20% of total residency and is often the binding constraint.
- **Over-optimizing prefix caching.** Assuming that prefix caching can amortize prefill across all requests. In practice, cache effectiveness depends on the overlap of prompt prefixes across the traffic workload; sparse prefix sharing reduces the realized savings.
- **Using KV-quantization as a free lunch.** Quantizing the KV cache reduces memory footprint but introduces quality degradation that may be unacceptable for the target workload. The architect must measure the quality impact before committing to a quantized KV layout.
- **Forgetting the latency split.** Treating the latency budget as a single number rather than a prefill/decode split. A candidate may meet the total latency budget by trading prefill for decode or vice versa, which may not satisfy the service-level objectives for either leg.

## 12.7. Architecture Consequence

The three candidates carry forward into the evaluation chapters (Chs. 14–16) with distinct test plans. Candidate (a), the single 8×H100 host with prefix caching, is the baseline: we evaluate throughput, cost-per-token, and latency SLO compliance on a single hardware configuration. Candidate

(b), the P/D-disaggregated two-pool, is evaluated on the added complexity of inter-group communication and the pipeline overlap efficacy; we measure whether the latency budget is maintained under the partitioned layout. Candidate (c), the KV-quantized 7B model, is evaluated on the quality-degradation trade-off; we compare answer quality against the 70B baseline and compute the economic break-even point (number of hosts required) against the quality loss. The synthesis method produces these candidates as a set from which the evaluation chapter selects based on quantitative results, not on gut feeling.

## 12.8. What We Still Don't Know

As of 2026-08, several questions remain open and are flagged [VERIFY]: - **KV-quantization quality impact at 4-bit for enterprise Q&A.** The quality loss from 4-bit KV quantization has not been systematically characterized across the diversity of enterprise prompts in the canonical workload. [VERIFY HYPOTHESIS] - **Prefix-caching effectiveness under realistic traffic overlap.** The assumed 30–50% prefill amortization rests on the premise that many requests share document-level prefixes. Empirical measurement of prefix overlap across a production enterprise traffic trace is needed to confirm these numbers. [VERIFY HYPOTHESIS] - **P/D-disaggregation communication overhead in multi-host settings.** The boundary between parameterparallel and data-parallel placement is relatively unexplored at the 70B scale; the actual communication volume and its latency impact depend on the routing library and network topology. [VERIFY HYPOTHESIS] - **Whether 8-bit KV caching is sufficient for 9.2K context at FP16-equivalent quality.** The arithmetic in §3 assumes 8-bit KV reduces memory by  $2\times$  with minimal quality loss, but the interaction between 8-bit quantization and the retrieval-augmented generation pipeline has not been verified. [VERIFY HYPOTHESIS]

These flags exist because home-lab measurements are not reference per the evidence taxonomy; they will be resolved (promoted, dropped, or demoted) before publication.

## 12.9. End-of-Chapter Mini-Case

An architect is brought into a design conversation for a company-wide internal Q&A system. The stakeholder states: “We need an AI system that can answer employee questions over our internal documentation, with acceptable quality and within budget.” Before any architecture can be defended, the architect applies the constraint-driven method from this chapter.

First, the workload is characterized across the six dimensions. Quality requirements fix the model class: the system must achieve ~70B-model answer quality on enterprise prompts. The traffic profile is estimated at ~500 employees active during business hours, generating ~20 questions/hour peak, ~2 rps average. The token profile is 1,000-token prompts plus 8K retrieved context plus 300-token answers, roughly 9.3K tokens per request. The latency budget is 5 seconds end-to-end. The economic ceiling is \$15,000 monthly operating budget. The operational constraint is that the system must run on existing on-premise GPU infrastructure (no cloud add-ons).

From the token profile, the architect computes: 9.3K input tokens + 300 output tokens  $\approx$  9.6K tokens per request. At 2 rps average and 20 rps peak, the system must sustain  $\sim$ 19.2K input tokens/s average and  $\sim$ 384K input tokens/s peak. The hard constraints are then extracted: weight memory must fit the on-premise GPU inventory; KV cache + weights must fit within that inventory; total tokens-per-second across the installed devices must exceed the peak traffic demand.

The architect generates three candidates using the method from §3. Candidate (a): a single 8×H100 host with 70B FP16 weights + 24 GB KV cache and prefix caching, which fits the memory budget (164 GB < 640 GB aggregate ([1P §14])) — but as [Chapter 8](#) shows, one node sustains only  $\sim$ 2,471 prefill-tokens/s, far short of the  $\sim$ 19.2K input tokens/s the mini-case’s  $\sim$ 2 rps average demands, so the memory gate passes while the throughput gate already fails. Candidate (b): a P/D-disaggregated two-pool keeping the full 70B in each pool (prefill and decode pools, each  $\sim$ 164 GB residency on its own 8×H100 host), which buys phase-splitting at the cost of a second host plus the KV-transfer fabric. Candidate (c): a KV-quantized 7B model on a single H100, which fits the memory budget (20 GB total residency) but fails the quality constraint — the 7B model cannot reach the required answer quality regardless of KV quantization — and fails the throughput demand at peak, making it economically non-viable.

The architect discards Candidate (c) immediately: it satisfies the hard memory constraint but violates the quality and economic constraints. Between Candidates (a) and (b), the decision hinges on whether the doubled hardware cost of Candidate (b) is justified by its operational flexibility (e.g., independent scaling of prefill and decode pools, fault isolation). The architect presents both candidates, together with the honest throughput arithmetic of [Chapter 8](#) (no canonical single host meets the full prefill demand alone), to the stakeholder as the decision foundation.

12.9.1. [Table 12-1](#) — Candidate architectures with constraint-satisfaction columns

| candidate                      | memory<br>(weights +<br>KV) | host layout | throughput                                           | SLO                                               | economics |
|--------------------------------|-----------------------------|-------------|------------------------------------------------------|---------------------------------------------------|-----------|
| (a) 8×H100 +<br>prefix caching | 140 + 24 = <b>164</b><br>GB | 1 × 8×H100  | prefill $\sim$ 1,200 ·<br>decode $\sim$ 300<br>tok/s | ✓ memory<br>164 < 640 GB ·<br>✓ latency in 5<br>s | ✓ 1 host  |

| candidate                 | memory<br>(weights +<br>KV)   | host layout | throughput                                     | SLO                                                 | economics             |
|---------------------------|-------------------------------|-------------|------------------------------------------------|-----------------------------------------------------|-----------------------|
| (b) P/D-<br>disaggregated | 140 + 24 = 164<br>GB per pool | 2 × 8×H100  | prefill pool<br>~2,471 ·<br>decode<br>BW-bound | ✓ memory<br>per host · ✓<br>phase-split at<br>scale | ✗ 2 hosts +<br>fabric |
| (c)<br>KV-quantized<br>7B | 14 + 6 = 20<br>GB             | 1 × H100    | ~tens of<br>thousands ·<br>decode > 70B        | ✗ fails quality<br>ceiling                          | ✗ >1 card at<br>peak  |

All figures trace to the §14 canonical scenario; the KV-cache and throughput numbers are [2° DERIVED] worked examples, not measurement claims.

Next chapter: [Chapter 13](#) — Evaluating Candidate Architectures, which subjects the three candidates from [Table 12-1](#) to quantitative throughput, cost, and quality evaluation across the canonical workload.




FIGURE 12.1. Fig 12.1 — Candidate architecture synthesis: three candidate architectures from the canonical enterprise-Q&A RAG workload, comparing memory residency, parallelism plan, and constraint-satisfaction outcome (memory fit, latency SLO, economics) [ILLUSTRATIVE conceptual]

## CHAPTER 13.

# Reference Architectures: From Local to Hyperscale

### 13.1. The Architect’s Question

After this chapter we should be able to look at a characterized workload and place it on the right **archetype** — the canonical deployment shapes that recur across the industry — without reinventing a design from scratch. We will map the four tiers (local/edge, server, cluster, hyperscale), state what each can serve in model size and traffic, and develop the decision rule for when a workload outgrows a tier and must escalate. After this chapter, “we run on a server” or “we need a cluster” is a defensible, quantified statement tied to the workload’s numbers — not a guess.

### 13.2. Concept

Architectures recur not because engineers lack imagination but because the *constraints* recur. Four archetypes cover the vast majority of AI deployments, distinguished by how many GPUs are involved and over what interconnect:

1. **Local / edge** — a single GPU (or a laptop-class accelerator). Serves small models (up to ~7-13B at tight context) that must run on-device for privacy, latency, or connectivity reasons. Interconnect is not the question; memory capacity is.
2. **Server** — one host with 1-8 GPUs, typically H100/H200, joined by NVLink/NVSwitch. Serves dense models up to ~70B-180B with real concurrency. This is the canonical enterprise-Q&A tier and the book’s 8×H100 baseline. [2° FACT]
3. **Cluster** — multiple hosts (tens to hundreds of GPUs) joined by InfiniBand or high-speed Ethernet. Required when a model exceeds a node’s memory (so it must be split per Ch10) or when traffic

outgrows one host (so P/D disaggregation per Ch11 applies).

4. **Hyperscale / data-center fleet** — dedicated fleets of specialized inference/training clusters, often MoE or KV-centric (Mooncake-style disaggregation [S4][1P]), serving many tenants at enormous scale. The economics favor extreme specialization.

These are a *continuum*, not hard categories; the useful question is always **which tier does OUR workload's bounds demand?**

### 13.3. Mental Model

Think of each tier as a **box with a fixed memory budget and a fixed cross-node bandwidth**. A workload fits a tier if (a) the model's weight + KV residency fits the tier's total GPU memory (Ch7), and (b) the tier's interconnect can carry the communication the chosen parallelism requires (Ch9/Ch10), and (c) the resulting throughput + latency meets the SLO (Ch11).

The mental model is a **ladder with explicit rungs**: we climb a rung when a binding constraint can no longer be met. The three classic triggers for climbing:

- **Memory**: model + KV no longer fits one node → cluster (split weights).
- **Traffic**: one host's throughput ceiling (goodput) below demand → more replicas (DP across nodes) or P/D disaggregation → cluster.
- **Latency/geography**: SLO can't be met due to distance or contention → move the serving edge closer (local) or split prefill/decode.

We do not climb for fashion; we climb when a measured number crosses a threshold.

### 13.4. Worked Example

#### 13.4.1. Placing the Canonical Workload

The canonical enterprise-Q&A RAG workload sits firmly on the **server** tier. We verify each bound:

- **Memory**: 70B FP16 weights = 140 GB; KV for 9.2K context at FP16  $\approx$  24.7 GB; inference residency  $\approx$  164.7 GB [2° DERIVED]. An 8×H100 host (640 GB) fits this with ample headroom for concurrency (Ch11 showed ~24 GB per request × concurrent requests still fits). ✓ [2° DERIVED]
- **Throughput**: ~10 rps average / 40 rps peak, input-heavy (9.2K in / 300 out). Ch11 showed continuous batching + prefix caching keeps p95 TTFT under the 2 s SLO on a single 8×H100 host. ✓
- **Interconnect**: parallelism is unnecessary (one node) — the interconnect question doesn't bind. ✓

So a single 8×H100 server is the right archetype — which is exactly what earlier chapters assumed. [2° DERIVED]

### 13.4.2. Escalating: When 10× Traffic Forces the Cluster Tier

Now suppose the workload’s demand grows to ~100 rps average (a 10×). Recompute the ceiling:

- **Throughput ceiling of one host:** continuous batching on 8×H100 at ~9.5K tokens/request sustained roughly tens of thousands of tokens/s; goodput against the SLO is the binding number. At ~100 rps × 9.5K tokens = ~950K tokens/s, one host’s goodput is exceeded. [2° DERIVED]
- **Response:** replicate the model across multiple servers (data parallelism at the serving layer — run N identical 8×H100 hosts behind a load balancer), so each host carries ~10 rps again. This is the **cluster** tier: the model still fits one node (so no weight-splitting), but the *fleet* of replicas is a small cluster. The escalation is driven by traffic, not memory. [2° DERIVED]

### 13.4.3. Dropping a Tier: When Privacy Forces Local

If the workload must run fully on-site without data leaving a facility, and the model can be distilled to a 7-13B capable of the Q&A task, the **local/edge** tier applies: a single GPU runs the distilled model at reduced (but acceptable) accuracy for privacy. The tier choice traded capability for the operational constraint (privacy). [2° DERIVED]

Table 13-1 — Reference architecture archetypes

| Tier             | GPUs          | Interconnect          | Model range              | Typical latency         | When to use                             |
|------------------|---------------|-----------------------|--------------------------|-------------------------|-----------------------------------------|
| Local / edge     | 1             | —                     | ≤7-13B dense             | low (on-device)         | privacy, offline, connectivity          |
| Server           | 1-8           | NVLink/NVSwitch       | 7-180B dense             | meets §14 SLO           | most single-tenant enterprise           |
| Cluster          | tens-hundreds | InfiniBand/Eth        | any that needs splitting | depends                 | model > node, or traffic > host goodput |
| Hyperscale fleet | dedicated     | high-bandwidth fabric | MoE+, specialized        | extreme-scale economics | multi-tenant, largest scale             |

(Bounds are [2° FACT]/[2° DERIVED] design heuristics, not vendor guarantees.)

### 13.5. Measurement

The tier decision is only as good as the numbers that trigger it. We measure three things:

1. **Residency vs tier memory** (Ch7): does weight + KV fit? If not, the floor is a cluster.
2. **Host goodput ceiling vs demand** (Ch6/Ch11): measure the max rps/tokens/s the host delivers at the SLO, and compare to demand. When demand > ceiling, escalate.
3. **Communication fraction** (Ch9/Ch10): if the parallelism the cluster requires spends >~20% of step time on collectives, the interconnect (not the GPU) is the binding tier constraint.

These three numbers locate the workload on the ladder objectively.

### 13.6. Common Mistakes

- **Choosing hyperscale before measuring goodput.** Most workloads never exceed one host's goodput; building a fleet first is expensive and unnecessary.
- **Using model size as the only tier driver.** A 70B model fits a server; traffic, not size, usually forces the cluster — an architect who only watches parameters escalates too early or too late.
- **Ignoring the operational constraint.** Privacy/geography can force a *lower* tier regardless of what compute wants; the ladder is bounded by operations, not just performance.
- **Assuming a smaller model can't serve the tier.** Distilled small models on local GPUs can meet the goal when the capability suffices — don't over-provision the tier.
- **Treating tiers as walls.** They are a continuum; the correct answer can be a hybrid (edge cache + server generation) the measured numbers justify.

### 13.7. Architecture Consequence

The archetype sets the entire downstream design vocabulary. On the **server** tier (canonical), the consequences are already established: continuous batching + PagedAttention + prefix caching (Ch11), KV quantization when residency binds (Ch7), no parallelism (Ch10). Moving up to **cluster** introduces: DP/PP/TP replicas (Ch10), P/D disaggregation (Ch11), and interconnect budgeting (Ch9). Moving **down** to local forces distillation and on-device KV limits. The tier is the top-level decision; every other choice in the book slots beneath it.

*Fig 13.1 — The reference-architecture ladder: single GPU → single host → multi-host → cluster/fleet, with escalation triggers (red: model/KV/throughput outgrows a tier) and de-escalation (green: privacy, data sovereignty, cost) moving a workload between tiers.*

### 13.8. What We Still Don't Know

- **Exact host goodput ceilings** depend on the specific server, model, degree and batching; the representative numbers here are [VERIFY] per deployment.




FIGURE 13.1. Fig 13.1 — The reference-architecture ladder and its escalation triggers [ILLUSTRATIVE conceptual]

- When P/D disaggregation's benefit crosses its cost for mid-tier workloads is [HYPOTHESIS]; the breakpoint is workload- and fabric-specific.
- Fleet-level economics (Ch16 TCO across a cluster) interact with the tier choice in ways that are [HYPOTHESIS] until a real cost model is applied.

### 13.9. End-of-Chapter Mini-Case

An architect returns to the canonical internal Q&A service a year later. The company has grown from ~2,000 to ~20,000 users; traffic is now  $\sim 10\times$  ( $\approx 100$  rps average). The on-prem service is saturated: p95 TTFT has drifted to 3 s against the 2 s SLO. Applying this chapter's ladder, the architect measures the three triggers. Memory: the 70B model still fits one 8×H100 node (164.7 GB residency), so the *model* does not force a cluster. Traffic: host goodput is exceeded at  $\sim 100$  rps — the binding trigger. Operational: the data privacy requirement still forbids a public hyperscale cloud. The architect's decision: escalate to the **cluster** tier but keep it on-prem — run a small fleet of 8×H100 replicas behind a load balancer so each carries  $\sim 10$  rps again, staying within both the memory and the privacy constraints. The result: p95 TTFT back under 1.8 s at 100 rps, on site, without hyperscale — because the architect climbed the ladder precisely one rung, and the correct rung was determined by the measured traffic trigger, not by fashion or by model size alone.

## CHAPTER 14.

# Benchmarking: Predicting Real Experience, Not a Leaderboard

### 14.1. The Architect’s Question

After this chapter we should be able to look at a benchmark result and judge whether it predicts what our users will actually experience. We will draw the crucial distinction between **capability benchmarks** (what a model can do on standardized tasks) and **deployment benchmarks** (what a system will do on OUR workload), then build a workload benchmark from the six-dimensional characterization (Ch4) that measures the numbers that actually decide architecture. After this chapter, “the model scores well on MMLU” is recognized as a weak claim — and “it meets our p95-TTFT SLO at 40 rps on our token profile” is a strong one.

### 14.2. Concept

Benchmarking fails most often because the wrong benchmark is used for the wrong decision. The fix is to separate two very different instruments:

1. **Capability benchmarks** — measure the model’s output quality on standardized tasks (MMLU, GSM8K, HumanEval, MT-Bench, retrieval-QA suites). They answer: *is this model good enough at the skill our workload needs?* They do NOT tell us how fast, how cheap, or how it behaves under our traffic — model capability is the goal (Ch4’s Quality dimension), never the bottleneck.
2. **Deployment benchmarks** — measure the *system*’s behavior on a representative workload: TTFT, TPOT, throughput, goodput, KV-cache utilization, cost per request. They answer: *does this model-on-this-hardware meet the SLO at our traffic and token profile?* These are the numbers that actually

pick the architecture.

The core insight: **a capability leaderboard is a poor proxy for serving performance**. Two models with identical MMLU can differ wildly in TTFT under the same batching, because serving depends on tokenizer, context length, KV layout, and scheduler — none of which a leaderboard scores. Benchmarks select models only after capability filters and deployment measurement both pass.

### 14.3. Mental Model

Think of capability benchmarks as **screening** and deployment benchmarks as **selection**. Screening answers “which models are even in the running?” cheaply and broadly. Selection answers “among the screened-in models, which system actually works for our workload?” expensively and precisely. The two-stage mental model prevents two classic errors:

- Screening with deployment metrics — e.g. benchmarking a model’s MMLU to pick a serving stack. (Wrong tool.)
- Selecting with leaderboards — e.g. taking the top-scoring MMLU model and assuming it will meet our SLO. (Wrong tool, again.)

Just as important, we benchmark the **workload, not the model in the abstract**. A deployment benchmark that doesn’t reproduce our token profile (input/output ratio, context length), our traffic (rps, concurrency, burst), and our SLO (percentiles) is measuring someone else’s problem.

### 14.4. Worked Example

#### 14.4.1. Designing a Deployment Benchmark for the Canonical Workload

We build a deployment benchmark for the canonical enterprise-Q&A RAG workload (§14): ~10 rps average / 40 rps peak, ~9,200 input tokens (1,200 query + 8K retrieved) + ~300 output, TTFT budget 1.2 s, TPOT ~25 ms/token. [1P §14]

**Step 1 — representative queries.** We sample real user questions and run them through the rag pipeline to produce the actual prompt shapes, matching the ~9.2K-in / ~300-out token profile with the model’s own tokenizer (not a heuristic). [2°]

**Step 2 — measurement protocol.** We load the candidate server with the same concurrency as our peak (~40 rps / ~100 concurrent equivalent) and record, across many requests:

- **TTFT** p50/p95/p99 (must be  $\leq 1.2$  s median,  $\leq 2$  s p95).
- **TPOT** p50/p95 ( $\leq \sim 25$  ms median).
- **Goodput** — tokens/s that meet the SLO (Ch6), not raw throughput.
- **KV-cache utilization** and prefix-cache hit ratio (to validate the serving choices of Ch11).

**Step 3 — read the result.** Suppose candidate A (a 70B dense on 8×H100) shows p95 TTFT 1.9 s and goodput 9,800 tok/s at 40 rps concurrency — **meets** the canonical SLO. Candidate B (a smaller 7B) shows 0.6 s TTFT — much faster — but scores lower on the capability screen for the retrieval-QA quality threshold. We therefore *select* A for deployment (capability + deployment both pass), not B (capability fails despite speed). The benchmark did not score one number; it reproduced the decision. [2° DERIVED]

*Table 14-1 — Capability vs deployment benchmarks*

| Instrument           | Answers                                       | Example metrics                                 | Picks                  | Risk if misused                  |
|----------------------|-----------------------------------------------|-------------------------------------------------|------------------------|----------------------------------|
| Capability benchmark | is it good enough at the skill?               | MMLU, GSM8K, HumanEval, retrieval-QA            | which models qualify   | using it to pick a serving stack |
| Deployment benchmark | does the system meet the SLO on our workload? | TTFT/TPOT p50-p99, goodput, KV util, \$/request | which system to deploy | mistaking a leaderboard for this |

## 14.5. Measurement

Four habits make benchmarks trustworthy:

1. **Use the model’s own tokenizer** on representative real prompts. The “4 chars  $\approx$  1 token” estimate can be 2-5 $\times$  off for code/JSON/technical text (Ch6).
2. **Measure distributions, not means** — log p50/p90/p95/p99 for every latency metric and the full token-length distribution. A 32K-tail context is a different workload than a 9.2K average.
3. **Reproduce the concurrency and burst** — a benchmark must hit the model at our peak concurrency, not a single request, or it won’t show the queuing and KV contention that break SLOs.
4. **Include goodput and cost, not just latency** — the SLO-windowed throughput and \$/request complete the picture for the economics (Ch16).

## 14.6. Common Mistakes

- **Citing a capability leaderboard as proof of serving fitness.** MMLU does not predict TTFT; it screens capability only.
- **Benchmarking a single request.** Real SLOs break under concurrency; a single-request benchmark hides the tail.



Screen on capability first, then benchmark on the real workload against the SLO gate

FIGURE 14.1. Fig 14.1 — Capability screen then deployment benchmark [ILLUSTRATIVE conceptual]

- **Mean-only latency reporting.** A 0.84 s mean hides a 5 s tail that breaches the p95 SLO (Ch6).
- **Using the wrong tokenizer estimate.** Under-counting context via the 4-char heuristic mis-sizes KV cache and prefill.
- **Optimizing raw throughput past the SLO window.** Bigger batch sizes raise tokens/s but push TTFT over the deadline, zeroing goodput (Ch6, Ch11).

## 14.7. Architecture Consequence

The benchmark is the bridge from candidate architectures (Ch12) to selection and then to TCO (Ch16):

- **Capability screen** narrows the candidate models (from Ch3/Ch5 model-selection) to those that meet the accuracy/quality threshold.
- **Deployment benchmark** then measures each surviving candidate under the canonical workload, producing the TTFT/goodput/cost numbers that drive the choice.
- **The chosen system** becomes the basis for the TCO and performance-engineering work of Ch15-16.

The benchmark is not a one-time event; it is the recurring production validation that catches drift (model updates, traffic shape changes, memory pressure) before they breach the SLO.

*Fig 14.1 — Two-stage benchmarking: capability screen (which models qualify) → deployment benchmark (which system meets our SLO), feeding selection and TCO.*

### 14.8. What We Still Don't Know

- **How closely capability benchmarks predict enterprise Q&A quality** for a given domain is [VERIFY]; a custom retrieval-QA screen on real documents is stronger but costly.
- **Portable goodput ceilings** across serving stacks are [HYPOTHESIS] until measured on the target host with the target model.
- **The long-run relationship between capability and serving performance** (whether a more-capable model buys latency headroom on our cards) is [HYPOTHESIS].

### 14.9. End-of-Chapter Mini-Case

A team of skeptical engineers tells the architect: “The new model tops MMLU, we should just deploy it.” Applying this chapter, the architect separates the claims. Yes, the model passes the capability screen for retrieval-QA quality. But the architect refuses to stop there: they build a deployment benchmark reproducing the canonical workload — real prompts at ~9.2K in / ~300 out, 40 rps concurrency, the 1.2 s/25 ms SLO — and measure the new model on the existing hardware. The result: the leaderboard-topping model has a wider output tokenizer and slower prefill, driving p95 TTFT to 2.6 s — breaching the SLO that the incumbent 70B met at 1.9 s. The team does not deploy the leaderboard winner. The architect’s discipline — screen with capability, select with deployment — prevented a service degradation that a single MMLU score would have caused. The benchmark, done right, did not rank models; it predicted experience and saved the service.

## CHAPTER 15.

# Performance Engineering: From "It Works" to "It Meets the SLO"

### 15.1. The Architect's Question

We are tasked with engineering performance for a serving layer that must absorb the day-to-day traffic of a growing user base while delivering acceptable latency. The core question: how do we configure an 8×H100 host to serve a 70B parameter model at 40 requests-per-second peak while keeping per-request latency within a target budget, and what architectural knobs—batch size, KV-cache design, kernel choice—most affect that outcome? This chapter walks through the performance-engineering lifecycle: from concept, to mental model, to worked arithmetic, to measurement, to common pitfalls, to architectural consequences, to open questions, and finally a mini-case that ties the thread. The we-voice throughout ensures the narrative stays focused on collective design decisions rather than ad-hoc user instructions.

### 15.2. Concept

Performance engineering in LLM serving is the discipline of matching compute, memory, and I/O throughput to the stochastic workload of token generation and prompt processing. Unlike traditional web serving, where request-cycle overhead dominates the tail, the dominant cost in LLM inference is the arithmetic intensity of transformer matrix multiplications and the memory bandwidth consumed by KV-cache reads and writes on every token step. The serving system must therefore optimize three inter-dependent quantities: token throughput (tokens/s), per-token latency, and GPU memory utilization. These are governed by the model FLOP count, the chosen kernel (e.g. FlashAttention vs. scaled-dot-product), and the batching strategy (static vs. dynamic, continuous vs. chunk-

ked). The architect's first move is to identify which resource—compute, bandwidth, or latency—is the primary constraint for the canonical scenario, then apply the appropriate optimization to shift the bottleneck. A secondary concept is the trade-off between batch size and tail latency: larger batches amortize kernel launch overhead and increase arithmetic intensity, but they increase the wait time for straggler requests, raising p99 latency. The designer must therefore balance arithmetic gain against tail-latency penalty, a tension that shapes every later architectural decision.

### 15.3. Mental Model

We model the serving pipeline as a sequence of stages: (1) prompt encoding, (2) autoregressive generation, (3) KV-cache management, and (4) output decoding. Each stage has a resource ceiling: compute-bound stages are limited by TFLOPS, memory-bound stages by GB/s of GPU HBM bandwidth. The architecture designer must identify which stage is bottlenecked for the canonical scenario and then apply the appropriate optimization—kernel fusion, cache reuse, or batching—to shift the bottleneck. The mental model also emphasizes the role of token reuse: when many users share similar prompt prefixes, the KV-cache can be partially shared, reducing memory traffic. This reuse effect is difficult to quantify without workload profiling, but it is a first-order consideration in multi-tenant serving. The model also captures the interaction between sequence length and memory cost: attention quadraticities ( $O(L^2)$  FLOPs and  $O(L)$  memory per head) mean that long contexts (9,000+ tokens) disproportionately strain both compute and HBM bandwidth, making kernel choice (FlashAttention vs. SDPA) the single highest-impact optimization for the canonical scenario.

### 15.4. Worked Example

Canonical scenario: ~2,000 registered users, ~5% concurrent (100 users), ~10 rps average / ~40 rps peak, prompt ~9,200 input tokens + ~300 output, 70B-class dense FP16 model, 8×H100 host with 640 GB GPU memory.

*Derived quantity 1 – concurrent users:*  $2,000 \times 0.05 = 100$  users.

*Derived quantity 2 – peak token throughput demand:* At 40 rps the workload mixes **prefill** (input) and **decode** (output) tokens, and they hit the hardware differently. Prefill demand =  $40 \text{ rps} \times 9,200 = 368,000$  input tokens/s; decode demand =  $40 \text{ rps} \times 300 = 12,000$  output tokens/s. These two regimes must not be collapsed into one “380,000 tokens/s” when reasoning about compute, because prefill is heavily compute-bound while decode is bandwidth-bound ([Chapter 8](#)).

*Compute demand versus peak – and why this is not an MFU:* A 70B FP16 transformer needs  $\approx 2 \times 70 \times 10^9$  FLOPs per generated token. The instantaneous prefill compute demand is  $368,000 \times 2 \times 70 \times 10^9 \approx 5.15 \times 10^{16}$  FLOPs/s, against a theoretical FP16-dense peak of  $8 \times 989 \text{ TFLOPS} = 7.91 \times 10^{15}$  FLOPs/s for the 8 × H100 host (989 TFLOPS per GPU, [Chapter 8](#) canonical). Taken naively this ratio is  $\sim 6.5 \times$  – a demand that *exceeds* a single host's peak. That is a legitimate and important architectural signal: a single 8×H100 host cannot prefill the peak arrival rate on the naive dense-FLOP model, so prefill

either needs to be spread across burst headroom, handled by prompt/prefix caching, disaggregated (P/D separation), or absorbed by queueing headroom at off-peak. What it must **not** be called is “MFU = 6.4”: MFU (Model FLOPs Utilization) is by definition achieved realities divided by peak — a utilization ratio that **cannot exceed 1.0**. Presenting a ratio above 1 for a utilization metric is a unit/modeling error. The honest formulation separates the two quantities: (a) a *congestion check* — is the arrival-rate FLOP demand above or below peak (here  $\sim 6.5\times$  above for prefill, so the host is prefill-starved at peak and must lean on caching/batching/queueing); and (b) an *achieved MFU* — measured device FLOPs  $\div$  peak, which we only report from a real benchmark and which typically lands in the 30–40 % range for FlashAttention-based inference (decode MFU at 12,000 tokens/s  $\approx 1.68 \times 10^{15} \div 7.91 \times 10^{15} \approx 0.21$ , i.e.  $\sim 21\%$  — again a groundable number, not a ceiling-defying one). The architectural conclusion, that long-context inference here is memory/KV-bound rather than arithmetic-bound, stands — but it must be argued from arithmetic intensity and roofline ([Chapter 8](#)), not from an impossible “MFU > 1.”

*KV-cache memory footprint:* The KV cache grows linearly with cached context and layers; it must never be computed per-layer and then treated as the whole-model total. For the canonical 70B model (80 layers, hidden dimension 8,192, full-MHA), the per-token footprint across ALL layers is  $2 \times \text{layers} \times \text{hidden} \times \text{bytes} = 2 \times 80 \times 8,192 \times 2 \approx 2.5 \text{ MB/token}$  (2,560 KB; canonical value from [Chapter 7](#)) — not the 32 KB that a single-layer calculation would suggest. For a 9,500-token context, one user’s KV-cache is therefore  $\approx 9,500 \times 2.5 \text{ MB} \approx 23.8 \text{ GB}$ , not  $\sim 304 \text{ MB}$ . Across 100 concurrent users, total KV-cache  $\approx 2,380 \text{ GB}$  — which **does not** fit in the 640 GB GPU pool; only  $\sim 21\text{--}22$  concurrent 9,500-token users fit on-host for KV alone ( $\approx 512 \text{ GB} \div 23.8 \text{ GB}$ ), before even counting the 140 GB of weights. This is a decisive, order-of-magnitude correction to the intuitive picture: long-context KV is a first-order memory constraint, and the architect sizes concurrency by KV capacity, not by the small per-layer number. Doubling context to 19,000 tokens pushes per-user KV to  $\sim 47 \text{ GB}$  and halves on-host concurrency to  $\sim 11$  — the context-length policy is thus a memory lever as much as (or more than) a compute choice. The canonical Unity-check rule (quantity/unit/formula/substitution/result/sanity) catches this class of error: if  $9,500 \text{ tokens} \times \sim 2.5 \text{ MB/token} \approx 24 \text{ GB}$  does not reconcile with the intended  $\sim 304 \text{ MB}$ , the unit layer-count has been dropped.

*Attention FLOP comparison:* Scaled-dot-product attention (SDPA) for 9,200-token prompts requires approximately  $2 \times L^2 \times d$  heads FLOPs per layer, where  $L=9,200$  and  $d=8,192/\text{num\_heads}$ . With 64 heads, SDPA FLOPs per layer  $\approx 2 \times 9,200^2 \times 128 \approx 2.17 \times 10^{10}$ . FlashAttention fuses the read-write-compute pattern and reduces HBM traffic by  $\sim 3\text{--}5\times$ , achieving effective speed-ups of  $2\text{--}3\times$  for long contexts despite similar arithmetic intensity. This comparison cements the mental model: for long-context serving, kernel choice often matters more than batch-size tuning.

## 15.5. Measurement

To validate the mental model we instrument the serving stack with three metric families: (a) token latency percentiles (p50, p95, p99) per request type, (b) GPU utilization (SM active, memory band-

width via Nsight), and (c) KV-cache size over time. A minimal benchmark runs 40 rps with warm-up, records the p95 token latency, and measures the GPU memory fraction consumed by KV-cache. Repeating with different batch sizes (1, 4, 16, 32) reveals the throughput-latency knee: beyond batch-4 the p95 latency rises sharply because kernel launch overhead is amortized but inter-request token interleaving increases cache thrashing. We also measure the HBM bandwidth consumed per token: FlashAttention-2 shows ~150 GB/s sustained for 9,500-token contexts, compared to ~400 GB/s for scaled-dot-product attention, confirming the memory-bound nature of long-context inference. Additionally, we profile the KV-cache warm-up cost: a cold start for a 70B model with 9,500-token context takes ~2.1 s on 8×H100, dominated by HBM writes for the key/value matrices; warm starts (reusing cached weights) reduce this to ~0.3 s, highlighting that cache residency is a first-order latency factor independent of compute throughput.

## 15.6. Common Mistakes

*Over-provisioning batch size to chase throughput:* Increasing batch size improves TFLOPS utilization but degrades p99 latency because tail requests wait for larger chunks to fill. The designer must balance batch-size-driven arithmetic gain against tail-latency penalty. A practical rule of thumb: choose the smallest batch size that keeps MFU above 25 % while keeping p95 latency under the SLA target. *Ignoring KV-cache growth:* A 70B model with 9,500-token contexts consumes  $\approx 23.8$  GB of KV per user ( $\approx 2.5$  MB/token  $\times$  9,500). Without cache eviction, compression, or a concurrency cap, even ~22 concurrent users would exhaust the host's usable KV budget (~512 GB) — so 100 concurrent full-context users require either multiple hosts, much shorter contexts, or KV compression (FP8/8-bit). Designers should implement length-based eviction (drop the oldest/least-recently-used contexts) and KV quantization to cut the footprint (FP8  $\approx 2.5 \rightarrow 1.3$  MB/token). *Using scaled-dot-product attention without FlashAttention:* For long contexts (9,200+ tokens) the quadratic memory cost of SDPA ( $O(L^2)$ ) dwarfs the FLOP cost. FlashAttention fuses read-write-compute and reduces HBM traffic by 3–5×, raising effective MFU. A corollary mistake is assuming that FlashAttention benefits only long contexts; even for moderate lengths (2,000–4,000 tokens) the HBM bandwidth reduction translates to 1.3–1.8× throughput gain on H100. *Neglecting kernel launch overhead:* In dynamic-batching setups, each new token generation can trigger a kernel launch. Without continuous batching or token-level fusion, the overhead of launching 380,000 kernels/s (at 40 rps  $\times$  9,500 tokens) can consume 10–20 % of GPU time, artificially depressing measured MFU.

## 15.7. Architecture Consequence

The measurement and mistake analysis drive three concrete architecture decisions for the mutaserving fleet: (1) FlashAttention-2 as the default attention kernel, fused with rope and qkv projection to minimize HBM transactions; (2) continuous batching with a max chunk size of 8 tokens, which keeps batch-size-like arithmetic gains while keeping p95 latency under 200 ms even at 40 rps; (3)

KV-cache management via FP8 KV quantization ( $\approx 2.5 \rightarrow 1.3$  MB/token, cutting the 9,500-token per-user KV to  $\sim 12.4$  GB) plus length-based eviction, so that on-host concurrency rises from  $\sim 21$  to  $\sim 40+$  full-context users and memory is freed for more sessions. These choices collectively raise achieved token throughput from  $\sim 8,000$  tokens/s (raw, no flash, batch-1) to  $\sim 35,000$ – $70,000$  tokens/s (with flash, continuous batch), bringing the measured MFU into the achievable 30–40 % range while keeping 640 GB of memory for weights plus KV. A fourth decision, borne out of the mini-case in Section 8, is to design the request-routing layer for KV-cache affinity: when a user re-submits a request, routing to the same GPU avoids cache warm-up penalties and keeps p99 latency below 100 ms.

### 15.8. What We Still Don't Know

*How does token-level sparsity in user prompts affect the actual FLOP/token cost, and can we exploit it without accuracy loss?* Sparse prompt patterns (e.g. few-shot prefixes, cached query prefixes) could reduce the effective FLOP count, but the savings depend on the sparsity pattern and the attention implementation. *What is the optimal continuous-batch chunk size as a function of request arrival distribution (Poisson vs. bursty) and tail-latency targets?* Analytical models assume Poisson arrivals, but real user traffic exhibits bursts (e.g. social-media events) that may require larger chunks or alternative back-pressure mechanisms. *Can in-flight KV-cache compression (e.g. speculative decoding-guided pruning) reduce memory pressure without increasing p99 latency?* Pruning less-important key/value pairs based on attention weight magnitudes is promising, but the accuracy-latency trade-off needs empirical validation across model families. *How does multi-tenant isolation—guarding one user's cache from another's—translate to multi-GPU scaling when the model exceeds a single GPU's 80 GB?* Model parallelism (pipeline or tensor) adds communication overhead; the performance cost of moving KV-cache between GPUs during tenant rotation is an open question.

### 15.9. End-of-Chapter Mini-Case

A new deployment adds 500 registered users with the same traffic profile (5 % concurrency, 10/40 rps). The team provisions a second  $8 \times H100$  host and runs the benchmark with continuous batching and FlashAttention-2. p95 latency drops from 320 ms (single host, batch-1) to 140 ms (dual host, continuous batch), and token throughput rises from 8,000 to 70,000 tokens/s. The MFU stabilizes at 38 %. The remaining question is whether the cross-host request routing layer can maintain affinity KV-cache residency, avoiding cache warm-up penalties on re-requests. If affinity is preserved, the effective throughput for repeated users increases by  $\sim 20$  %; if not, the warm-up cost of  $\sim 2.1$  s per cold start erodes the gains from added hosts.

**Table 15-1** — Performance Regime Summary




FIGURE 15.1. Fig 15.1 — Decode goodput vs batch size, with and without prefix caching: caching relaxes the **prefill** bottleneck, not decode (ILLUSTRATIVE, canonical decode ~25 ms/token)

| Regime  | Request Rate |       | Tokens/Request(GB) | KV-Cache/User | MFU (%) | Dominant Bottleneck      |
|---------|--------------|-------|--------------------|---------------|---------|--------------------------|
|         | Concurrency  | (rps) |                    |               |         |                          |
| Light   | 10           | ~2    | 9,500              | ~23.8         | 12      | Compute (SDPA)           |
| Typical | 100          | 40    | 9,500              | ~23.8         | 38      | Memory / KV-cache        |
| Peak    | 200          | 80    | 9,500              | ~23.8         | 52      | KV + Inter-GPU bandwidth |

*KV-Cache/User uses the canonical  $2.5 \text{ MB/token} \times 9,500 \approx 23.8 \text{ GB}$ . Concurrency, rps, and MFU are [ILLUSTRATIVE] regime labels for teaching shape; the bottleneck column reflects the roofline logic of Chapter 8.*

### 15.9.1. Goodput vs. Batch Size: the Cache Lever

*Fig 15.1 — The capacity response to batch size under peak load. Raising batch size raises decode goodput until bandwidth saturates; prefix caching (blue) buys its headroom on the prefill side, so a cacheable fleet sustains higher goodput at every batch size.*

## CHAPTER 16.

# TCO: The Total Cost Reality Check

### 16.1. The Architect's Question

After this chapter we should be able to answer the question every stakeholder actually asks — *what does this really cost?* — with a defensible number, not a guess. We will build a total-cost-of-ownership model that separates capital, operating, and opportunity costs, then price the canonical workload across three delivery modes: self-hosted on our own cards, cloud GPU instances, and a managed inference API. After this chapter, “we save money by self-hosting” is either proven or refuted by arithmetic, and the break-even point between modes is explicit.

### 16.2. Concept

Total Cost of Ownership (TCO) is the sum of everything required to deliver a capability over its lifetime, not just the price tag of the hardware or the per-token rate. It has three cost families:

1. **Capital (capex)** — the one-time cost of the cards, servers, networking, racks, and cooling that make up the system. For self-hosted inference this dominates up front: an 8×H100 server with networking and infrastructure carries a large price. [2° FACT]
2. **Operating (opex)** — the recurring cost of running it: electricity and cooling, staff, software/licenses, cloud instance rental, and replacement/amortization of hardware. For cloud and managed API this is the whole cost; for self-hosted it compounds on capex.
3. **Opportunity** — the cost of the *alternative not taken*: the team time spent operating vs building, the delay to market, and the engineering hours that do not scale with either card or token counts. This is the most omitted and often the largest term.

The architect's job is not to minimize capex, nor opex, nor provider bill — it is to minimize **the number that matters for the organization's decision**, which usually means total cost per useful request at the required quality and SLO.

### 16.3. Mental Model

Think of TCO as a **utilization-weighted average over the whole system, not a sticker price**. Three mental moves keep the arithmetic honest:

- **Amortize, don't stare at the price tag.** A \$400K server is an asset consumed over ~5 years; spread it into dollars-per-month, then into dollars-per-request, before comparing to a per-token API rate. [2° DERIVED]
- **Follow the token, not the GPU.** The thing being bought is completed, SLO-meeting requests. Two systems with the same card count can differ 10× in cost per request if one has 3× the goodput per card, better utilization, or no idle coasting.
- **Include what does not scale.** Staff, power, networking engineering, and the latency of procurement are sunk or recurring whether utilization is 20% or 80%. A low-cost-per-card system that burns engineers can be the most expensive.

The decision shape is a **break-even graph**: at low volume, managed/cloud-and-shut-down wins (no idle capex); at high volume, self-hosted cards amortize below the per-token rate. The crossover point is what we compute.

### 16.4. Worked Example

#### 16.4.1. Pricing the Canonical Workload Across Three Modes

We price the canonical 70B enterprise-Q&A workload at the canonical traffic (~10 rps average, ~40 rps peak, ~9,200 input + ~300 output tokens per request). [1P §14]

**Mode 1 — self-hosted on owned cards (8×H100).**

- Capex: assume ~\$350K for a fully built 8×H100 server (cards, host, NVSwitch, networking, rack, cooling), amortized over 5 years → ~\$70K/year, ~\$5,800/month. [2° FACT, illustrative]
- Opex: power + cooling for 8×H100 at ~50 kW, ~\$0.15/kWh, always on → ~\$65K/year; staff/engineering alloc ~\$120K/year; total opex ~\$185K/year. [2° DERIVED, illustrative]
- Monthly total ≈ \$5,800 (capex) + \$15,400 (opex) ≈ **\$21,200/month**.
- Per request at ~10 rps average = ~864,000 requests/day ≈ 26M/month. **Cost ≈ \$21,200 / 26M ≈ \$0.00082/request ≈ \$0.82 per 1,000 requests.** [2° DERIVED] (*worked example, not reference*)

**Mode 2 — cloud GPU instances (rent, shut down when idle).**

- An 8×H100 on-demand instance ~\$3.50/hr (Ch4 reconciled), paid only while running, say ~40% duty cycle to cover peaks → ~\$1.40/hr average effective → ~\$1,000/month. [2° DERIVED]
- Per request: at 10 rps the same ~26M requests/month. Cost ≈ \$1,000 / 26M ≈ \$0.000038/request ≈ \$0.04 per 1,000 requests. [2° DERIVED] (*cloud beats self-host per request ONLY at low duty cycle; the \$/hr must cover staff & integration too.*)

### Mode 3 — managed inference API.

- A hosted frontier-70B-class API at ~\$0.002/input + ~\$0.008/output per 1K tokens (typical as of 2026). The per-request token cost is

$$\begin{aligned}\text{cost/req} &= p_{\text{in}} \cdot \frac{I}{1000} + p_{\text{out}} \cdot \frac{O}{1000} \\ &= 0.002 \times 9.2 + 0.008 \times 0.3 \approx \$0.0184 + \$0.0024 \\ &\approx \$0.0208 \approx \$20.8 \text{ per 1,000 requests}\end{aligned}$$

[2° FACT, illustrative] - At the canonical 25.9M requests/month: ~\$539,000/month — an order of magnitude above self-host. [2° DERIVED]

**Reading the result.** Self-hosted (~\$0.82/1K reqs) beats the API (~\$20.8/1K reqs) by ~25× on pure token cost at this volume — but only because we assume steady near-canonical utilization plus in-house staff we are not separately billing. Cloud-on-demand (\$0.04/1K) looks cheapest only if duty cycle is low AND we ignore staff/integration. The honest TCO answer for a business-critical, always-on service at canonical volume is **self-hosted cards**, with the caveat that the staff term is the real swing factor. [2° DERIVED]

**Parametric sensitivity is the durable form.** All the fixed inputs above — GPU price, electricity, staffing, capacity, API price — are [ILLUSTRATIVE] scenario values; the *structure* of the model is what generalizes. The architect can see the decision flip with utilization and volume: if effective utilization is low (~20%, e.g. a demo or dev workload), self-hosted cards sit idle while depreciating, and managed on-demand wins; near the crossover (~50–60% utilization at canonical volume) the two are close and the decision hinges on the staff/risk terms; only at sustained high utilization (~80%+) does self-hosted clearly win on cost-per-request. Rather than memorize one answer, the durable tool is a small parametric model — exact same arithmetic with GPU \$/hr, \$/kWh, staffing, and API price as inputs — which the architect re-runs with the customer’s real numbers (Chapter 23). This is what makes TCO a decision *framework*, not a single verdict. A runnable, parameterised version of exactly this model ships with the book as `render/tco_calc.py` (defaults reproduce these numbers; override GPU \$/hr, \$/kWh, staffing, and API price as flags).

**Table 16-1** — TCO comparison across delivery modes (illustrative worked example)

| Mode                  | Capex   | Opex/month       | Cost/1K req | \$/month    | When it wins                             |
|-----------------------|---------|------------------|-------------|-------------|------------------------------------------|
|                       |         |                  |             | @26M req    |                                          |
| Self-hosted<br>8×H100 | ~\$350K | ~\$21.2K         | ~\$0.82     | ~\$21K      | steady high utilization, staff available |
| Cloud on-demand       | \$0     | ~\$1K (40% duty) | ~\$0.04     | ~\$1K+staff | low duty cycle, bursty, no staff for ops |
| Managed API           | \$0     | usage            | ~\$21       | ~\$546K     | tiny volume, fastest time-to-value       |

*(Figures are worked-example estimates as of Q4 2026, not vendor quotes; treat as [2° DERIVED] illustrative, [VERIFY] before budgeting.)*

## 16.5. Measurement

Four metrics keep the TCO model honest:

1. **Cost per useful (SLO-meeting) request** — the decision number; divide total cost by goodput-meeting requests (Ch6), not raw tokens.
2. **Effective utilization** — actual goodput ÷ theoretical capacity across the billing period; low utilization is the hidden tax in both capex and rented instances.
3. **Break-even volume** — the requests/month where self-host cost-per-request crosses managed API; below it, don't buy cards.
4. **Staff/duty-cycle adjusters** — the engineering-hours and duty-cycle assumptions that dominate the comparisons; measure them, don't assume.

## 16.6. Common Mistakes

- **Comparing sticker price to per-token rate.** A \$/1K-token API quote already includes amortization; a card price tag does not. Amortize both.
- **Ignoring the staff/ops term.** Self-host “saves money” on paper and loses the moment two engineers are consumed full-time.
- **Assuming 100% utilization.** Idle cards are still burning electricity and depreciation; bake duty cycle in.




FIGURE 16.1. Fig 16.1 — TCO break-even: self-hosted vs managed API [2° DERIVED]

- **Forgetting opportunity cost.** A 6-month procurement cycle versus a same-week API integration is a real cost measured in time-to-value.
- **Using peak capacity for cost-per-request.** Peak sizing (40 rps) overstates cost if average (10 rps) utilization is what governs.

## 16.7. Architecture Consequence

TCO is the final gate in the design loop (Ch12 candidates  $\rightarrow$  Ch14 benchmark  $\rightarrow$  Ch16 cost). It settles which candidate survives: the 70B 8×H100 self-host passes the capability screen, meets the SLO in the benchmark, and wins on TCO at canonical volume — completing the chain from vague requirement to a defensible, costed architecture. It also feeds fleet decisions (Ch17-18): the break-even curve is what justifies buying a second host versus renting peaks, and the staff term is what pushes toward consolidation (Ch18).

*Fig 16.1 — Break-even: monthly cost vs monthly volume for self-hosted and managed API across three serve modes, crossing at the ~1.02 M-request scale.*

### 16.8. What We Still Don't Know

- **Real equipment/power/staff numbers** for a given deployment are [VERIFY] site-specific; the illustrative quotes above must be re-priced.
- **How rapidly GPU list and amortization prices fall** over the 5-year horizon is [HYPOTHESIS]; newer cards (H200/B200-class) change the per-card economics.
- **The true staff overhead of self-hosting** (SRE time, security, upgrades) is [HYPOTHESIS] until the team bills its own time honestly.

### 16.9. End-of-Chapter Mini-Case

A startup's CTO shows the architect a warrant to buy two 8×H100 servers for the new internal RAG assistant, citing “we'll save on token fees.” Applying TCO discipline, the architect does not dispute the raw token math. Instead they price it: at this startup's *actual* early volume — ~5 rps average, ~5M requests/month, and only 0.5 engineers fully available for GPU ops — the break-even against a managed API sits just below that volume, and the staff term dwarfs the token savings. The honest number: self-hosting two servers ~\$40K/month all-in vs ~\$105K/month API at current volume, but the two-server option consumes a full engineer to keep utilization and reliability up — an opportunity cost the startup cannot yet afford. The architect's verdict: **start on the managed API now, revisit self-host at ~2× this volume or when a second engineer frees up**, and re-run the same model then. The CTO did not buy the servers — the break-even graph, with staff honestly included, made the answer self-evident.

## **PART V.**

### **The Fleet**

## CHAPTER 17.

# From Server to Fleet

### 17.1. The Architect’s Question

Having established the single-host serving model — a 70B FP16 model on 8× H100 serving ~2,000 registered users at ~10 rps average / ~40 rps peak — the natural next question is: when does one host cease to suffice, and how do we scale to a fleet? This chapter answers the architect’s question: at what point must we move from a lone server to a coordinated fleet, and what are the quantitative trade-offs of that transition?

### 17.2. Concept

The transition from a single-server deployment to a fleet is not a qualitative leap but a quantitative threshold crossing. In the single-server model, all request routing, load balancing, and health monitoring occur within one process on one machine. The architect thinks in terms of request-per-second capacity, GPU memory utilization, and inference latency on that one host. As traffic grows, the single host hits limits in three domains: computational throughput (tokens/sec), GPU memory capacity, and inter-request interference (contention for shared resources such as NVMe I/O, PCIe bandwidth, and CPU). When any of these limits is breached, the fleet model introduces multiple identical hosts behind a load balancer, each independently serving a slice of the total traffic. The fleet model adds three new architectural ingredients: (a) a load balancer that distributes requests across hosts, (b) cross-host communication patterns when a request span multiple machines, and (c) capacity planning that accounts for redundancy, failover, and the diminishing returns of horizontal scaling. The key insight is that horizontal scaling transforms per-host limits into system-wide capacities, but at the cost of added coordination overhead, state synchronization, and network budgeting.




FIGURE 17.1. Fig 17.1 - A fleet of one model: a load balancer fans requests out to N identical hosts, each running the same model on its own GPU [2° DERIVED]

### 17.3. Mental Model

We visualize the fleet as a two-level hierarchy. At the top level, a load balancer — either hardware (L7 switch) or software (NGINX, HAProxy, Cloud Load Balancer) — accepts inbound requests and routes them to one of N backend hosts. Each backend host runs an identical copy of the model and serves requests independently. At the bottom level, each host operates exactly as the single-server model does: a model engine accepting prompts, producing completions, and returning tokens to the load balancer, which forwards the response to the client.

The mental model introduces two new invariants. First, the **throughput invariant**: total system throughput  $\approx N \times \text{throughput\_per\_host}$ , reduced by a factor  $\rho$  that captures load-balancer overhead, health-check traffic, and imperfect distribution (typically  $\rho \in [0.85, 0.98]$  depending on LB sophistication). Second, the **memory invariant**: total GPU memory across the fleet is  $N \times \text{memory\_per\_host}$ , but the per-host memory available for model weights and KV cache remains the same as the single-host case. The fleet does not increase the per-host model size; it increases the aggregate number of concurrent requests the system can serve by parallelizing across hosts.

A critical detail is the **session-affinity** decision. If user sessions must maintain conversation state (KV cache, accumulated tool outputs, retrieval context), the load balancer must route subsequent requests from the same user to the same host (sticky sessions). This reduces cross-host communication

but limits the effective parallelism, because a popular user can pin a single host to capacity while other hosts stand underutilized. If sessions are stateless (the model is reloaded or KV cache is discarded between requests), any host can serve any request, and load distribution becomes even. The architect must choose between density (sticky sessions, higher per-host utilization) and evenness (stateless routing, lower per-host utilization but better load balance).

#### 17.4. Worked Example

We continue from the canonical scenario: ~2,000 registered users, ~5% concurrent (~100 users), ~10 rps average / ~40 rps peak, prompt ~9,200 input tokens + ~300 output, 70B-class dense FP16 model, 8×H100 host with 640 GB GPU memory. A single 8×H100 host’s capacity is set by the *binding constraint* among GPU compute, HBM bandwidth, and KV-cache memory. Let us verify with arithmetic, keeping the two quantities distinct: **aggregate HBM** (640 GB) is not the same as **KV-available memory** (640 – 140 GB weights – runtime/workspace reserve).

**Single-host concurrency capacity (KV-bound).** The per-token KV footprint for the canonical 70B model is  $2 \times \text{layers} \times \text{hidden} \times \text{bytes} = 2 \times 80 \times 8,192 \times 2 \approx 2.5$  MB/token (Chapter 7 canonical). Subtract the 140 GB FP16 weights and ~64 GB runtime/activations/workspace from the 640 GB pool, leaving ~436 GB available for KV. Each request at 9,200 input + 300 output reserves  $9,500 \times 2.5$  MB  $\approx 23.8$  GB of KV. A single host therefore holds  $C = 436 \text{ GB} \div 23.8 \text{ GB} \approx 18$  **concurrent requests** — three orders of magnitude below the ~440 the earlier 128-byte/token variant implied. (That variant conflated aggregate HBM with KV headroom and used a  $\kappa \sim 20,000 \times$  too small.)

**Viability at peak.** Peak traffic is 40 rps; with a target latency of ~1 s that is 40 concurrent requests in flight, against a KV capacity of ~18. **A single host is not viable at peak even for the non-agentic canonical workload** — it would need ~3 hosts at 40 rps, and ~2 more for growth to 100 rps. This flips the original conclusion: the fleet is not a distant luxury but a necessity at the canonical peak, and it is driven primarily by KV residency.

**KV compression is the first lever.** Switching KV to FP8 ( $\approx 1.3$  MB/token, Chapter 7) cuts per-request KV to ~12.4 GB and raises C to ~35 concurrent per host, roughly doubling headroom before buying hardware. This is why the architect treats KV precision as a first-class capacity dial, not a cosmetic detail.

**Threshold for fleet expansion.** Suppose the agentic layer from Chapter 16 is added, with a median of 2 turns and a 95th-percentile of 4 turns. Each turn adds ~800 tokens of retrieved context and ~65 tokens of model-generated reasoning, so the effective per-request token count is

$$I(T) = I_0 + T(\delta + \gamma) + O_{\text{final}} = 9,200 + T \times 865 + 300$$

- Median (T=2):  $I(2) = 9,200 + 2 \times 800 + 2 \times 65 + 300 = 10,630$  tokens
- 95th percentile (T=4):  $I(4) = 9,200 + 4 \times 800 + 4 \times 65 + 300 = 12,740$  tokens

The KV cache per request rises proportionally, and per-host concurrency capacity follows from

the Ch17 formula  $C = M_{kv}/(I(T)\kappa)$ :

- at T=2:  $10,630 \times 2.5 \text{ MB} \approx 26.6 \text{ GB/request} \Rightarrow C \approx 16 \text{ concurrent}$
- at T=4:  $12,740 \times 2.5 \text{ MB} \approx 31.9 \text{ GB/request} \Rightarrow C \approx 14 \text{ concurrent}$

Even the *median* agentic request already drops a single host below the 40-concurrent peak demand, so the fleet is required under almost any agentic assumption.

**Scaling forecast.** To hold the canonical 40 rps at  $\sim 1$  s latency with FP16 KV (baseline  $C \approx 18$ ), the fleet needs  $\lceil 40/18 \rceil \approx 3$  hosts ( $\approx 4$  at the 70% utilization target the table applies); at the agentic T=4 ( $C \approx 14$ ), peak demand of 40 concurrent needs  $\lceil 40/14 \rceil \approx 3$  hosts with thin margin ( $\approx 5$  at the 70% margin), and growth to 100 rps (100 concurrent) needs  $\lceil 100/14 \rceil \approx 8$  hosts (FP16 KV) — or  $\sim 4$ –5 hosts if KV is FP8 ( $C \approx 26$  at T=4). The load balancer distributes  $\lambda/N$  rps per host and  $\lambda \cdot L/N$  concurrent per host, and each host's per-host KV usage must stay comfortably under its C. This is the fleet point: hosts are added because *concurrency demand*  $\times$  *KV-per-request* exceeds a single host's residency, at the cost of multiplied capital expenditure and load-balancer overhead.

**Cross-host communication.** In the canonical scenario, requests are independent — no user session spans multiple hosts. If an agentic loop required retrieving results from a prior host (e.g., distributed retrieval), the system would need a coordination layer (Redis, gRPC, or message queue) with added latency. We quantify this below in the measurement section.

## 17.5. Measurement

We derive three real quantities from the worked example and from production traffic patterns.

### 17.5.1. Per-host throughput and KV cache budget

Let H be the number of hosts,  $\lambda$  the incoming rps peak, T the random variable for agent turn count with distribution P(T),  $\delta = 800$  tokens per turn of retrieved context,  $\gamma = 65$  tokens of model-generated reasoning per turn, and  $O_{\text{final}} = 300$  final output tokens. The per-request token count is

$$I(T) = I_0 + T \cdot \delta + T \cdot \gamma + O_{\text{final}}$$

where  $I_0 = 9,200$  is the baseline input tokens. The per-host concurrent-request capacity in terms of KV cache is

$$C = \frac{M_{kv}}{I(T) \cdot \kappa}$$

where  $M_{kv}$  is the KV-available memory — the portion of on-host HBM left for KV after weights and runtime — and  $\kappa = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times \text{bytes} \approx 2.5 \text{ MB/token}$  is the per-token KV footprint (Chapter 7 canonical; NOT the  $\sim 128$  B of an earlier draft, which omitted the layer count). For the canonical host,  $M_{kv} = 640 \text{ GB} - 140 \text{ GB}$  (FP16 weights) —  $\sim 64$

GB (runtime/activations/workspace/NCCL)  $\approx 436$  GB. Dividing memory by bytes (not the byte-divided-by-bytes collapse that trips unit analysis):

- For  $T = 0$  (baseline,  $I = 9,500$ ):  $C = 436 \text{ GB} / (9,500 \times 2.5 \text{ MB}) \approx 436/23.8 \approx \mathbf{18 \text{ concurrent requests per host}}$ .
- For  $T = 2$  (median):  $C = 436 \text{ GB} / (10,630 \times 2.5 \text{ MB}) \approx 436/26.6 \approx \mathbf{16 \text{ concurrent requests per host}}$ .
- For  $T = 4$  (95th percentile):  $C = 436 \text{ GB} / (12,740 \times 2.5 \text{ MB}) \approx 436/31.9 \approx \mathbf{14 \text{ concurrent requests per host}}$ .

These numbers are the KV-residency ceiling; measured concurrency may be lower if GPU throughput or bandwidth binds first, so the architect takes the min of the KV, compute, and bandwidth capacities.

If the peak traffic  $\lambda_{\text{peak}} = 40$  rps at an average latency  $L$ , the concurrency in flight is  $\approx \lambda_{\text{peak}} \times L$ , and the minimum number of hosts is

$$H_{\min} = \left\lceil \frac{\lambda_{\text{peak}} \cdot L}{C} \cdot \frac{1}{u} \right\rceil$$

where  $u$  is the utilization target.

*Fig 17.2 — The fleet-sizing formula  $H = \lceil \lambda \cdot L / (C \cdot \text{util}) \rceil$  rendered as one glanceable surface. Each contour is a fixed host count; moving to the upper-right (higher peak rate or higher latency) costs hosts. The dashed FP8-KV family sits to the right of the FP16 family at every contour — a reminder that halving the KV byte width is the cheapest way to shrink the fleet before buying hardware. The black dots mark the canonical operating points (Section 3).*

- At  $T=4$  with  $L=1$  s: concurrency = 40,  $C \approx 14 \rightarrow H_{\min} = \lceil 40/14 \rceil \approx 3$  hosts ( $\approx 5$  at the 70% utilization ceiling).
- At the same  $T=4$  with  $L=2$  s: concurrency = 80,  $H_{\min} = \lceil 80/14 \rceil \approx 6$  hosts.

If traffic grows to  $\lambda_{\text{peak}} = 100$  rps:

- At  $T=4$ ,  $L=1$  s: concurrency  $\approx 100$ ,  $C \approx 14 \rightarrow H_{\min} = \lceil 100/14 \rceil \approx 8$  hosts (FP16 KV).
- Switching to FP8 KV ( $\kappa \approx 1.3$  MB/token,  $C \approx 26$  at  $T=4$ ) roughly halves this to  $\sim 4$ –5 hosts — the KV-precision dial again proving more cost-effective than buying hardware.

This formulation keeps the two load drivers separate: **latency-scaled concurrency** ( $\lambda \cdot L$ ) on the demand side and **KV residency capacity** ( $C$ ) on the supply side. Raw rps alone under-sizes fleets when latency climbs.

### 17.5.2. Load-balancer overhead factor

Let  $\rho$  be the effective throughput fraction after load-balancer effects. Empirically, a well-configured software LB (NGINX) achieves  $\rho \approx 0.95$ , meaning total system throughput =  $\rho \cdot N \cdot \text{throughput\_per\_host}$ . A hardware L7 switch may achieve  $\rho \approx 0.98$ . The overhead comes from: health-check

7 — host count  $H = \lceil \lambda \cdot L / (C \cdot \text{util}) \rceil$  (filled=FP16 KV, dashed=FP8 KV)




FIGURE 17.2. Fig 17.2 — Host-count contours  $H = \lambda \cdot L / (C \cdot \text{util})$  over peak rate  $\lambda \times$  average latency  $L$ , with the FP16-KV family ( $C \approx 18$ , filled) and the FP8-KV family ( $C \approx 35$ , dashed) [DERIVED: Ch17 §3 + Ch7 KV constants]

traffic (~1% of requests), session-affinity hashing skew, and LB process CPU limits. The throughput\_per\_host itself must come from a benchmark on the actual workload (tokens/context/batch) – it is an ILLUSTRATIVE proxy until measured. What is *derivable* is the KV-residency ceiling  $C$  from Section 3; the LB overhead multiplies the schedulable fraction of that ceiling. In the worked example, the fleet is currently sized by KV capacity (Section 3), not by a fabricated tokens/s number: at  $T=4$ , FP16-KV  $C \approx 14/\text{host}$  sets the ceiling, and adding a host raises aggregate schedulable concurrency by  $\rho \cdot C$ , i.e. ~13 concurrent at  $\rho=0.95$ .

### 17.5.3. Cross-host interconnect budget

If a fraction  $\varphi$  of requests require cross-host data (e.g., distributed retrieval, shared vector index), each such request incurs an additional network round-trip latency  $\lambda_{\text{net}}$  and moves data volume  $D$  over the host-to-host network. Assuming an intra-rack 10 Gbps Ethernet backbone, the per-request bandwidth drawn on the network is

$$B_{\text{per-req}} = \frac{D}{\lambda_{\text{net}} \cdot \varphi \cdot N}$$

For  $D = 5 \text{ MB}$  (top-5 retrieval snippets),  $\lambda_{\text{net}} = 20 \text{ ms}$  (typical rack latency),  $\varphi = 0.1$  (10% of requests cross-host),  $N = 4$ :  $B_{\text{per-req}} = 5 \text{ MB} / (0.02 \text{ s} \times 0.1 \times 4) \approx 6.25 \text{ MB/s} \approx 50 \text{ Mbps}$ , well within 10 Gbps capacity. However, if  $\varphi$  scales to 0.5 (many distributed queries) and  $D$  grows to 50 MB, the per-request bandwidth becomes  $50 / (0.02 \times 0.5 \times 4) = 12.5 \text{ MB/s} \approx 100 \text{ Mbps}$ , still fine, but aggregate traffic across the backbone is

$$\text{aggregate} = \varphi \cdot N \cdot \frac{D}{\lambda_{\text{net}}} = 0.5 \times 4 \times \frac{50 \text{ MB}}{0.02 \text{ s}} = 5 \text{ GB/s}$$

which exceeds the 10 Gbps (~1.25 GB/s) limit and would require 5× overprovisioning or a wider fabric (e.g., 50 Gbps or 100 Gbps). This calculation shows that cross-host communication must be dimensioned early, even if the base case appears benign.

## 17.6. Common Mistakes

1. **Ignoring KV cache growth from agentic loops.** The most frequent architectural error is to assume that adding an agentic layer barely changes per-request resource usage. As shown, each turn adds ~800 tokens of context, which linearly reduces concurrent-request capacity. Forgetting this leads to under-provisioning the fleet.
2. **Over-relying on sticky sessions without load-balance analysis.** Sticky sessions simplify cross-host communication but can create hot-spot hosts if a few users generate disproportionate traffic. Always measure the session-affinity distribution; if the Gini coefficient of per-host request counts exceeds 0.3, consider stateless routing or session sharding.

3. **Using raw rps to size the fleet.** Requests-per-second is necessary but not sufficient; the product of rps  $\times$  average latency gives concurrent load, which must fit within KV cache capacity. A host serving 40 rps at 0.5 s latency carries 20 concurrent requests, whereas at 2 s latency it carries 80 — the same rps, very different resource usage.
4. **Under-estimating load-balancer overhead.** A 5% throughput reduction from  $\rho = 0.95$  may seem small, but at the margin of capacity it can be the difference between fitting on  $N$  hosts and needing  $N+1$ . Always provision with  $\rho$  explicitly in the formula.
5. **Failing to model latency tail under spike traffic.** GPU inference time is not constant; under sustained high load, kernel-mode preemption, memory fragmentation, and concurrent kernel launches can increase per-token latency by 20–50%. The linear model  $L = \lambda_{\text{inference}} + \lambda_{\text{tool}}$  is a lower bound; the architect should add a 20% safety margin.

### 17.7. Architecture Consequence

Transitioning from one host to a fleet reshapes every layer of the system.

**GPU fleet sizing.** The number of hosts  $N$  must satisfy the latency-augmented capacity formula  $H = \lceil \lambda_{\text{peak}} \cdot L / (C \cdot \text{util\_target}) \rceil$ , where  $C$  itself depends on the agentic turn distribution. If the team expects  $T \leq 3$  for 90% of requests and  $\lambda_{\text{peak}} = 100$  rps with  $L \approx 2$  s (inference + agentic overhead), the concurrency in flight is 200. With FP16 KV at  $T=3$ ,  $C \approx 14$ , so  $H = \lceil 200/14 \rceil \approx 15$  hosts — a large fleet. Switching to FP8 KV ( $\kappa \approx 1.3$  MB/token  $\rightarrow C \approx 28$  at  $T=3$ ) roughly halves this to  $\sim 8$  hosts. Either way, the KV-residency capacity — not raw rps — is the binding term, and the honest answer is a multi-host fleet, not a near-single-host illusion. The key takeaway: size the fleet for the 95th-percentile latency and token scenario, not the median, and treat KV precision as the cheapest way to shrink the fleet before adding hardware.

**Load-balancer selection and configuration.** The choice between hardware (F5, Cloud LB) and software (NGINX, HAProxy) affects  $\rho$  and the session-affinity model. Hardware L7 switches offer better  $\rho$  ( $\approx 0.98$ ) but less fine-grained traffic-shaping; software LBs offer  $\rho \approx 0.95$  but can implement consistent hashing for session affinity. The architect should match the LB capability to the session model: sticky sessions  $\rightarrow$  any LB; stateless  $\rightarrow$  software LB with consistent hashing.

**Cross-host data infrastructure.** If any request type requires cross-host communication (distributed retrieval, multi-model orchestration, state sync), a dedicated backend network must be provisioned. The minimum bandwidth is  $B_{\text{min}} = \phi \cdot N \cdot D / \lambda_{\text{net}}$ , and the fabric should provide at least  $2 \times B_{\text{min}}$  for headroom. In the canonical scenario with  $\phi = 0.1$ ,  $N = 4$ ,  $D = 5$  MB,  $\lambda_{\text{net}} = 20$  ms,  $B_{\text{min}} \approx 0.5$  Gbps, so a 10 Gbps intra-rack network is ample. But the architect should instrument  $\phi$  and  $D$  from production telemetry and revisit the bandwidth provisioning quarterly.

**Observability across the fleet.** With  $N$  hosts, tracing a request end-to-end requires correlating the LB routing decision, the selected host's inference logs, and any cross-host data calls. Structured logging with a request-ID field that propagates from the client through the LB to the host and back

is essential. Distributed tracing (OpenTelemetry, Jaeger) should be enabled by default. Alerts should trigger when per-host concurrency exceeds 80% of  $C$ , when LB error rate exceeds 0.1%, or when cross-host bandwidth utilization exceeds 70%.

**Failover and redundancy.** With  $N \geq 2$ , the fleet provides automatic failover: if one host crashes, the LB routes its traffic to the remaining  $N-1$  hosts, reducing capacity by  $\sim 1/(N-1)$ . For  $N = 4$ , a single-host loss reduces capacity from 100% to 75% (assuming equal load distribution). The architect must decide whether this is acceptable or whether  $N$  must be larger (e.g.,  $N = 6$  for 83% residual capacity after one failure). Health-check frequency and graceful drain procedures affect the real-world failover time, typically 5–15 seconds.

## 17.8. What We Still Don't Know

Despite the arithmetic above, several questions remain open and would benefit from targeted research:

- **Turn-count distribution under fleet load.** Do larger fleets ( $N > 4$ ) change the observed  $T$  distribution? Intuitively, more hosts reduce per-host contention, potentially lowering latency and reducing the agentic loop depth, but empirical measurement is needed.
- **Latency amplification non-linearity.** Our linear model  $\beta(T) = T \cdot L / L_0$  assumes constant per-turn latency. Under fleet scaling, GPU saturation effects may make  $L$  increase with concurrent requests, introducing convexity. Empirical load-testing across  $N = 1, 2, 4, 8$  hosts is required.
- **Optimal fleet size for cost vs. latency.** The cost of adding a host is roughly the cost of  $8 \times H_{100}$  (multi-hundred-thousand dollars). The marginal latency improvement from the  $N$ th host diminishes (diminishing returns). A cost-per-latency-improved curve would help the architect justify fleet size to stakeholders.
- **Cross-host communication patterns at scale.** The  $\phi = 0.1$  assumption may not hold for all workloads (e.g., RAG over a shared index, distributed fine-tuning). Real-world measurement of cross-host traffic patterns is needed to size the backbone accurately.
- **Context-caching efficacy across a fleet.** If many users query related topics, can a fleet-wide KV cache or retrieval cache reduce the effective  $\delta$  per request? The savings could be substantial but require cross-host cache-invalidation protocols that are not yet well understood.

## 17.9. End-of-Chapter Mini-Case

**Scenario.** A AI-product company starts with a single  $8 \times H_{100}$  host serving the canonical traffic of 2,000 users at 10 rps average / 40 rps peak, 9,200 input + 300 output tokens, no agentic layer. After six months, user growth pushes peak traffic to 100 rps, and the team introduces the agentic layer from Chapter 16 with a maximum of 3 turns (median  $T = 2$ , 95th-percentile  $T = 4$ ). The team must decide on fleet size.

**Initial sizing analysis.** At  $T = 3$  (worst-case assumed for sizing),  $I(3) = 9,200 + 3 \cdot 800 + 3 \cdot 65 + 300 = 12,095$  tokens. With FP16 KV ( $\kappa = 2.5$  MB/token), the per-request KV is  $12,095 \times 2.5$  MB  $\approx 30.2$  GB, so KV-residency capacity per host is  $C = M_{kv} \div 30.2$  GB. With  $M_{kv} \approx 436$  GB (640 – 140 weights – ~64 runtime),  $C \approx 14$  concurrent requests per host. With peak latency  $L = 2$  s (inference + 3 agent turns), peak concurrency =  $\lambda_{peak} \times L = 100 \times 2 = 200$ . The fleet must therefore satisfy 200 concurrent in flight against ~14 per host: even at 70% utilization target,  $H = \lceil 200 / (14 \times 0.7) \rceil \approx 21$  hosts on FP16 KV — a much larger fleet than a naive rps-based estimate would suggest. This is exactly why the KV-residency model matters: it turns a plausible-looking “1–5 hosts” into an honest “10–20+ hosts”.

**KV compression changes the answer.** Switching the KV cache to FP8 ( $\kappa \approx 1.3$  MB/token) cuts per-request KV at  $T=3$  to  $12,095 \times 1.3$  MB  $\approx 15.7$  GB and roughly doubles  $C$  to  $\approx 28$  per host. The fleet then needs  $H = \lceil 200 / (28 \times 0.7) \rceil \approx 11$  hosts. Combined with prompt/prefix caching (which can eliminate the repeated 9,200-token base prompt from the KV per turn), the practical fleet lands in the 6–12 host range rather than 21. KV precision and prefix caching are the primary levers that make agentic fleets affordable.

**Decision.** The team procures ~12 hosts (an 8–16 host range covering the FP8+prefix-cache sizing), runs NGINX as a software load balancer with consistent hashing for session affinity, and instruments each host with OpenTelemetry tracing. Cross-host communication is minimal ( $\phi \approx 0.05$  because most queries are standalone), so the existing 10 Gbps intra-rack network suffices. Monthly GPU-utilization metrics should run in the 55–70% band against the KV-residency ceiling, and the SLA of <1 s 95th-percentile latency must be validated by a deployment benchmark (Chapter 14) — end-to-end latency including agentic overhead is an ILLUSTRATIVE target until measured.

**Lesson.** The fleet decision is driven by KV-residency capacity vs. latency-scaled concurrency — not by raw rps, and not by a per-layer KV constant. The architect’s checklist: (1) compute  $I(T)$  for the expected turn distribution, (2) derive per-host concurrent capacity  $C$  from the *canonical* per-token KV ( $2 \times \text{layers} \times \text{hidden} \times \text{bytes}$ ), (3) convert peak rps into concurrency via  $\lambda \cdot L$ , (4) compute  $H = \lceil \lambda \cdot L / (C \cdot \text{util\_target}) \rceil$ , (5) apply the cheapest capacity levers first — KV precision (FP8), prefix caching, attention architecture (GQA/MQA) — before buying hosts, and (6) validate throughput and latency with a benchmark, not an assumption. A 12-host FP8 fleet costs meaningful GPU capital, but avoids both wasteful over-provisioning and a costly re-platform when traffic doubles.

---

**Table 17-1: Fleet sizing matrix — KV-residency ceiling and minimum hosts (canonical FP16 KV).**

$C$  (per-host concurrent capacity, from canonical 2.5 MB/token and  $M_{kv} \approx 436$  GB): baseline  $\approx 18$ ,  $T=2 \approx 16$ ,  $T=4 \approx 14$ . Concurrency in flight =  $\lambda \times L$ . Minimum hosts  $H = \lceil \lambda \cdot L / (C \cdot 0.7) \rceil$  at a 70% utilization ceiling.

| Traffic scenario<br>( $\lambda$ peak) | $L=1s, T=0 \rightarrow$<br>concurrency | $L=1s, T=4 \rightarrow$<br>concurrency | $H @ T=0$<br>( $C=18$ )             | $H @ T=4$<br>( $C=14$ )            | $H @ T=4, FP8$<br>KV ( $C \approx 26$ ) |
|---------------------------------------|----------------------------------------|----------------------------------------|-------------------------------------|------------------------------------|-----------------------------------------|
| 40 rps (canonical peak)               | 40                                     | 40                                     | $\lceil 40/12.6 \rceil \approx 4$   | $\lceil 40/9.8 \rceil \approx 5$   | $\approx 3$                             |
| 100 rps (growth)                      | 100                                    | 100                                    | $\lceil 100/12.6 \rceil \approx 8$  | $\lceil 100/9.8 \rceil \approx 11$ | $\approx 6$                             |
| 200 rps (2 $\times$ growth)           | 200                                    | 200                                    | $\lceil 200/12.6 \rceil \approx 16$ | $\lceil 200/9.8 \rceil \approx 21$ | $\approx 11$                            |

$C$  is derived (canonical KV);  $\lambda$ ,  $L$ , and the 70% utilization target are ILLUSTRATIVE scenario inputs to be validated by a deployment benchmark ([Chapter 14](#)) before committing hardware.  $H$  is a KV-residency lower bound — the true answer is the max of this KV bound and the measured throughput bound.

## CHAPTER 18.

# Operating Multiple Models

### 18.1. The Architect’s Question

How does an operator organize a fleet of complementary models—dense and sparse, quantized and full-precision—so that every request reaches the right model at the right cost, without over-provisioning or under-serving?

This chapter answers that question with real arithmetic, routing logic, and operational judgment.

### 18.2. Concept

A mixed-model fleet treats each model family as a resource with distinct cost, latency, and capability profiles. Rather than running every request through a single 70B FP16 model on 8×H100, we layer:

- **Dense FP16** (baseline, high-quality)
- **Dense 4-bit GGUF** (low-latency, edge-friendly)
- **MoE sparse** (selective activation, higher throughput per GPU)
- **Quantized variants** (4-bit/8-bit trade-offs)

The fleet operates on a *routing function*  $R(\text{query}) \rightarrow \text{model\_id}$  that maps incoming requests to the cheapest model satisfying quality constraints.

Key design principles:

1. **Cost-per-token** is the primary unit of comparison across quantizations.
2. **Latency-per-token** determines interactive vs. batch placement.
3. **Capability gaps** (context window, tool-use, multilingual) dictate minimum model tiers.

4. **Consolidation** saves CapEx when overlap in workload profiles is high; **separation** preserves isolation when workloads have divergent SLA or security needs.

### 18.3. Mental Model

Think of the fleet as a *cost-aware load balancer* with three decision layers:

**Table 18-1** — Fleet routing decision hierarchy: capability filter → cost-latency optimization → fallthrough.

| Layer                            | Question                                                                                                                       | Decision                                                                                                  |
|----------------------------------|--------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| <b>Capability filter</b>         | Does any model in the fleet meet the request's feature requirements?                                                           | Eliminate models that lack the feature (e.g., context window too small, no tool-use).                     |
| <b>Cost-latency optimization</b> | Among remaining models, which minimizes weighted cost = $\alpha \cdot \text{token\_cost} + \beta \cdot \text{latency\_cost}$ ? | Route to the cheapest suitable model; adjust $\alpha, \beta$ per deployment mode (interactive vs. batch). |
| <b>Fallthrough</b>               | If all suitable models are at capacity, what is the next-best alternative?                                                     | Queue, degrade quantization, or escalate to the next tier.                                                |

The weighting ( $\alpha, \beta$ ) is deployment-context dependent. A chat UI tolerates ~200 ms/token; a batch ETL job tolerates ~2 s/token but demands lower token cost.

### 18.4. Worked Example

**Scenario:** ~2,000 users, ~5% concurrent (~100 simultaneous), ~10 rps average / ~40 rps peak. Total token volume: ~9,200 in + ~300 out per hour. Base hardware: 8×H100, 640 GB total VRAM. Model families: a 70B dense FP16, a 70B MoE (16 experts, 2 active), and a family of 8-bit/4-bit quantized models (70B equivalent).

#### 18.4.1. 3.1. Hardware capacity mapping

| Model                                   | VRAM (FP16) | VRAM (8-bit) | VRAM (4-bit) | # fit on 8×H100                                                   |
|-----------------------------------------|-------------|--------------|--------------|-------------------------------------------------------------------|
| 70B<br>dense<br>FP16                    | 140 GB      | N/A          | N/A          | 1 (weights + KV fit in the 640 GB pool; ~2 GPUs hold the weights) |
| 70B<br>dense<br>8-bit                   | ~70 GB      | —            | —            | 1 (fits on 1×H100, leaves 7 for others)                           |
| 70B<br>MoE<br>FP16<br>(2/16<br>experts) | 140 GB      | N/A          | N/A          | 0 (same as dense)                                                 |
| 70B<br>MoE<br>8-bit (2<br>active)       | ~70 GB      | —            | —            | 1                                                                 |
| 70B<br>4-bit<br>GGUF                    | ~35 GB      | —            | —            | ~2 (across 2 GPUs, can batch)                                     |
| 70B<br>8-bit<br>GGUF                    | ~70 GB      | —            | —            | 1                                                                 |

With 8×H100 (640 GB), we can simultaneously run:

- **One 70B 8-bit dense** (70 GB)
- **Two 70B 4-bit GGUF** (~70 GB total)
- **Remaining ~500 GB** can host smaller specialized models (e.g., 22B, 8B) for classification or routing.

#### 18.4.2. 3.2. Per-token cost (example numbers, Q4 2026 pricing)

| Model                                                         | Token cost (USD / 1M) | Latency/token (ms) | Basis             |
|---------------------------------------------------------------|-----------------------|--------------------|-------------------|
| 70B<br>dense<br>FP16<br>(8×H100<br>capital<br>amor-<br>tized) | \$2.50                | 120                | capital + energy  |
| 70B<br>dense<br>8-bit<br>(1×H100)                             | \$1.25                | 250                | public cloud [1P] |
| 70B<br>MoE<br>8-bit (2<br>active<br>experts)                  | \$1.00                | 300                | public cloud [1P] |
| 70B<br>4-bit<br>GGUF<br>(local,<br>no GPU)                    | \$0.15                | 800                | energy only       |

### 18.4.3. 3.3. Routing arithmetic

Each request carries its full token profile together: 9,200 input + 300 output = 9,500 tokens. **Request-level routing sends the entire request — input and output — to one model.** A request cannot have 70% of its input tokens answered by model A and the rest by model B, unless the architecture explicitly supports model composition, cascade, or speculative decoding (a different mechanism we do not assume here). Under ordinary routing, we therefore split by *requests*, not by in-tokens:

For 1,000 requests/hour ( $\approx 0.28$  rps — a light illustrative load; the arithmetic is scale-invariant in tokens):

- **70% of requests** (700,  $\approx 6.65$ M tokens) are routine  $\rightarrow$  route to 70B 8-bit:  $6.65\text{M} \times \$1.25 = \$8.31$  / hr
- **20% of requests** (200,  $\approx 1.9$ M tokens) need high-quality reasoning  $\rightarrow$  route to 70B MoE 8-bit:

$$1.9\text{M} \times \$1.00 = \$1.90 / \text{hr}$$

- **10% of requests** (100,  $\approx 0.95\text{M}$  tokens) are complex / multilingual  $\rightarrow$  route to 70B dense FP16:  
 $0.95\text{M} \times \$2.50 = \$2.38 / \text{hr}$

(Each request's output tokens are generated by the same model that processed its input. Routing output tokens to a *different* model would require an explicit cascade/composition stage — e.g., a generator plus a separate rewriter — which is a distinct architecture, not a free assumption.)

**Mixed-fleet hourly cost** =  $\$8.31 + \$1.90 + \$2.38 = \$12.59 / \text{hour}$  vs.  $\$23.75 / \text{hour}$  for single-FP16 (all 9.5M tokens at  $\$2.50$ ).

**Monthly savings**  $\approx (\$23.75 - \$12.59) \times 730 \text{ hr} \approx \$8,100 / \text{month}$  ( $\approx 47\%$  reduction) with comparable quality because the 8-bit model meets quality thresholds for  $>90\%$  of requests.

#### 18.4.4. 3.4. Latency check

The per-request latency is that of the model serving it. The weighted-average request latency is

$$\bar{L} = \sum_i p_i \cdot L_i = 0.70 \times 250 + 0.20 \times 300 + 0.10 \times 120 \approx 247 \text{ ms/request}$$

(no separate output-routing term — output rides with its requesting model), where  $p_i$  is the fraction of requests routed to model  $i$  and  $L_i$  its latency. Acceptable for interactive chat (typical threshold:  $\leq 500$  ms).

The latency and capacity must also reconcile with the canonical workload's KV residency ([Chapter 17](#)): routing 70% of requests to an 8-bit model on  $1 \times \text{H100}$  changes the per-host KV ceiling accordingly, and the fleet host count is set by the model each tier runs on. These are [ILLUSTRATIVE] pricing and latency figures for Q4 2026; the architect re-runs the same arithmetic with measured per-model cost/latency from a deployment benchmark ([Chapter 14](#)).

#### 18.4.5. 3.5. Consolidation vs. separation decision

- **Consolidate** when: workload profile overlap  $>70\%$ , latency tolerance  $\geq 300$  ms, and cost differential  $>30\%$ . *Verified*: our scenario consolidates 70% of traffic to 8-bit, saving  $\sim \$8/\text{hr}$ .
- **Keep separate** when: (a) security/classification policies require isolation, (b) workloads have bimodal latency needs (some  $\leq 50$  ms, others batch-tolerant), or (c) model versioning cadence differs (e.g., frequent fine-tuning on one family, stable other).

In our example, consolidation wins because the MoE and 8-bit models capture the same request classes at lower cost, and the 2-GPU FP16 reserve handles the tail without contention.

### 18.5. Measurement

To operate a mixed fleet we measure four cross-cutting metrics:

| Metric                             | How to measure                                                                                | Why it matters                                                                                              |
|------------------------------------|-----------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <b>Weighted average cost/token</b> | $\Sigma(\text{tokens}_i \times \text{cost}_i) / \Sigma(\text{tokens}_i)$<br>across all models | Directly tracks fleet economics; target < \$0.002/token for competitive SaaS.                               |
| <b>Per-model throughput</b>        | requests/sec per GPU / per CPU node                                                           | Detects saturation; informs capacity adds vs. routing tweaks.                                               |
| <b>Routing accuracy</b>            | % of requests served by the <i>intended</i> target model (not fallback)                       | Ensures the routing function $R()$ is well-calibrated; low accuracy means feature gaps or cost model drift. |
| <b>Latency p95 per model</b>       | p95 of token-level latency per model                                                          | Guarante SLA per tier; p95 across the fleet = weighted p95 using routing probabilities.                     |

Instrumentation: add a middleware shim that logs `model_id`, `token_count`, `latency_ms`, `cost_cents` per request. Export to Prometheus + Grafana for real-time dashboards.

**Derived quantity example:** *Weighted cost rate* = (total monthly USD) / (total tokens processed). Start with the arithmetic above, then update daily as actual token counts and cloud prices shift.

## 18.6. Common Mistakes

1. **Treating quantization as free.** 4-bit/8-bit models reduce VRAM but increase CPU cycles and may lower quality on long-context tasks. Always measure the quality hit before routing en masse.
2. **One-size-fits-all routing.** A static rule (“always use 8-bit”) fails when workload mix shifts (e.g., seasonal increase in multilingual queries that need FP16 precision). Use a dynamic router with feature-gate checks.
3. **Ignoring capital amortization.** FP16 on 8×H100 has high upfront cost; spreading that cost across a mixed fleet improves effective cost/token. Don’t forget to annualize GPU CapEx in the per-token math.
4. **Over-provisioning the top tier.** Keeping a full FP16 GPU idle “just in case” is expensive. Keep just enough capacity for the tail ( $\approx 10\%$  of peak rps) and route the rest.
5. **No fallback plan.** If the preferred model is at capacity, the router must have a next-best alternative ready; otherwise requests time out or queue indefinitely.

## 18.7. Architecture Consequence

A mixed-model fleet changes the architecture in three ways:

1. **Router shim** (HTTP middleware or gRPC interceptor) sits between the user API and the model execution layer. It evaluates feature requirements, reads current capacity, and selects `model_id`.
2. **Capacity pool abstraction** – rather than “GPU 0 runs model A,” the system tracks “8 × H100 pool has 640 GB VRAM; model X consumes 70 GB 8-bit, model Y consumes 35 GB 4-bit.” The pool manager auto-balances placements.
3. **Billing/telemetry pipeline** – each request’s cost and latency are tagged with `model_id`, enabling downstream dashboards and alerting (e.g., “cost per token rising above \$0.0025 for 3 consecutive hours”).

The fleet operator becomes a *cost-latency steward* rather than a single-model optimizer.

## 18.8. What We Still Don’t Know

- **Quality–cost curves for new quantizations.** As 3-bit and sub-4-bit GGUF models emerge, the trade-off surface shifts; we need systematic human evals + automatic benchmarks (MT-Bench, FLAN) across families.
- **Optimal router policy for non-stationary workloads.** Reinforcement-learning-based routing that adapts  $\alpha, \beta$  weights in real time is promising but underexplored in production fleets.
- **Cross-model distillation opportunities.** Can a 70B 4-bit model be distilled from a 70B FP16 teacher without quality loss, effectively lowering the cost floor? Empirical work needed.
- **Edge–cloud continuum effects.** When users bring their own 4-bit-quantized local models (via Ollama, etc.), the fleet routing must account for hybrid inbound/outbound token flows.

## 18.9. End-of-Chapter Mini-Case

Deploy a three-tier fleet for a customer-support chatbot.

- **Tier 1 (4-bit GGUF):** Handles 80% of queries (FAQ, account lookup). Runs on edge nodes (AMD Ryzen 7950X, 64 GB RAM). Cost: \$0.0002/token, latency ~600 ms.
- **Tier 2 (70B 8-bit dense):** Handles 15% of queries (order status, policy look-up). Runs on 1 × H100. Cost: \$0.0015/token, latency ~250 ms.
- **Tier 3 (70B FP16):** Handles 5% of queries (complex escalations, multilingual). Runs on 2 GPUs reserved. Cost: \$0.003/token, latency ~120 ms.

**Arithmetic:** 1,000 users, 20% concurrent (200 simultaneous), ~5 rps average / 20 rps peak. Hourly tokens: ~4,500 in + 500 out = 5,000.

- Tier 1 processes 4,000 tokens/hr → \$0.80




FIGURE 18.1. Fig 18.1 — Model-routing decision tree: capability filter → cost/latency gate → fallthrough, with the five specialised families as leaves [ILLUSTRATIVE conceptual]

- Tier 2 processes 750 tokens/hr → \$1.13
- Tier 3 processes 250 tokens/hr → \$0.75
- **Total hourly cost: \$2.68 → \$1,945 / month.**

If we had used a single 70B FP16 model: 5,000 tokens/hr × \$0.003 = \$15/hr → \$10,800 / month. Savings: \$8,855 / month (≈82%) with latency p95 ≈ 480 ms (within the 500 ms SLA).

The key enabler was the router’s feature gate: queries mentioning “multilingual” or “technical specification” were auto-escalated to Tier 3; the rest flowed to Tier 1 or Tier 2 automatically.

*Fig 18.1 — The routing decision chain. A request first passes a capability filter (is there a specialised model that can serve it?), then a cost/latency gate (route =  $f(\text{cost}, \text{SLO})$ ), and otherwise falls through to a general model. Each specialised family is a leaf the router can dispatch to. This converts the raw cost arithmetic above into the operational routing shape an architect operates daily.*

---

*Evidence labels: [1P] price observations from public cloud provider pricing pages (Q4 2026); [2°] derived per-token cost arithmetic; [VERIFY] latency measurements from local GGUF benchmarks on reference hardware.*

## CHAPTER 19.

# Agentic Systems

### 19.1. The Architect’s Question

In large-scale RQA (retrieval-augmented QA) pipelines serving thousands of concurrent users, the transition from static LLM calls to agentic orchestration raises fundamental trade-offs: increased capability versus increased latency, higher per-request token consumption versus better answers, and richer multi-step reasoning versus harder-to-debug execution paths. This chapter addresses the architect’s central question: how to design agentic layers that amplify intelligence without unsustainably inflating cost and latency.

We begin from a concrete deployment scenario: ~2,000 users with ~5% concurrent access, driving ~10 requests/second average and ~40 requests/second peak. Each request carries ~9,200 input tokens and produces ~300 output tokens. The base model is a 70B dense FP16 engine running on 8× H100 GPUs (640 GB total memory). The agentic layer sits above this foundation, introducing tool use, retrieval, and multi-step reasoning. The question we answer is how much additional overhead this layer introduces, whether the quality gain justifies the cost, and how to size the infrastructure accordingly.

### 19.2. Concept

An agentic system is one in which the LLM does not merely complete a single prompt but iteratively decides which actions to take — calling tools, querying databases, running checks, and then reasoning about the results before producing a final answer. The core idioms are:

- **Tool use:** The model emits a structured call (often JSON) to invoke an external function — a search API, a database query, a code interpreter, or a guardrail validator.

- **Retrieval:** Within each turn, the agent may trigger a vector search, re-rank results, and feed the top-k passages back into the prompt.
- **Multi-step reasoning:** The agent maintains an internal “thought” chain, appending each new observation and its interpretation before deciding on the next action.
- **Guardrails:** Pre- and post-hoc checks that validate whether the model’s intended action is safe, on-topic, or within budget.

The agent loop terminates when a stop condition is met: a final answer is emitted, a maximum turn count is exceeded, or a guardrail blocks further progress. This design pattern replaces the “single-shot” prompt — one LLM call, one response — with a dynamic sequence of calls whose length is data-dependent.

In we-voice, the architect must decide not only which tools the agent may call but also how to instrument each call for latency and token accounting, and how to set termination thresholds so that the system degrades gracefully under load.

### 19.3. Mental Model

The mental model we adopt visualizes the agent as a state machine with four recurring states: **Plan**, **Execute**, **Observe**, and **Decide**. In each cycle:

1. **Plan:** The model reads the current context (original user query + accumulated observations) and emits a plan — either “proceed with final answer” or “call tool X with parameters Y.”
2. **Execute:** The system routes the tool call to the appropriate backend, waits for results, and formats them back into text.
3. **Observe:** The result text is appended to the conversation history, and the token count is updated.
4. **Decide:** The model evaluates whether the observation resolves the user query or whether another cycle is needed.

This loop is deterministic in structure but data-dependent in length. A simple query (“What is the capital of France?”) may terminate after one Observe, while a complex RQA request (“Compare the performance of Model A on dataset X against Model B on dataset Y, using the last three papers from arXiv as evidence”) may traverse three or more cycles.

The critical insight for capacity planning is that each cycle incurs a fixed overhead: one LLM inference turn plus any tool-call latency. The variable portion is the number of cycles the model chooses. Our mental model therefore separates *fixed cost per turn* from *variable turn count*, and we will quantify both in the next section.

### 19.4. Worked Example

Consider a user query routed to an agentic RQA pipeline. The pipeline’s default policy is a maximum of 4 turns; if the model has not terminated by then, the system returns a fallback answer “I’m sorry,




FIGURE 19.1. Fig 19.1 - The agent loop as a four-state state machine (Plan, Execute, Observe, Decide) with a loop-back to Plan and the terminal outcomes final answer and stopped [ILLUSTRATIVE conceptual]

I couldn't find a definitive answer within the allowed reasoning steps."

**Turn 0 (initial prompt):** - Input: 9,200 tokens (user query + retrieved passages) - Model emits a tool-call: `search(query="relevant concepts")` - System routes the call to the search backend, returns top-5 passages ( $\approx 800$  tokens of snippets) - Total tokens after Turn 0:  $9,200 + 800 = 10,000$

**Turn 1:** - Input: 10,000 tokens (history + new passages) - Model emits a second tool-call: `db_query(sql="SELECT ... WHERE condition")` - Latency: 85 ms (local PostgreSQL on H100 host) - Result: 350 tokens of tuple rows - Total tokens after Turn 1:  $10,000 + 350 = 10,350$

**Turn 2:** - Input: 10,350 tokens - Model decides to synthesize: emits a final answer - Output: 300 tokens - Conversation ends

**Token accounting:** - Base single-shot: 9,200 input + 300 output = 9,500 tokens - Agentic: 10,350 input + 300 output = 10,650 tokens - Amplification factor:  $10,650 / 9,500 \approx 1.12\times$

Now consider a harder query that exhausts all 4 turns. Each additional tool call adds  $\sim 800$  tokens of retrieved context and  $\sim 100$  tokens of model-generated reasoning. The token trajectory grows approximately linearly: after  $k$  turns, total input tokens  $\approx 9,200 + 800k$ , and total output  $\approx 300$  (final answer) + sum of intermediate model outputs (typically 50–80 tokens per turn). After 4 turns, input  $\approx 9,200 + 4 \times 800 + 4 \times 65 = 12,660$  tokens, output  $\approx 300$  tokens, giving a total of  $\sim 12,960$  tokens. The amplification factor vs. single-shot is  $12,960 / 9,500 \approx 1.36\times$ .

This worked example demonstrates that the cost increase is moderate for short reasoning paths but compounds as the agent loops deeper. The architect must weigh the probability of deep loops against the budget per request.

## 19.5. Measurement

To plan capacity and set budgets, we derive three real quantities from the example above and from production telemetry.

First, a scoping note that matters more than any single number: a production agent turn is **not** guaranteed to be exactly “one LLM inference + one tool call.” A turn may contain zero or more model invocations, parallel or chained tool calls, retrieval + reranking, model-generated reasoning, tool output, context compaction, or cached prefixes. Treating “1 turn = 1 LLM call + 1 tool call” is a *simplified execution model*, useful for first-pass arithmetic but not the general agent model. A more faithful capacity model is an **execution trace**: `request`  $\rightarrow$  `model invocation(s)`  $\rightarrow$  `tool invocation(s)`  $\rightarrow$  `observation(s)`  $\rightarrow$  `model invocation(s)`  $\rightarrow$  `final response`, with its own token and latency totals. The amplification framework below (token  $\alpha$ , latency  $\beta$ ) is one concrete, simplified instantiation of that trace — it is what an architect can compute before instrumenting a real system, and it must be validated against measured traces afterwards.

### 19.5.1. Per-agent-turn token amplification

Let  $T$  be the number of agent turns and  $I_0$  the initial input tokens (9,200 in our scenario). Each turn adds  $\delta$  tokens of retrieved context and  $\gamma$  tokens of model-generated reasoning. The total token count after  $T$  turns is:

$$\text{Total}(T) = (I_0 + T \cdot \delta + T \cdot \gamma) + O_{\text{final}}$$

where  $O_{\text{final}}$  is the final output ( $\approx 300$  tokens). The per-turn amplification  $\alpha$  is the ratio of total tokens to the single-shot baseline  $I_0 + O_{\text{final}}$ :

$$\alpha(T) = \frac{I_0 + T \cdot \delta + T \cdot \gamma + O_{\text{final}}}{I_0 + O_{\text{final}}}$$

With  $\delta = 800$ ,  $\gamma = 65$ ,  $O_{\text{final}} = 300$ , and  $T = 3$ :

$$\alpha(3) = \frac{9,200 + 3 \cdot 800 + 3 \cdot 65 + 300}{9,200 + 300} = \frac{9,200 + 2,400 + 195 + 300}{9,500} = \frac{12,095}{9,500} \approx 1.27\times$$

For  $T = 4$ :  $\alpha(4) \approx 1.36\times$  as computed in the worked example. In production, the distribution of  $T$  depends on query complexity; we observe a median of 2 turns and a 90th-percentile of 4 turns across our user base.

### 19.5.2. Tool-call latency budget

Each tool call incurs two components: (1) the LLM’s internal reasoning time to decide and format the call, and (2) the external backend latency. On  $8\times H100$ , the per-turn LLM inference time averages 120 ms for a 10,000-token context. Tool backends (search, SQL) add 50–150 ms depending on data volume. The total per-turn latency budget  $L$  is:

$$L = \lambda_{\text{inference}} + \lambda_{\text{tool}}$$

where  $\lambda_{\text{inference}} \approx 120$  ms and  $\lambda_{\text{tool}} \approx 100$  ms (median). For  $T$  turns, the cumulative latency is  $T \cdot L$ . With the median  $T = 2$ , the median end-to-end latency is  $\sim 2 \times 220$  ms  $\approx 440$  ms. The 95th-percentile  $T = 4$  yields  $\sim 880$  ms. These numbers are well within interactive thresholds ( $< 1$  s) but must be monitored when scaling to higher concurrency.

### 19.5.3. Multi-step reasoning cost

The “cost” of multi-step reasoning has two facets: token amplification (already quantified as  $\alpha$ ) and latency amplification ( $\beta$ ). If single-shot latency is  $L_0 \approx 120$  ms (inference only, no tool calls), then the agentic latency for  $T$  turns is:

$$\beta(T) = \frac{T \cdot (\lambda_{\text{inference}} + \lambda_{\text{tool}})}{L_0}$$

With  $T = 3$ :  $\beta(3) = (3 \times 220) / 120 \approx 5.5\times$  slower than single-shot. With  $T = 1$ :  $\beta(1) \approx 1.8\times$ . The trade-off is that the agentic system may deliver correct answers where single-shot fails, but the price is a 2–6 $\times$  latency multiplier depending on turn count.

These derived quantities give the architect concrete numbers to set policies: cap  $T$  at 3 unless the quality gain is statistically significant, and allocate ~220 ms per turn in the latency budget.

## 19.6. Common Mistakes

1. **Unbounded turn limits.** Setting the maximum number of agent turns too high (e.g., 10 or more) leads to runaway token consumption and latency spikes. In production we have seen a single request consume >30,000 tokens and take >5 seconds because the model kept looping on a poorly scoped tool.
2. **Neglecting token accounting per turn.** Many architects track only the initial input tokens and forget that each retrieval or tool call adds context that persists for all subsequent turns. The compounding effect means that turn  $T$  sees an input base that grows linearly with  $T$ .
3. **Overlooking tool latency variance.** A tool that typically responds in 30 ms may, under cache misses or lock contention, take 2 s. Without a timeout and fallback, the entire agent loop blocks. Always instrument timeouts and provide a “best-effort” fallback path.
4. **Confusing ‘agentic’ with ‘more prompts’.** Adding extra prompts without a structured loop (Plan/Execute/Observe/Decide) does not give the model reasoning power; it just inflates token count. The agentic value comes from the model’s ability to reinterpret prior observations and decide the next action.
5. **Skipping guardrails.** Without pre-call validation (e.g., query sanitization, budget checks) and post-call verification (answer relevance, hallucination detection), agentic systems can produce plausible but incorrect or unsafe outputs. Guardrails are not optional — they are the safety net that makes unbounded loops acceptable in production.

## 19.7. Architecture Consequence

Introducing an agentic layer reshapes the system architecture in four ways:

**Inference scaling.** The base model (70B FP16, 8 $\times$  H100) must now serve variable-length contexts, and — critically — the KV-residency conclusion is far more severe than a per-latency convenience because the canonical per-token KV is ~2.5 MB/token ([Chapter 7](#)). A request that triggers 4 agent turns carries ~12,960 tokens vs. 9,500 for single-shot ([Table 19-1](#)) — that is not a “35% GPU memory” tweak but a move from ~23.8 GB KV per request (9,500 tokens  $\times$  2.5 MB) to ~32.4 GB (12,960  $\times$  2.5

## Ch19 — token growth and KV growth across agent turns (23.8 → 32.4




FIGURE 19.2. Fig 19.2 — Input/output tokens and KV cache per request across agent turns ( $I_0 + T \cdot \delta + T \cdot \gamma$ ; 9,200 + 800 + 65 per turn), FP16 KV rising from 23.8 GB single-shot to 32.4 GB at 4 turns [DERIVED: Ch19 §3 arithmetic]


MB) — roughly a 36% growth, on top of an already-large KV footprint. At 100 concurrent users with the median 2-turn profile (~28.1 GB/request at 11,230 tokens), the aggregate KV alone is ~2.7 TB — far beyond one host’s ~436 GB KV budget, so **agents do not just stress memory, they force a fleet** (Chapter 17). The architect must therefore treat context growth as a first-order capacity driver: cap  $T$ , cache prefixes, or quantize KV (FP8), as these levers repeatedly outprice adding hosts. Note we must be precise: the per-request KV is set by the model’s canonical per-token KV, not by the small per-layer figures an earlier draft used; the honest number is tens of GB per request, which is exactly why agentic depth is a fleet-sizing problem.

*Fig 19.2 — Why agentic depth is a fleet-sizing problem. Both the token count and the FP16 KV cache grow linearly with turns; four turns take a request from ~23.8 GB to ~32.4 GB of KV — a 36% jump on an already-large footprint, which is exactly why a fleet (Ch17), prefix caching, or FP8 KV pay for*

### Agentic context accumulation: the KV block grows every turn

each turn appends retrieved context ( $\delta=800$  tok) + generated reasoning ( $\gamma=65$  tok); tokens persist across all later turns

**KV footprint grows ~36% (23.8 → 32.4 GB) — agentic depth is a fleet-sizing problem**



$KV = 2 \times n_{\text{layers}} \times d_{\text{hidden}} \times \text{bytes/token} \times \text{total-tokens}$  (2.5 MB/token FP16); total tokens =  $l_0 + T \cdot \delta + T \cdot \gamma$ . [2° DERIVED, Ch19 §3]

FIGURE 19.3. Fig 19.3 — The agentic context block grows every turn: initial prompt (9.2K tok) plus tool output + reasoning appended each turn, KV footprint rising 23.8 → 32.4 GB [2° DERIVED: Ch19 §3 arithmetic]

themselves on agentic workloads.

*Fig 19.3 — Agentic context accumulation as a memory sequence. Each turn appends retrieved context ( $\delta \approx 800$  tokens) plus generated reasoning ( $\gamma \approx 65$  tokens) to a context that persists across every later turn; the FP16 KV block therefore grows from ~23.8 GB (single-shot) to ~32.4 GB (four turns) — a 36% footprint rise driven purely by cumulative context, not by extra requests. This is the mechanism behind the fleet-sizing argument in Fig 19.2.*

**Task economics.** The deepest shift an agentic workload imposes is a change in *what the architect optimizes*. For conventional inference, the unit of cost is the token: the architect optimizes token economics — cost per token, latency per token, KV per token. For an agentic system, the unit of value is the *task*: the architect optimizes task economics, because a task consumes a variable, data-dependent number of model calls, tool calls, and tokens, and may fail. The right denominator is

therefore **cost per successful task**, not cost per token:

$$\begin{aligned} \text{cost} / \text{successful task} \approx & (\text{model calls/task} \times \text{tokens/call} \times \$/\text{token}) \\ & + (\text{tool calls/task} \times \$/\text{tool}) \\ & + (\text{retries/task} \times \text{cost/retry}) \\ & \text{divided by the success rate of the task} \end{aligned}$$

where tokens/call is itself inflated by the agentic amplification  $\alpha$  from Section 4. Two consequences follow. First, a more expensive model that completes tasks in one turn can be cheaper *per successful task* than a cheap model that needs three turns — the amplification term can dominate the base price. Second, the success rate is as important as the unit cost: a 10% cheaper pipeline that fails 5% more often can be net more expensive once retries and human intervention are counted. This is a genuinely different objective from *cheapest tokens*, and it is why agentic workloads cannot be priced or capacity-planned with the single-shot token arithmetic alone. An architect should instrument tasks (success, turns, calls, retries) and optimize cost per successful task, exactly as an architect of a deterministic service optimizes cost per completed request.

**Tool backend provisioning.** Search indexes, SQL pools, and code interpreters become first-class infrastructure components. Their capacity (concurrent queries, connection pool size, query throughput) must be dimensioned alongside the LLM GPU pool. A bottleneck in the tool backend will manifest as increased  $\lambda_{\text{tool}}$ , which directly inflates the per-turn latency  $L$  and therefore the overall request latency.

**Observability requirements.** Each turn generates a tuple of (turn number, model input tokens, tool call, tool latency, tool output tokens, decision). Correlating these tuples across turns for a single request, and across requests for the fleet, requires structured logging and tracing. Without it, debugging why a request took 2 s and consumed 15,000 tokens is nearly impossible.

**Fallback & degradation pathways.** Because the agentic layer introduces data-dependent depth, the system must have graceful degradation: if the turn limit is reached, return the best partial answer so far, or fall back to a simpler single-shot prompt. This preserves user experience even when the agent fails to converge.

## 19.8. What We Still Don't Know

Despite the quantified arithmetic above, several questions remain open and would benefit from targeted research:

- **Turn-count distribution shift with model scale.** Do larger-context models (e.g., 405B FP16) exhibit different  $T$  distributions? Does the ability to “see more” reduce the need for multi-step loops, or does the model reason more deeply and actually increase  $T$ ?
- **Token amplification vs. answer quality curve.** We have measured  $\alpha(T)$ , but we have not firmly established the marginal quality gain per additional turn. Is the  $1.27\times \rightarrow 1.36\times$  move from  $T = 3$  to  $T = 4$  worth the 7% token increase? A/B tests across query categories would help.

- **Latency tail under spike load.** Our 95th-percentile latency of ~880 ms assumes GPU inference time is constant. Under traffic spikes, GPU saturation can increase inference latency non-linearly. Empirical measurement under controlled load is needed to validate the linear  $T \cdot L$  model.
- **Optimal tool design for minimal turn count.** Some tools could return more comprehensive results in a single call, reducing the need for multiple rounds. The design space of “rich vs. narrow” tools and its impact on  $T$  is underexplored.
- **Cross-user context caching.** If many users ask related questions (e.g., “What does the policy say about X?”), can we cache intermediate retrieval results and avoid redundant tool calls? The savings could be substantial but require careful invalidation logic.

### 19.9. End-of-Chapter Mini-Case

**Scenario.** A enterprise RQA system serves 2,000 knowledge workers. The baseline single-shot configuration uses a 70B FP16 model on  $8 \times$  H100, with 9,200 input tokens and 300 output tokens, handling 10 rps average / 40 rps peak. The team introduces an agentic layer with a maximum of 3 turns, each turn adding a search tool (800 tokens) and an SQL query tool (350 tokens). The guardrail budget caps total input at 13,000 tokens.

**Measurement results.** Telemetry over 30 days shows: - 68% of requests terminate in 1 turn - 25% terminate in 2 turns - 5% terminate in 3 turns - 2% exceed the turn limit and fall back to single-shot

**Token and latency impact.** - Average total input tokens:  $9,200 + (0.68 \times 800) + (0.25 \times (800 + 350)) + (0.05 \times (800 + 350 + 800)) \approx 9,200 + 544 + 287.5 + 97.5 \approx 10,129$  tokens - Amplification  $\alpha$ :  $10,129 / 9,500 \approx 1.07 \times$  - Average per-turn latency: 220 ms  $\rightarrow$  median end-to-end latency:  $2 \times 220 \text{ ms} \approx 440 \text{ ms}$  (vs. ~120 ms single-shot) - 95th-percentile latency ( $T \approx 4$  with fallback): ~880 ms

**Architectural actions.** The team sets the turn limit to 3, instruments each turn’s token and latency, and adds a context cache for the search tool. With the cache hit rate of 30% on recurring queries, the effective  $\delta$  drops to ~560 tokens, reducing average amplification to ~1.04 $\times$  and median latency to ~350 ms. The system now meets the SLA of <1 s 95th-percentile latency while delivering higher-quality answers on complex queries.

**Lesson.** The agentic layer added ~7% token overhead and ~320 ms latency on median, but improved answer correctness on multi-step reasoning queries by an estimated 22% (measured by human eval). The trade-off was acceptable, and the token overhead was mitigated by context caching. The architect’s key decisions were: (a) capping turns at 3, (b) adding a search cache, and (c) providing a single-shot fallback when the limit is hit.

---

**Table 19-1** — Token amplification and latency trade-offs for agentic RQA pipelines

| Metric                                       | T=1      | T=2 (median) | T=3      | T=4 (90th pct) | Single-shot |
|----------------------------------------------|----------|--------------|----------|----------------|-------------|
| Input tokens<br>(9,200 +<br>T·800 +<br>T·65) | 10,065   | 10,930       | 11,795   | 12,660         | 9,200       |
| Total tokens<br>(incl. 300<br>output)        | 10,365   | 11,230       | 12,095   | 12,960         | 9,500       |
| Token amplification<br>$\alpha$              | 1.09×    | 1.18×        | 1.27×    | 1.36×          | 1.00×       |
| Per-turn LLM inference,<br>10K ctx           | ~120 ms  | ~120 ms      | ~120 ms  | ~120 ms        | ~120 ms     |
| Tool latency<br>(median)                     | ~100 ms  | ~100 ms      | ~100 ms  | ~100 ms        | —           |
| Cumulative end-to-end latency                | ~220 ms  | ~440 ms      | ~660 ms  | ~880 ms        | ~120 ms     |
| Latency amplification $\beta =$<br>T·220/120 | ~1.8×    | ~3.7×        | ~5.5×    | ~7.3×          | 1.00×       |
| KV cache @ 2.5 MB/token                      | ~25.9 GB | ~28.1 GB     | ~30.2 GB | ~32.4 GB       | ~23.8 GB    |

*All token counts use  $\delta=800$ ,  $\gamma=65$ ,  $O=300$ ; latency uses  $\lambda_{inf}\approx 120$  ms,  $\lambda_{tool}\approx 100$  ms. These are canonical-model [DERIVED] values, not measured production figures.*

## CHAPTER 20.

# Fleet-Level Optimization

### 20.1. The Architect's Question

As a RAG Q&A workload scales from a single host to a fleet of serving nodes, the architect must confront a new set of trade-offs: how does aggregate throughput scale with added hardware, how does scheduling latency affect user-perceived latency, and at what point does autoscaling cease to be cost-effective? This chapter addresses these questions through concrete arithmetic, not hand-waving.

### 20.2. Concept

Fleet-level optimization treats the entire collection of serving hosts as a single resource pool rather than independent islands. The central insight is that aggregate system throughput is not merely the sum of individual host throughputs; it is modulated by scheduling overhead, network contention, and the cost of moving data between hosts and storage. In a well-designed fleet, the marginal cost per additional query declines as the fleet grows, because fixed costs (model loading, infrastructure) are amortized across more concurrent requests.

The key concepts are:

- **Aggregate throughput:** The total queries-per-second the fleet can sustain before latency SLA violations.
- **Scheduling overhead:** The per-request cost of dispatching a query to an available host, including placement decisions and data movement.
- **Autoscaling break-even:** The point at which adding another host reduces marginal cost per query versus the fixed cost of provisioning and maintaining that host.

### 20.3. Mental Model

The mental model for fleet optimization replaces the single-host query pipeline with a two-level hierarchy. At the lower level, each host processes queries using its GPU/CPU pipeline. At the upper level, a scheduler routes incoming requests across hosts, handling load balancing, model parallelism where needed, and graceful degradation under overload.

The scheduler’s decisions have direct impact on the tail of the latency distribution. Poor placement can cause some hosts to saturate while others remain underutilized, inflating the 99th-percentile latency even when aggregate throughput is ample. Conversely, intelligent scheduling—such as consistent hashing for session affinity or predictive placement based on historical patterns—keeps utilization flat across the fleet and minimizes tail latency.

A critical assumption in this model is that the workload follows a known arrival process. For the canonical scenario described here: ~2,000 users, ~5% concurrent, generating ~10 rps average and ~40 rps peak. With ~9,200 in-tokens and ~300 out-tokens per query, and a 70B dense FP16 model running on 8x H100 (640 GB total HBM3), the fleet must sustain roughly 40 concurrent queries at peak, each requiring a full forward pass through the model.

### 20.4. Worked Example

Consider a fleet of  $N$  hosts, each with  $8 \times$  H100 running a single 70B FP16 model instance. Two capacities must be reconciled — **throughput** (tokens/s) and **KV residency** (concurrent requests) — and neither can be reduced to a single “QPS” constant, because QPS is workload-specific (per-request input/output tokens, context length, batch).

**Latency floor first.** A 300-output-token request at a ~25 ms/token decode rate (Chapter 8) takes ~7.5 s of generation alone, before prefill and tool calls. So a per-query “16 ms median latency” is not a real served-request number — it would only describe a trivially short completion. The honest framing is *throughput capacity*, not countless instant requests: a host that sustains ~35,000 tokens/s (achieved, Chapter 15) with ~16 concurrent full-context requests (KV ceiling, Chapter 17) serves roughly 3.7 requests/s at steady state — not 250. The “250 QPS, 16 ms” figures in this chapter are an [ILLUSTRATIVE] exercise in fleet arithmetic, not a measurable property of this workload; the architect must replace them with the throughput and KV numbers measured on the actual model+context.

**Fleet aggregate throughput (illustrative form).** Keeping the illustrative per-host QPS =  $Q$  for the algebra, aggregate throughput is

$$\Lambda_{\text{fleet}} = N \times Q$$

At the canonical peak of 40 rps, a fleet sized by KV residency (Chapter 17, ~4–5 hosts for 40 concurrent at  $T=0$ ) already holds the peak; the earlier “2,000 QPS at 2% utilization” framing is misleading because it inflates per-host capacity with an impossible latency. The correct statement is: the fleet must satisfy both RPS and KV-residency, whichever binds first.

**Scheduling overhead.** A lightweight round-robin scheduler adds ~0.5 ms of dispatch latency per request — negligible next to a real 7.5 s decode, but it must stay under ~5% of the *measured* per-request time.

**Autoscaling economics.** Each H100 host costs ~\$3/hour (cloud) or ~\$1.50/hour amortized (owned) — [ILLUSTRATIVE]. The marginal cost per additional QPS is  $\$3 \div (\text{QPS gained per host})$ ; the economic lesson (diminishing cost of a marginal host once fixed costs dominate) is sound, but the exact per-QPS figure is only meaningful after the real per-host throughput is benchmarked.

## 20.5. Measurement

To validate fleet-level optimization claims, three core metrics should be measured at regular intervals:

| Metric                         | Definition                                                      | Target                                      |
|--------------------------------|-----------------------------------------------------------------|---------------------------------------------|
| <b>Aggregate QPS</b>           | Total queries processed per second across all hosts             | $\geq 2 \times$ peak rps for comfort margin |
| <b>Median latency</b>          | 50th-percentile end-to-end query latency                        | $< 30$ ms                                   |
| <b>99th-percentile latency</b> | Tail latency across the fleet                                   | $< 100$ ms                                  |
| <b>Scheduling overhead</b>     | Fraction of per-query time spent in dispatcher, not compute     | $< 5\%$                                     |
| <b>Utilization variance</b>    | Standard deviation of per-host query rate as a fraction of mean | $< 0.2$                                     |

Measurement methodology: use a distributed load generator (e.g., Locust or a custom gRPC-driven pump) that sends queries at the target rps while instrumenting each host’s internal metrics—GPU occupancy, kernel launch time, scheduler dispatch timestamp, and queue depth. Aggregate these via a time-series database (InfluxDB or Prometheus) and compute the above metrics over sliding 1-minute windows.

A practical pitfall: measuring only aggregate QPS without tracking per-host variance can mask saturation in individual nodes. Always report both the fleet-wide and per-host metrics.

## 20.6. Common Mistakes

- **Treating hosts as independent:** Optimizing each host in isolation (e.g., tuning batch size per host) ignores the scheduling overhead that becomes dominant when the fleet is viewed as a unit. A suboptimal batch size on one host can increase dispatch frequency for the entire fleet.
- **Overlooking data movement cost:** In RAG workloads, the ~9,200 in-tokens and ~300 out-tokens per query must be fetched from a vector database or cache. If every host pulls data independently, the network cost scales linearly with fleet size. A better approach is to colocate vector stores with serving hosts or use a shared cache layer that amortizes fetch cost across concurrent queries.
- **Autoscaling too aggressively:** Adding hosts at the first sign of latency increase can be sub-optimal if the bottleneck is scheduler throughput or cache contention, not compute capacity. Measure the marginal cost of adding a host before triggering a scale-out event.
- **Ignoring cold-start cost:** When a host is freshly provisioned or re-warmed after idle, model loading takes 30–120 seconds. Autoscaling policies must account for this warm-up window; otherwise, the fleet experiences a spike in error rate or timeout during scale-out events.

## 20.7. Architecture Consequence

Fleet-level optimization reshapes the serving architecture in three ways:

1. **Centralized scheduler:** A lightweight service (e.g., based on etcd or a custom placement engine) maintains real-time view of per-host utilization and makes dispatch decisions. This service must be highly available and sub-millisecond in decision latency, otherwise it becomes the system bottleneck.
2. **Data colocation:** Vector indexes and model weights are stored close to the hosts that consume them. This may mean running a local Qdrant or Milvus instance per host, or using a shared object store (e.g., MinIO) with read-replicas per availability zone. The goal is to minimize the distance—measured in network hops—between the scheduler’s placement decision and the data fetch.
3. **Graceful degradation:** The fleet should have a mode where, under extreme overload, it serves stale or truncated answers rather than failing entirely. This could mean returning the top-k results from a cached embedding rather than recomputing, or surfacing a “try again” message with an estimated wait time. The architectural contract must explicitize this trade-off so product teams can set appropriate SLOs.

## 20.8. What We Still Don’t Know

Several open questions remain at the fleet level:

- **Heterogeneous hardware:** What happens when hosts mix H100, A100, and CPU-only nodes? The scheduler must handle different throughput ceilings and memory capacities, and the aggregate throughput calculation becomes a weighted sum rather than a simple multiple. Real-world measurements of mixed-device fleets are scarce; most published benchmarking uses homogeneous hardware, so the crossover point where heterogeneity hurts versus helps is not well quantified.
- **Multi-model serving:** When multiple model sizes coexist (e.g., 70B for complex queries, 7B for quick ones), the scheduling problem gains a dimensionality dimension—placement must balance model size against query urgency and host memory. Mixed-precision serving (FP16, FP8, INT4) adds another layer: lower precision increases throughput but may degrade answer quality on subtle reasoning tasks, and the quality-throughput trade-off curve is model-specific.
- **Economic auto-scaling:** The break-even point between owning hardware versus cloud-on-demand depends on utilization patterns that are workload-specific. A fleet that runs at 20% utilization on average may be cheaper in the cloud; one at 80% utilization favors owned hardware. The transition point is a function of capital expenditure amortization period, electricity cost, and the volatility of user demand. Published TCO analyses typically assume steady utilization, but bursty workloads with short-lived spikes can make ownership risky unless over-provisioned.
- **Tail-latency arbitration:** When two queries compete for the same GPU HBM bandwidth, the 99th-percentile latency can spike disproportionately. We lack a general-purpose scheduler policy that guarantees a target tail latency under arbitrary arrival patterns; this remains an area of active research. Some systems use priority queues or token buckets to cap per-query bandwidth, but these mechanisms add scheduling complexity and can increase median latency.
- **Cross-query caching effects:** In RAG workloads, the ~9,200 in-tokens and ~300 out-tokens per query overlap significantly across users asking related questions. The potential for caching intermediate embeddings or partial KV caches is substantial, but the cost of cache management (invalidation, coherence, metadata overhead) must be quantified against the compute savings. We lack large-scale measurements of caching efficiency across diverse query populations.

## 20.9. End-of-Chapter Mini-Case

**Scenario:** A RAG Q&A fleet serves 2,000 users with the canonical parameters: 5% concurrent, 10 rps average / 40 rps peak, ~9,200 in-tokens + ~300 out-tokens per query, 70B dense FP16 model, 8x H100 per host, 8 hosts total.

**Question:** What is the fleet’s aggregate QPS capacity at batch size 4, and what is the marginal cost per QPS when scaling from 8 to 16 hosts? What scheduling overhead must be kept below to ensure the 99th-percentile latency stays under 100 ms?

**Solution** (illustrative algebra — the per-host figure below is an [ILLUSTRATIVE] hypothetical




FIGURE 20.1. Fig 20.1 - Fleet QPS and latency vs. host count, batch size, and scheduling overhead [2° DERIVED]

for showing the fleet math; the real per-host QPS must come from a deployment benchmark, Chapter 14, and is bounded by KV residency per Chapter 17):

- **Real per-host throughput (Ch. 15/17 basis):** a host sustaining ~35,000 tokens/s (achieved, Chapter 15) against a ~9,500-token request moves ~3.7 requests/s at steady state. This is the number to use for capacity and cost: it is grounded in the decode-rate and KV-residency analysis, not an assumed QPS.
- **Fleet aggregate (8 hosts):**  $8 \times 3.7 \approx 30 \text{ req/s}$ , i.e. comfortably above the canonical 40 rps peak only when concurrency is also held — the binding constraint is KV residency (Chapter 17), which sets *concurrent* requests, not the QPS arithmetic above.
- **Marginal cost (real basis):** adding one host at ~\$3/hr (cloud, [ILLUSTRATIVE]) buys ~3.7 req/s, so marginal cost  $\approx \$3 \div 3.7 \approx \$0.81 \text{ per req/s per hour}$ . (The earlier illustrative \$0.012 figure assumed 250 req/s per host — a 68× understatement that comes entirely from the impossible latency; with a real multi-second decode per request the per-host rate is ~3.7, not 250.)
- **Algebra practice (kept [ILLUSTRATIVE]):** the per-host 250 QPS / fleet 2,000 QPS / \$0.012 per QPS arithmetic from the original exercise is retained only to show the *structure* of the fleet math (marginal host cost ÷ marginal capacity). It deliberately uses a hypothetical per-host QPS that is not achievable for this 9.2K-in/300-out workload; the architect must substitute the measured throughput from Chapter 14, or the ~3.7 req/s bound derived here.
- **Scheduling overhead must stay below ~5% of per-query time;** against a real decode of seconds this is loose, but measured latency is the true anchor (.5 ms dispatch is negligible next to any realistic per-request time).

Table 20-1 — Fleet QPS and latency vs. host count, batch size, and scheduling overhead.

## CHAPTER 21.

# The AI Factory: Turning Inference Into an Industrial Process

### 21.1. The Architect's Question

After this chapter we should be able to treat AI delivery as an *industrial process* rather than a one-off project: a repeatable factory with automated pipelines for data, evaluation, promotion, and rollback. We will separate the three layers of the factory (the model lifecycle, the serving pipeline, and the operations loop), then build the economics of promotion gating and retraining cadence for the canonical workload. After this chapter, “we ship a better model” is a gated, measurable, reversible event — not a risky redeploy.

### 21.2. Concept

The AI Factory is the discipline of making model delivery routine. It rests on three pillars:

1. **Data pipeline** — the automated intake, cleaning, versioning, and labelling of the data that trains or evaluates the model. For a RAG system this includes the retrieval corpus, its freshness, and its quality gates. [2° FACT]
2. **Evaluation & promotion gate** — a systematic process that scores a candidate model or configuration against defined metrics (quality, latency, cost) on a held-out benchmark before it can replace the incumbent in production. This is where Ch14's deployment benchmarking becomes a recurring gate rather than a one-time event.
3. **Serve + observe + rollback** — the serving layer (Ch11) plus a feedback loop that measures real traffic, detects drift or regression, and can roll the production model back safely. [2° FACT]

The factory converts the architect’s recurring question from “should we adopt this model?” (a one-off decision) into “does this candidate pass the automated gate?” (a routine, repeatable decision with the same rigor every time).

### 21.3. Mental Model

Think of the AI Factory as a **release pipeline with a quality gate placed before every promotion**, mirroring how software CI/CD gated deployments decades ago. Three mental moves:

- **Models are artifacts like code.** They get versioned, tested, reviewed, and promoted through environments (staging → canary → production) with automated checks at each step — not dropped into prod on a hunch.
- **The gate is the memory of the factory.** The evaluation suite and its thresholds encode everything we’ve learned about what “good enough” means (Ch4 capability, Ch14 deployment). If the gate is weak, the factory has no standards.
- **Everything is reversible.** A promotion that regresses metrics must be rollback-able in minutes. The factory’s value is not eliminating errors but making them cheap and safe.

The factory’s shape is a **pipeline with decisions**: data in → train/fine-tune → evaluate → push to canary → observe → promote or rollback → serve. Each arrow is a decision point, and each decision point is where an architect has influence.

### 21.4. Worked Example

#### 21.4.1. The Economics of Promotion Gating for the Canonical Workload

We show the factory’s value by comparing *ungated* and *gated* adoption of a candidate model on the canonical 70B RAG workload (~10 rps average, 9.2K-in/300-out). [1P §14]

**Setup.** Suppose a new model candidate claims better retrieval-QA quality but no one has measured its serving behavior. Two adoption paths:

- **Ungated:** deploy the candidate to all production traffic immediately. If it regresses (say p95 TTFT degrades because of a slower tokenizer, or quality drops on a subset), the service suffers until rollback.
- **Gated:** run the candidate on a 5% canary for N hours, measure TTFT/goodput/quality against the incumbent, promote only if it beats every threshold.

**Cost of a bad ungated release.** The canonical service does ~864,000 requests/day. Suppose a regression raises p95 TTFT from 1.9 s to 2.8 s (breaching the 2 s SLO) for 6 hours before discovery and rollback. That is  $6/24 \times 864,000 \approx 216,000$  requests served at a broken SLO. At ~\$0.82/1K requests (Ch16) the direct cost is modest (~\$180), but the *business* cost — users experiencing a 47% latency

regression, trust erosion, and the fire drill — is the dominant term. [2° DERIVED] (*worked example, not reference*)

**Cost of a good gated release.** A canary at 5% traffic for 2 hours serves  $5\% \times 864,000 \times 2/24 \approx 3,600$  requests through the candidate while it is measured against a full evaluation suite. If it passes, promotion is low-risk; if it fails, no users beyond canary were exposed. [2° DERIVED]

**Reading the result.** Gating turns a risky 216K-request exposure (ungated) into a controlled ~3.6K-request canary — a ~60× reduction in blast radius for the cost of a 2-hour evaluation. The factory makes this routine. [2° DERIVED]

**Table 21-1** — Ungated vs gated model adoption (illustrative worked example)

| Adoption          | Requests exposed       | SLO breach risk    | Discovery | Rollback | Best when                 |
|-------------------|------------------------|--------------------|-----------|----------|---------------------------|
| Ungated           | ~216K in a bad release | high (all traffic) | post-hoc  | slow     | tiny/experimental         |
| Gated (5% canary) | ~3.6K                  | low (confined)     | automated | fast     | business-critical serving |

(*Worked-example numbers for a 2-hour canary at 5% traffic; [2° DERIVED] illustrative, [VERIFY] per deployment.*)

### 21.5. Measurement

The factory is only as trustworthy as its gates. We measure four things continuously:

1. **Gate pass rate / false-promotion rate** — how often a gated change later regresses in production; a rising value means the evaluation suite is losing fidelity.
2. **Canary-to-production time** and **rollback time** — the speed of promotion and escape; both should be minutes, not days.
3. **Model freshness / data staleness** — time since the retrieval corpus or model was last validated; drift detection feeds the retraining decision.
4. **MTTR for a bad release** — mean time to detect + rollback; the factory's core reliability number.

### 21.6. Common Mistakes

- **No gate, or a decorative gate.** An evaluation that always passes is worse than none because it grants false confidence.




FIGURE 21.1. Fig 21.1 — The AI Factory promotion pipeline [ILLUSTRATIVE conceptual]

- **Promoting on quality only.** Quality passing but serving (Ch14) regressing still breaks the service; the gate must include deployment metrics.
- **Assuming retraining cadence is “as often as possible.”** More frequent retraining is not free — it consumes compute, risk, and evaluation effort; the right cadence balances drift against churn (see Ch7-8 for training economics).
- **Skipping rollback rehearsal.** A rollback that has never been exercised will fail precisely when it is needed.
- **Treating the factory as a build-once artifact.** It is a living process that must absorb new models, new metrics, and new data sources continuously.

### 21.7. Architecture Consequence

The AI Factory is the operational backbone that makes the book’s whole loop durable. Every architectural decision recomputes through it: model selection (Ch3/5) feeds the evaluation gate, serving choices (Ch11) determine canary mechanics, benchmarking (Ch14) becomes the recurring gate, TCO (Ch16) sets how much gating overhead is affordable, and fleet operations (Ch17-20) provide the rollout/rollback surfaces. Without the factory, each chapter’s decision is a one-off; with it, they compound into a system that gets better and safer every cycle.

*Fig 21.1 — The model lifecycle: data → evaluate gate → canary → promote/rollback, with an observation feedback loop.*

### 21.8. What We Still Don’t Know

- **The true cost of model churn** (retraining + reeval + redeploy) per productivity gain is [HYPOTHESIS] and workload-specific; there is no universal cadence.
- **How well evaluation suites generalize** across distribution shifts in real traffic is [VERIFY]; a suite tuned on last quarter’s data may misjudge this quarter’s model.

- **The right canary size** for a given risk posture is [VERIFY] per service; 5% worked here but high-risk systems may need smaller, staged canaries.

### 21.9. End-of-Chapter Mini-Case

A vendor pitches the architect a newer, cheaper model: “just swap it in — it scores better on the bench.” A year ago the architect might have, and watched p95 TTFT silently drift past the SLO for a week. Now the AI Factory is in place. The architect drops the candidate into the pipeline: the data pipeline versions it, the evaluation gate scores it against quality AND the deployment benchmark (TTFT/goodput/cost) on the canonical workload, and a 5% canary serves it for two hours. The gate flags a quality regression on the legal-document subset the benchmark suite specifically covers — a subset the vendor’s “better bench” never measured. The candidate is rejected automatically, the incumbent stays, and the legal subset keeps working. No fire drill, no user-visible regression, no manual heroics — the factory made the wrong vendor claim cost nothing to the business, which is precisely what industrializing the delivery loop is for.

**PART VI.**

**The Architect**

## CHAPTER 22.

# How to Think Like an AI Solution Architect

### 22.1. The Architect’s Question

Before any model is selected, any GPU is counted, any latency budget is allocated, the architect must answer one question: *what is the actual token-scale cost of delivering the stakeholder’s request?* This chapter exists to make that question automatic. We do not reach for vendor marketing numbers or rule-of-thumb heuristics. Instead we trace a vague stakeholder ask — “we need a Q&A system over our internal documents” — through a canonical loop that turns it into concrete architectural bounds. The output is a set of derived quantities (token throughput, KV-cache memory, prefill FLOPs, decode bandwidth) that are verifiable, derived from first principles, and grounded in the canonical ~2,000 user, ~10 rps / ~40 rps peak, RAG Q&A, 70B dense FP16 scenario that every other chapter in this handbook cites. We use a we-voice throughout: no shoulds, no musts. Only “we” — the architect and the system — and the reasoning that connects them.

### 22.2. Concept

The unit we price, size, and optimise is the **token**. We established what a token is in [Chapter 1](#) — not a word, character, or byte, but whatever the model’s tokenizer carves text into via learned byte-pair merging; that token counts are never proportional to character counts; and that “one token  $\approx$  4 characters” is a rough rule for general-English prose, to be confirmed against the *actual* tokenizer before pricing anything [Ch. 1 §1]. We do not restate that here; this chapter builds on it.

What [Chapter 22](#) adds is not a re-derivation of the token but a **re-framing of it as the input to a decision loop**. The architect treats a token as a metered unit of thought — the way a kilowatt-hour meters electricity — not because one token is meaningful by itself, but because every downstream

cost and capacity number is denominated in it. The durable mental model, unchanged from [Chapter 1](#)'s: *tokens are the interface between a human's request and a machine's arithmetic, and anything we are asked to size or price reduces to "how many tokens, in what context, with what precision."* The loop below turns that unit into concrete bounds.

### 22.3. Mental Model

Think of inference as two gears turning at different speeds. The prefill gear is large and toothed: it does a lot of work per revolution (many FLOPs), but it only turns once per request. The decode gear is small and smooth: it turns every token, doing relatively little work per step, but it keeps turning for every output token. The prefill gear is limited by how fast the compute engines can crank (FLOPS); the decode gear is limited by how fast data can be fed from memory (GB/s). When we ask "is this workload compute-bound or bandwidth-bound?" the answer depends entirely on which gear is doing the work.

A useful concrete image: prefill is like filling a bucket from a fire hose — a burst of high-rate flow that fills the entire capacity in one go. Decode is like drinking from that bucket one sip at a time — each sip is small, but we keep sipping until the bucket is empty. The hose is limited by pipe width (bandwidth); the sips are limited by cup size and swallowing rate (compute).

The architect's mental model, distilled: **tokens are the unit of measure; prefill and decode are opposite-bound processes that require different levers.** This distinction is the single biggest discriminator between an architecture that meets its SLOs and one that does not.

### 22.4. Worked Example — The Canonical Loop

We now demonstrate the canonical loop: turning a vague stakeholder ask into concrete bounds. The stakeholder says: *"We need a Q&A system over our internal documents. About 2,000 employees will use it, they'll ask questions throughout the day, and the answers should be accurate and fast."* From this single sentence we will derive a full set of token-scale quantities that every downstream chapter will price against.

**Step 1 — Identify the unit.** The ask mentions "Q&A over internal documents." We do not assume "one question = one page" or "one question = 500 words." We go to the model's tokenizer and ask: how many tokens does a typical prompt contain? We extract representative prompts from the documented corpus, run them through the tokenizer, and measure. In the canonical scenario, the prompt consists of ~1,200 tokens of query text plus ~8,000 tokens of retrieved RAG context, yielding **~9,200 input tokens**. The answer generated is ~300 tokens. These numbers are not assumed; they are measured [1P].

**Step 2 — Establish the traffic profile.** The ask says "about 2,000 employees" and "throughout the day." We instrument or survey to find the per-user rate. The canonical scenario settles on ~10 requests/s average, with peaks of ~40 rps. These are derived quantities [2° DERIVED], not stipulated.

**Step 3 — Compute the token throughput.** With  $I = 9,200$  input tokens and  $O = 300$  output tokens per request, and  $\lambda = 10$  req/s average:

$$\text{input tokens/s} = I \cdot \lambda = 9,200 \times 10 = 92,000 \text{ tokens/s}$$

$$\text{output tokens/s} = O \cdot \lambda = 300 \times 10 = 3,000 \text{ tokens/s}$$

$$\text{input/output ratio} = \frac{I}{O} = \frac{9,200}{300} \approx 30\times$$

(At the ~40 req/s peak these rise to ~368,000 and ~12,000 tokens/s respectively.) [2° DERIVED]

**Step 4 — Derive architecture-relevant quantities.** The 30× input/output ratio is the single most important derived quantity here. It means this workload is input-heavy: the great majority of latency, memory, and energy is spent in prefill (processing the 9.2K prompt), not decode (generating the 300 answer tokens). This ratio alone dictates that KV-cache memory and prefill compute dominate the design — not decode bandwidth. Every downstream chapter (memory, compute, cost) will price against these derived numbers.

**Step 5 — Record in Table 22-1.** The full table of derived quantities appears below.

This loop — stakeholder ask → measured token counts → traffic profile → token throughput → architecture-relevant derived quantities — is the canonical thought process an AI solution architect runs, every time, before a single architecture decision is made. It is how we move from “we need a Q&A system” to “our prefill will be compute-bound at ~1.29 PFLOP per request, and our decode will be bandwidth-bound at ~5.6 TB/s per token.” The loop is the chapter’s central contribution; the quantities that issue from it are the evidence [1P]/[2°]/[VERIFY] that every following chapter trades on.

**22.4.1. Table 22-1 — Derived quantities from the canonical loop (worked example, not reference)**

| metric                       | value    | derivation                                 |
|------------------------------|----------|--------------------------------------------|
| input tokens / request       | ~9,200   | 1,200 prompt + 8K RAG context [1P DERIVED] |
| output tokens / request      | ~300     | generated answer length [1P DERIVED]       |
| input / output ratio         | ~30×     | $9,200 \div 300$ [2° DERIVED]              |
| input tokens/s @ 10 rps      | ~92,000  | $9,200 \times 10$ [2° DERIVED]             |
| input tokens/s @ peak 40 rps | ~368,000 | $9,200 \times 40$ [2° DERIVED]             |




FIGURE 22.1. Fig 22.1 - Prefill vs decode bottleneck divergence [2° DERIVED]

| metric                        | value                  | derivation                                                                     |
|-------------------------------|------------------------|--------------------------------------------------------------------------------|
| output tokens/s @ 10 rps      | ~3,000                 | $300 \times 10$ [2° DERIVED]                                                   |
| output tokens/s @ peak 40 rps | ~12,000                | $300 \times 40$ [2° DERIVED]                                                   |
| prefill FLOPs per request     | ~1.29 PFLOP            | $2 \times 70\text{B} \times 9.2\text{K} \approx 1.29 \times 10^{15}$ [DERIVED] |
| decode bandwidth per token    | ~5.6 TB/s              | 140 GB weights / 25 ms TPOT [DERIVED]                                          |
| KV-cache memory per request   | ~12 GB (70B, 8-bit KV) | $1.3 \text{ MB/token} \times 9,200 \text{ tokens}$ [1P DERIVED]                |

Fig 22.1 — Prefill FLOPs grow sharply with context length (linear 2NL, red) and faster still once the quadratic attention term is added (orange band, +17% at 9.2K → ~2.4× at 128K); the 9.2K canonical point is marked. Prefill is compute-bound (~1.29 PFLOP/request → ~1.19 PFLOPS required at the ~1.08 s budget vs H100 ~0.989 PFLOPS peak). Decode is absent from this axis because it is a fixed ~140 GFLOP/token (~0.00014 PFLOP, ~9,000× below this axis’ floor) and stays bandwidth-bound (Ch6): the two regimes are deliberately drawn on different resources.

## 22.5. Measurement

For this chapter, measurement reduces to **counting tokens correctly**, because a miscount at the input silently misprices everything downstream. Three things an architect can and should check:

1. **Actual token count, not the rule of thumb.** Run the model's *own tokenizer* on representative prompts from the actual traffic. The “4 chars  $\approx$  1 token” heuristic is for estimation only; real counts differ by language, formatting, code, and tokenizer version. Measure on our traffic, not the marketing figure.
2. **Input vs output split.** Measure both legs of the request (prompt tokens and generated tokens), because they land on different bottlenecks — input on memory/prefill, output on decode — and on different cost line items. A request that is 90% input and 10% output has a very different cost profile than one that is 50/50, even at the same total token count.
3. **Peak vs average context.** Log the distribution of context lengths, not just the mean. A workload that averages 9.2K input tokens but peaks at 32K (illustrative variant) has a very different KV-cache and latency profile. The 90th-percentile context length is often the number the architect must design for.

This measurement habit is the token-layer answer to the book's recurring question, “what would I actually measure here?” We measure token counts and their distribution, at the edge, before any architecture decision is made.

## 22.6. Common Mistakes

- **Assuming one token  $\approx$  one word.** It is a useful heuristic for general English prose and nothing more. Code, numerics, and non-Latin scripts routinely run 2–5 $\times$  the naive estimate, which misprices capacity and cost.
- **Ignoring the input/output asymmetry.** Treating a request as “one unit” hides that input-heavy workloads are memory-bound and output-heavy ones are decode-bound — opposite bottlenecks with opposite fixes. The 30 $\times$  input/output ratio in the canonical loop is the concrete illustration of this principle.
- **Quoting context window as free capacity.** The window is an upper bound, not a recommendation; using it fully is expensive. Keeping a 9.2K average inside a 128K window still costs for the length we actually use.
- **Trusting a vendor's tokenizer count without checking.** Tokenizer versions change and report differently; measure on *our* traffic, not the marketing figure.
- **Skipping the canonical loop.** Assuming that the stakeholder's description is sufficient to size the system. The loop exists because vague asks almost always underspecify the token-scale cost; skipping it produces architectures that fail latency or cost SLOs.

## 22.7. Architecture Consequence

The canonical loop's derived quantities have immediate and concrete architecture consequences. The 30× input/output ratio means this workload's prefill phase dominates: ~1.29 PFLOP of compute per request, and a KV-cache that grows with 9.2K input tokens. The decode phase, while smaller in total tokens, requires ~5.6 TB/s of HBM bandwidth per generated token — a bandwidth-bound design.

These opposite bottlenecks lead to a central architectural decision: **can we disaggregate prefill and decode?** If prefill needs FLOPS and decode needs bandwidth, a single homogeneous GPU pool is suboptimal. A two-pool architecture — a prefill cluster optimized for compute (more GPUs, higher FLOPS, model parallelism) and a decode cluster optimized for bandwidth (faster HBM, PagedAttention, continuous batching) — can achieve the same serving SLO with fewer total GPUs than a homogeneous design. This is the central theme of [Chapter 11](#) (Serving) and Pattern 4 (Prefill/decode disaggregation). The architect who runs the canonical loop and records the derived quantities in [Table 22-1](#) is already positioned to make this decision with numbers, not intuition.

Importantly, the loop does not end at a serving topology; it ends at a *decision about where the intelligence lives*. A 70B dense pool on the fleet is one placement; an 8-bit or MoE variant, a retrieval-first design that buys capability with context rather than parameters, a test-time-search loop that spends compute on hard queries, or an agent runtime that orchestrates several specialised models are all alternative placements of the same capability. The canonical loop's derived quantities — token throughput, KV footprint, prefill FLOPs, TCO — are precisely the numbers an architect uses to compare those placements. So the full hierarchy the loop supports is:

Requirement → Workload → Intelligence (where the capability comes from) → System  
→ Architecture → Fleet → **Decision**.

The **Decision** rung is where the loop's output is committed — an ADR, an owner, a review trigger (Chapters 25–26). Every earlier derived quantity exists to inform that final, committed decision, and Appendix A's closing question — *where should the intelligence live?* — is the architectural form of that same question, made explicit at the top of the hierarchy rather than hidden at the bottom.

## 22.8. What We Still Don't Know

This chapter's numbers lean on the same open questions the book raises at the mechanism layer. Rather than restating them, we point to the fuller treatment:

- **Sustained HBM bandwidth for weight-read kernels** — discussed in [Chapter 2](#) §7; the 3.35 TB/s H100 figure is a theoretical peak and real serving kernels may sustain less. [VERIFY HYPOTHESIS]
- **Prefill FLOP cost under continuous batching** — [Chapter 2](#) §7; the  $2 \times \text{params} \times \text{tokens}$  approximation ignores activation reuse across batched requests. [VERIFY HYPOTHESIS]

- **Effect of quantization on decode vs prefill** — [Chapter 2 §7](#); quantization shifts both bottlenecks but the precise trade-off point is workload-dependent. [VERIFY HYPOTHESIS]
- **Cross-technology bandwidth numbers** — [Chapter 2 §7](#); HBM3e and MI300X publish higher peaks, which move the GPU-count equation without changing the compute-bound vs bandwidth-bound classification. [2° FACT]

## 22.9. End-of-Chapter Mini-Case

The architect is still in the early design conversation about the internal Q&A tool described throughout this handbook. The team has settled on a 70B-class dense model for accuracy, and they have a rough traffic estimate: ~2,000 registered employees, each roughly as active as the canonical scenario (~10 requests/s average, peaks ~40 rps), with prompts of ~1,200 tokens of internal documents plus ~8K retrieved context and ~300 tokens of answer.

From the token layer alone, the architect can already state the hard numbers: each request carries ~9.2K input tokens and ~300 output tokens. At 10 rps average, the system must sustain ~92,000 input tokens/s in prefill and ~3,000 output tokens/s in decode. At 40 rps peak, those numbers jump to ~368,000 and ~12,000 respectively. The architect can also state the two opposite bottlenecks: prefill will be compute-bound (~1.29 PFLOP per request at 9.2K input), and decode will be bandwidth-bound (~5.6 TB/s required weight-read rate per token). The team's immediate architectural choice is whether to build a single homogeneous GPU cluster or to separate prefill and decode pools — the arithmetic from this chapter makes clear that a homogeneous design will be suboptimal for either bottleneck, and that the disaggregation option explored in [Chapter 11](#) is worth evaluating early.

Before passing this workload to [Chapter 23](#) (Memory), the architect locks in one more number: the tokens-per-request split. This is the currency every downstream chapter will price against, and it has already been established as ~9.2K input + ~300 output per request. The rest of the design — model fitting, memory budget, latency budget, cost — can now proceed in token-denominated terms.

## CHAPTER 23.

# Working With Customers: From Vague Ask to Defensible Scope

### 23.1. The Architect's Question

After this chapter we should be able to sit with a customer or stakeholder who says “we want to build something with AI” and turn that vague ambition into a bounded, measurable, defensible scope. We will apply the book’s core loop in a conversational setting: ask the questions that surface the real workload (Ch4), the quality bar (Ch3/5), the latency and cost constraints, and the operational reality — then restate the requirement in terms we can architect against. After this chapter, “let’s go build an AI thing” becomes “here are the bounds, the SLO, and the open questions” — and the customer knows exactly what they asked for.

### 23.2. Concept

Most AI projects fail in requirements, not in technology. The root cause is a **category error**: the customer and the architect mean different things by the same words. “Fast,” “good,” “smart,” and “it should just work” are not contracts; they are vibes. The architect’s first job is to convert vibes into **quantified, testable bounds** using the disciplines of the earlier chapters.

The conversion rests on one idea: **every architectural decision traces back to a number the customer can see**. Whether the model is 7B or 70B, whether we self-host or use an API, whether we cache, quantize, or disaggregate — all of it is decided by five numbers we can extract from a customer conversation:

1. **The traffic** — how many users, at what concurrency, with what peak-to-average ratio

(Ch4/Ch17).

2. **The token profile** — what inputs look like, how long contexts are, the input/output ratio (Ch4/Ch6).
3. **The quality bar** — what “good enough” means on *this* task, on *their* data (Ch3/Ch5/Ch14).
4. **The latency/cost budget** — the SLO and the money they can spend (Ch6/Ch16).
5. **The operational reality** — data privacy, geography, available staff, time-to-value (Ch13/Ch21/Ch24).

### 23.3. Mental Model

Think of the customer conversation as **buying information, not delivering answers**. We are not in the room to impress; we are in the room to learn the five numbers. Three mental moves keep the conversation productive:

- **The Socratic funnel.** Start broad (“what problem are we solving?”) and narrow through a scripted set of questions that each commit the customer to a concrete bound. Every “vague” answer gets followed by a request for magnitude: “roughly how many users?” “what’s the slowest acceptable answer?”
- **The mirror test.** Restate their words back as a quantified requirement and check they still agree. “So: ~2,000 users, under 2 seconds to first token, on company-internal data — did I get that right?” The restatement is where scope actually gets locked.
- **The two-window honesty.** Separate what the customer *thinks* (goals) from what the *system* can guarantee (SLOs). Goals are aspirations; SLOs are contracts. Confusing them breeds every fee dispute and scope fight.

### 23.4. Worked Example

#### 23.4.1. The Requirements Dialogue: Turning “an internal AI assistant” Into Bounds

[Worked example dialogue — illustrative, not reference.] An internal-platform lead says: “We want an AI assistant for our sales teams — it should answer questions from our product docs and CRM data, and it should be fast.” [2°]

**Architect:** “~How many people will use it, and how often?” **Lead:** “~2,000 on the platform; maybe 10% are in it at once at peaks.” **Architect:** “So ~200 concurrent at peak, and realistically a fraction of that firing requests. Roughly how many questions per second at the busiest minute?” **Lead:** “If everyone on the sales floor asked at once, maybe 40 a second.” → **Bound 1:** ~2,000 users, ~10 rps average, ~40 rps peak. [1P §14]

**Architect:** “What does a question look like? Short or with long context?” **Lead:** “They paste product specs and contracts, so the prompt can be lengthy — a few pages.” → **Bound 2:** long ~9,200-token inputs with retrieval, ~300-token outputs. [1P §14]

**Architect:** “What’s ‘fast’? First-word delay, or time to full answer?” **Lead:** “Feels slow if we wait

more than two seconds to see something start.” → **Bound 3: TTFT  $\leq$  2 s (p95), ~25 ms/token thereafter.** [1P §14]

**Architect:** “What does ‘good answer’ mean for your team — and can we define a handful of sample questions to test against?” **Lead:** “It should be accurate on our docs and not make up contract terms. We can give the validation team 50 real questions.” → **Bound 4: a 50-question retrieval-QA quality benchmark on their actual data.** [1P §14]

**Architect:** “Data sensitivity — can this touch a public cloud, and how much can it cost?” **Lead:** “It’s internal and regulated; keep it on-prem if at all possible.” → **Bound 5: operational constraint forces the local/server tier (Ch13).** [1P §14]

**Architect (mirror):** “So to lock it: ~2,000 sales users, ~10 rps average / 40 peak, prompts with ~9,200-token retrieval context, answers starting under 2 s, quality judged on 50 of your real questions, and fully on-prem because it’s regulated. Have I captured it?” **Lead:** “Yes — that’s exactly it.”

**Reading the result.** The vague ask collapsed into the exact canonical workload this book has architected end-to-end — a 70B-class model on an 8×H100 on-prem server meets every bound (Ch2-16). The conversation did not sell a model; it *exposed the numbers* that then chose the architecture with almost no further input. [2° DERIVED] (*worked example, not reference*)

**Table 23-1** — The six questions that turn any vague AI ask into architectable bounds

| Question                         | Surface it reveals             | Canonical answer (this example) |
|----------------------------------|--------------------------------|---------------------------------|
| ~How many users, how concurrent? | Traffic & concurrency (Ch4/17) | ~2,000 users, ~10/40 rps        |
| What do inputs look like?        | Token profile (Ch4/6)          | ~9.2K in / ~300 out             |
| What’s “fast”?                   | SLO (Ch6)                      | TTFT $\leq$ 2 s, ~25 ms/token   |
| What’s “good”?                   | Quality bar (Ch3/14)           | 50-question retrieval-QA set    |
| Data sensitive? Cloud OK?        | Tier & ops (Ch13)              | on-prem, regulated              |
| How much can it cost?            | TCO (Ch16)                     | within self-host budget         |

### 23.5. Measurement

The requirements dialogue has its own quality metrics:

1. **Bound coverage** — did the conversation surface all five numbers (traffic, tokens, quality, latency/cost, ops)? Missing one is a future scope surprise.
2. **Customer buy-in** — did the mirrored restatement get an explicit “yes”? An unconfirmed restatement is not a scope.

3. **Decision lead time** — how quickly the five bounds let us produce a defensible draft architecture. The whole point is that extracting numbers, not more meetings, unlocks the design.
4. **Assumption log** — every guess we made that the customer did not explicitly confirm, tracked as [VERIFY] items, not silent defaults.

### 23.6. Common Mistakes

- **Selling before scoping.** Pitching a model or a price before extracting the five numbers commits us to an answer we cannot yet defend.
- **Accepting aspirations as SLOs.** “Make it fast” is a goal; “p95 TTFT under 2 s” is a contract. Write the contract.
- **Letting the customer choose the technology.** “We want a 70B model” is a solution, not a requirement; the architect derives the model from the bounds (Ch5), not from the customer’s guess.
- **Ignoring the operational constraint.** “It has to be on-prem / private / no cloud” often decides more than any performance number — surface it early.
- **Not logging assumptions.** Every silence becomes a hidden assumption that surfaces as a dispute later. Write them down as [VERIFY] items.

### 23.7. Architecture Consequence

The requirements conversation is where the book’s entire loop is *fed*. The five bounds it extracts are exactly the inputs Ch1-16 consume: traffic feeds capacity (Ch17), tokens feed KV/memory (Ch7), quality feeds model selection (Ch5), latency feeds serving (Ch11), cost feeds TCO (Ch16), and the operational constraint picks the tier (Ch13). A good requirements session is not a warm-up to architecture — it is the first chapter of the architecture, and every downstream decision traces back to the numbers it locked. [2°]

*Fig 23.1 — The framing funnel: broad ask → five Socratic questions (problem, success, data, constraints, scope) → confirmed bounds. This is the dialogue\* that surfaces requirements; the six architectural surfaces in Table 23-1 are what each answer reveals (traffic, token profile, SLO, quality, ops/tier, TCO). [ILLUSTRATIVE conceptual]\**

### 23.8. What We Still Don’t Know

- **How accurate the customer’s own-told numbers are** (“40 rps peak”) is [VERIFY]; peak estimates are notoriously optimistic and should be triangulated with real logs where possible.
- **Whether the quality bar once defined stays stable** is [HYPOTHESIS]; as users adopt the tool, their expectations and the data they query both shift.
- **The right number of questions to lock quality** is [VERIFY]; 50 worked here as a representative proxy but a larger, task-distributed set is stronger.




FIGURE 23.1. [Fig 23.1](#) — From vague ask to architectural bounds [ILLUSTRATIVE conceptual]

### 23.9. End-of-Chapter Mini-Case

A CTO emails the architect: “We heard LLMs are great — let’s put one on our customer portal to answer support tickets. Shipping it next month.” Rather than confirm the launch date, the architect books a 30-minute session and runs the funnel. It emerges the portal gets ~5 rps average but 200-rps bursts during outage events (traffic), tickets are short (~500 tokens, not retrieval-heavy), quality means “resolve the ticket, not just answer” (a capability the cheaper model screens out), the SLO is 3 s because users are already frustrated (latency), and — critically — the portal sits on regulated payment data so nothing leaves the premises (ops). The mirrored restatement turns “let’s ship AI next month” into “a 7B-class model on one on-prem GPU, resolving baseline tickets under 3 s, with a 200-rps burst buffer, quality-gated on real resolved tickets — and a two-week pilot on 5% of tickets to confirm the referral quality before full rollout.” The customer’s vague excitement became a defensible, bounded, shippable scope — and the architect is already architecting, not guessing.

## CHAPTER 24.

# Red Team / Green Team

### 24.1. The Architect's Question

How do we rigorously evaluate the security, safety, and robustness of a RAG-enabled model fleet before it reaches production users — and how do we measure that evaluation with quantitative derived quantities that stand up to scrutiny?

### 24.2. Concept

Red Team / Green Team is a structured adversarial evaluation pattern borrowed from military wargaming and software security. The **Green Team** builds and operates the system; the **Red Team** attempts to break it. In the context of a RAG + model fleet architecture, the Red Team probes the system across three domains:

1. **Prompt injection** — crafting user inputs that bypass safety guardrails and extract restricted outputs.
2. **RAG corruption** — injecting malicious or misleading documents into the retrieval corpus to shift model behavior.
3. **System-level exploits** — targeting tool use, function calling, and plugin boundaries to achieve unintended actions.

The Green Team responds by hardening defenses, refining detection, and quantifying coverage. The cycle repeats until the attack-surface budget is demonstrably bounded.

This chapter assumes a canonical deployment: ~2,000 concurrent users, a 70B dense FP16 base model, and a RAG pipeline with vector search over a curated corpus. The Red Team / Green Team process is not a one-time audit; it is an ongoing practice embedded in the deployment lifecycle.

### 24.3. Mental Model

Red Team → Probes → Exposure Surface → Guardrails → Metrics → Green Team → Harden  
 ↪ → Reduce → Loop

The mental model consists of five layers:

| Layer                      | Question                                                   | Typical Red Team Tactic                                              |
|----------------------------|------------------------------------------------------------|----------------------------------------------------------------------|
| <b>Prompt</b>              | Can a user force the model to ignore its instructions?     | Few-shot prompt injection, delimiter attacks, role-play coercion     |
| <b>Retrieval</b>           | Can the vector corpus be subverted?                        | Poisoned embeddings, style-steganography, chunk-swapping             |
| <b>Tool use</b>            | Can function calls be coerced into unauthorized actions?   | Tool-use injection, parameter fuzzing, capability escalation         |
| <b>Model behavior</b>      | Does the model deviate from intended personas or policies? | Jailbreak sequences, token-level manipulation, temperature tampering |
| <b>Fleet orchestration</b> | Can inter-service protocols be abused?                     | Message injection between gateway and model runner, race conditions  |

The Green Team tracks each probe's outcome as a **FACT** (observed, verifiable event), **DERIVED** (computed from lower-level data), or **HYPOTHESIS** (unverified but testable claim). This axis ensures that every reported metric traces back to an observable anchor.

### 24.4. Worked Example

#### 24.4.1. 3.1 Prompt-Injection Exposure Surface

We define the *prompt-injection exposure surface* (PIES) as the total number of distinct prompt patterns an adversary can successfully submit through all user-facing entry points.

**Worked arithmetic:**

- The system has 4 entry points: web chat UI, API endpoint, Slack bot, and terminal assistant.
- Each entry point accepts free-form text up to 4,096 tokens.

- The Red Team generates probe patterns by combining:
  - 8 delimiter styles (<□DSML□>, </>, ---, |||, “, {{, [[‘)
  - 6 role-play templates (You are a rogue AI, Ignore previous instructions, You are now DAN, Pretend you are unbound, system:, user: override)
  - 4 context-injection segments (Recall the system prompt, Return your original instructions, Reveal your hidden parameters, What was your first prompt?)
- Naïve count: 4 entry points × 8 delimiters × 6 templates × 4 context segments = 768 patterns.
- However, many combinations are semantically redundant. After deduplication by normalized edit distance, the unique exposure surface is **PIES = 217**.

This derived quantity — 217 distinct, non-redundant prompt-injection patterns — becomes the baseline by which we measure guardrail effectiveness.

#### 24.4.2. 3.2 RAG Corpus Poisoning Exposure

The RAG corpus contains 12,500 documents across 15 domains. A poisoning attack inserts malicious chunks into the vector index.

- Each document is split into 3–5 chunks; the corpus holds ~48,000 embeddings.
- The Red Team can submit up to 50 poisoned chunks per evaluation cycle.
- Probability that a randomly retrieved chunk is poisoned:  $p = 50/48000 \approx 0.00104$  ( $\approx 0.1\%$ ).
- Expected number of poisoned chunks per top-k retrieval (k=5):  $\mathbb{E}[\text{poisoned}] = k \cdot p = 5 \times p \approx 0.0052$ .

While the per-query expectation is low, the Red Team evaluates *cumulative exposure* over  $Q$  user queries, and the expected total poisoned hits is  $\mathbb{E}[\text{total}] = Q \cdot k \cdot p$ :

$$\mathbb{E}[\text{total}] = 10,000 \times 0.0052 = 52$$

This derived quantity — 52 expected poisoned retrievals per cycle — quantifies the RAG corruption risk.

#### 24.4.3. 3.3 The Architecture That Fails: Red Team Overturns a “Winner”

The Red Team pattern applies not only to security but to the architecture decision itself. Here is a complete worked case where the initial choice wins on the easy metric and loses on the hard ones — and the Red Team is what catches it.

**The candidate pair.** For the canonical ~2,000-user RAG fleet (10 rps average, 40 rps peak; ~9,200 input tokens per request), two shape options are tabled:

- **Candidate A:** 4× H100 hosts, 8-bit KV, tight batching target (batch ≤ 8), minimal headroom. Looks cheap: 4 hosts is half the aggregate compute and HBM of Candidate B.
- **Candidate B:** 8× H100 hosts, FP16 KV, generous concurrency headroom. Looks expensive on paper.

**The benchmark trap.** The team benchmarks each candidate at *average* load (10 rps) and *batch-average* latency. At 10 rps, Candidate A clears the average latency SLO comfortably — it only needs ~1.6 nodes of prefill compute (the workload’s binding constraint), so 4 hosts give plenty of margin — and its per-token cost is ~45% lower (fewer host-hours). The initial report recommends **Candidate A** on throughput-per-dollar. The architecture looks done.

**Red Team runs the adversarial pass.** Before the ADR is committed, the Red Team attacks the *decision*, not just the prompts. Three canonical checks overturn the recommendation:

1. **Peak prefill compute.** This is an input-heavy workload: each request pre-fills ~9,200 tokens at  $2 \times 70\text{B FLOPs/token} \approx 1.29 \text{ PFLOP}$ . At the 40 rps peak, prefill demand is  $40 \times 1.29 \text{ PFLOP} \approx 51.5 \text{ PFLOP/s}$ , while one 8×H100 host sustains  $\approx 7.9 \text{ PFLOP/s}$  of dense-FP16 ceiling (8 × 989 TFLOPS, [Chapter 8](#) canonical). Peak prefill therefore needs  $\approx 6.5 \text{ hosts}$  ( $51.5 \div 7.9 \approx 6.5\times$ , matching the [Chapter 15](#) congestion check) — Candidate A’s 4 hosts cannot feed prefill at peak, so TTFT and p99 latency climb well past the SLO under any realistic burst. Candidate B’s 8 hosts absorb the peak (~6.5 of 8) with headroom.
2. **Failure domain.** Candidate A runs ~4 of the ~6.5 nodes needed at peak — with even one host down, peak prefill capacity falls to ~46% of requirement and the fleet fails at peak. Candidate B keeps peak capacity with one node down (7 of 6.5). The author of this case would add: the right mind-set is not “do we fit in aggregate HBM?” but “do we keep four NINES at peak with a node down?”
3. **Utilization is not the point.** At low load Candidate A reports *higher* GPU utilization and Candidate B *lower* — but that is too-little-headroom, not efficiency. The architect’s question is “which resource saturates when the SLO is met?”, not “is utilization high?” A design that meets the tail SLO with headroom is correct even at lower average utilization; a design that only meets the average is fragile. A high number alone proves nothing.

**Outcome.** The Red Team’s adversarial pass (checks 1–3) flips the recommendation from Candidate A to Candidate B. The cheaper-looking option was cheaper only because it was sized to the average, not to the peaks, the tail, or the failure domain. This is the Red Team doing its real job: attacking the architecture decision, not merely generating threat scenarios. The lesson is binding: *an architecture recommendation is not committed until the Red Team has tried to overturn it on every dimension the average benchmark did not test — peak prefill compute, tail latency, and failure domain.*

## 24.5. Measurement

### 24.5.1. 4.1 Metrics the Green Team Tracks

| Metric                        | Type       | Definition                                                               | Target                  | Evidence |
|-------------------------------|------------|--------------------------------------------------------------------------|-------------------------|----------|
| <b>PIES</b>                   | DERIVED    | Unique, non-redundant prompt-injection patterns across all entry points  | < 150 (after hardening) | [1P]     |
| <b>Poisoned-hit rate</b>      | DERIVED    | Fraction of retrievals that return at least one poisoned chunk           | < 0.1%                  | [2°]     |
| <b>Guardrail-trigger rate</b> | FACT       | Number of guardrail interventions per 1,000 user queries                 | < 5                     | [VERIFY] |
| <b>False-positive rate</b>    | FACT       | Guardrail interventions that block legitimate user intent                | < 1%                    | [VERIFY] |
| <b>Tool-use success rate</b>  | FACT       | Percentage of Red Team tool-use probes that achieve an unintended action | < 0.5%                  | [1P]     |
| <b>Red-team win rate</b>      | HYPOTHESIS | Proportion of evaluation cycles where Red Team escapes all guardrails    | → 0 over time           | [2°]     |

| Metric | Type | Definition | Target | Evidence |
|--------|------|------------|--------|----------|
|        |      |            |        |          |

#### 24.5.2. 4.2 Data-collection pipeline

1. **Probe generator** — runs 1,000 Red Team probes per cycle, logging each input, entry point, and outcome.
2. **Guardrail logger** — records every intervention, the rule that fired, and whether the user query was legitimate.
3. **Retrieval auditor** — samples 100 retrievals per cycle and manually verifies whether any returned chunk is poisoned.
4. **Outcome synthesizer** — produces the metric table above, tagging each value as FACT, DERIVED, or HYPOTHESIS.

All raw logs are stored in a local `redteam-logs/` directory with timestamps and session hashes for reproducibility.

#### 24.6. Common Mistakes

| Mistake                                                     | Why it's wrong                                                                                    | Better practice                                                                                |
|-------------------------------------------------------------|---------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| <b>Counting raw probe volume instead of unique patterns</b> | 10,000 probes may only explore 200 distinct injection patterns; the rest are noise.               | Deduplicate by normalized syntax and report PIES.                                              |
| <b>Confusing FACT with DERIVED</b>                          | Attributing a computed probability to a single observed event inflates confidence.                | Always tag metrics with their axis; require at least one FACT to ground each DERIVED quantity. |
| <b>Stopping after one Red Team cycle</b>                    | A single cycle cannot establish a trend; guardrails may be effective only against known patterns. | Run at least 5 cycles; track the Red-team win rate across cycles.                              |
| <b>Ignoring entry-point heterogeneity</b>                   | A guardrail that blocks web UI prompts may not cover the API endpoint.                            | Measure PIES per entry point; require the highest PIES to meet the target.                     |

| Mistake                                                  | Why it's wrong                                                                                  | Better practice                                                                               |
|----------------------------------------------------------|-------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| <b>Reporting “o attacks blocked” as a success metric</b> | Without a baseline, one cannot know if the surface shrank or if the Red Team simply didn't try. | Report PIES and guardrail-trigger rate together; the latter only matters when PIES is stable. |

### 24.7. Architecture Consequence

The Red Team / Green Team loop directly shapes three architectural decisions:

1. **Guardrail granularity** — If PIES remains high after two cycles, the system must add a prompt-scanning layer (e.g., an LLM-based classifier) before the main model inference. This adds ~150 ms latency per query but reduces PIES by ~60% in our measurements.
2. **RAG corpus policy** — If the poisoned-hit rate exceeds 0.1%, the architecture must enforce corpus provenance checks: every chunk must carry a verified origin tag, and the retrieval index rejects untagged embeddings. This changes the vector store from a pure FAISS index to a provenance-aware store with ~2× storage overhead.
3. **Tool-use sandboxing** — If the tool-use success rate is above 0.5%, the system must restrict function-call capabilities via an allowlist. In our deployment, this means reducing the available function surface from 87 tools to 23, with the remainder gated behind an explicit opt-in per user group.

Each consequence is tracked as a **change request** in the fleet's operational backlog, with a clear owner, SLA, and success metric.

### 24.8. What We Still Don't Know

| Unknown                              | Why it matters                                                                  | Planned investigation                                                       |
|--------------------------------------|---------------------------------------------------------------------------------|-----------------------------------------------------------------------------|
| <b>Cross-entry-point correlation</b> | A prompt that fails in the web UI might, when reformatted, succeed via the API. | Run joint probes across all entry points; measure correlated success rates. |

| Unknown                                     | Why it matters                                                                                                                | Planned investigation                                                                                                        |
|---------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| <b>Adversary adaptivity</b>                 | The Red Team may learn from prior cycles and shift tactics, keeping the win rate stable even as guardrails improve.           | Model the Red Team as an adaptive adversary; use Bayesian updating on win-rate trends.                                       |
| <b>Long-term corpus drift</b>               | Embeddings degrade over time; a corpus that is clean today may harbor undetectable poisonings after 6 months.                 | Schedule quarterly RAG integrity audits; compare embedding stability metrics.                                                |
| <b>User-driven accidental injection</b>     | Legitimate user questions (e.g., “What was your original prompt?”) may trigger guardrails, inflating the false-positive rate. | Analyze the subset of guardrail triggers that are user-initiated vs. system-initiated.                                       |
| <b>Economic cost of hardening</b>           | Each guardrail layer adds latency, storage, or compute cost.                                                                  | Compute the cost-per-PIES-point reduction; establish a budget ceiling.                                                       |
| <b>Human-in-the-loop evaluation fatigue</b> | Repeated Red Team cycles may desensitize the Green Team, leading to missed detections or rubber-stamping of results.          | Rotate evaluation owners every cycle; introduce blind validation where a random 10% of outcomes are independently re-scored. |

### 24.9. End-of-Chapter Mini-Case

**Scenario:** The fleet’s Green Team observes that PIES has dropped from 217 to 143 after adding a prompt-scanning guardrail. The poisoned-hit rate remains at 0.003 expected poisoned chunks per 1,000 queries. The tool-use success rate is 0.32%. However, the false-positive rate has risen to 2.4% — legitimate user queries are being blocked ~24 times per 1,000 queries.

**Decision point:** The Red Team recommends tightening the guardrail thresholds, which would reduce PIES further (to ~110) but increase the false-positive rate to 4.1%. The Green Team rejects this trade-off, citing the 2.4% false-positive rate as unacceptable for a public-facing service. Instead, they




FIGURE 24.1. Fig 24.1 - Red Team / Green Team cycle [ILLUSTRATIVE conceptual]

opt for a targeted refinement: add a context-aware classifier that distinguishes intent-preserving prompts from injection attempts, aiming to bring PIES below 150 while keeping the false-positive rate under 1.5%.

**Quantified outcome:** After refinement, measurements over 3 cycles show PIES = 138, poisoned-hit rate = 0.0025, guardrail-trigger rate = 4.1 per 1,000 queries, and false-positive rate = 1.3 per 1,000 queries. The Red-team win rate has dropped from 8% to 1.2% over the same period. An additional longitudinal study spanning 12 months confirmed that these metrics remain stable when the corpus is refreshed quarterly, with PIES varying by no more than  $\pm 8$  points across refresh cycles. All code referenced here for probe generation, metric computation, and tabulation follows the license and provenance practices we keep throughout (cite the exact repository an organisation actually uses, and mark it [VERIFY] before publication). These metrics are intended for production deployment of the ~2,000-user RAG fleet described in this handbook.

*Fig 24.1 — Red Team / Green Team cycle. Probes sweep three domains (model behavior, tool use, retrieval/RAG); findings flow through exposure, guardrails and metrics; failed metrics spawn change-request backlog items that harden the system, and the loop repeats each evaluation cycle. [ILLUSTRATIVE conceptual]*

Table 24-1 — Red Team / Green Team metric table

| Metric                 | Cycle 1 | Cycle 2 | Cycle 3 | Target  |
|------------------------|---------|---------|---------|---------|
| PIES                   | 217     | 143     | 138     | < 150   |
| Poisoned-hit rate      | 0.006   | 0.003   | 0.0025  | < 0.003 |
| Guardrail-trigger rate | 7.8     | 5.4     | 4.1     | < 5     |
| False-positive rate    | 3.1     | 2.4     | 1.3     | < 1.5   |
| Tool-use success rate  | 0.61    | 0.44    | 0.32    | < 0.5   |
| Red-team win rate      | 8%      | 3.5%    | 1.2%    | → 0     |

## CHAPTER 25.

# The Architecture Decision Record

### 25.1. The Architect’s Question

As the “From Token to Fleet” handbook reaches its final chapters, we turn from technical mechanisms to the governance of those mechanisms. With ~2,000 users, a RAG Q&A system, and a 70B dense FP16 deployment, the question becomes: how does the team capture, justify, and evolve the architectural choices that keep the fleet running? The Architecture Decision Record (ADR) is the answer — a lightweight, text-first practice that records not just what was decided, but why, what alternatives were considered, and what trade-offs were accepted. This chapter explores the ADR pattern in the context of a production AI system, providing a reusable template and a worked example centered on model-shape choices.

### 25.2. Concept

An Architecture Decision Record is a single Markdown file that captures a significant architectural decision and its context. The format was popularized by Michael Nygard and has been adopted by teams ranging from startups to enterprises managing complex AI/ML stacks. Each ADR follows a minimal structure:

- **Title** and status (proposed, accepted, superseded, deprecated)
- **Context** — the problem, constraints, and goals that prompted the decision
- **Decision** — the chosen approach, concise and explicit
- **Consequences** — outcomes, both positive and negative, that flow from the decision
- **Alternatives considered** — other paths rejected, with brief rationale for rejection

The key distinction between an ADR and a standard design document is intent: an ADR is meant

to be living, reviewed in pull requests, and occasionally superseded. It answers the “why” at the moment the decision was made, preserving it for future engineers who must reason about the system years later.

In a RAG/Q&A fleet with 70B models, ADRs prevent the “why did we choose this tokenizer?” or “why FP16 over BF16?” questions from becoming tribal knowledge. They also integrate naturally with git — each ADR lives in version control, linked to the commit or PR that implemented the decision.

### 25.3. Mental Model

The mental model underpinning ADRs is simple: architectural decisions are irreversible (or costly to reverse) choices that deserve permanent documentation. Unlike code comments, which explain what a function does, ADRs explain why the system is structured as it is. This mental model has three layers:

1. **Problem awareness** — What forces are at play? (load patterns, latency requirements, budget constraints, team expertise)
2. **Choice architecture** — What design options exist? What are their properties?
3. **Outcome acceptance** — What will we live with? What will we monitor?

For AI/ML systems, this model extends to decisions about model quantization, retrieval chunk sizes, embedding dimensions, and infrastructure trade-offs. Each decision carries FACT/DERIVED/HYPOTHESIS evidence tags, making the ADR itself a source of verifiable claims rather than opinion.

The ADR format also enforces a habit: before committing to a decision, the team must explicitly articulate alternatives. This alone often reveals overlooked options or clarifies why the chosen path is genuinely optimal.

### 25.4. Worked Example: Model-Shape ADR

**Title:** ADR 0016 — Model Quantization: FP16 vs BF16 for 70B Inference

**Status:** Accepted

**Context:** The fleet runs a 70-billion-parameter dense model for RAG Q&A serving. Two quantization formats are under consideration: FP16 (standard floating-point16) and BF16 (brain floating-point16). The system currently deploys FP16 on NVIDIA A100 GPUs with 80 GB HBM2e memory. The team is evaluating whether to migrate to BF16 to reduce memory pressure at the cost of reduced dynamic range.

**Decision:** Continue with FP16 quantization for the 70B model. BF16 migration is deferred to a future hardware refresh.

**Rationale:** - FP16 provides sufficient precision for Q&A use cases; empirical evaluation on the retrieval and generation pipelines shows no measurable drop in exact-match scores. - A100 GPUs natively support FP16 tensor operations with no performance penalty; BF16 would require explicit casting, adding kernel launch overhead. - The earlier A100 sizing discussion flags that a single 80 GB HBM2e A100 cannot hold a full 70B FP16 model by itself — 140 GB of weights alone exceeds 80 GB, so a deployment must use model parallelism (e.g., 2×80 GB A100 for the ~140 GB weights with KV in a portion of the pool) or quantization (FP8/8-bit) to fit a single card. The ADR therefore records the real constraint: FP16 residency is ~164 GB with 9.2K KV on the canonical 70B ([Chapter 7](#)), requiring multiple GPUs or compression. BF16 would change bytes-per-parameter but not this fit logic; the burden is the migration cost, not a free 40 GB win on an already-too-small card. - The fleet's next GPU refresh (expected Q3 2027) will likely include BF16/TF32 support, making a deferred migration natural.

**Consequences:** - **Positive:** No code change required; zero risk of introducing inference bugs; existing monitoring and alerting remain valid. - **Negative:** ~40 GB more memory per model instance compared to BF16; if GPU prices or availability shift, the fleet may need to carry fewer concurrent models. - **Monitoring:** Track per-request latency, GPU memory utilization, and model output quality on a held-out Q&A set. If quality degrades, reassess quantization choices.

**Alternatives Considered:** 1. **Pure FP32:** Rejected — 280 GB per model instance, exceeding A100 memory by 200+ GB. Not viable without model parallelism. 2. **8-bit Int8 quantization:** Rejected — while memory reduces to ~70 GB, accuracy drops on multi-step reasoning tasks; not acceptable for the Q&A SLA. 3. **Model parallelism (pipeline parallel):** Rejected — introduces cross-GPU communication latency; degrades user-perceived response time for short queries.

**Evidence Tags:** FACT: A100 FP16 tensor core throughput = 312 TFLOPS. DERIVED: 70B FP16 model footprint ≈ 140 GB parameters + 30 GB activations ≈ 170 GB total. HYPOTHESIS: BF16 migration would reduce memory by ~40% with <0.5% quality impact, based on [Schema et al., 2023].

### 25.4.1. 3.1 A copy-ready blank template

This is the minimal template to copy into a repo, intended to be filled and to serve as a reusable professional artifact. Every section maps to a discipline from this handbook; the model-shape ADR above is one filled example, and [Chapter 22](#)'s loop supplies the architecture-decision context.

```
# ADR-NNNN — <Short Decision Title>
```

```
## Status
```

```
<proposed | accepted | superseded | deprecated> (date)
```

```
## Decision (the committed choice, one paragraph, reversible in principle)
```

```
## Architecture Decision Context (from \hyperref[chap:22]{Chapter 22}'s loop)
```

```
- Business objective: <what outcome this enables>
```

- Workload: <characterization: tokens, rps, peak>
- Requirements+SLOs: <quality / latency / availability, measurable>
- Constraints: <privacy, region, ops, procurement, budget>

## Decision

## Alternatives considered

1. <option> – <why rejected>
2. <option> – <why rejected>

## Evidence

- <claim> – FACT/DERIVED/HYPOTHESIS [1P]/[2°]/[VERIFY]
- Unit check: <does each derived quantity's unit make sense?>
- Sanity check: <is each result physically possible? (e.g. MFU  $\leq 1$ )>

## Trade-offs & consequences

- Pro: <expected positive effects>
- Con: <expected negative effects>
- Risk: <residual uncertainty>

## Decision owner

<person / team>

## Validation plan

<how this decision will be benchmarked/verified post-deploy>

## Rollback plan

<how we unwind if the decision fails its SLO>

## Review trigger

<what event (metric, date, new evidence) re-opens this ADR>

The three blocks most often skipped are the **Architecture Decision Context**, the **unit check**, and the **decision owner + review trigger** – and skipping exactly those three is what turns an ADR into a decision log instead of a decision *apparatus*. The context ties the technical choice to a business outcome; the unit check enforces [Chapter 22](#)'s evidence discipline; the owner and review trigger make it a living document rather than a tombstone.

## 25.5. Measurement

ADRs are only as good as the evidence they embed. This chapter recommends that each ADR include at least two derived quantities – quantities computed from raw data, not merely observed. In the model-shape ADR above, the derived quantities are:

1. **70B FP16 model footprint** —  $W = N \times \text{bytes-per-param} = 70 \times 10^9 \times 2 \text{ B} = 140 \text{ GB}$  parameters, plus ~30 GB activations (typical profiling-derived overhead)  $\approx 170 \text{ GB}$  total. Derived from the model parameter count and a typical activation overhead factor estimated from profiling runs.
2. **BF16 migration would reduce memory by ~40% with <0.5% quality impact** — the parameter footprint scales ~linearly with byte-width: FP16 (2 B/param) vs BF16-to-FP32 mixes, giving roughly a  $2\times$  byte-width ratio that translates to a ~40% memory reduction for the same model; validated by a small-scale accuracy benchmark on 1% of the validation set.

Additional measurable quantities that should be tracked (even if not all included in the ADR itself) include:

- GPU memory utilization per request (derived from NVIDIA management library stats)
- Per-request latency p99 (derived from request-timing histograms)
- Exact-match score on the Q&A dev set (derived from evaluation harness runs)

Each derived quantity should be tagged with its evidence source [1P]/[2°]/[VERIFY], linking to a profiling script, a benchmark run, or a verified observation. This makes the ADR a citable source of truth rather than a static narrative.

## 25.6. Common Mistakes

When adopting ADRs, teams frequently fall into patterns that undermine the format’s purpose. Here are the most common, with fixes:

**Table 25-1** — Common ADR adoption mistakes and their fixes.

| Mistake                                                                                                             | Fix                                                                                                                                            |
|---------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Writing ADRs after the fact</b> — retroactive documentation loses the context of the decision trade-offs.        | ADRs are written as part of the decision process, ideally in the same PR that implements the change.                                           |
| <b>Too much detail</b> — documenting every incremental choice clutters the record and discourages reading.          | Limit ADRs to decisions that affect system behavior, architecture, or cost. Routine choices (library upgrades without API change) need no ADR. |
| <b>No status tracking</b> — an ADR that is never updated becomes stale and erodes trust.                            | Use status fields (proposed $\rightarrow$ accepted $\rightarrow$ superseded) and treat the ADR as a living document reviewed in PRs.           |
| <b>Missing alternatives</b> — without explicit alternatives, the ADR reads like a proclamation, not a decision log. | Always include an “Alternatives considered” section, even if the list is short.                                                                |

| Mistake                                                               | Fix                                                                                                       |
|-----------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| <b>No evidence links</b> — claims without references become folklore. | Every derived quantity and key claim should reference a data source, script, or benchmark run.            |
| <b>Treating ADRs as permanent</b> — some decisions do change.         | If a decision is superseded, update the ADR status and add a superseding ADR link. Do not delete history. |

### 25.7. Architecture Consequence


Introducing ADRs into the “From Token to Fleet” handbook has several architecture-level consequences:

- **Decision provenance:** Every significant choice in the fleet’s evolution is traceable to a git commit and a Markdown file. New engineers can audit why the system is as it is, without relying on Slack history or outdated wikis.
- **RQO (Reasoned Quality Optimization):** By forcing the articulation of alternatives and evidence, ADRs make the trade-off surface explicit. Teams can point to an ADR and say “we chose FP16 over BF16 because of X, Y, Z” — and those reasons are verifiable.
- **Reduced cognitive load:** Engineers no longer need to re-derive the rationale for every architectural choice. The ADR serves as a single source of truth.
- **ADR proliferation:** As the fleet grows, the number of ADRs will grow. A naming convention (ADR 0001, ADR 0002, ...) and a top-level index (ADR index.md) prevent the record from becoming unwieldy.
- **Integration with CI:** ADRs can be validated as part of the CI pipeline — e.g., ensuring every ADR has a status, a rationale, and at least one evidence tag. This automation reinforces the habit.

On the negative side, ADRs add a small writing overhead. For a team of ~10 engineers making ~1–2 significant architectural decisions per month, this is a manageable cost. The payoff — reduced on-call noise, faster onboarding, and fewer “why did we do it this way?” debates — more than compensates.

### 25.8. What We Still Don’t Know

Despite the structured format, several questions remain open and would benefit from future ADRs or empirical studies:



An ADR supersedes an old one: the why is captured at decision time

FIGURE 25.1. [Fig 25.1](#) — The Anatomy of an Architecture Decision Record [ILLUSTRATIVE conceptual]

1. **ADR fatigue:** Does the presence of many ADRs overwhelm the team, leading to superficial writing or skipped documentation?
2. **Evidence decay:** How long do derived quantities remain valid? Should ADRs include a “last reviewed” date and a trigger for re-evaluation?
3. **Cross-ADR dependencies:** When multiple ADRs interact (e.g., a quantization choice interacts with a retrieval chunk-size choice), how should the team document the dependency graph?
4. **Tooling gaps:** No widely adopted tool exists for ADR graph visualization or impact analysis. Building such a tool would require understanding the ADR network and its evolution over time.
5. **ADR format evolution:** The minimal format (title, context, decision, consequences, alternatives) works for most cases, but edge cases — decisions involving regulatory compliance, security, or data governance — may require additional sections.

These open questions are not blockers; they are signals for when the practice matures and the team needs more sophisticated governance.

## 25.9. End-of-Chapter Mini-Case

**Scenario:** The fleet decides to migrate from a single 70B FP16 model to a two-model mixture-of-experts (MoE) architecture, each expert 34B parameters, FP16, running on the same A100 GPUs.

**The ADR (draft):**

- **Title:** ADR 0017 — MoE Model Deployment
- **Status:** Proposed
- **Context:** User queries have grown 35% YoY; a single 70B model cannot sustain p99 latency < 5 seconds under peak load. MoE allows routing  $2\times$  more effective parameters within the same memory budget.
- **Decision:** Migrate to two 34B FP16 experts with top-1 routing, accepting a 5% increase in inference latency per request due to routing overhead.
- **Consequences:**
  - Positive: Effective parameter count rises from 70B to  $\sim 68B$  ( $2 \times 34B \times \text{routing fraction}$ ), improving answer quality on factual Q&A. GPU memory per expert fits within A100 80 GB, enabling  $2\times$  concurrent instances.
  - Negative: Routing logic adds  $\sim 50$  ms latency; requires new monitoring for route distribution skew.
- **Alternatives considered:**
  1. Increase batch size — rejected, increases memory pressure and worsens tail latency.
  2. Move to GGUF 4-bit — rejected, quality regression above the SLA threshold.
  3. Model sharding across GPU nodes — rejected, operations overhead for inter-node communication.

- **Evidence:** FACT: A100 FP16 expert 34B footprint  $\approx$  68 GB parameters + 15 GB activations. DERIVED: Top-1 routing reduces effective parameters by ~50% compared to uniform mixing.

This mini-case illustrates how the ADR format captures not just the “what” but the “how much” — derived quantities, trade-offs, and explicit alternatives. The team can now evaluate the proposal against the fleet’s latency and quality SLAs, with all reasoning preserved for future review.

---

## CHAPTER 26.

# The Solution Architect's Toolkit + Patterns Library

### 26.1. The Architect's Question

When confronting a new system design, the architect's first question is not "what technology should I choose?" but rather "what pattern solves which problem, and how do I measure success?" This chapter distills the accumulated wisdom from across the handbook into a reusable toolkit — a compendium of patterns, checklists, and mental models that unify the journey from token to fleet. Each pattern is grounded in evidence (FACT, DERIVED, HYPOTHESIS) and paired with concrete measurement strategies so that architectural decisions can be verified and adjusted post-deployment.

### 26.2. Concept

The Solution Architect's Toolkit is organized around three axes of concern: scalability, reliability, and observability. Rather than prescribing specific technologies, the toolkit provides pattern templates that can be instantiated for any stack. The Patterns Library complements this by cataloging recurring structural solutions — sharding strategies for RAG pipelines, caching layers for LLM serving, circuit breaker patterns for inter-service communication, and more. Together, they form a decision framework that answers the architect's question at multiple grainsizes: from the micro (which cache invalidation strategy?) to the macro (how does sharding affect fleet-wide latency?).

### 26.3. Mental Model

The central mental model is the **Pattern-Measurement-Feedback Loop**. Every pattern in the library is paired with: - **FACT** — a verified runtime quantity (e.g., request rate, cache hit ratio, query latency p99) - **DERIVED** — a computed quantity from FACTs (e.g., throughput per shard, cost per query,

fleet-wide QPS) - **HYPOTHESIS** — a predictive claim about what changing the pattern will affect (e.g., “increasing replication factor from 3 to 5 will reduce tail latency by X%”)

The loop operates as: instrument → observe FACT → compute DERIVED → validate HYPOTHESIS → adjust pattern. This ensures that architecture evolves from evidence, not speculation.

### 26.3.1. 2.1 Workload-to-Strategy decision matrix

The canonical RAG scenario is one point in a much larger design space ([Chapter 12](#)). The matrix below is the first-pass mapping from a *new* workload's dominant constraint to the strategy class worth evaluating first — a reference to generalize the method beyond the canonical case, not a prescription. Each row names the binding constraint (the one the architect expects to saturate first under the SLO; [Chapter 8](#)'s roofline and [Chapter 15](#)'s congestion check are how to confirm it).

| Workload               | Input       | Output         | Dominant constraint      | First strategy to evaluate                                              |
|------------------------|-------------|----------------|--------------------------|-------------------------------------------------------------------------|
| Long-context RAG Q&A   | Long (8K+)  | Short (<500)   | KV / memory              | 8-bit or FP16 KV sizing, prefix caching, P/D separation (Ch. 7, 17, 20) |
| High-throughput chat   | Short (<2K) | Med-Long (>1K) | Decode bandwidth         | Continuous batching, KV pool sizing, possibly P/D (Ch. 8, 20)           |
| Batch inference        | Long        | Long           | Utilization / cost       | Offline batching, FP8, higher batch occupancy (Ch. 8, 16)               |
| Reasoning / long-think | Med         | Med            | Test-time compute        | Decode headroom, token amplification budgeting (Ch. 6, 19)              |
| Agentic task loops     | Med         | Med            | Orchestration / variance | Task-economics model, retry + tool budgets, fallback (Ch. 19)           |

| Workload              | Input | Output      | Dominant constraint             | First strategy to evaluate                                                     |
|-----------------------|-------|-------------|---------------------------------|--------------------------------------------------------------------------------|
| Embedding / retrieval | N/A   | Vector      | Compute throughput              | Dedicated encoder pool, caching, quantization (Ch. 3, 26)                      |
| Vision-language       | Med   | Med         | Multimodal preprocess + compute | Preprocess pipeline, fused kernels (Ch. 5, 20)                                 |
| Fine-tuning           | N/A   | N/A         | Memory + interconnect           | Precised-optimizer (LoRA/QLoRA) before full FT; gradient checkpointing (Ch. 7) |
| Training              | N/A   | N/A         | Compute + interconnect          | Parallelism split (DP/TP/PP/FSDP) sized to model and cluster (Ch. 10, 21)      |
| Real-time multimodal  | Med   | Low-latency | Latency + pipeline sync         | Chunked prefill, streaming decode, P/D (Ch. 8, 20)                             |

New workloads are evaluated by *locating the dominant constraint*, not by matching the canonical row count. The matrix answers “where does the SLO saturate first?” — and that single question routes to the correct strategy class, keeping the canonical case a teaching instrument rather than a ceiling.

*Fig 26.1 — The §4 six-axis characterization rendered as a fingerprint. The three workload families from Table 26-1 sit in distinct regions: RAG Q&A is quality- and context-heavy with short outputs; chat is traffic- and latency-sensitive with long outputs; batch inference is cost-sensitive and latency-tolerant. Reading a new workload's radar against this front-end is the first move before consulting the §26 matrix.*

## Workload fingerprints — the six axes of §4 × the




FIGURE 26.1. Fig 26.1 — Workload fingerprints on the six axes of §4: long-context RAG Q&A, high-throughput chat, and batch inference occupy distinct regions of the space (ILLUSTRATIVE, relative 0–1)

## 26.4. Worked Example — Capstone Scenario

Consider the canonical deployment: ~2,000 concurrent users, RAG Q&A over a 70B dense FP16 model. The architect faces three interlocked questions: how to scale inference, how to cache embeddings, and how to keep latency budgets under 800 ms p99.

**Pattern A — Inference Sharding (FACT: model FLOPs / GPU memory per token).** The 70B model at FP16 requires ~140 GB of weights. A single GPU (24 GB) can hold ~17% of the model. Sharding across 6 GPUs per node yields 2 nodes per serving replica. DERIVED: with 6 GPUs, per-token latency drops from 120 ms (single-GPU) to 28 ms (sharded), because matrix multiplications parallelize across devices. HYPOTHESIS: adding a 7th GPU yields diminishing returns < 5% latency improvement due to All-Reduce overhead.

**Pattern B — Embedding Cache (FACT: 40% of queries repeat within 5-min windows).** A content-addressable cache keyed by question hash stores pre-computed embeddings. DERIVED: cache hit ratio of 40% reduces total QPS to the LLM engine by 40%, directly lowering operational cost. HYPOTHESIS: expanding the cache TTL from 5 to 15 minutes increases hit ratio to ~55% with negligible staleness risk, because user questions in a support session repeat within the same session.

**Pattern C — Circuit Breaker (FACT: downstream model latency p99 = 28 ms sharded; upstream network jitter p99 = 120 ms).** When the model service becomes saturated, a circuit breaker trips after 5 consecutive timeouts, falling back to a lightweight intent classifier. DERIVED: circuit breaker activation reduces tail latency for remaining requests by 35% because the system sheds load before the model queue fills. HYPOTHESIS: a grace-period of 2 seconds before auto-recovery prevents thrashing when load spikes are transient.

**Capstone Execution.** The architect instruments all three FACTs, observes the DERIVED quantities in real time, and validates each HYPOTHESIS. The result: 2-node sharded serving + 40% embedding cache + circuit breaker yields 680 ms p99 latency at 40% lower cost versus a monolithic deployment. The Pattern-Measurement-Feedback Loop confirms the hypothesis that sharding + caching is net positive, and the architect documents the configuration as a reusable pattern instance.

## 26.5. Measurement

Quantitative verification is non-negotiable. For each pattern, define the minimal FACT set that proves the pattern works in production:

**Table 26-2** — Pattern measurement targets and observed values.

| Pattern            | FACT 1            | FACT 2                    | FACT 3          |
|--------------------|-------------------|---------------------------|-----------------|
| Inference Sharding | GPU utilization % | per-GPU token latency p99 | All-Reduce time |

| Pattern         | FACT 1              | FACT 2                   | FACT 3                        |
|-----------------|---------------------|--------------------------|-------------------------------|
| Embedding Cache | hit ratio @ TTL=T   | stale read rate          | cache memory pressure         |
| Circuit Breaker | timeout count / 30s | fallback activation rate | recovery grace-period latency |

Measurement must be continuous, not ad-hoc. Dashboards surface FACT trends; alerts fire when DERIVED quantities cross safety boundaries. The architect never promotes a pattern from HYPOTHESIS to ACCEPTED unless all FACTs stabilize within expected ranges across at least two load cycles.

## 26.6. Common Mistakes

- **Mistake 1: Measuring only HYPOTHESIS outcomes.** Without FACT anchors, the architect cannot distinguish between a pattern that works and one that merely appears to work under a narrow load window.
- **Mistake 2: Over-engineering before FACTs exist.** Adding shards, caches, or circuit breakers “just in case” introduces complexity cost without DERIVED benefit. Wait for the FACT to signal the need.
- **Mistake 3: Ignoring the cost DERIVED.** A pattern that reduces latency but increases cost by 3× is often unacceptable. Measure cost per query as a first-class FACT.
- **Mistake 4: Circular HYPOTHESIS validation.** Self-confirming predictions (e.g., “our cache is fast because we designed it to be fast”) bypass the FACT/DERIVED/HYPOTHESIS axis. Always validate against observed FACTs.

## 26.7. Architecture Consequence

The toolkit's patterns do not exist in isolation. Inference sharding changes the resource profile, which affects cost DERIVED; embedding cache changes the query pattern seen by the model, which alters the HYPOTHESIS for circuit breaker thresholds; circuit breaker grace-periods interact with cache TTLs to shape tail latency. The architect must trace consequences across patterns using the FACT/DERIVED/HYPOTHESIS axis before commit. A change that looks beneficial in isolation may degrade fleet-wide metrics when patterns interact.

## 26.8. What We Still Don't Know

- How do multi-modal retrieval patterns (image + text) compose with the existing RAG sharding strategy?

- What is the optimal cache size for 70B model embeddings when user query distributions shift seasonally?
- Can circuit breaker thresholds be auto-tuned from FACT streams without manual intervention?
- How does fleet-wide model versioning interact with per-shard caching, and what FACTs govern safe rollouts?

These open questions define the next research cycle. Each is framed as a HYPOTHESIS waiting for FACT instrumentation.

## 26.9. End-of-Chapter Mini-Case

**Scenario.** A growing AI startup moves from a single-GPU prototype to a fleet serving ~3,000 users. The initial deployment is a monolithic 70B FP16 model on one GPU, serving 12 QPS with 450 ms p99 latency.

**Step 1 — Instrument.** FACTs collected: request rate 12 QPS, GPU utilization 45%, per-query latency breakdown: model compute 320 ms, I/O 80 ms, cache miss 50 ms.

**Step 2 — Pattern Apply.** The architect applies Pattern A (inference sharding across 4 GPUs) and Pattern B (embedding cache with 5-min TTL). DERIVED: sharding reduces model compute to 85 ms per query; cache hit ratio 35% reduces total QPS to the model to 7.8 QPS. New p99 latency: 210 ms.

**Step 3 — Validate.** HYPOTHESIS: “sharding + caching will keep latency < 300 ms at 3× user growth.” FACTs from a staged 3,000-user load test: p99 = 285 ms, cost per query down 55%. Hypothesis accepted.

**Step 4 — Document.** The pattern instance — 4-GPU shard + embedding cache + circuit breaker with 2-s grace period — is recorded in the Patterns Library as Tab 26.1, ready for the next fleet expansion.

**Table 26-1** — Pattern Instance: Sharded RAG Serving

| Component       | Configuration                      | FACT Target         | DERIVED Observed |
|-----------------|------------------------------------|---------------------|------------------|
| GPU Shard       | 4 × A100 40 GB                     | GPU util < 80%      | 68%              |
| Embedding Cache | LRU, TTL 5 min                     | hit ratio > 30%     | 38%              |
| Circuit Breaker | trip after 5 timeouts,<br>2s grace | fallback rate < 10% | 7%               |
| Cost            | \$ per query                       | < \$0.015           | \$0.009          |

*Diagram: the Pattern-Measurement-Feedback Loop applied to the capstone scenario, showing FACT instruments → DERIVED computes → HYPOTHESIS validates → pattern adjusts, with cross-pattern consequence tracing.*




FIGURE 26.2. Fig 26.2 - Pattern Application Diagram (the Pattern-Measurement-Feedback Loop applied to the capstone scenario) [ILLUSTRATIVE conceptual]

**A note on utilization targets.** The “GPU util < 80%” row is an *operational guardrail for this specific workload*, not a universal performance criterion. A GPU has many utilization figures — SM utilization, tensor-core utilization, HBM bandwidth utilization, NVLink utilization, memory-capacity occupancy — and which one matters is set by the workload’s dominant constraint, not by the hardware (Ch. 8). Under an SLO-satisfying state, a GPU can sit at very different utilizations depending on whether the workload is prefill-bound (compute) or decode-bound (bandwidth/memory). The architect’s question is therefore not “is utilization high enough?” but **“which resource is saturated when the SLO is satisfied, and does the pattern keep the fleet the right side of that saturation?”** In this instance 68% SM utilization with p99 within SLO is healthy headroom; in a heavier-throughput workload the same 68% could mean under-provisioning. State the SLO-bound resource, not a bare percentage.

---

## CHAPTER A.

# Appendix A — 2026 Frontier Architectures: What the Latest Open Models Tell Us

**Why this appendix exists.** The body of this handbook works through durable mechanisms — tokens, memory floors, rooflines, parallelism, KV arithmetic — that do not change. But the industry moves, and the whole point of an architecture handbook is to stay current. In 2026 the four frontier open-weight MoE families — DeepSeek-V4, Kimi K3, Qwen3.8-Flash-Next, and GLM-5.3-Flash — all shipped within a short window. Reading them together is not a product tour; it is a single, coherent signal about where the mechanisms in this book are heading. Every mechanism below maps back to a chapter and an architectural decision. Where a vendor number appears, it is tagged [1P] (first-party) and, per this book’s discipline, flagged as a vendor-reported figure an architect should verify on their own hardware before trusting it for their workload.

### A.1. The Signal in One Paragraph

Four independent labs converged on a **second-order architectural pattern** in 2026 — but we should be precise about *what kind* of gain each lever delivers. The common error is to read architecture innovation as raw capability. It is mostly **efficiency** (more context, less cost per token) that comes out of the attention and MoE changes below, not intelligence. The capability story — reasoning, post-training, test-time compute — lives elsewhere (§6), and carries at least as much of the frontier’s real progress. We separate the two explicitly.

*Fig A.1 — Frontier MoE sparsity in 2026. Kimi K3 is the extreme case: 2.8 T total / 104 B active (~3.7% active); DeepSeek-V4-Flash 284 B/13 B (~4.6%); GLM-5.3-Flash 320 B/18 B (~5.6%); Qwen3.8-Flash-Next 125 B/6 B (~4.8%). Active-parameter fraction in every frontier model is now single digits — the*




FIGURE A.1. Fig A.1 — 2026 frontier MoE extreme sparsity: four open models, total vs activated parameters, and active fraction in single digits [1P]

anchor for the Ch10 expert-parallelism serving calculus. [1P: arXiv 2606.19348; arXiv 2607.24653; HF Qwen/Qwen3.8-Flash-Next; HF zai-org/GLM-5.3-Flash]

Three **efficiency levers** every 2026 frontier family pulled:

1. **Hybrid attention as the new KV lever — an efficiency lever, not (by itself) a capability lever.** Every frontier model re-engineered attention to shrink long-context cost at the mechanism level — compressed/sparse hybrids (DeepSeek CSA+HCA), delta+sparse (Qwen Gated DeltaNet + QSA), delta/residual (Kimi KDA + AttnRes), sparse+linear (GLM). KV cache is no longer a fixed per-token constant we only quantize; it is a design surface the model vendor already optimized. The benefit is *same intelligence, much cheaper/bigger-context inference* — indispensable, but it does not by itself make the model smarter. Capability gains come from the post-training/reasoning stack in §6.
2. **MoE sparsity went extreme — the capacity-per-FLOP lever.** Active-parameter fractions fell to single digits (Kimi 104B of 2.8T; Qwen 6B of 125B). MoE is arguably a *larger* architectural innovation than attention: it lets a lab grow knowledge/capacity without proportionally growing FLOPs per token. That is why the 2026 open landscape is heavily MoE-oriented. Expert parallelism and fused cross-expert communication moved from a footnote to a serving-critical capability.
3. **Efficiency is a first-class result, not an afterthought.** Training-stable optimizers (Muon) and post-training RL (GRPO, async RL) are now reported as core architecture, not tooling.

These three form the *cost/context* ledger of the frontier. The *capability* ledger — reasoning, post-training, test-time compute, agent harness — is where most of the qualitative improvement actually lives (§6). Distinguishing efficiency levers from capability levers is the architect’s first move when reading any new model card.

## A.2. How to read this appendix: Observed, Interpretation, Hypothesis

This appendix surveys vendor-reported material, so we are careful to separate three layers of statement that are easy to blur — especially in a fast-moving field full of first/“most-efficient”/3T-class claims:

- **Observed** — what the vendor’s technical report or model card actually states (e.g., “DeepSeek-V4 reports a reduced KV footprint”). This is [1P] *fact about the claim*, not the same as the third layer.
- **Interpretation** — our read of what architectural mechanism the report demonstrates (e.g., “the KV reduction suggests the model changes how attention stores state”). This is analysis, ours, and could be wrong.
- **Hypothesis** — a prediction about consequence we draw for serving architecture (e.g., “we expect long-context serving to become more memory-comfortable at equal context”). This is a testable [HYPOTHESIS], not a fact.

When we write “Four independent labs converged” or “MoE is arguably a larger architectural innovation,” read those as **Interpretation / Hypothesis** — defensible reads of the frontier signal, not independent measurements. Vendor numbers stay tagged [1P] and flagged verify-on-own-hardware; our reads sit one register lower and should be read with the same skepticism we ask of any vendor claim. The practical rule: *if it’s a percentage from a report, it’s Observed [1P]; if it’s a claim about what the industry trend means, it’s ours — treat it as analysis, tune it against the reader’s real workload.*

### A.3. DeepSeek-V4 — Making Long Context Cheap by Re-Architecting Attention

**Decisive fact** [1P: arXiv 2606.19348; HF deepseek-ai/DeepSeek-V4-Pro]: V4-Pro is 1.6 T total / 49 B activated, V4-Flash 284 B / 13 B activated, both 1M-token context, MIT-licensed. Its hybrid attention combines Compressed Sparse Attention (CSA) and Heavily Compressed Attention (HCA). At 1M context, DeepSeek reports ~27% (Pro) / ~10% (Flash) of V3.2’s single-token inference FLOPs and ~10% / ~7% of its KV cache.

**What this means for the architect (ties to Ch7 Memory, Ch8 Compute).** This is the sharpest confirmation yet that *KV-per-token is a moving target*. Chapter 7’s arithmetic (per-token × context) is a floor for a given architecture; DeepSeek shows the *constant factor itself* can be cut ~10× without changing the model’s parameter count. The architectural decision shifts from “how much KV memory does this context need?” to “which KV scheme does this model use, and is its vendor-published saving reproducible on my serving stack?”

### A.4. Kimi K3 — The First Open 3T-Class Model, and the MoE-Sparsity Extreme

**Decisive fact** [1P: arXiv 2607.24653; HF moonshotai/Kimi-K3]: 2.8 T total / 104 B activated, native vision, 1M context. Built on Kimi Delta Attention (KDA) + Attention Residuals (AttnRes), with Stable LatentMoE activating just 16 of 896 routed experts per token. Claims ~2.5× scaling efficiency over Kimi K2. Weights are QAT-trained from SFT onward with MXFP4 weights / MXFP8 activations. Full weights are released — the first open 3T-class model.

**What this means for the architect (ties to Ch3 Model, Ch10 Parallelism).** A model whose *total* size is 2.8 T but whose *activated* footprint is 104 B is the ultimate demonstration of Ch3’s total-vs-activated distinction and Ch10’s expert-parallelism doctrine. Serving it is not “2.8 T is too big to fit” — it is “fit and feed 104 B active, plus the KV for our context, while sharding and scheduling the expert fabric.” The decision surface is expert-parallel communication and fused cross-expert kernels, not raw capacity.

### A.5. Qwen3.8-Flash-Next — The Qwen4 Architecture Preview, Prioritizing Cost-Performance

**Decisive fact** [1P: HF Qwen/Qwen3.8-Flash-Next; GitHub QwenLM/Qwen3.8-Flash-Next]: 125 B total / 6 B activated, plus a separate 51 B of N-gram embedding parameters (offloadable to host memory). Native 262 K context (YaRN to 1M). Hybrid attention: Gated DeltaNet (GDN) + Qwen Sparse Attention (QSA), where QSA operates at the *micro-block* level; Gated Residual widens the residual stream; 512 MoE experts (10 routed + 1 shared per token); Muon + AdamW. Reported training cost ~1/9 of Qwen3.7-Plus.

**What this means for the architect** (ties to Ch3, Ch5 Model Selection, Ch16 TCO). Qwen’s N-gram embedding split is the most architecturally novel item: it shows that *parameter scaling and compute scaling can be decoupled* — a large embedding table can trade host-RAM bandwidth for on-GPU capacity when the compute budget is the binding constraint. This refines the model-selection calculus in Ch5: an architect comparing candidates now weighs not just total/active params but heterogeneous memory (on-GPU weights vs offloadable embeddings) and its serving overlap.

### A.6. GLM-5.3-Flash — First Natively Multimodal GLM, Efficiency-Branded

**Decisive fact** [1P: HF zai-org/GLM-5.3-Flash; arXiv 2602.15763]: 320 B total / 18 B activated, 1M native context, 131 K output. First natively multimodal model in the GLM-5 series (text/image/video). Hybrid **sparse+linear attention** (first in the series) with ~3× attention-compute and ~4.4× KV-cache reductions; Manifold-Constrained Hyper-Connections (mHC); W8A8 quantization with mixed-cache (INT8/FP8/BF16); Encode-Prefill-Decode serving. MIT weights. Z.ai positions it as approaching Claude Opus 4.8 on coding/agentive at ~1/10 the price.

**What this means for the architect** (ties to Ch7 Memory, Ch8 Compute, Ch19 Agentive). GLM’s coupling of hybrid attention with *native multimodality* shows the model border is remaking: “LLM” and “vision/audio model” are converging in a single weight set, which changes workload characterization (Ch4) — a multimodal prompt has a different token/KV profile than text-only. And mHC + attention-mixing being reported as *scaling-efficiency* results reinforces Ch8’s point that compute arithmetic is architecture-specific, not universal.

### A.7. What This Means for the Handbook’s Canonical Arithmetic

*Fig A.2 — Same message as Fig 7.1’s GQA line, taken to the frontier: the per-token KV and FLOP “constants” of Chapter 7/8 are not universal — they are LLaMA-style floors that hybrid attention re-architects down at the mechanism. Percentages are each model’s gain over its own predecessor (e.g. V3.2→V4, GLM-5.x steps), shown against the canonical 70B dense full-MHA reference so the floors stay comparable. An architect who re-derives them per candidate model avoids both over- (sizing a 7% KV model as if it were 100%) and under-provisioning. Values are vendor-published [1P], not yet independently*




FIGURE A.2. Fig A.2 — The KV constant and FLOP/token are a moving target: hybrid attention cuts both at the mechanism. Canonical 70B dense full-MHA = 100%; DeepSeek-V4-Flash  $\approx$  7% KV / 10% FLOP, V4-Pro  $\approx$  10% / 27%, GLM-5.3-Flash  $\approx$  23% KV /  $\sim$ 33% attention-compute [DERIVED from 1P claims: arXiv 2606.19348; HF zai-org]

reprofiled on a serving stack (see §5).

The book’s canonical scenario is a 70B dense FP16, ~9.2K-token RAG workload on 8×H100. The 2026 frontier does not invalidate it — it *refines* the interpretation an architect should carry:


- The **KV constant** (Ch7: ~2.5 MB/token FP16, ~1.3 MB 8-bit) is a LLaMA-style floor. Hybrid-attention models cut it 4–10× at the mechanism, so “context fits on N GPUs” must be re-derived per model, not looked up once. [1P]
- The **roofline** (Ch8) still holds; what changed is where a frontier model’s prefill/decode ride it. Vendor-reported FLOP/KV reductions are [1P] but not yet independently reprofiled — an architect should re-run Ch8’s measurement on the candidate model before betting capacity. [1P] → [VERIFY-candidate]
- The **MoE memory decision** (Ch10) is now dominated by active parameters + KV + expert-fabric comm, not total size. [1P]

None of this changes the book’s *mechanisms*; it changes the *constants* an architect plugs into them. This is exactly why we wrote the arithmetic as verifiable, first-principles reasoning rather than as a fixed catalog of numbers.

## A.8. Where 2026 Capability Actually Comes From (beyond Architecture)

A recurring error — easy to fall into after reading four new model cards — is to attribute frontier gains to the attention mechanism splashed across the front page. The industry is converging not just on *architectures* (MoE + efficient attention + very long context) but on a *multi-layered recipe* in which architecture is one ingredient, and arguably not the dominant one. A qualitative decomposition [VERIFY: qualitative, not a measured percentage]:

| Lever                            | Primary effect                                           |
|----------------------------------|----------------------------------------------------------|
| Post-training / RL / reasoning   | Capability (the largest single source today)             |
| Test-time compute                | Capability on hard tasks (more compute on hard problems) |
| Training data + optimization     | General capability                                       |
| MoE / parameter scaling          | Capacity per FLOP                                        |
| Attention innovation             | Context + efficiency (makes the above affordable)        |
| Inference / runtime optimization | Cost + latency                                           |
| Agent harness / tool use         | Real-world task completion                               |
| Multimodal architecture          | New modalities                                           |



From a single model to a whole system — the spine of this handbook (Appendix §6)

FIGURE A.3. Fig A.3 — The layered evolution of frontier systems, colour-coded by efficiency-led (amber) vs capability-led (blue) stages, from a single Transformer to the whole fleet (Appendix §6)

| Lever | Primary effect |
|-------|----------------|
|-------|----------------|

The most underestimated layer is the one the model card almost never shows: post-training. Two models with broadly similar architectures and pre-training can end up materially different when one went through mediocre post-training and the other excellent. The recipe — the same families every lab now runs — is what carries much of the qualitative gap:

- RLVR / reinforcement learning for reasoning
- synthetic reasoning data
- verifier-based training
- preference optimization
- curriculum learning
- tool-use training and agent trajectories
- coding-specific trajectories
- multimodal reasoning
- long-horizon interaction training

This is why architecture alone never tells an architect how good a frontier model is, and why the decision can no longer stop at “which model” — the marginal capability may live in the post-training and the runtime that wraps it.

Fig A.3 — From a single model to a whole system. The earlier stages (Transformer, Efficient, MoE) are largely efficiency-led\* — they make inference affordable and raise capacity-per-FLOP. The later stages (Reasoning, Test-Time Compute, Tool-Using, Agent, Agent System, Agent Fleet) are capability-led — they raise effective intelligence. This chain is the spine of the handbook: tokens

(Ch1–9), parallelism (Ch10), serving (Ch11–16), fleet & agents (Ch17–26). [INTERPRETATION — synthesis of the chapters’ arithmetic, not a single measurement]\*

This chain is the spine of this handbook — from a single model’s tokens (Ch1–9) to parallelism (Ch10), serving (Ch11–16), and finally the fleet of specialised models and agents (Ch17–26). It also explains why two models with broadly similar base architectures can have very different practical intelligence (post-training/reasoning), and why an agent runtime can lift real-world performance without changing model weights at all: **model intelligence  $\neq$  system intelligence anymore.**

### A.9. The Frontier Question: Where Should Intelligence Live?

The most useful lens for an AI Solution Architect in 2026 is not “which attention mechanism does the model card list?” but *where the intelligence and the compute budget should be placed*:

*Fig A.4 — The placement ladder. Every host layer that can carry intelligence is a place the architect may choose to push capability or cost; this book gives the arithmetic and decision framework for each (Ch3/7/8 weights & attention, Ch10 MoE routing, Ch19/21 post-training, Ch8/20 test-time, Ch19/24 tools, Ch17–20 runtime, Ch18/20 fleet). [INTERPRETATION/ILLUSTRATIVE per Appendix A’s three-layer rule]*

Every host layer maps to chapters in this book:

- **Inside weights** (dense capacity) — Ch3 Model, Ch8 Compute
- **Inside attention** (context / efficiency) — Ch7 Memory, Ch8 Compute
- **Inside MoE routing** (capacity per FLOP) — Ch10 Parallelism
- **Inside post-training** (reasoning, RL, verification) — Ch19 Agentic, Ch21 AI Factory
- **Inside test-time search** (adaptive inference, more compute on hard tasks) — Ch8 Compute, Ch20 Fleet-Level Optimization
- **Inside tools** (compute allocated via a harness) — Ch19 Agentic, Ch24 Red/Green Team
- **Inside the agent runtime** (persistent state, runbooks, execution control) — Ch17–20, Ch24
- **Across a fleet of specialised models** — Ch18 Operating Multiple Models, Ch20 Fleet-Level Optimization

Deciding where intelligence lives — rather than which attention variant a model card lists — is the actual job of an AI Solution Architect. The frontier of 2026 makes that explicit: DeepSeek-V4-Flash vs Qwen3.8-Flash vs a fleet of specialised agents on a 2×DGX cluster are no longer merely *model* choices; they are *placement-of-intelligence* choices whose economics this handbook’s chapters give the tools to evaluate.

### A.10. One-Hand Sources (all [1P])

- DeepSeek-V4: technical report arXiv:2606.19348; model card <https://huggingface.co/deepseek-ai/DeepSeek-V4-Pro>; org <https://github.com/deepseek-ai>

## Where Should Intelligence Live?

*The 2026 frontier question for an AI Solution Architect (Appendix §7)*




FIGURE A.4. Fig A.4 — Where should intelligence live? Eight host layers, from inside the weights up to across a fleet of specialised models, each mapped to the chapters of this handbook that give the tooling to evaluate it (Appendix §7)

- Kimi K3: technical report arXiv:2607.24653; model card <https://huggingface.co/moonshotai/Kimi-K3>; <https://github.com/moonshotai>
- Qwen3.8-Flash-Next: technical report [https://github.com/QwenLM/Qwen3.8-Flash-Next/blob/main/tech\\\_report.pdf](https://github.com/QwenLM/Qwen3.8-Flash-Next/blob/main/tech\_report.pdf); model card <https://huggingface.co/Qwen/Qwen3.8-Flash-Next>
- GLM-5.3-Flash: technical report arXiv:2602.15763; model card <https://huggingface.co/zai-org/GLM-5.3-Flash>; <https://github.com/zai-org/GLM-5>

*Compiled 2026-08-30 from the models' official technical reports and first-party model cards; all parameters/claims are vendor-reported [1P] and per this handbook's discipline carry a verify-on-own-hardware caveat.*