

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Q1.1 tokens

Q1.2 decoder: what's active

Q1.3 six-dim fingerprint

↓ per-token which reads & feeds decision tier

Part II — The Workload

Q1.4 capacity: Greenoodput

Q1.5 which model family

↓ measure & gate

Part III — The System

Q1.6 KV/cache: no fork to connect

Q1.7 partitioning, cache, P-D

↓ GPUs ↓ servering stack

Part IV — The Architecture

Q1.8 can't tie an LLM to a two-stage bmark

Q1.9 compiler: too many MFUes

↓ synthesis & deployment / reject

Part V — The Fleet

Q1.10 H := $\hat{H} + \hat{C} + \hat{V}$ / KV per cost

Q1.11 fleet size: multi-tenant routing

↓ capacity

Part VI — The Architect

Q1.12 both have a KV cache for green probe

Q1.13 AC H12: what's a 2026 strategy

↓ where to cope the possible design