

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

↓ per-token price is not fixed

↓ what to characterize

Part II – The Workload

capacitance
↓ which is a measure & gate

Part III – The System

↓ fit / compile memory-bound

↓ GPUs ↓ **sharing stack**

Part IV — The Architecture

↓ synthesis deployment / reject

↓ fleet saturation/cloud/API

Part V — The Fleet

↓ fleet size multi-dept routing capacity

file → n2 + rollback
↓ QPS vs. host rollout

Part VI – The Architect

both ways → design problem
↓ where to place the probe

↓ record the determined anchors