

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

1.1.1

token = cost

→ price & KV floor

1.1.2

decode vs prefill

→ which binds

1.1.3

total vs active

→ model tier

1.1.4

six-dim fingerprint

→ to characterize

Part II — The Workload

1.2.1

capability screen

→ model family

1.2.2

SLOs + goodput

→ measure & gate

Part III — The System

1.3.1

KV/token × ctx

→ fit & concurrency

1.3.2

roofline vs ridge

compute/mem-bound

1.3.3

interconnect

→ comm cost

1.3.4

parallelism

→ GPUs & split

1.3.5

batching/cache/P-D

→ serving stack

Part IV — The Architecture

1.4.1

candidate arch.

→ synthesis & score

1.4.2

tier ladder

→ deploy tier

1.4.3

two-stage bmark

→ accept / reject

1.4.4

congestion & MFU

→ fleet saturated?

1.4.5

TCO · 3 modes

→ host / cloud / API

Part V — The Fleet

1.5.1

$H = \lambda L / (C \cdot \text{util})$

→ fleet size

1.5.2

route by cap/cost

multi-model routing

1.5.3

KV per agent turn

→ depth × capacity

1.5.4

fleet req/s

→ QPS vs host

1.5.5

gate + rollback

→ safe rollout

Part VI — The Architect

1.6.1

bottleneck diverge

→ where to invest

1.6.2

vague ask → bounds

→ scope it

1.6.3

red/green probe

→ stress-test design

1.6.4

ADR template

→ record decision

1.6.5

workload → strategy

→ pick pattern

1.6.6

2026 constants

→ re-derive anchors