

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)  
each chapter supplies a derived quantity, which feeds a decision**

### Part I — The Token

- 📊 token = cost → price & KV floor
- 📊 decode vs prefill → which binds
- 📊 total vs active → model tier
- 📊 six-dim fingerprint → to characterize

### Part II — The Workload

- 📊 capability screen → model family
- 📊 SLOs + goodput → measure & gate

### Part III — The System

- 📊 KV/token × ctx → fit & concurrency
- 📊 roofline vs ridge → compute/mem-bound
- 📊 interconnect → comm cost
- 📊 parallelism → GPUs & split
- 📊 batching/cache/P-D → serving stack

### Part IV — The Architecture

- 📊 candidate arch. → synthesis & score
- 📊 tier ladder → deploy tier
- 📊 two-stage bmark → accept / reject
- 📊 congestion & MFU → fleet saturated?
- 📊 TCO · 3 modes → host / cloud / API

### Part V — The Fleet

- 📊  $H = \lambda L / (C \cdot \text{util})$  → fleet size
- 📊 route by cap/cost → multi-model routing
- 📊 KV per agent turn → depth × capacity
- 📊 fleet req/s → QPS vs host
- 📊 gate + rollback → safe rollout

### Part VI — The Architect

- 📊 bottleneck diverge → where to invest
- 📊 vague ask → bounds → scope it
- 📊 red/green probe → stress-test design
- 📊 ADR template → record decision
- 📊 workload → strategy → pick pattern
- 📊 2026 constants → re-derive anchors