

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

↓ per-token price for full resolution

↓ what to characterize

Part II — The Workload

↓ which model families

Part III — The System

↓ fit / column / column-bound

↓ GPUs ↓ **swapping stack**

Part IV — The Architecture

↓ synthesis deployment, reject

↓ fleet saturation/cloud/API

Part V — The Fleet

$H = \frac{1}{n} \sum_{i=1}^n \frac{1}{K_i} \text{Vapor cost}$
 ↓ fleet size ↓ multi-depth routing

file → log + rollback
↓ QPS vs. host rollout

Part VI — The Architect

both ways → green probe
↓ where to open the possible design

AC H12 w/ Apr 2026 strategy