

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Q17 tokens-decode: what's active
↓ per-token which reads & feeds decision tier

Q17 six-dim fingerprint
↓ what to characterize

Part II — The Workload

Q17 capacity: by Greenoodput
↓ which model family measure & gate

Part III — The System

Q17 KV/cache: no KV-cache connect
↓ fit / compute: owns memory-bound

Q17 partitioning, cache, P-D
↓ GPUs ↓ serving stack

Part IV — The Architecture

Q17 can't tie multi-stage bmark
↓ synthesis: deployment / reject

Q17 compute: TOOS MFules
↓ fleet saturation: cloud/API

Part V — The Fleet

Q17 H := $\hat{H} + \hat{C} + \hat{V}$ / KV-cache / agent turn
↓ fleet size: multi-indep. routing capacity

Q17 fleet: $\hat{H}$ / gate + rollback
↓ QPS vs. host rollout

Part VI — The Architect

Q17 both: $\hat{H}$ / $\hat{C}$ / $\hat{V}$ / KV-cache / green probe
↓ where to cope: the possible design

Q17 AC $\hat{H}$ / $\hat{C}$ / $\hat{V}$ / 2026 strategy
↓ record the decision: pattern anchors