

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Ch1 token, Ch2 decode, Ch3 total vs active
↓ per-token price & KV model dissection tier

Ch4 six-dim fingerprint
↓ what to characterize

Part II — The Workload

Ch5 capacity, Ch6 SLOs + goodput
↓ which model family measure & gate

Part III — The System

Ch7 KV/to, Ch8 xrofit, Ch9 siting connect
↓ fit / concurrency compute vs memory bound

Ch10 para, Ch11 batching, cache, P-D
↓ GPUs / whisksys stack

Part IV — The Architecture

Ch12 cand, Ch13 tier, Ch14 two-stage bmark
↓ synthesis & deployment accept / reject

Ch15 cong, Ch16 TCOM, Ch17 nodes
↓ fleet saturation, cloud/API

Part V — The Fleet

Ch18 H = , Ch19 mobility, Ch20 ap/ops
↓ fleet size, multi-model deployment capacity

Ch21 fleet, Ch22 gate + rollback
↓ QPS vs host safe rollout

Part VI — The Architect

Ch23 bottl, Ch24 dig, Ch25 red, Ch26 ops probe
↓ where to locate the probe, best design

Ch27 ADR, Ch28 work, Ch29 ApA > 2026
↓ record the pickles, parse & derive anchors