

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

↓ per-token price & resource selection tier

ch4 six-dim fingerprint

↓ what to characterize

Part II — The Workload

↓ which model train

Part III – The System

↓ fit / concurrent-compute-bound

parameters: matching, cache, P-D
↓ GPUs / which poolstack

Part IV – The Architecture

↓ synthesis is dependent on α / reject

Ch15 con Ch16 TCQ BIF modes

↓ fleet saturation? cloud/API

Part V — The Fleet

↓ fleet size multi-modal trucking capacity

floor gate + rollback

↓ OPS-vs-hosts rollout

Part VI — The Architect

↓ where to invest the press test design

ADF: [workflows](#) → [patterns](#) → [record the code](#) → [anchors](#)