

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)  
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

|     |                   |     |                   |     |                 |     |                           |
|-----|-------------------|-----|-------------------|-----|-----------------|-----|---------------------------|
| Ch1 | token = cost unit | Ch2 | decode vs prefill | Ch3 | total vs active | Ch4 | six-dimension fingerprint |
|-----|-------------------|-----|-------------------|-----|-----------------|-----|---------------------------|

- ↳ per-token price & KV floor
- ↳ which resource binds
- ↳ model selection tier
- ↳ what to characterize

Part II — The Workload

|     |                   |     |                |
|-----|-------------------|-----|----------------|
| Ch5 | capability screen | Ch6 | SLOs + goodput |
|-----|-------------------|-----|----------------|

- ↳ which model family
- ↳ what to measure & gate

Part III — The System

|     |                    |     |                       |     |                             |      |                         |      |                          |
|-----|--------------------|-----|-----------------------|-----|-----------------------------|------|-------------------------|------|--------------------------|
| Ch7 | KV/token × context | Ch8 | roofline: AI vs ridge | Ch9 | interconnect (time-vs-size) | Ch10 | parallelism composition | Ch11 | batching / caching / P-D |
|-----|--------------------|-----|-----------------------|-----|-----------------------------|------|-------------------------|------|--------------------------|

- ↳ does it fit / concurrency
- ↳ compute- vs memory-bound
- ↳ comm cost
- ↳ how many GPUs / which split
- ↳ serving stack

Part IV — The Architecture

|      |                         |      |                       |      |                     |      |                  |      |                 |
|------|-------------------------|------|-----------------------|------|---------------------|------|------------------|------|-----------------|
| Ch12 | candidate architectures | Ch13 | reference tier ladder | Ch14 | two-stage benchmark | Ch15 | congestion & MFU | Ch16 | TCO three modes |
|------|-------------------------|------|-----------------------|------|---------------------|------|------------------|------|-----------------|

- ↳ synthesis & score
- ↳ deployment tier
- ↳ accept / reject
- ↳ is the fleet saturated
- ↳ self-host vs cloud vs API

Part V — The Fleet

|      |                                               |      |                          |      |                         |      |                       |      |                           |
|------|-----------------------------------------------|------|--------------------------|------|-------------------------|------|-----------------------|------|---------------------------|
| Ch17 | hosts $H = \lambda L / (C \cdot \text{util})$ | Ch18 | route by capability/cost | Ch19 | token/KV per agent turn | Ch20 | fleet req/s @3.7/host | Ch21 | promote via gate+rollback |
|------|-----------------------------------------------|------|--------------------------|------|-------------------------|------|-----------------------|------|---------------------------|

- ↳ fleet size
- ↳ multi-model routing
- ↳ depth × capacity
- ↳ QPS-vs-host
- ↳ safe rollout

Part VI — The Architect

|      |                       |      |                    |      |                 |      |              |      |                            |     |                             |
|------|-----------------------|------|--------------------|------|-----------------|------|--------------|------|----------------------------|-----|-----------------------------|
| Ch22 | bottleneck divergence | Ch23 | vague ask → bounds | Ch24 | red/green probe | Ch25 | ADR template | Ch26 | workload → strategy matrix | ApA | 2026 constants vs canonical |
|------|-----------------------|------|--------------------|------|-----------------|------|--------------|------|----------------------------|-----|-----------------------------|

- ↳ where to invest
- ↳ scope the problem
- ↳ stress-test design
- ↳ record the decision
- ↳ pick the pattern
- ↳ re-derive anchors