

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Ch1 token = cost unit **Ch2** decode vs prefill **Ch3** total vs active **Ch4** six-dimension fingerprint

↳ per-token price & KV floor ↳ which resource binds ↳ model selection tier ↳ what to characterize

Part II — The Workload

Ch5 capability screen **Ch6** SLOs + goodput

↳ which model family ↳ what to measure & gate

Part III — The System

Ch7 KV/token × context **Ch8** roofline: AI vs ridge **Ch9** interconnect (time-vs-size) **Ch10** parallelism composition **Ch11** batching / caching / P-D

↳ does it fit / concurrency ↳ compute- vs memory-bound ↳ comm cost ↳ how many GPUs / which split ↳ serving stack

Part IV — The Architecture

Ch12 candidate architectures **Ch13** reference tier ladder **Ch14** two-stage benchmark **Ch15** congestion & MFU **Ch16** TCO three modes

↳ synthesis & score ↳ deployment tier ↳ accept / reject ↳ is the fleet saturated ↳ self-host vs cloud vs API

Part V — The Fleet

Ch17 hosts $H = \lambda L / (C \cdot \text{util})$ **Ch18** route by capability/cost **Ch19** token/KV per agent turn **Ch20** fleet req/s @3.7/host **Ch21** promote via gate+rollback

↳ fleet size ↳ multi-model routing ↳ depth × capacity ↳ QPS-vs-host ↳ safe rollout

Part VI — The Architect

Ch22 bottleneck divergence **Ch23** vague ask → bounds **Ch24** red/green probe **Ch25** ADR template **Ch26** workload → strategy matrix **EpA** 2026 constants vs canonical

↳ where to invest ↳ scope the problem ↳ stress-test design ↳ record the decision ↳ pick the pattern ↳ re-derive anchors