

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

- 1.1.1 token = cost → price & KV floor
- 1.1.2 decode vs prefill → which binds
- 1.1.3 total vs active → model tier
- 1.1.4 six-dim fingerprint → to characterize

Part II — The Workload

- 1.2.1 capability screen → model family
- 1.2.2 SLOs + goodput → measure & gate

Part III — The System

- 1.3.1 KV/token × ctx → fit & concurrency
- 1.3.2 roofline vs ridge compute/mem-bound
- 1.3.3 interconnect → comm cost
- 1.3.4 parallelism → GPUs & split
- 1.3.5 batching/cache/P-D → serving stack

Part IV — The Architecture

- 1.4.2 candidate arch. → synthesis & score
- 1.4.3 tier ladder → deploy tier
- 1.4.4 two-stage bmark → accept / reject
- 1.4.5 congestion & MFU → fleet saturated?
- 1.4.6 TCO · 3 modes → host / cloud / API

Part V — The Fleet

- 1.5.7 $H = \lambda L / (C \cdot \text{util})$ → fleet size
- 1.5.8 route by cap/cost → multi-model routing
- 1.5.9 KV per agent turn → depth × capacity
- 1.5.10 fleet req/s → QPS vs host
- 1.5.11 gate + rollback → safe rollout

Part VI — The Architect

- 1.6.2 bottleneck diverge → where to invest
- 1.6.3 vague ask → bounds → scope it
- 1.6.4 red/green probe → stress-test design
- 1.6.5 ADR template → record decision
- 1.6.6 workload → strategy → pick pattern
- 1.6.7 2026 constants → re-derive anchors