

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Ch1 token = cost unit Ch2 decode vs prefill Ch3 total vs active Ch4 six-dimension fingerprint

↳ per-token price & KV floor ↳ which resource binds ↳ model selection tier ↳ what to characterize

Part II — The Workload

Ch5 capability screen Ch6 SLOs + goodput

↳ which model family ↳ what to measure & gate

Part III — The System

Ch7 KV/token × context Ch8 roofline: AI vs ridge Ch9 interconnect (time-vs-size) Ch10 parallelism composition Ch11 batching / caching / P-D

↳ does it fit / concurrency ↳ compute- vs memory-bound ↳ comm cost ↳ how many GPUs / which split ↳ serving stack

Part IV — The Architecture

Ch12 candidate architectures Ch13 reference tier ladder Ch14 two-stage benchmarking Ch15 congestion & MFU Ch16 TCO three modes

↳ synthesis & score ↳ deployment tier ↳ accept / reject ↳ is the fleet saturated ↳ self-host vs cloud vs API

Part V — The Fleet

Ch17 hosts $H = \lambda L / (C \cdot util)$ Ch18 route by capability/cost Ch19 token/KV per agent turn Ch20 fleet req/s @3.7/host Ch21 promote via gate+rollback

↳ fleet size ↳ multi-model routing ↳ depth × capacity ↳ QPS-vs-host ↳ safe rollout

Part VI — The Architect

Ch22 bottleneck divergence Ch23 vague ask → bounds Ch24 red/green probe Ch25 ADR template Ch26 workload → strategy matrix Ch27 2026 constants vs canonical

↳ where to invest ↳ scope the problem ↳ stress-test design ↳ record the decision ↳ pick the pattern ↳ re-derive anchors