

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

📊 token = cost → price & KV floor

📊 decode vs prefill → which binds

📊 total vs active → model tier

📊 six-dim fingerprint → to characterize

Part II — The Workload

📊 capability screen → model family

📊 SLOs + goodput → measure & gate

Part III — The System

📊 KV/token × ctx → fit & concurrency

📊 roofline vs ridge → compute/mem-bound

📊 interconnect → comm cost

📊 parallelism → GPUs & split

📊 batching/cache/P-D → serving stack

Part IV — The Architecture

📊 candidate arch. → synthesis & score

📊 tier ladder → deploy tier

📊 two-stage bmark → accept / reject

📊 congestion & MFU → fleet saturated?

📊 TCO · 3 modes → host / cloud / API

Part V — The Fleet

📊 $H = \lambda L / (C \cdot \text{util})$ → fleet size

📊 route by cap/cost → multi-model routing

📊 KV per agent turn → depth × capacity

📊 fleet req/s → QPS vs host

📊 gate + rollback → safe rollout

Part VI — The Architect

📊 bottleneck diverge → where to invest

📊 vague ask → bounds → scope it

📊 red/green probe → stress-test design

📊 ADR template → record decision

📊 workload → strategy → pick pattern

📊 2026 constants → re-derive anchors