

Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision

Part I — The Token

Q17

tokens = doc + 16 * address file active

↓

per-token which reads 64KB from 16KB read buffer

Q17

six-dim fingerprint

↓

what to characterize

Part II — The Workload

Q17

capacity by Greenoodput

↓

which model family measure & gate

Part III — The System

Q17

KV/cache on 16KB token connect

↓

fit / compute vs memory-bound

Q17

partitioning, cache, P-D

↓

GPUs / servering stack

Part IV — The Architecture

Q17

can't tie a multi-stage bmark

↓

synthesis deployment / reject

Q17

compute TOOS MFUes

↓

fleet saturated cloud/API

Part V — The Fleet

Q17

$H := \frac{1}{N} \sum_{i=1}^N (KV)_{i, \text{avg cost}}$

↓

fleet size multi-tenancy routing capacity

Q17

fleet waste + rollback

↓

QPS vs host rollout

Part VI — The Architect

Q17

both have 16KB KV cache for probe

↓

where to cope with the possible design

Q17

AC H17 work to 2026 strategy

↓

record the design pattern